

System Overview

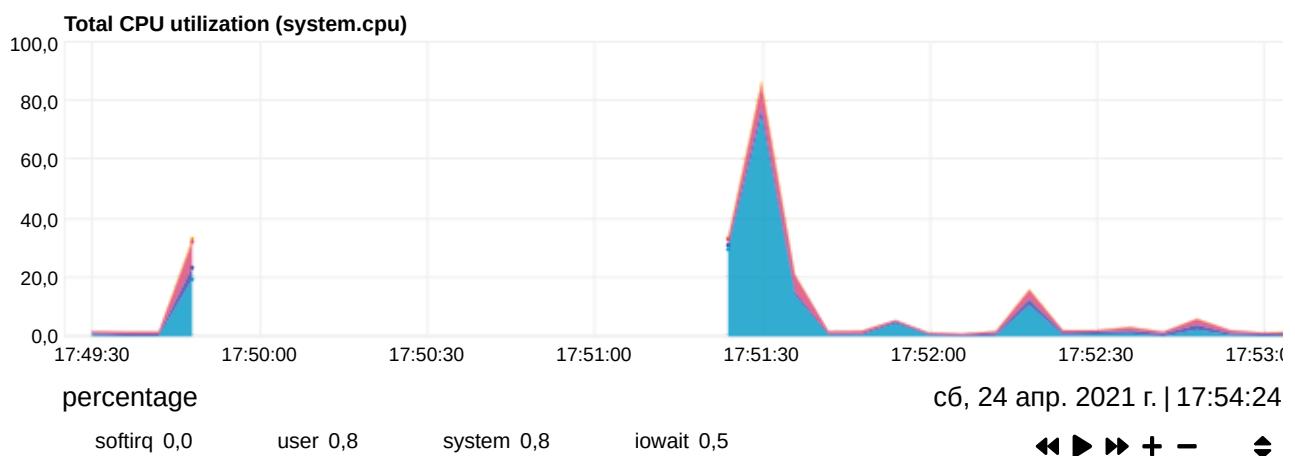
Overview of the key system metrics.

CPU

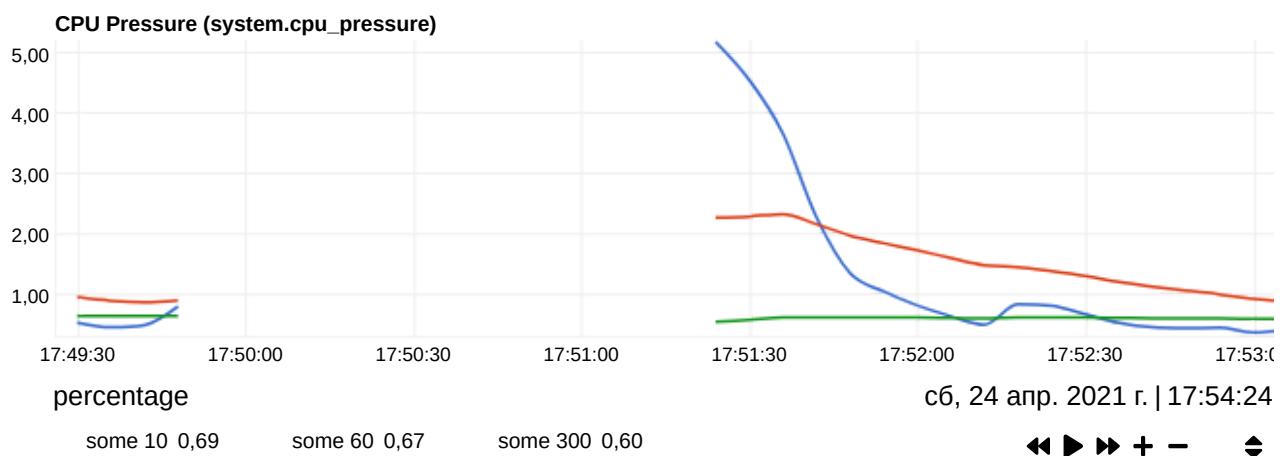
Total CPU utilization (all cores). 100% here means there is no CPU idle time at all. You can get per core usage at the CPUs section and per application usage at the Applications Monitoring section.

Keep an eye on **iowait** (2,5%). If it is constantly high, your disks are a bottleneck and they slow your system down.

An important metric worth monitoring, is **softirq** (0,00%). A constantly high percentage of softirq may indicate network driver issues.

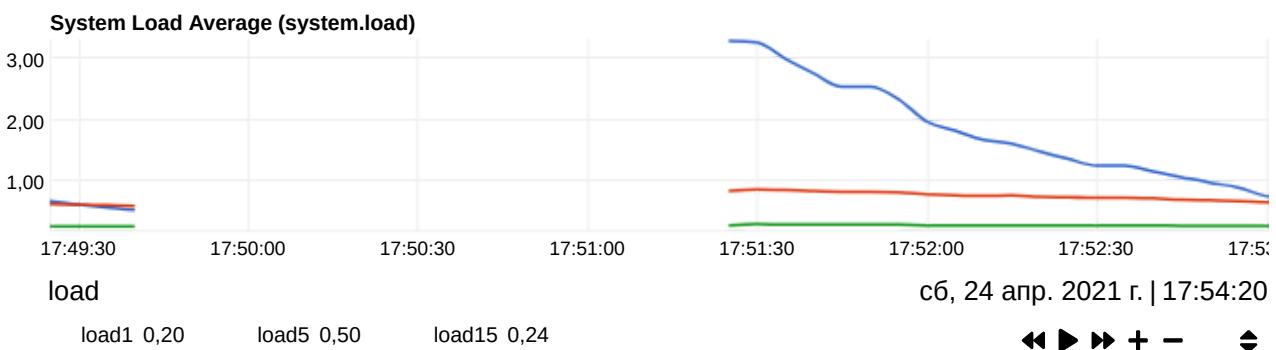


Pressure Stall Information (<https://www.kernel.org/doc/html/latest/accounting/psi.html>) identifies and quantifies the disruptions caused by resource contentions. The "some" line indicates the share of time in which at least **some** tasks are stalled on CPU. The ratios (in %) are tracked as recent trends over 10-, 60-, and 300-second windows.



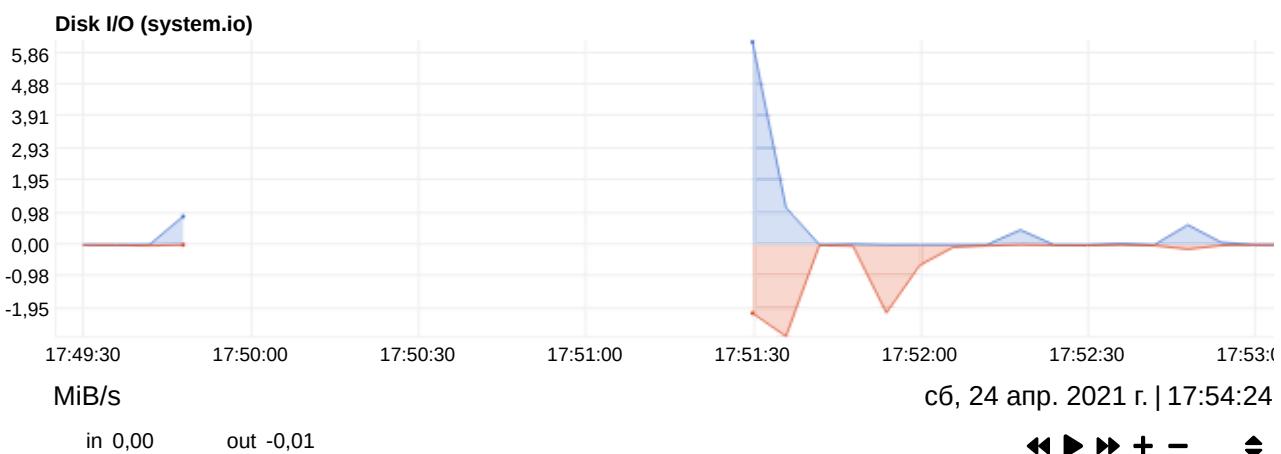
load

Current system load, i.e. the number of processes using CPU or waiting for system resources (usually CPU and disk). The 3 metrics refer to 1, 5 and 15 minute averages. The system calculates this once every 5 seconds. For more information check this wikipedia article ([https://en.wikipedia.org/wiki/Load_\(computing\)](https://en.wikipedia.org/wiki/Load_(computing))).



disk

Total Disk I/O, for all physical disks. You can get detailed information about each disk at the Disks section and per application Disk usage at the Applications Monitoring section. Physical are all the disks that are listed in `/sys/block`, but do not exist in `/sys/devices/virtual/block`.

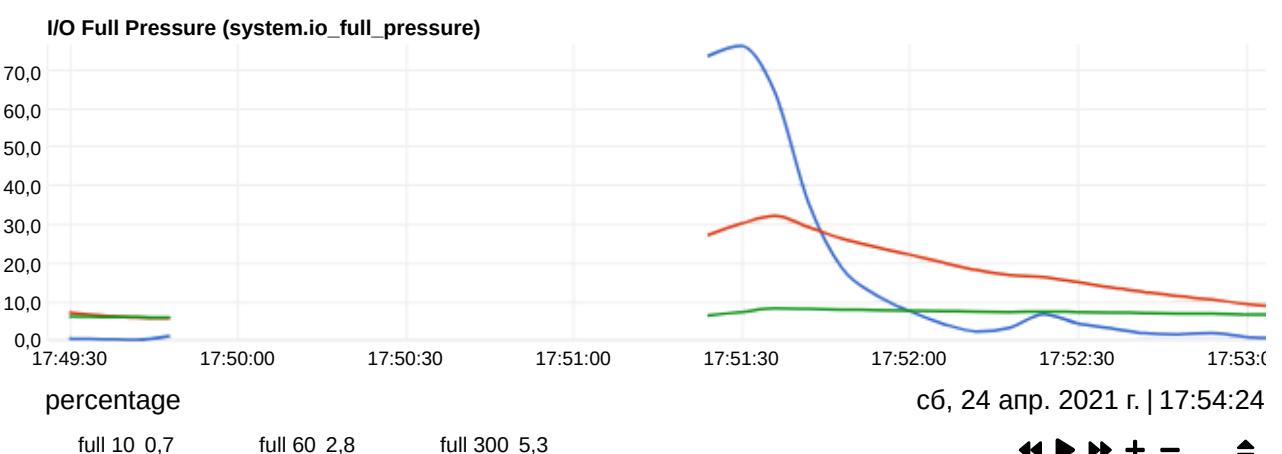
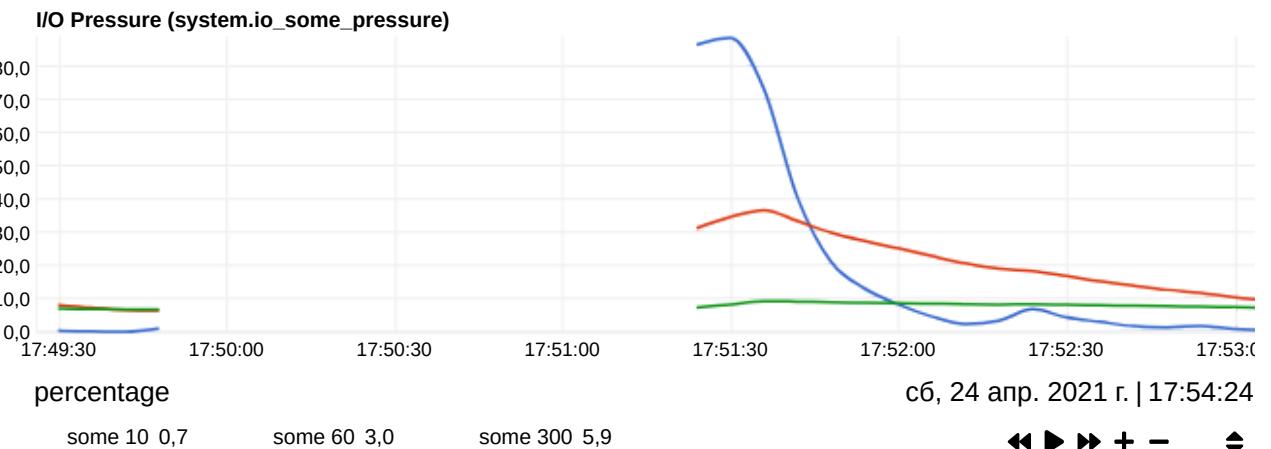


Memory paged from/to disk. This is usually the total disk I/O of the system.

Memory Paged from/to disk (system.pgpgio)



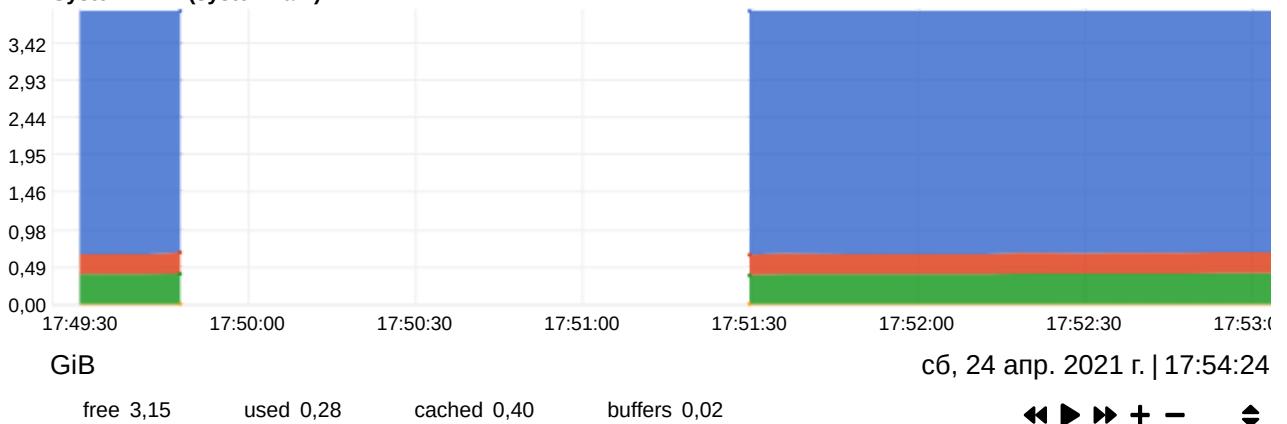
Pressure Stall Information (<https://www.kernel.org/doc/html/latest/accounting/psi.html>) identifies and quantifies the disruptions caused by resource contentions. The "some" line indicates the share of time in which at least **some** tasks are stalled on I/O. The "full" line indicates the share of time in which **all non-idle** tasks are stalled on I/O simultaneously. In this state actual CPU cycles are going to waste, and a workload that spends extended time in this state is considered to be thrashing. The ratios (in %) are tracked as recent trends over 10-, 60-, and 300-second windows.



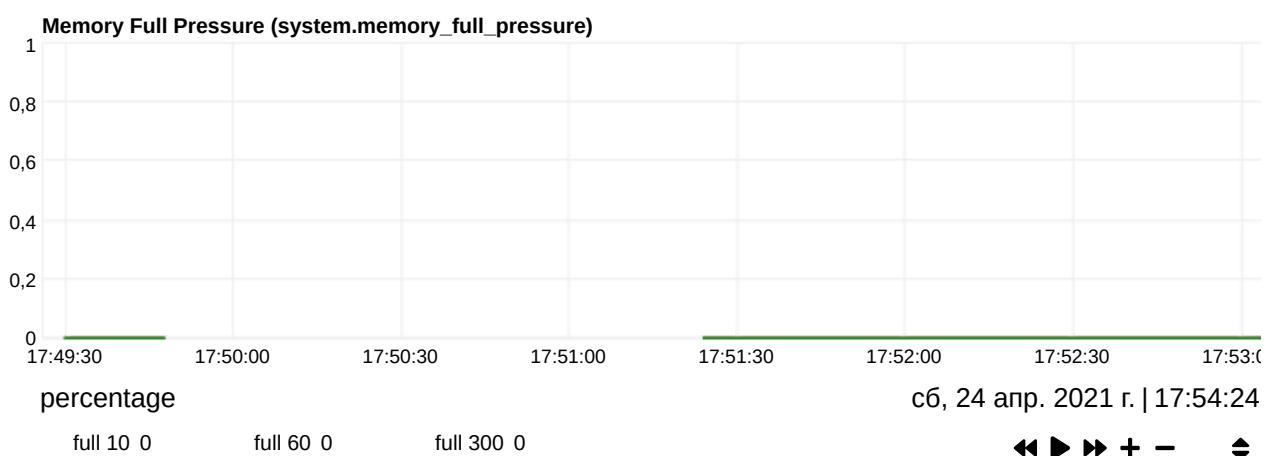
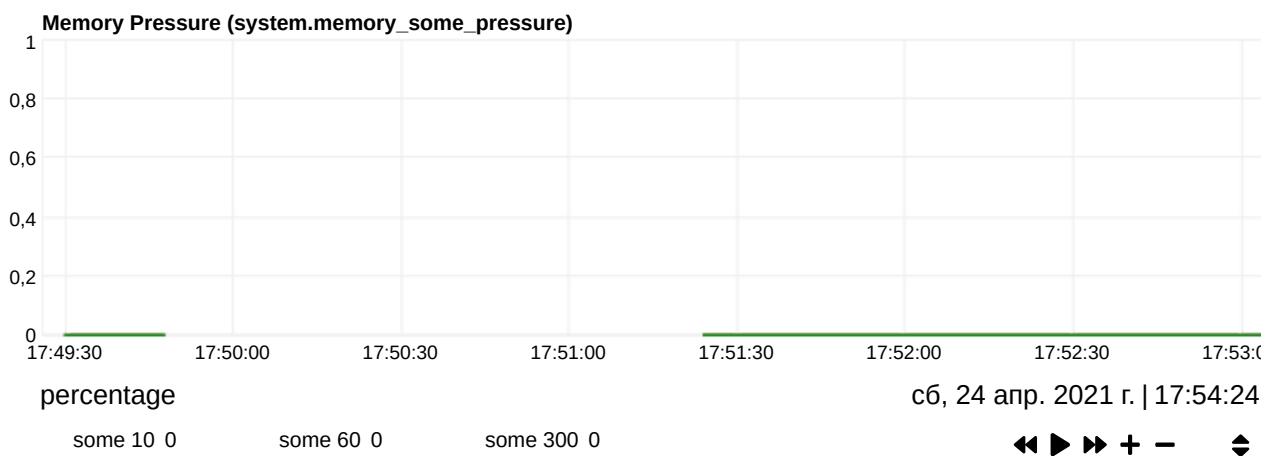
ram

System Random Access Memory (i.e. physical memory) usage.

System RAM (system.ram)



Pressure Stall Information (<https://www.kernel.org/doc/html/latest/accounting/psi.html>) identifies and quantifies the disruptions caused by resource contentions. The "some" line indicates the share of time in which at least **some** tasks are stalled on memory. The "full" line indicates the share of time in which **all non-idle** tasks are stalled on memory simultaneously. In this state actual CPU cycles are going to waste, and a workload that spends extended time in this state is considered to be thrashing. The ratios (in %) are tracked as recent trends over 10-, 60-, and 300-second windows.



Swap

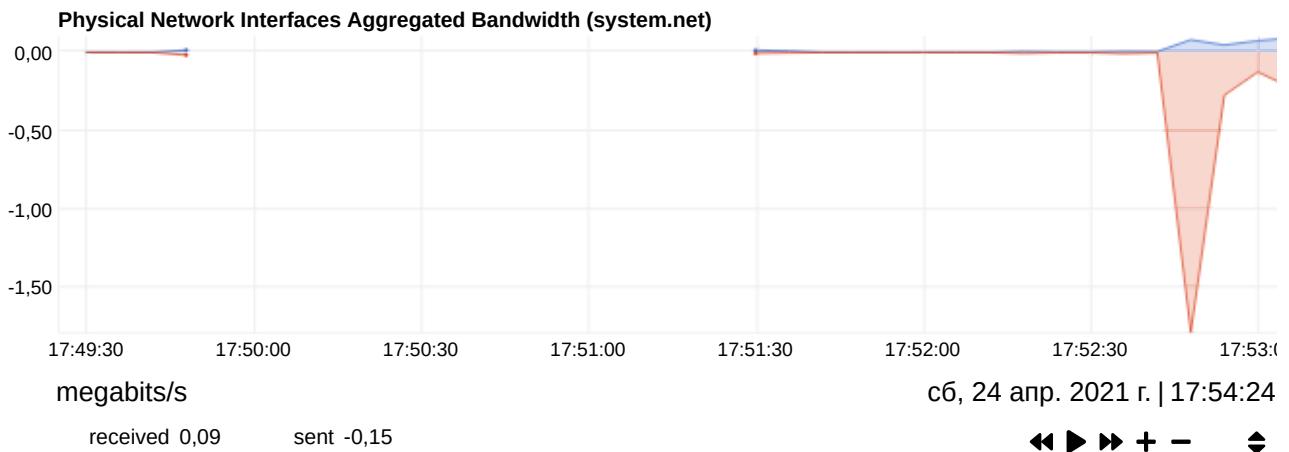
System swap memory usage. Swap space is used when the amount of physical memory (RAM) is full. When the system needs more memory resources and the RAM is full, inactive pages in memory are moved to the swap space (usually a disk, a disk partition or a file).

System Swap (system.swap)

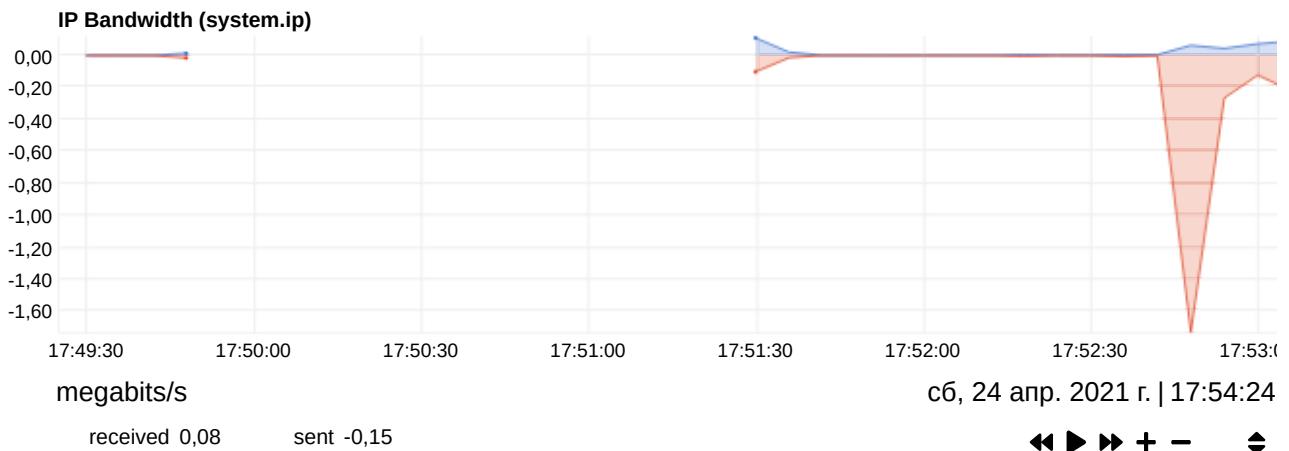


network

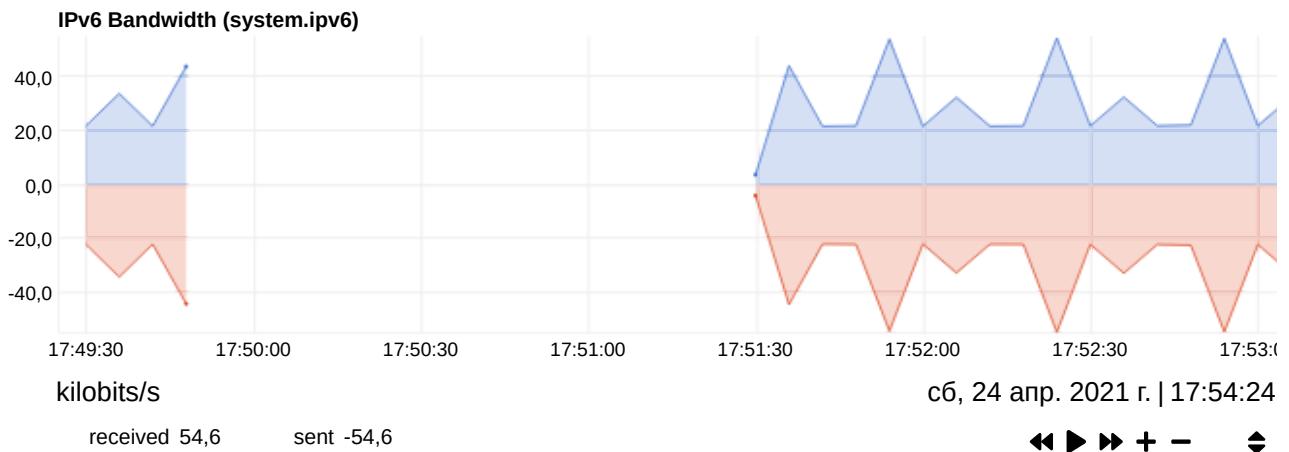
Total bandwidth of all physical network interfaces. This does not include `lo`, VPNs, network bridges, IFB devices, bond interfaces, etc. Only the bandwidth of physical network interfaces is aggregated.
 Physical are all the network interfaces that are listed in `/proc/net/dev`, but do not exist in `/sys/devices/virtual/net`.



Total IP traffic in the system.



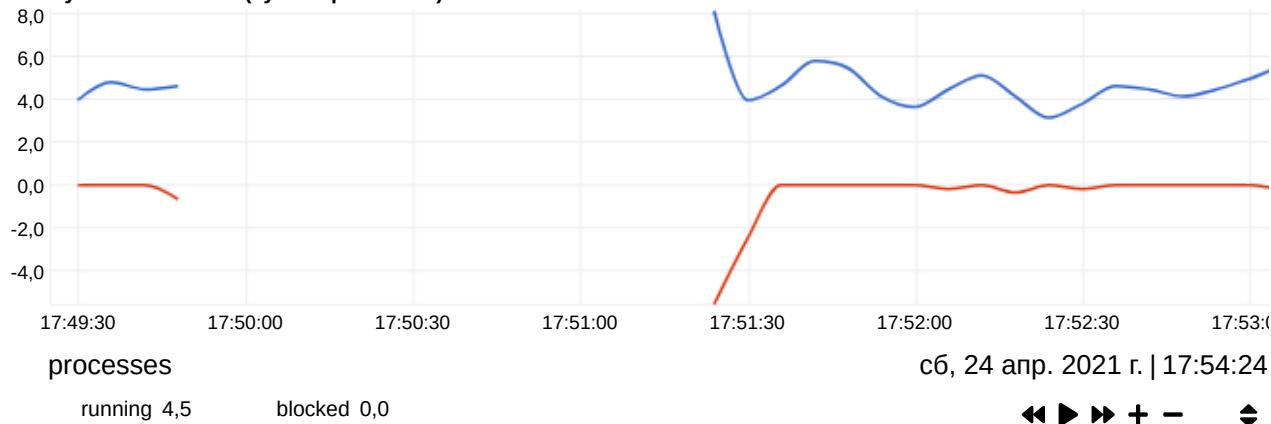
Total IPv6 Traffic.



processes

System processes. **Running** are the processes in the CPU. **Blocked** are processes that are willing to enter the CPU, but they cannot, e.g. because they wait for disk activity.

System Processes (system.processes)



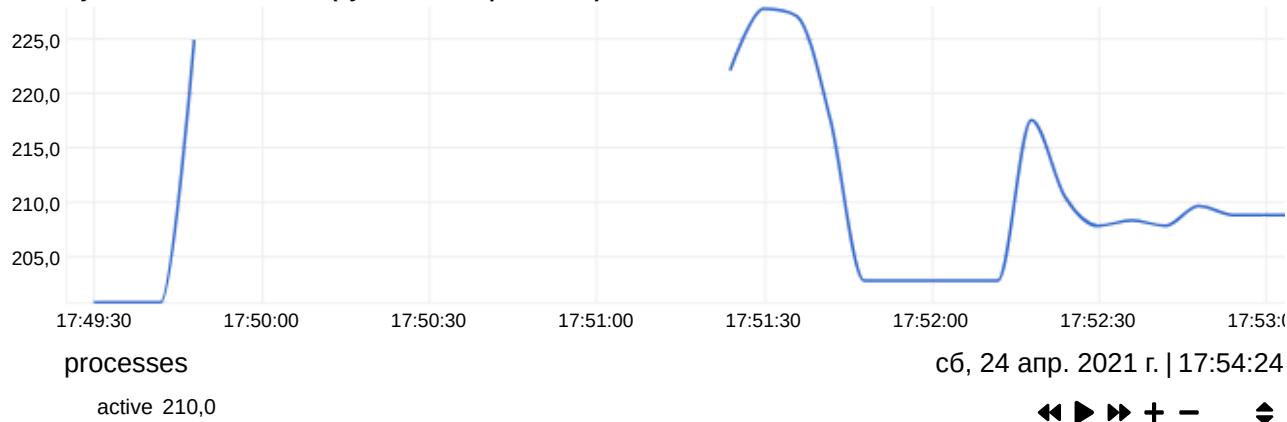
Number of new processes created.

Started Processes (system.forks)

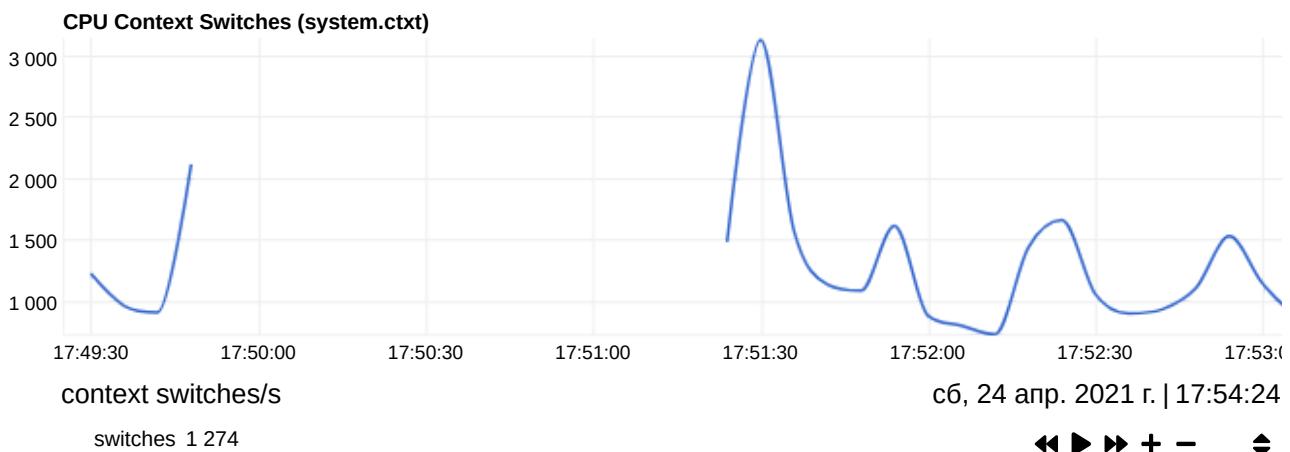


All system processes.

System Active Processes (system.active_processes)

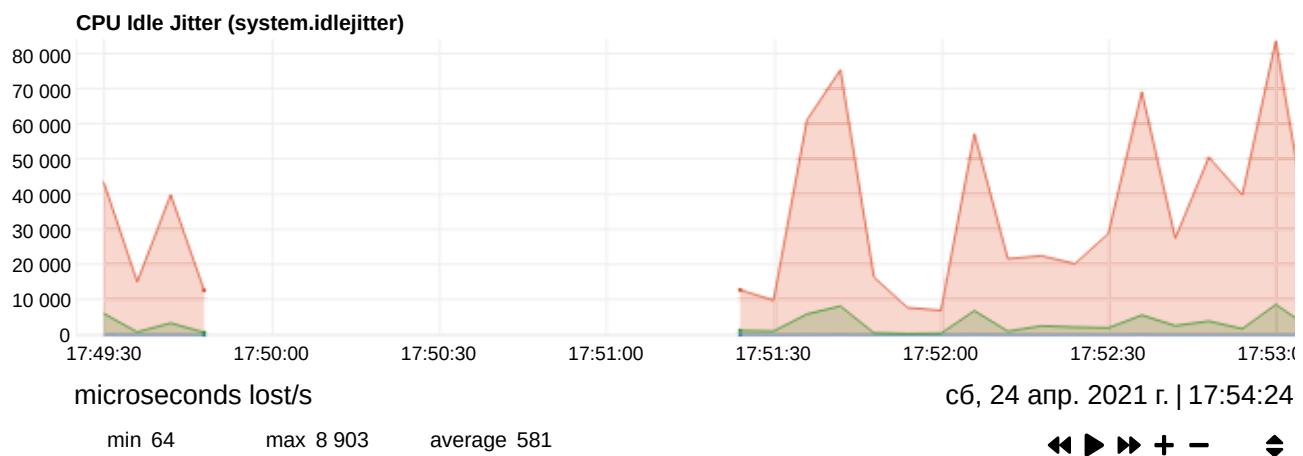


Context Switches (https://en.wikipedia.org/wiki/Context_switch), is the switching of the CPU from one process, task or thread to another. If there are many processes or threads willing to execute and very few CPU cores available to handle them, the system is making more context switching to balance the CPU resources among them. The whole process is computationally intensive. The more the context switches, the slower the system gets.



idlejitter

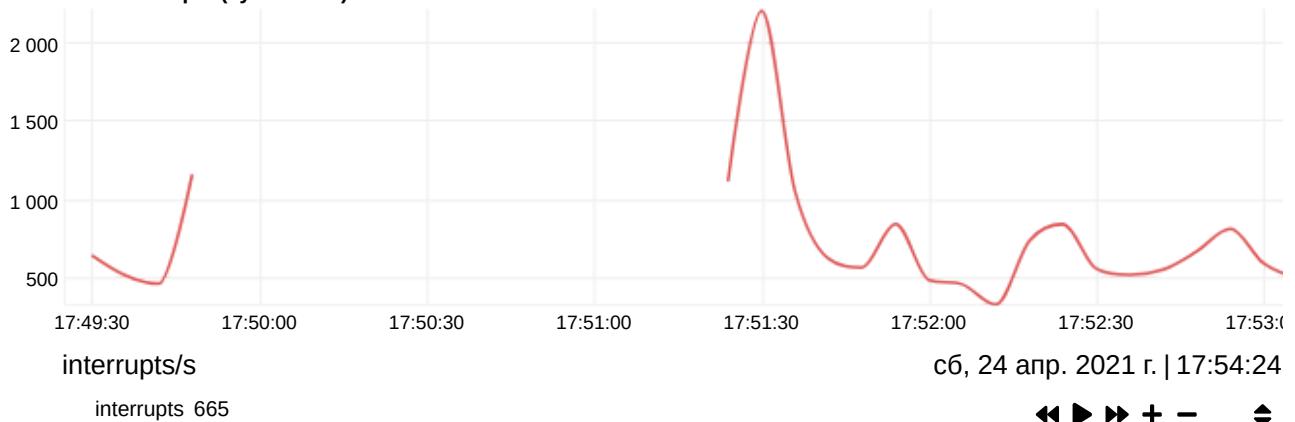
Idle jitter is calculated by netdata. A thread is spawned that requests to sleep for a few microseconds. When the system wakes it up, it measures how many microseconds have passed. The difference between the requested and the actual duration of the sleep, is the **idle jitter**. This number is useful in real-time environments, where CPU jitter can affect the quality of the service (like VoIP media gateways).



interrupts

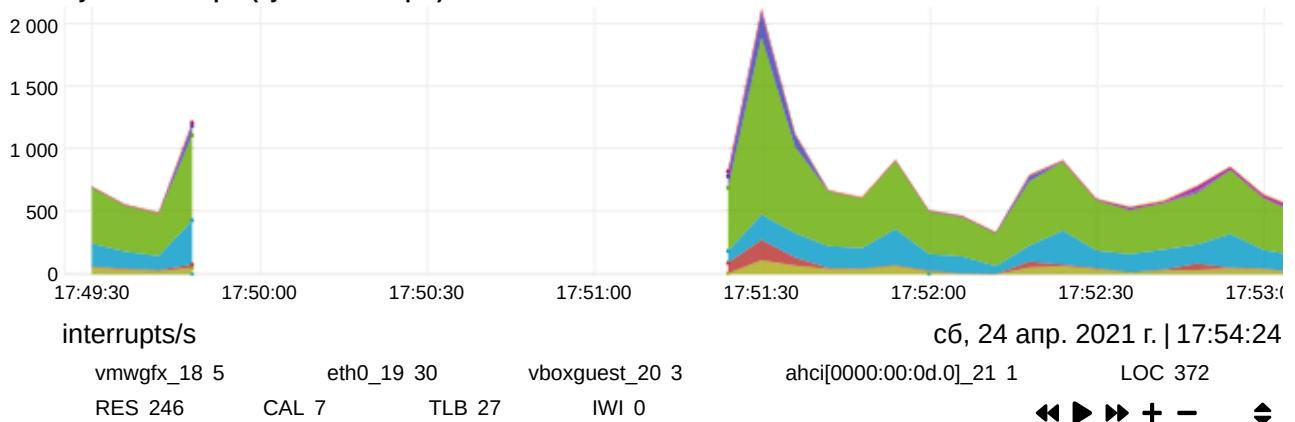
Total number of CPU interrupts. Check `system.interrupts` that gives more detail about each interrupt and also the CPUs section where interrupts are analyzed per CPU core.

CPU Interrupts (system.intr)



CPU interrupts in detail. At the CPUs section, interrupts are analyzed per CPU core.

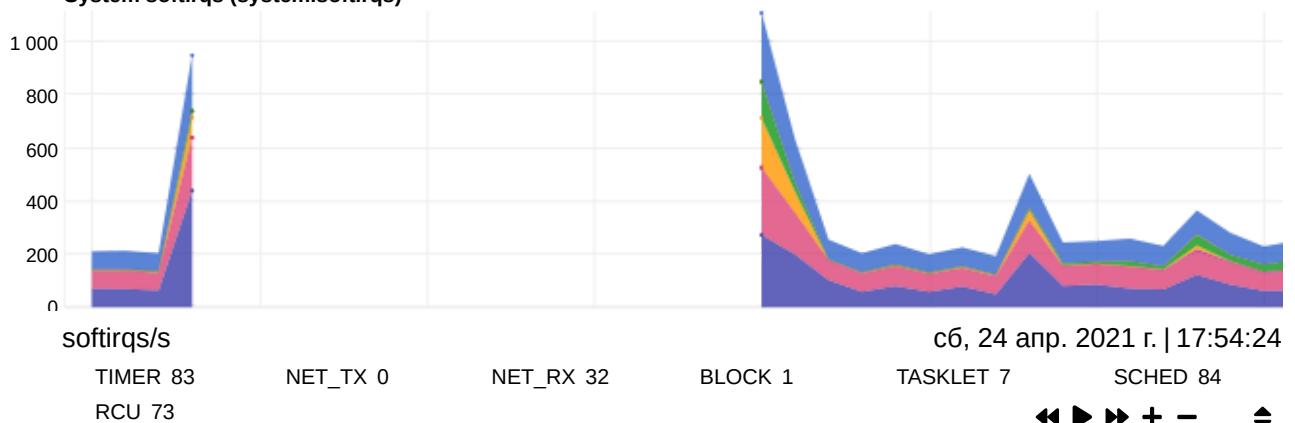
System interrupts (system.interrupts)



softirqs

CPU softirqs in detail. At the CPUs section, softirqs are analyzed per CPU core.

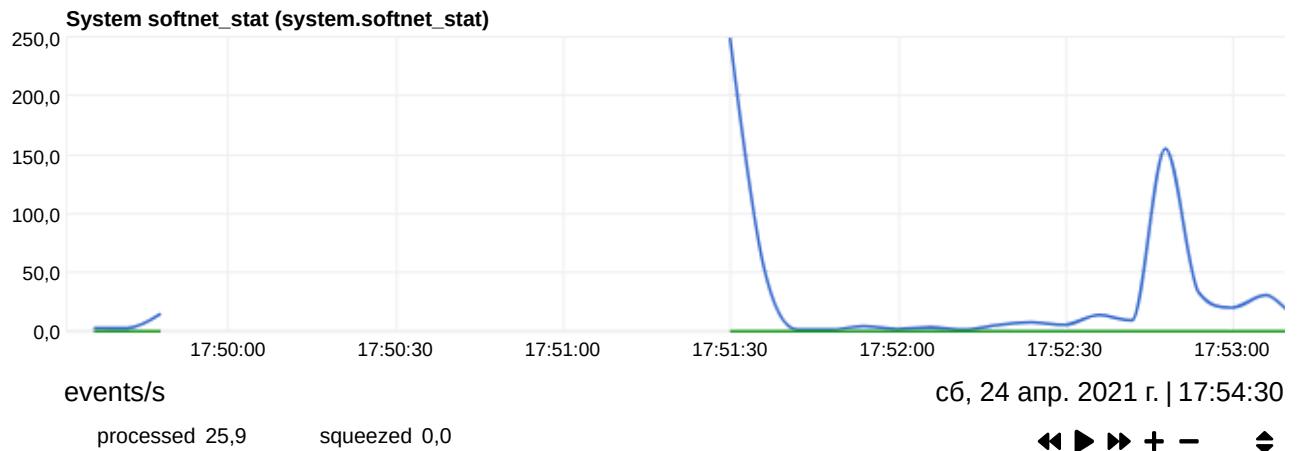
System softirqs (system.softirqs)



softnet

Statistics for CPUs SoftIRQs related to network receive work. Break down per CPU core can be found at CPU / softnet statistics. **processed** states the number of packets processed, **dropped** is the

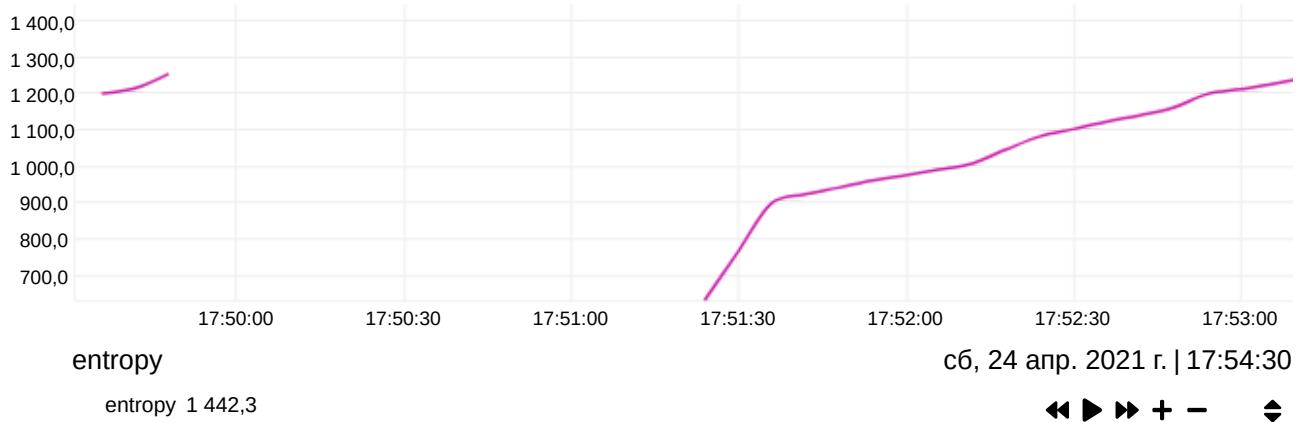
number packets dropped because the network device backlog was full (to fix them on Linux use `sysctl` to increase `net.core.netdev_max_backlog`), **squeezed** is the number of packets dropped because the network device budget ran out (to fix them on Linux use `sysctl` to increase `net.core.netdev_budget` and/or `net.core.netdev_budget_usecs`). More information about identifying and troubleshooting network driver related issues can be found at Red Hat Enterprise Linux Network Performance Tuning Guide (https://access.redhat.com/sites/default/files/attachments/20150325_network_performance_tuning.pdf).



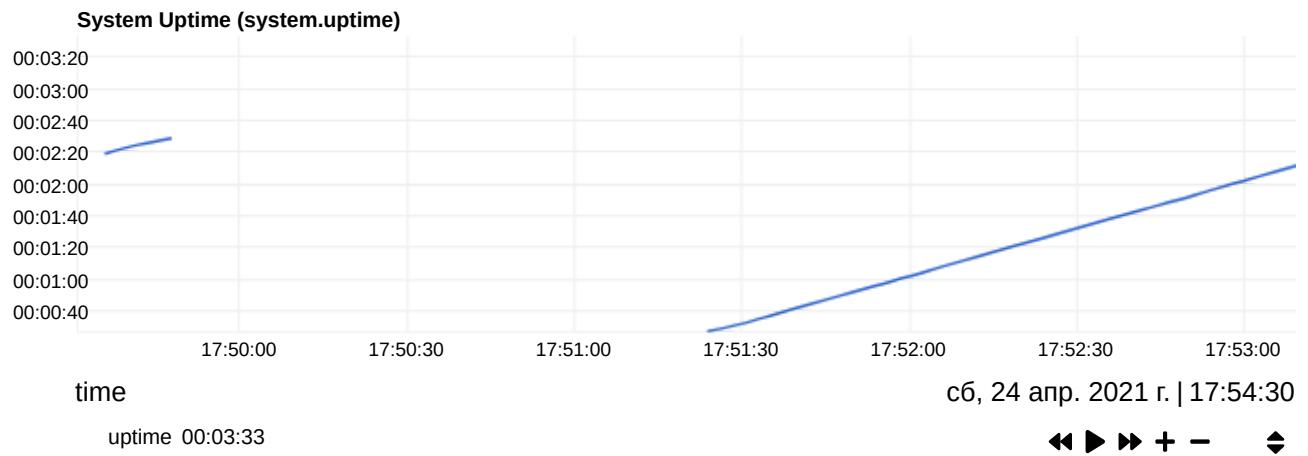
entropy

Entropy ([https://en.wikipedia.org/wiki/Entropy_\(computing\)](https://en.wikipedia.org/wiki/Entropy_(computing))), is a pool of random numbers (`/dev/random` (<https://en.wikipedia.org/wiki//dev/random>)) that is mainly used in cryptography. If the pool of entropy gets empty, processes requiring random numbers may run a lot slower (it depends on the interface each program uses), waiting for the pool to be replenished. Ideally a system with high entropy demands should have a hardware device for that purpose (TPM is one such device). There are also several software-only options you may install, like `haveged`, although these are generally useful only in servers.

Available Entropy (system.entropy)

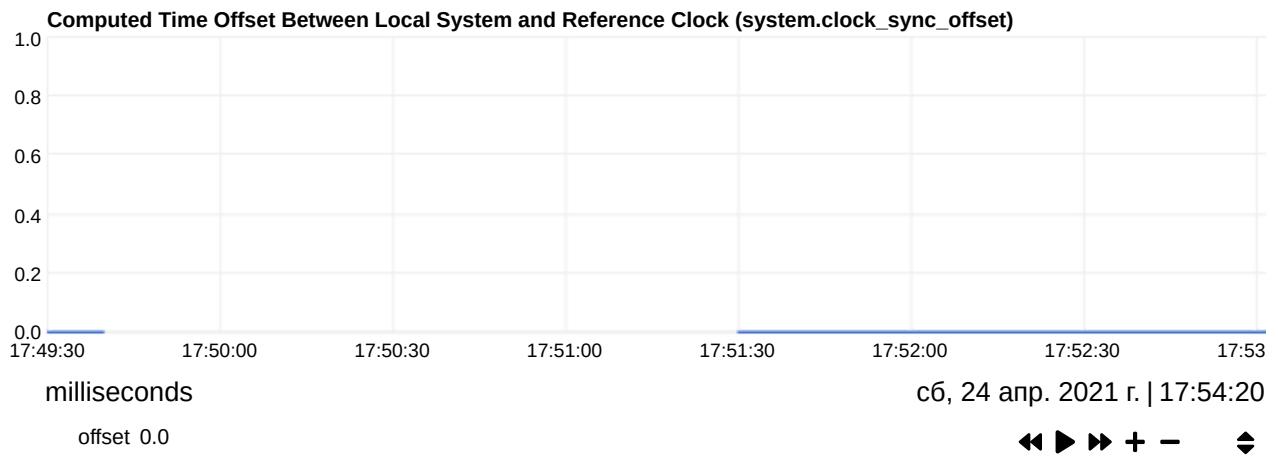
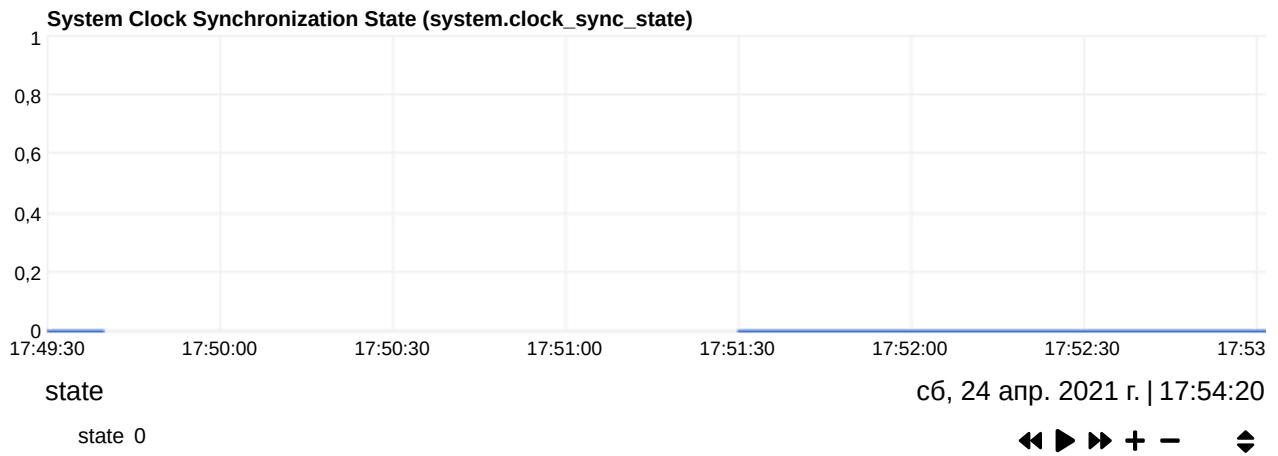


uptime

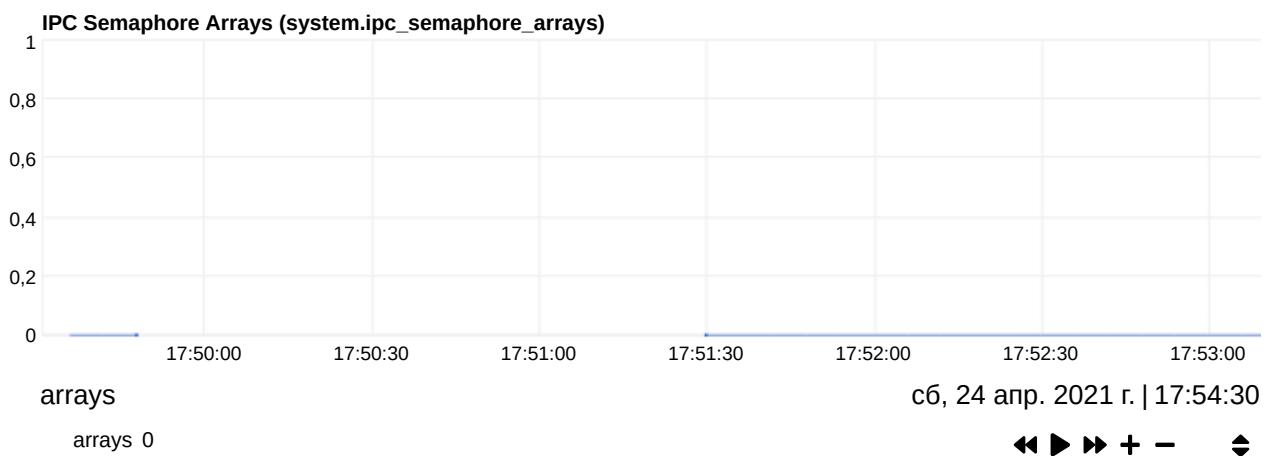
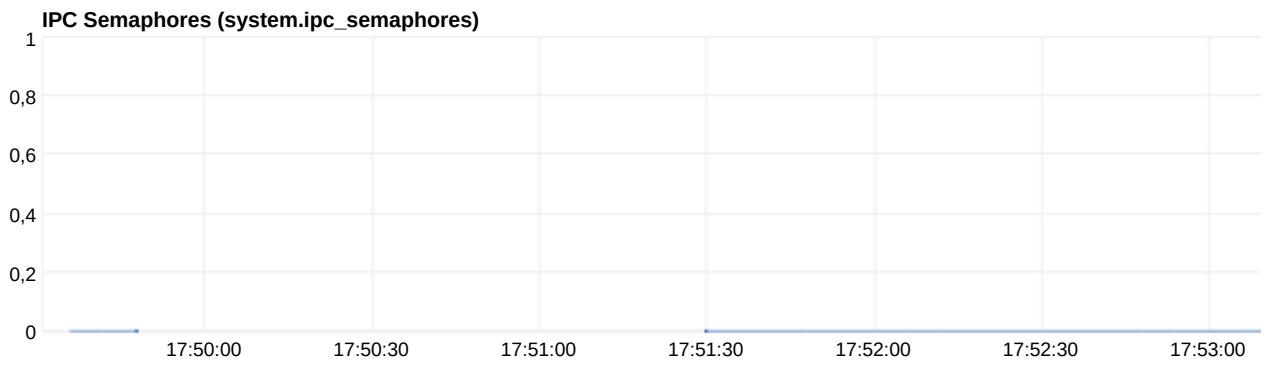


clock synchronization

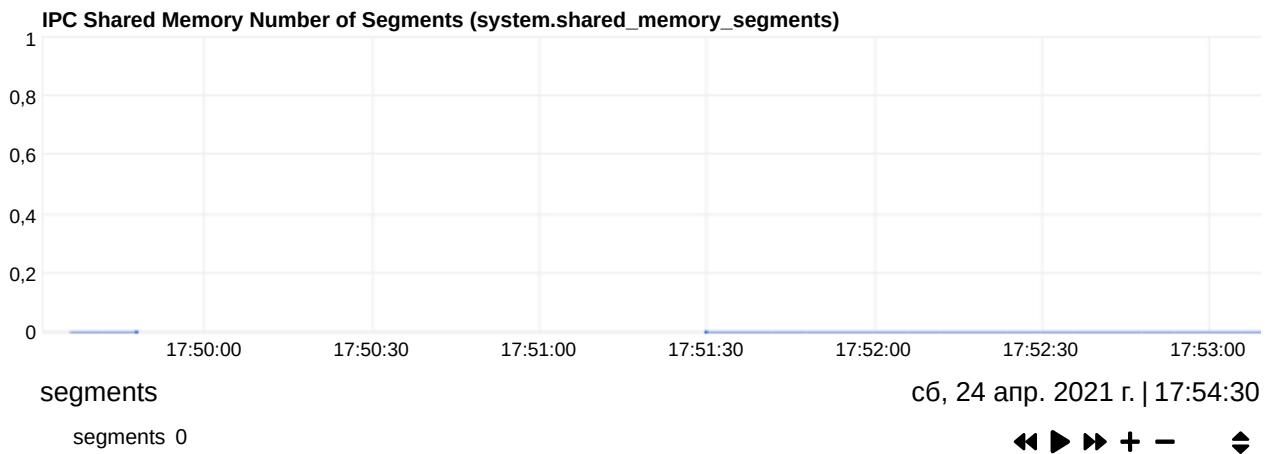
State map: 0 - not synchronized, 1 - synchronized

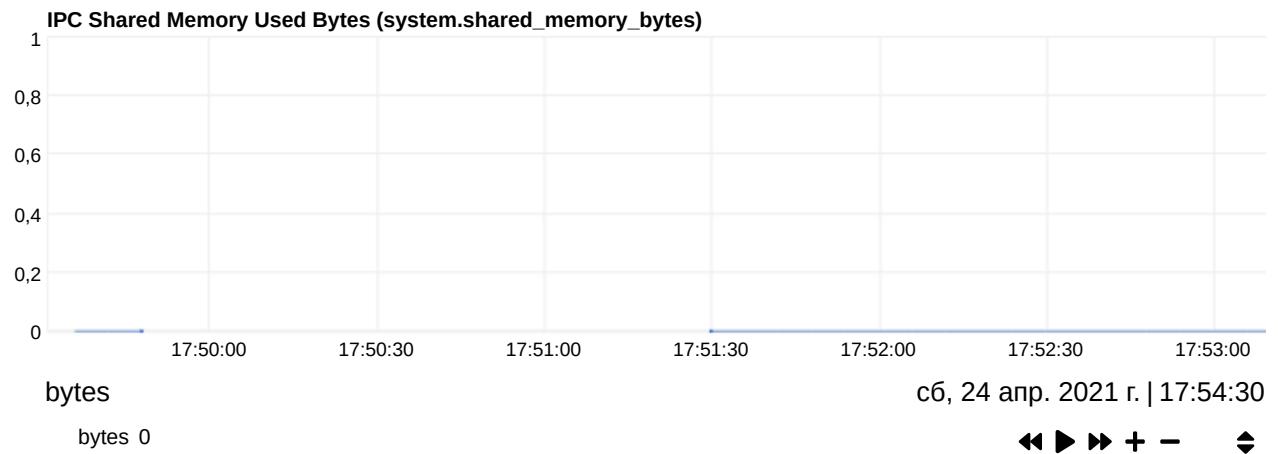


ipc semaphores



ipc shared memory

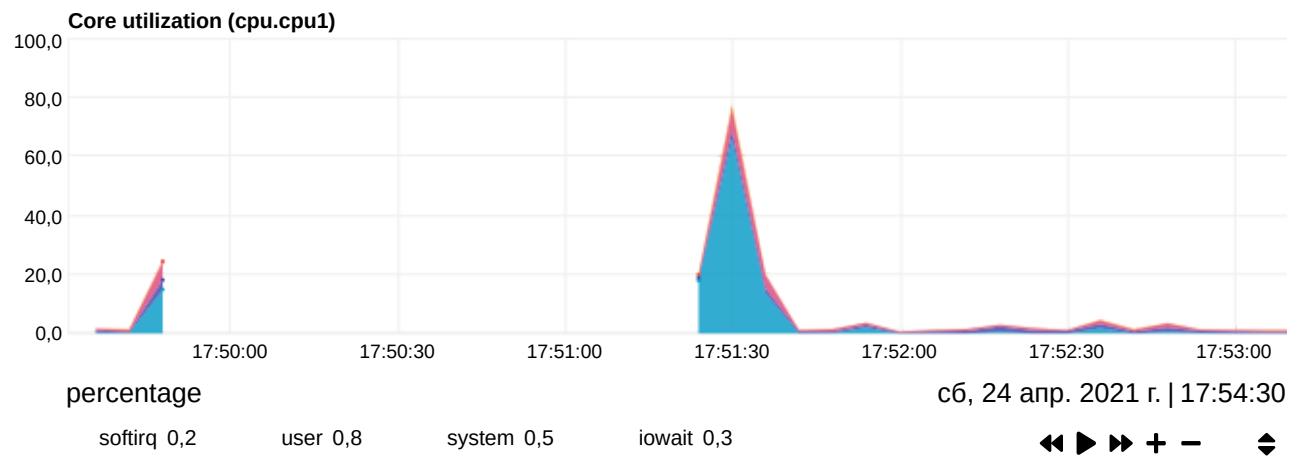
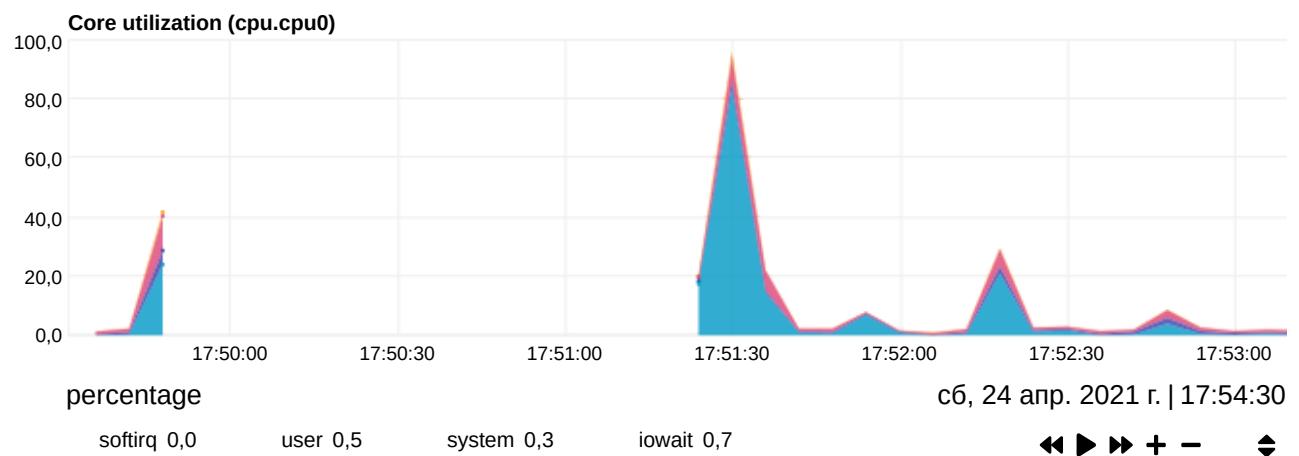




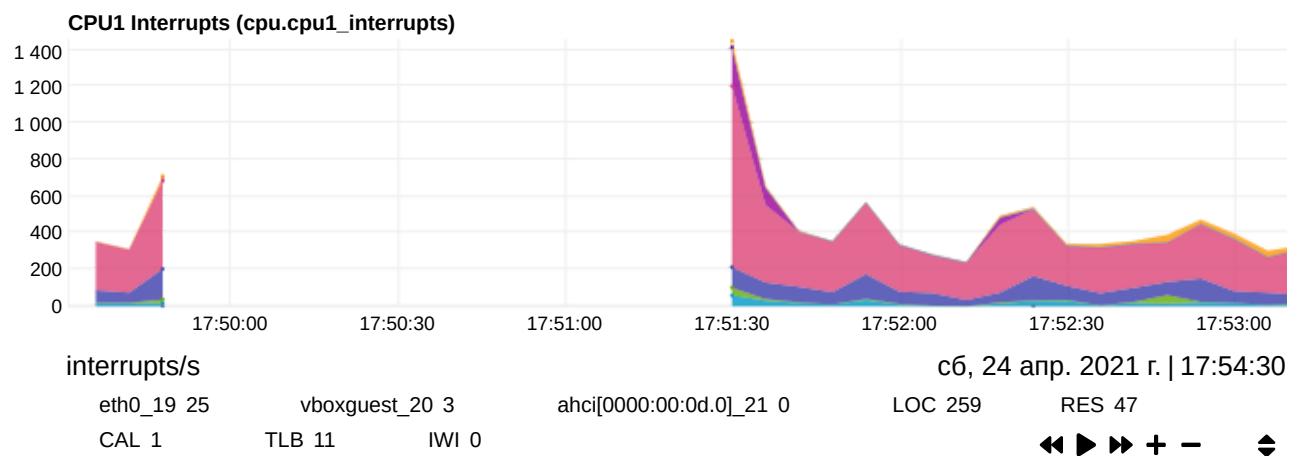
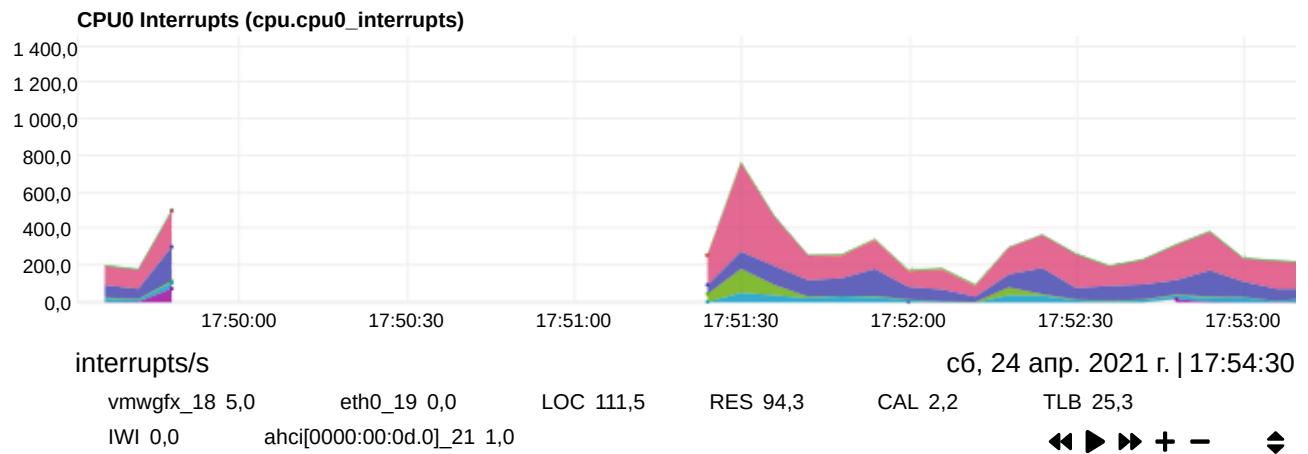
⚡ CPUs

Detailed information for each CPU of the system. A summary of the system for all CPUs can be found at the System Overview section.

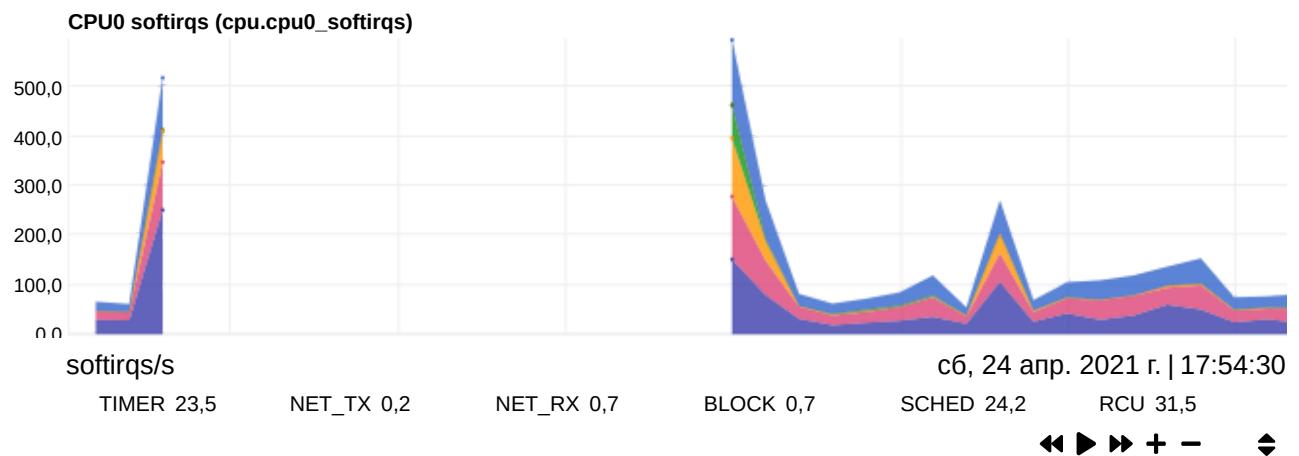
utilization

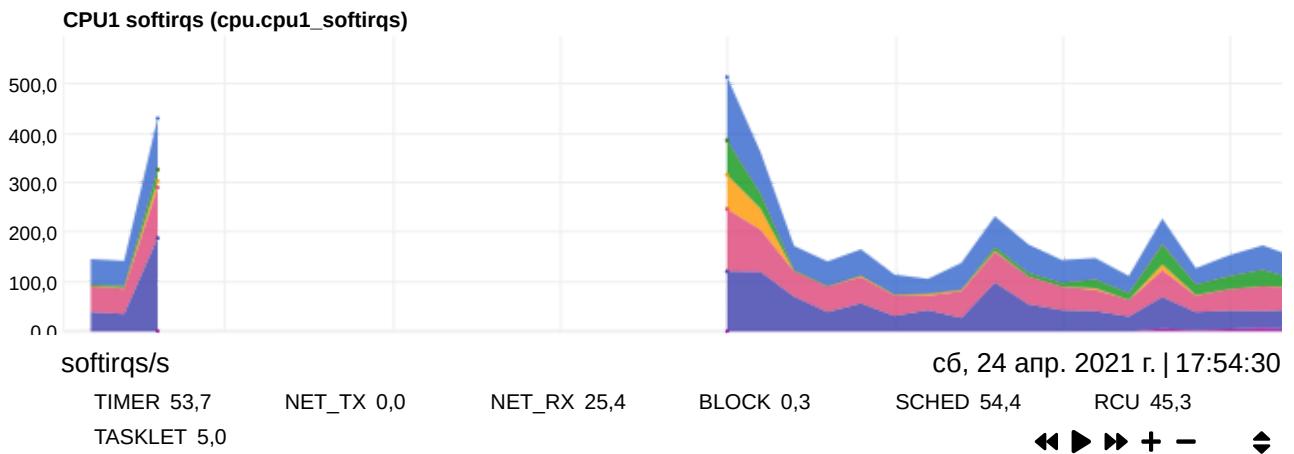


interrupts



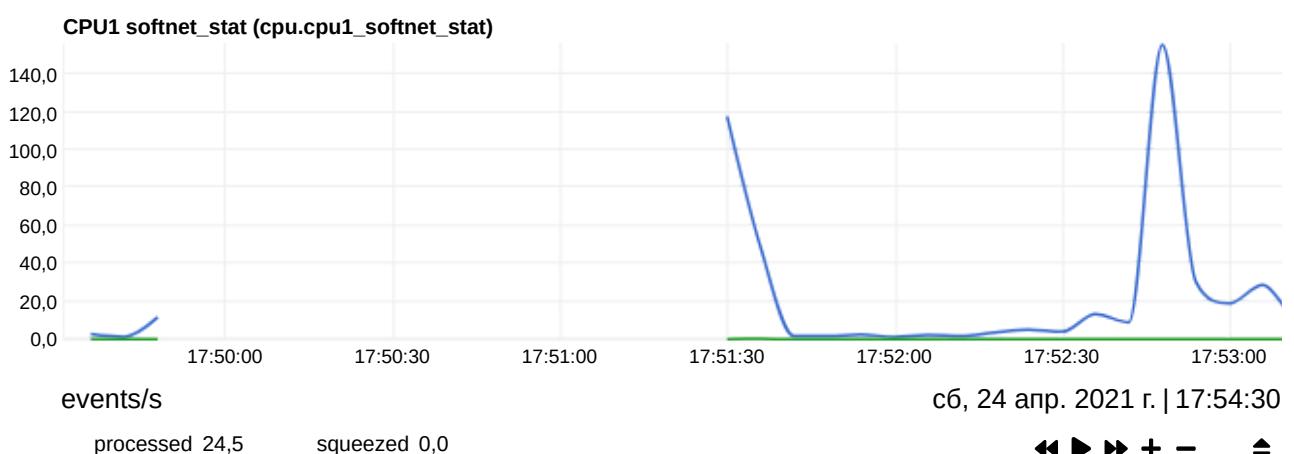
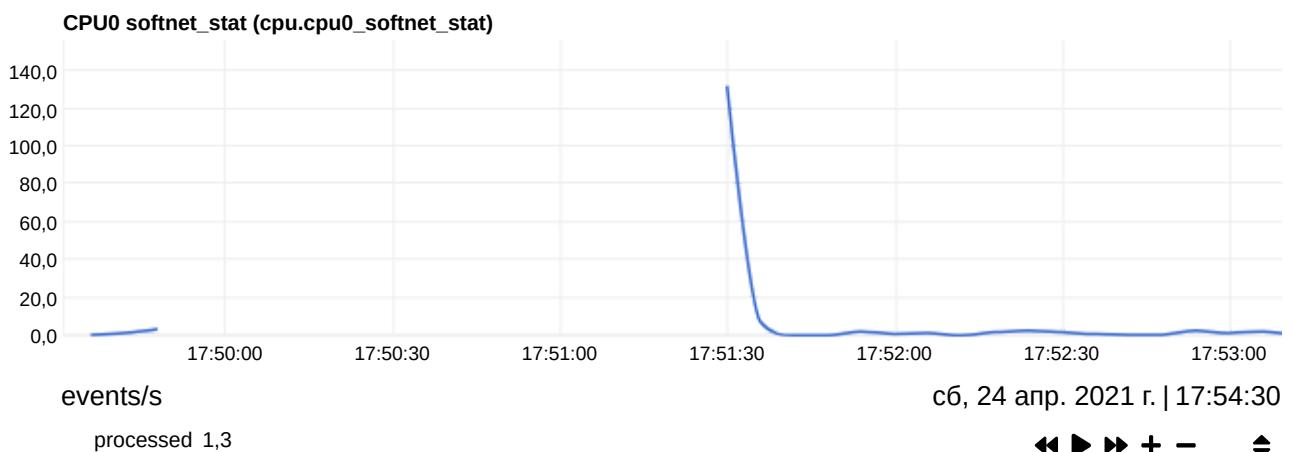
softirqs



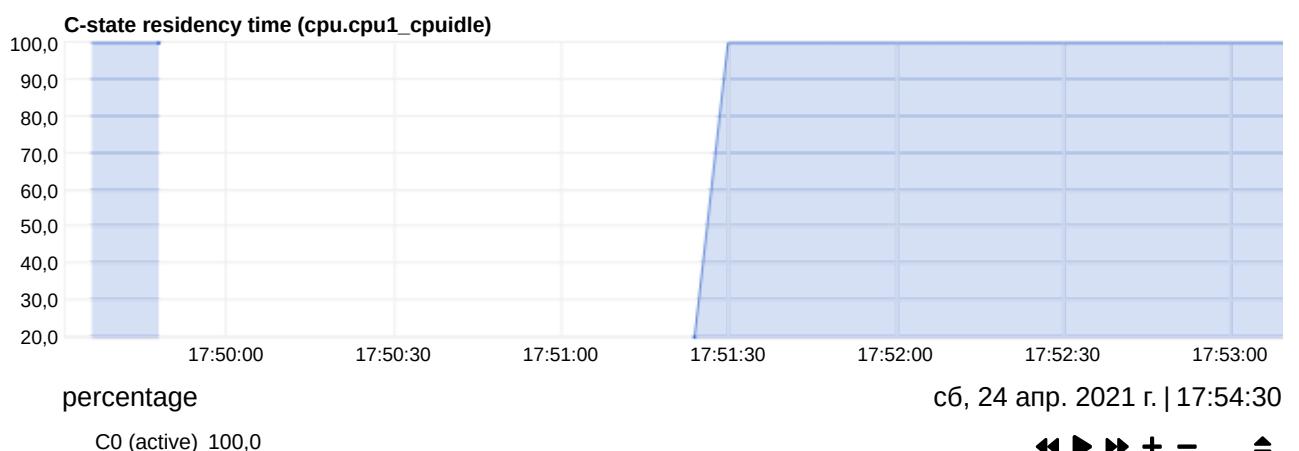
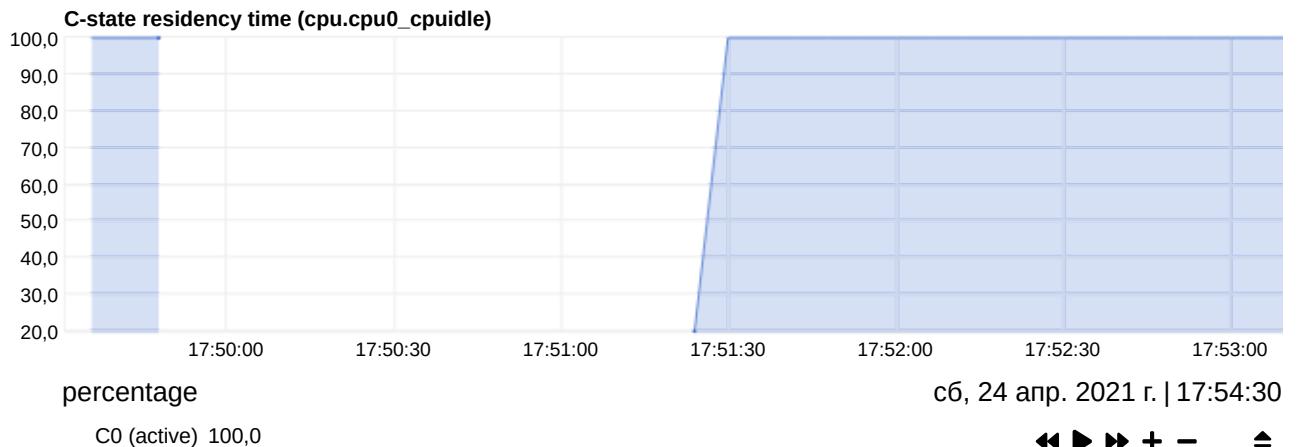


softnet

Statistics for per CPUs core SoftIRQs related to network receive work. Total for all CPU cores can be found at System / softnet statistics. **processed** states the number of packets processed, **dropped** is the number packets dropped because the network device backlog was full (to fix them on Linux use `sysctl` to increase `net.core.netdev_max_backlog`), **squeezed** is the number of packets dropped because the network device budget ran out (to fix them on Linux use `sysctl` to increase `net.core.netdev_budget` and/or `net.core.netdev_budget_usecs`). More information about identifying and troubleshooting network driver related issues can be found at Red Hat Enterprise Linux Network Performance Tuning Guide (https://access.redhat.com/sites/default/files/attachments/20150325_network_performance_tuning.pdf).



cpuidle



Memory

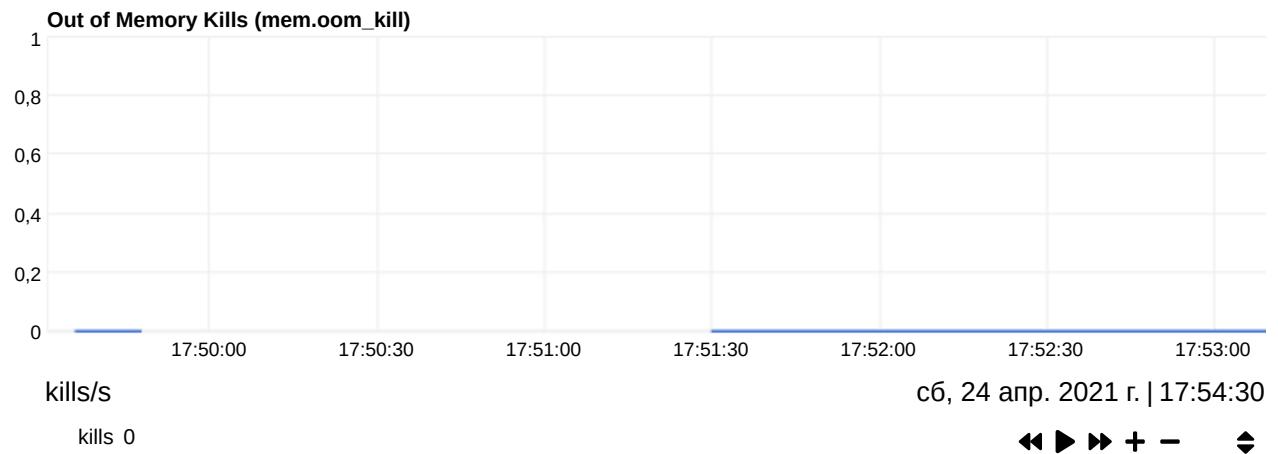
Detailed information about the memory management of the system.

system

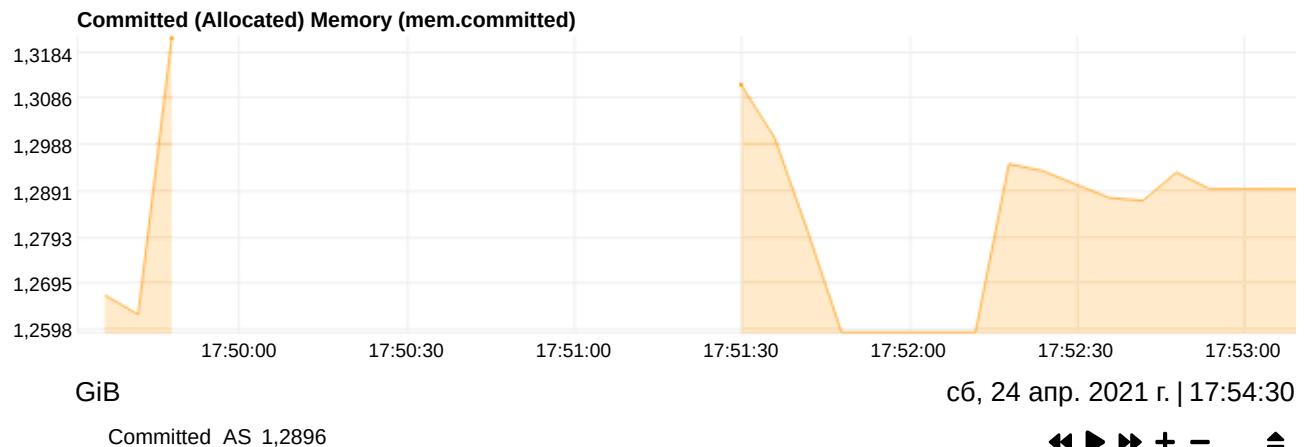
Available Memory is estimated by the kernel, as the amount of RAM that can be used by userspace processes, without causing swapping.

Available RAM for applications (mem.available)

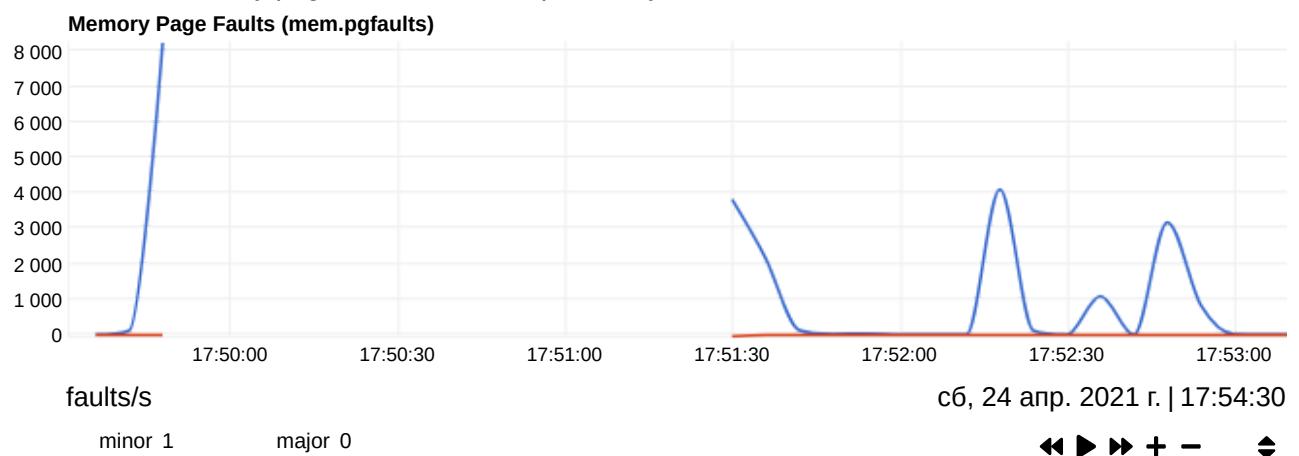




Committed Memory, is the sum of all memory which has been allocated by processes.



A page fault (https://en.wikipedia.org/wiki/Page_fault) is a type of interrupt, called trap, raised by computer hardware when a running program accesses a memory page that is mapped into the virtual address space, but not actually loaded into main memory. If the page is loaded in memory at the time the fault is generated, but is not marked in the memory management unit as being loaded in memory, then it is called a **minor** or soft page fault. A **major** page fault is generated when the system needs to load the memory page from disk or swap memory.



kernel

Dirty is the amount of memory waiting to be written to disk. **Writeback** is how much memory is actively being written to disk.

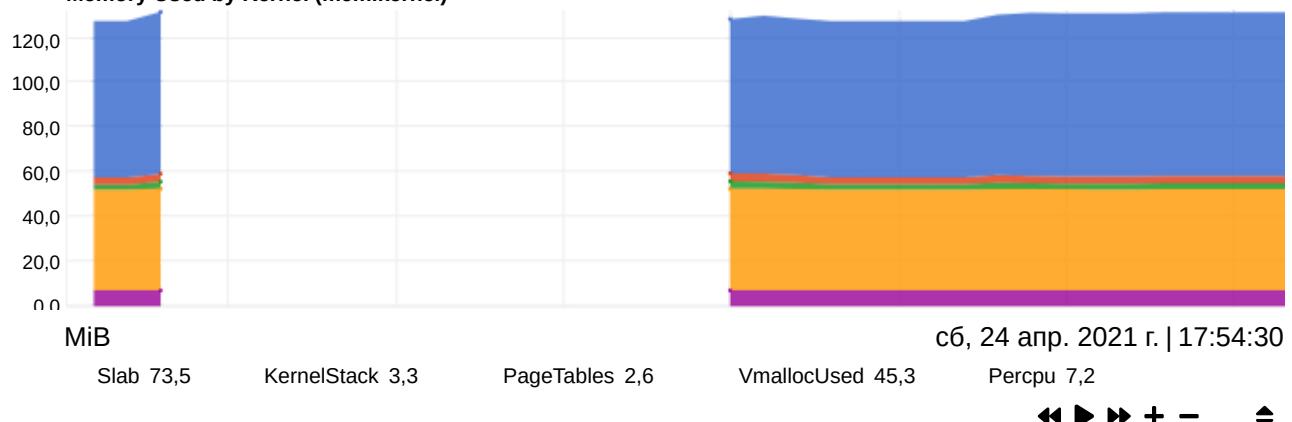
Writeback Memory (mem.writeback)



The total amount of memory being used by the kernel. **Slab** is the amount of memory used by the kernel to cache data structures for its own use. **KernelStack** is the amount of memory allocated for each task done by the kernel. **PageTables** is the amount of memory dedicated to the lowest level of page tables (A page table is used to turn a virtual address into a physical memory address).

VmallocUsed is the amount of memory being used as virtual address space.

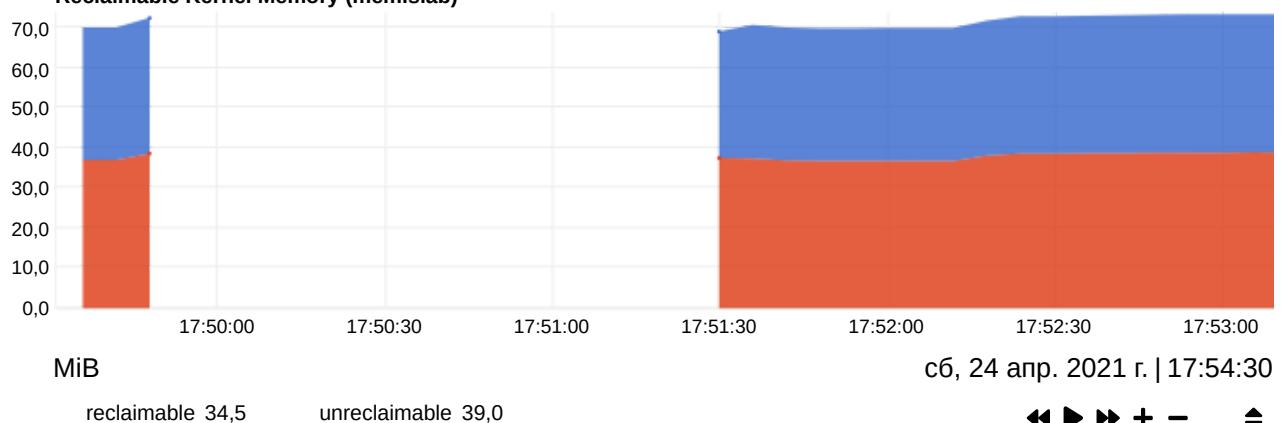
Memory Used by Kernel (mem.kernel)



slab

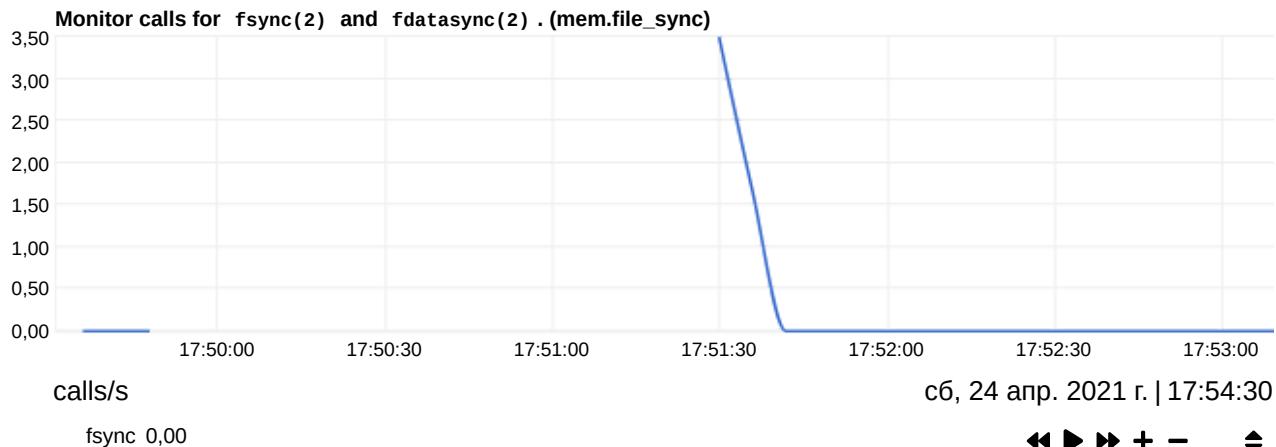
Reclaimable is the amount of memory which the kernel can reuse. **Unreclaimable** can not be reused even when the kernel is lacking memory.

Reclaimable Kernel Memory (mem.slab)

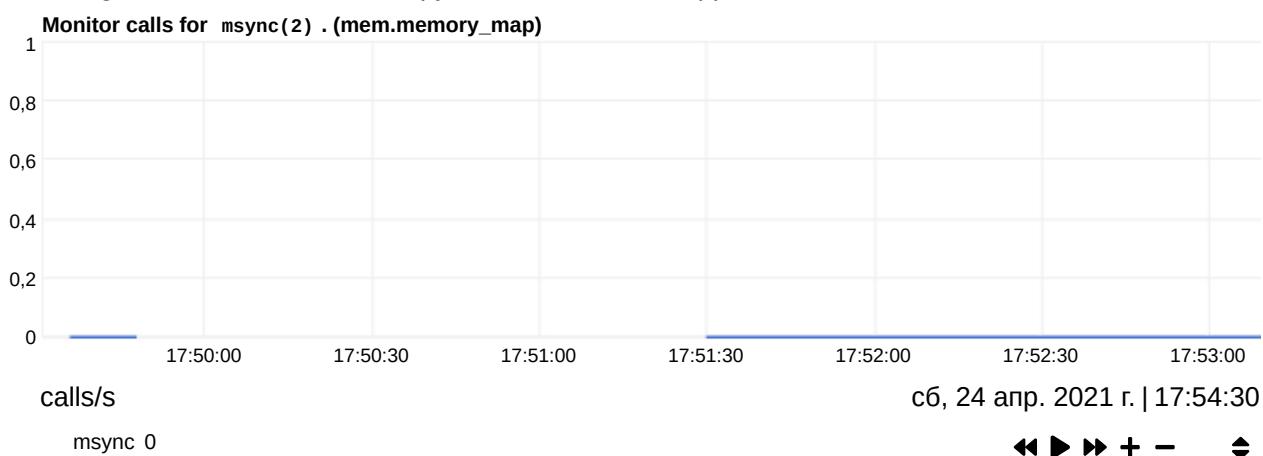


synchronization (eBPF)

System calls for fsync() and fdatasync() (<https://man7.org/linux/man-pages/man2/fsync.2.html>) transfer all modified page caches for the files on disk devices. These calls block until the device reports that the transfer has been completed.



System calls for msync() (<https://man7.org/linux/man-pages/man2/msync.2.html>) which flushes changes made to the in-core copy of a file that was mapped.

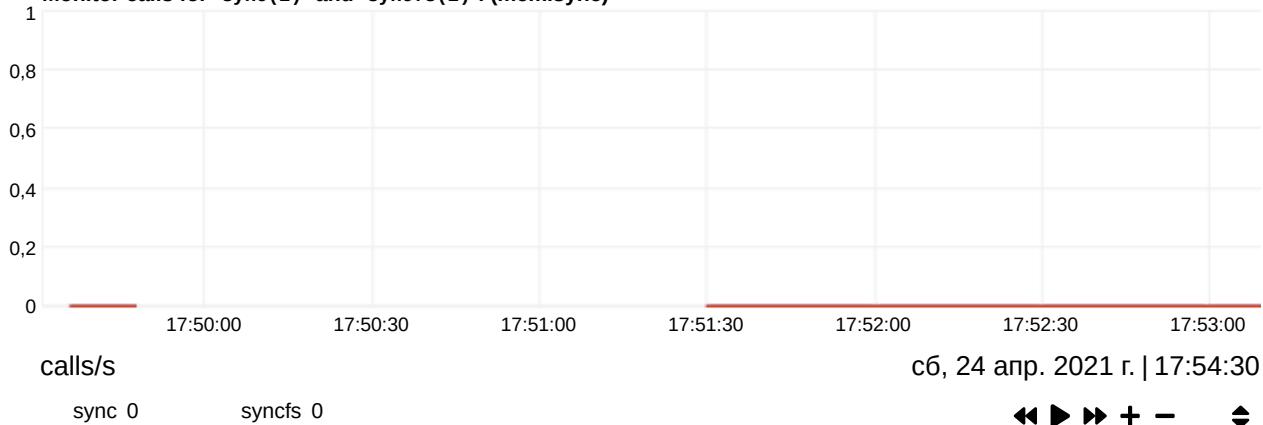


System calls for sync() and syncfs() (<https://man7.org/linux/man-pages/man2/sync.2.html>) which flush the file system buffers to storage devices. Performance perturbations might be caused by these calls.

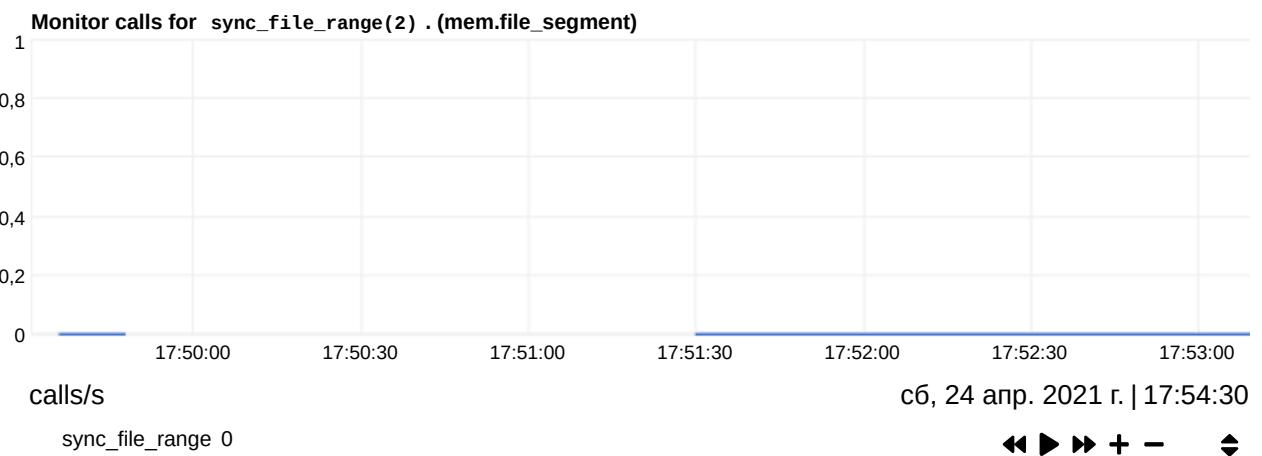
The sync() calls are based on the eBPF syncsnoop

(<https://github.com/iovisor/bcc/blob/master/tools/syncsnoop.py>) from BCC tools.

Monitor calls for sync(2) and syncfs(2) . (mem.sync)



System calls for sync_file_range() (https://man7.org/linux/man-pages/man2/sync_file_range.2.html) permits fine control when synchronizing the open file referred to by the file descriptor fd with disk. This system call is extremely dangerous and should not be used in portable programs.

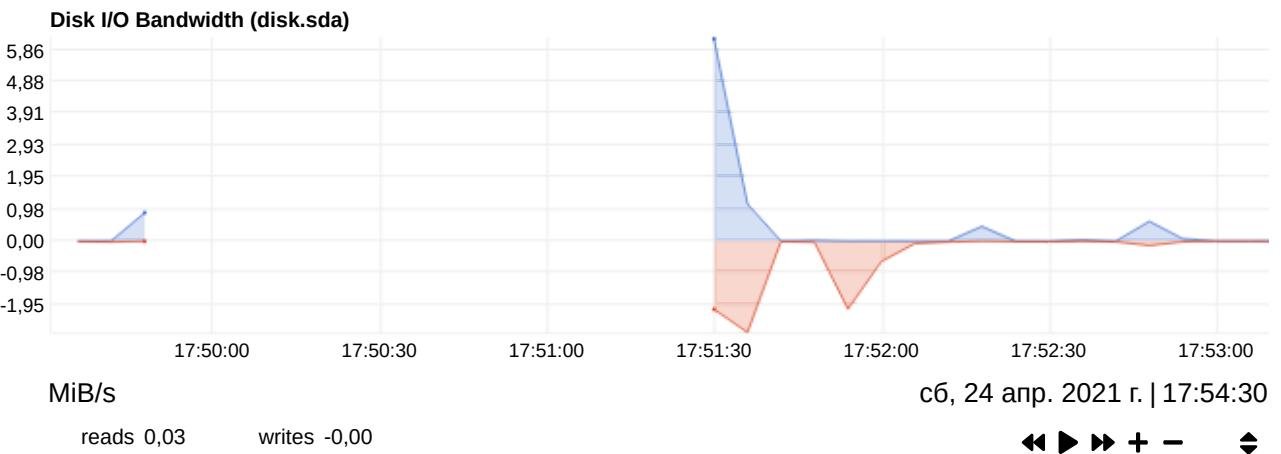


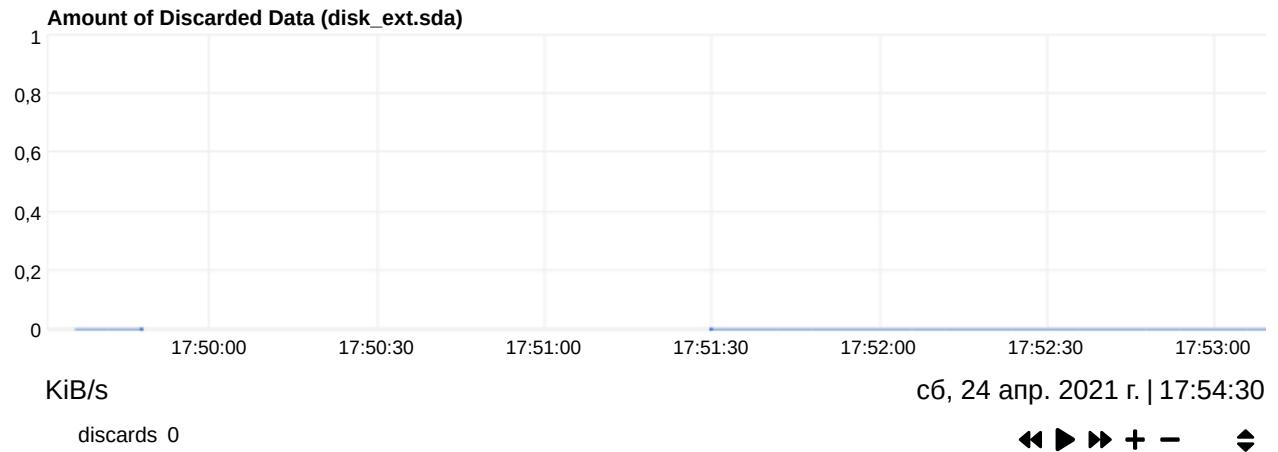
Disks

Charts with performance information for all the system disks. Special care has been given to present disk performance metrics in a way compatible with `iostat -x`. netdata by default prevents rendering performance charts for individual partitions and unmounted virtual disks. Disabled charts can still be enabled by configuring the relative settings in the netdata configuration file.

sda

Amount of data transferred to and from disk.



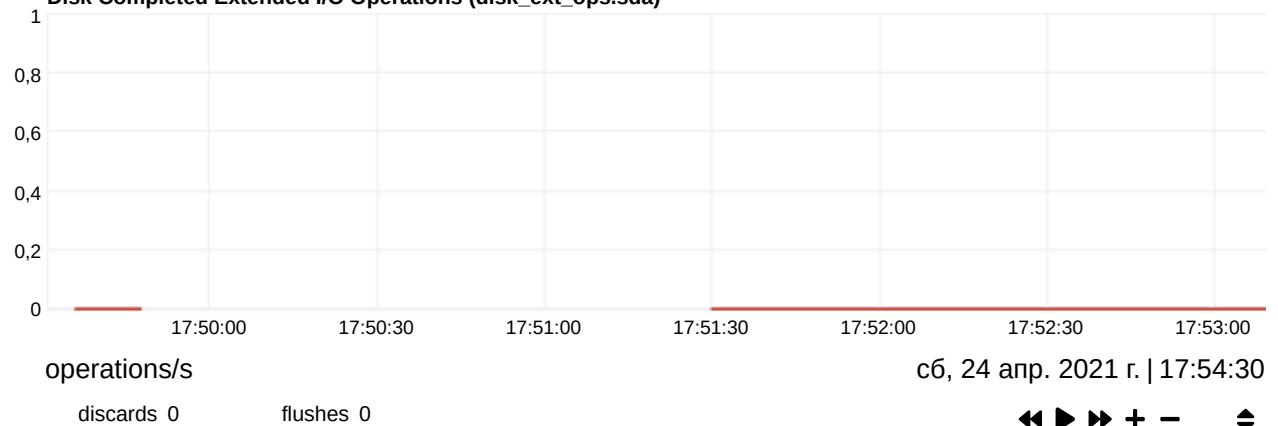


Completed disk I/O operations. Keep in mind the number of operations requested might be higher, since the system is able to merge adjacent to each other (see merged operations chart).

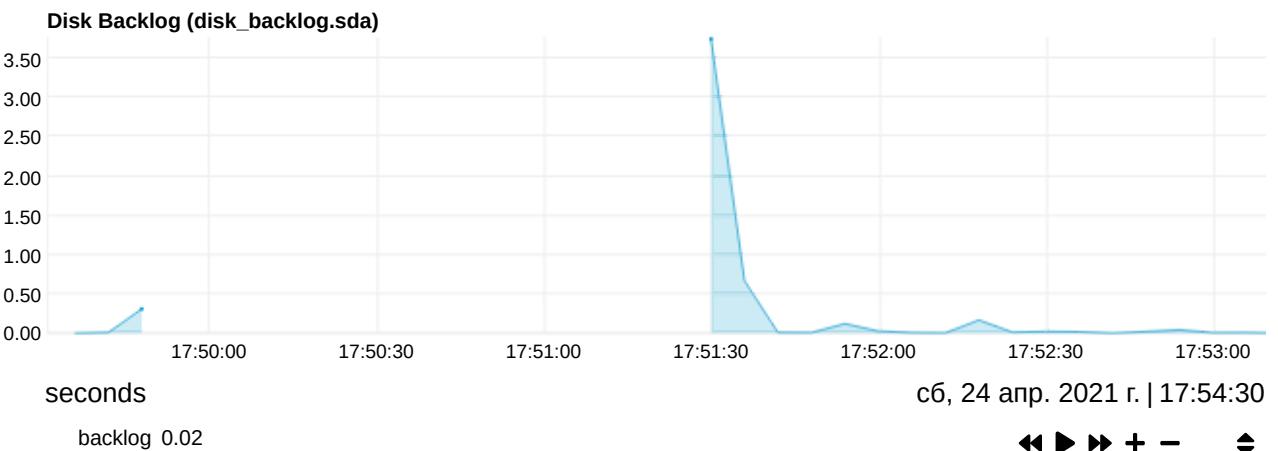
Disk Completed I/O Operations (disk_ops.sda)



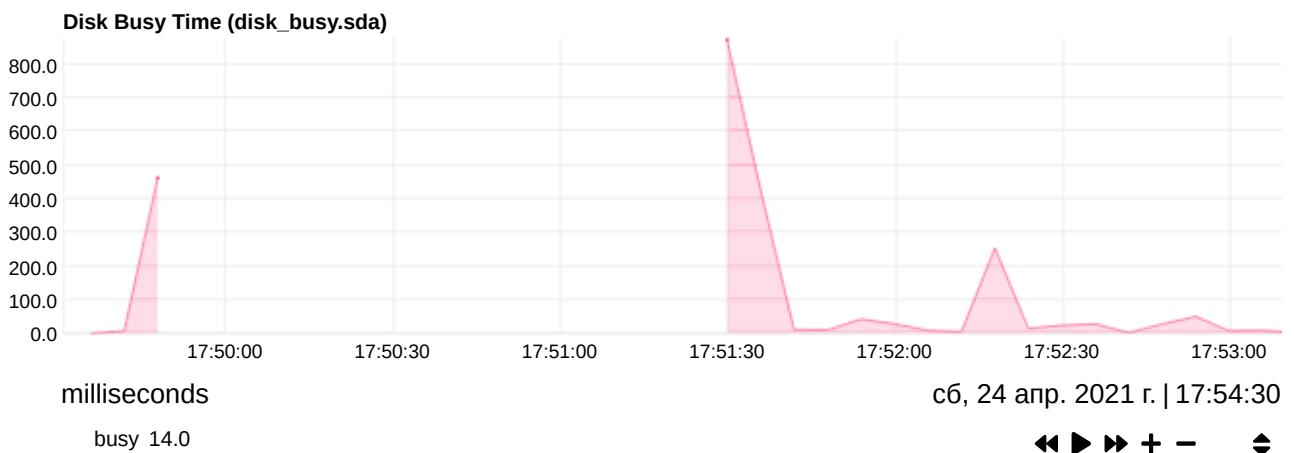
Disk Completed Extended I/O Operations (disk_ext_ops.sda)



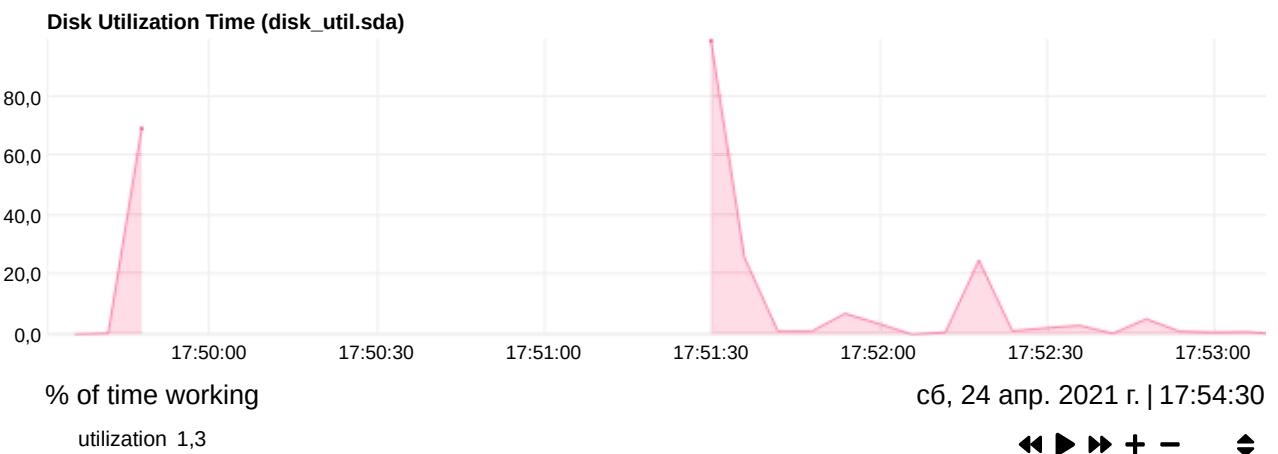
Backlog is an indication of the duration of pending disk operations. On every I/O event the system is multiplying the time spent doing I/O since the last update of this field with the number of pending operations. While not accurate, this metric can provide an indication of the expected completion time of the operations in progress.



Disk Busy Time measures the amount of time the disk was busy with something.



Disk Utilization measures the amount of time the disk was busy with something. This is not related to its performance. 100% means that the system always had an outstanding operation on the disk. Keep in mind that depending on the underlying technology of the disk, 100% here may or may not be an indication of congestion.



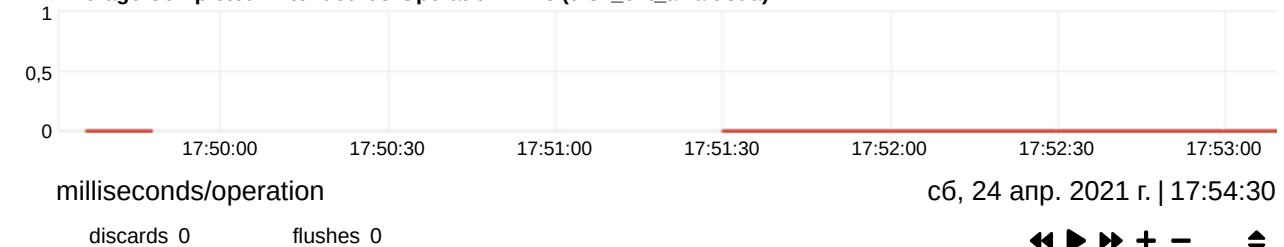
The average time for I/O requests issued to the device to be served. This includes the time spent by the requests in queue and the time spent servicing them.

Average Completed I/O Operation Time (disk_await.sda)



The average time for extended I/O requests issued to the device to be served. This includes the time spent by the requests in queue and the time spent servicing them.

Average Completed Extended I/O Operation Time (disk_ext_await.sda)



The average I/O operation size.

Average Completed I/O Operation Bandwidth (disk_avgsz.sda)



Average Amount of Discarded Data (disk_ext_avgsz.sda)



The average service time for completed I/O operations. This metric is calculated using the total busy time of the disk and the number of completed operations. If the disk is able to execute multiple parallel operations the reporting average service time will be misleading.

Average Service Time (disk_svctm.sda)



The number of merged disk operations. The system is able to merge adjacent I/O operations, for example two 4KB reads can become one 8KB read before given to disk.

Disk Merged Operations (disk_mops.sda)



Disk Merged Discard Operations (disk_ext_mops.sda)

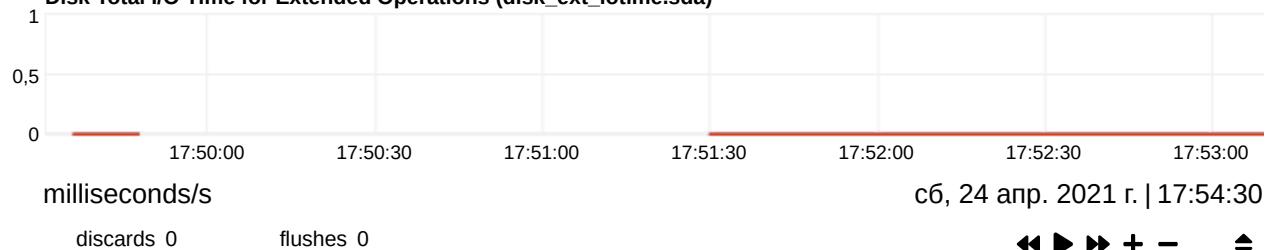


The sum of the duration of all completed I/O operations. This number can exceed the interval if the disk is able to execute I/O operations in parallel.

Disk Total I/O Time (disk_iotime.sda)

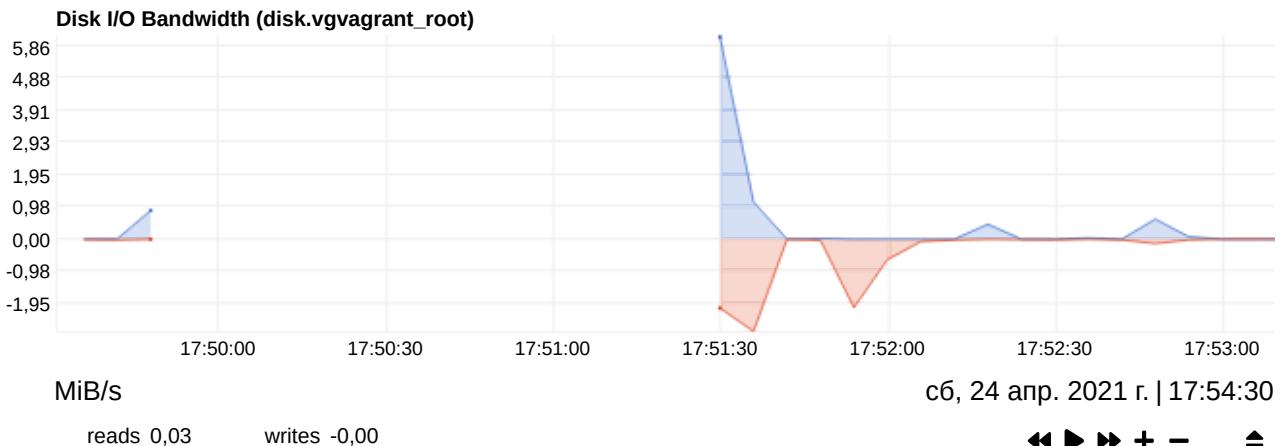


Disk Total I/O Time for Extended Operations (disk_ext_iotime.sda)

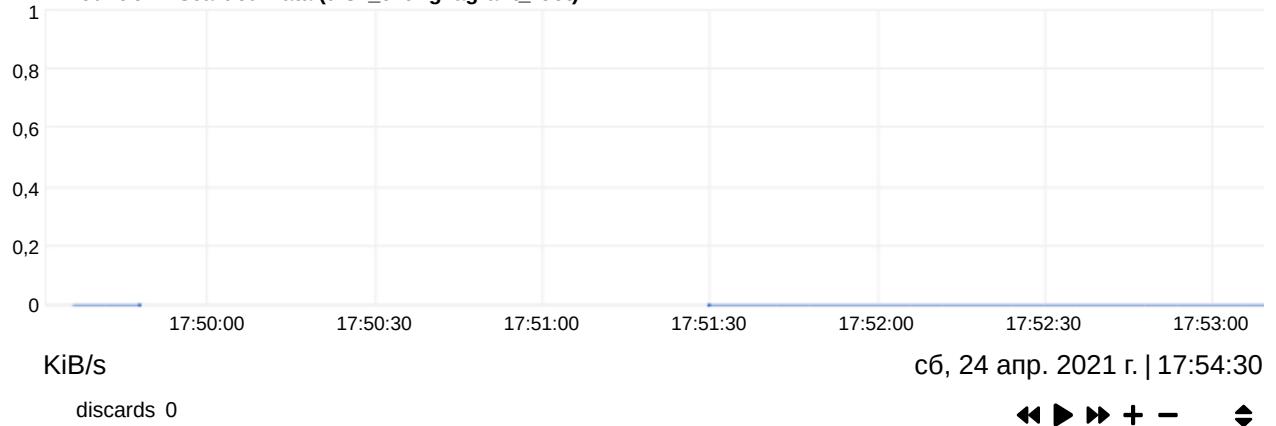


vgvagrant-root

Amount of data transferred to and from disk.



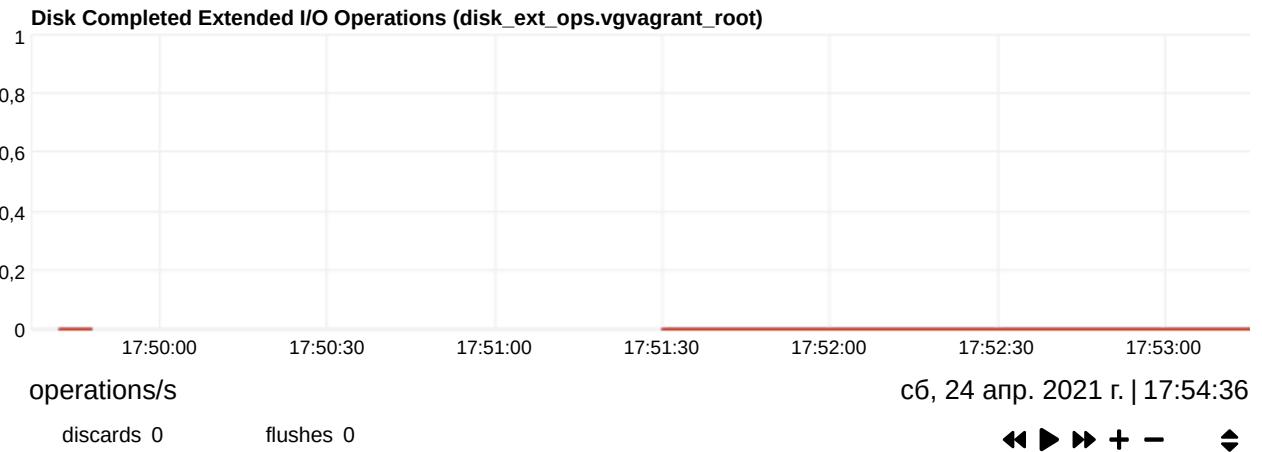
Amount of Discarded Data (disk_ext.vgvagrant_root)



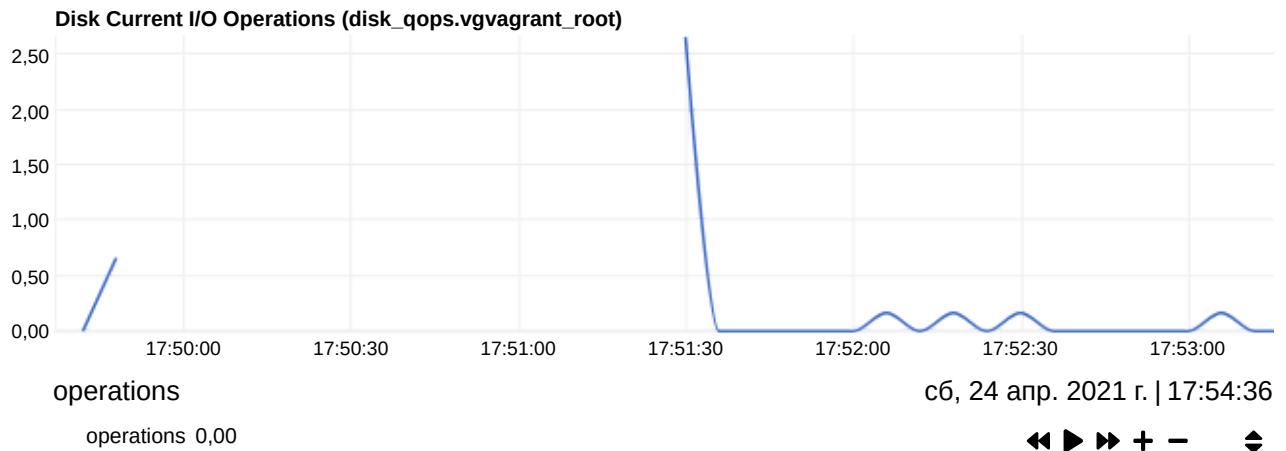
Completed disk I/O operations. Keep in mind the number of operations requested might be higher, since the system is able to merge adjacent to each other (see merged operations chart).

Disk Completed I/O Operations (disk_ops.vgvagrant_root)

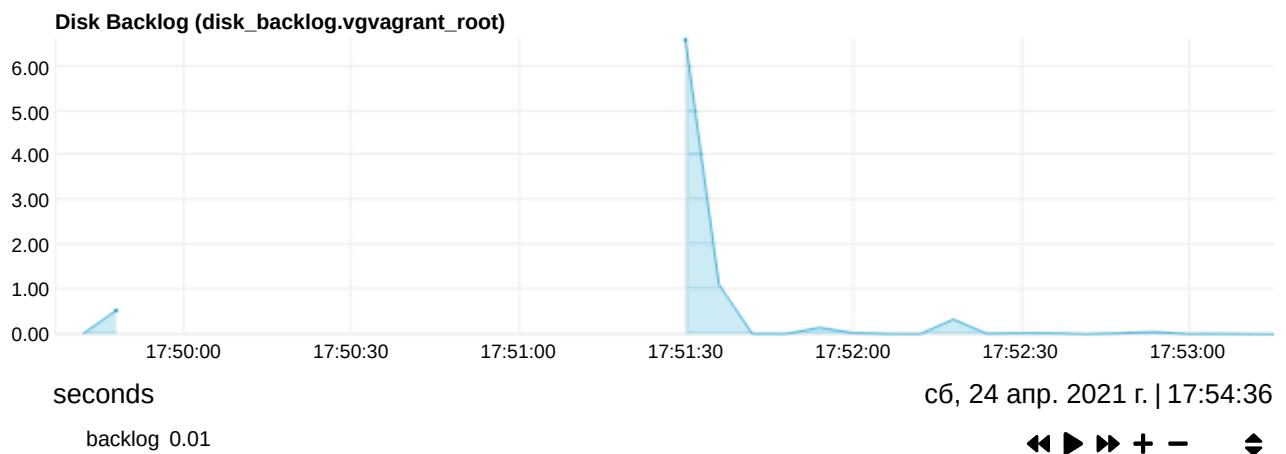




I/O operations currently in progress. This metric is a snapshot - it is not an average over the last interval.

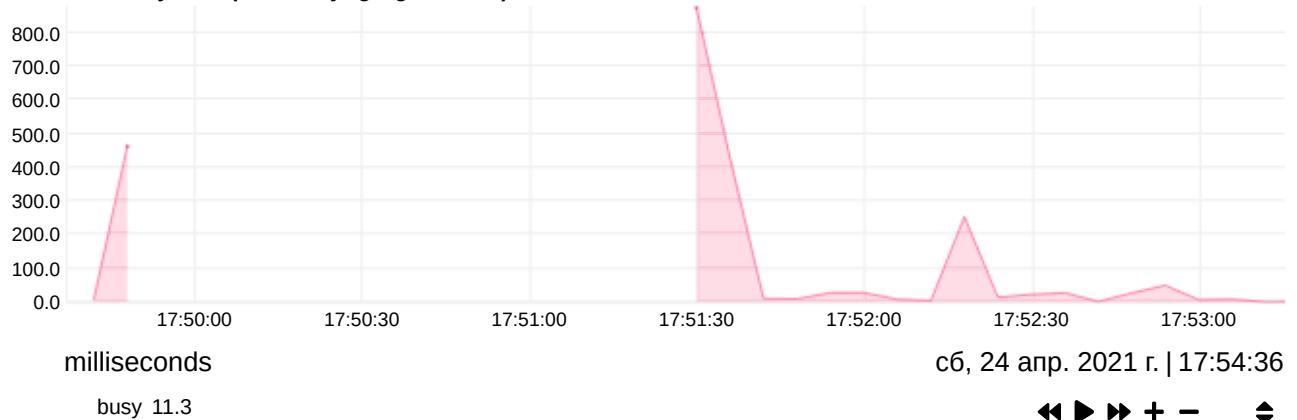


Backlog is an indication of the duration of pending disk operations. On every I/O event the system is multiplying the time spent doing I/O since the last update of this field with the number of pending operations. While not accurate, this metric can provide an indication of the expected completion time of the operations in progress.



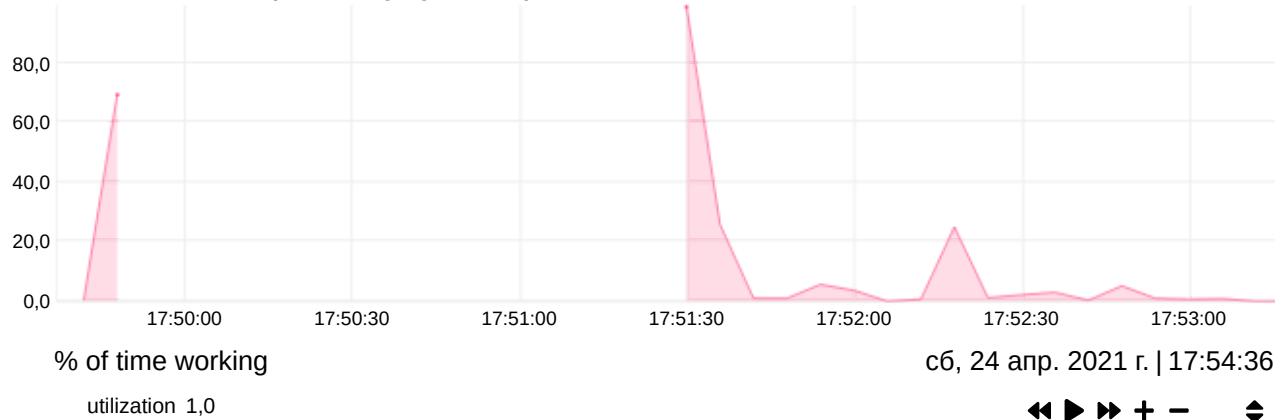
Disk Busy Time measures the amount of time the disk was busy with something.

Disk Busy Time (disk_busy.vg vagrant_root)



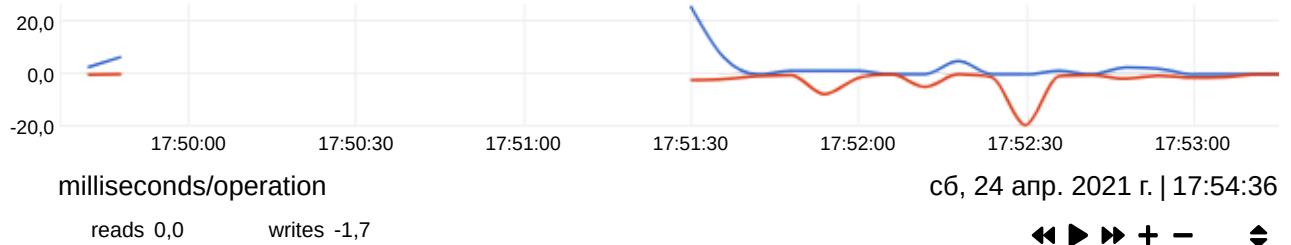
Disk Utilization measures the amount of time the disk was busy with something. This is not related to its performance. 100% means that the system always had an outstanding operation on the disk. Keep in mind that depending on the underlying technology of the disk, 100% here may or may not be an indication of congestion.

Disk Utilization Time (disk_util.vg vagrant_root)



The average time for I/O requests issued to the device to be served. This includes the time spent by the requests in queue and the time spent servicing them.

Average Completed I/O Operation Time (disk_await.vg vagrant_root)



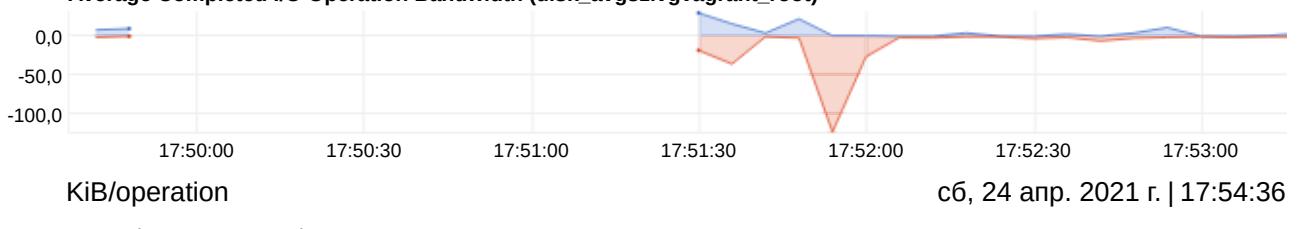
The average time for extended I/O requests issued to the device to be served. This includes the time spent by the requests in queue and the time spent servicing them.

Average Completed Extended I/O Operation Time (disk_ext_await.vgvagrant_root)



The average I/O operation size.

Average Completed I/O Operation Bandwidth (disk_avgsz.vgvagrant_root)



Average Amount of Discarded Data (disk_ext_avgsz.vgvagrant_root)



The average service time for completed I/O operations. This metric is calculated using the total busy time of the disk and the number of completed operations. If the disk is able to execute multiple parallel operations the reporting average service time will be misleading.

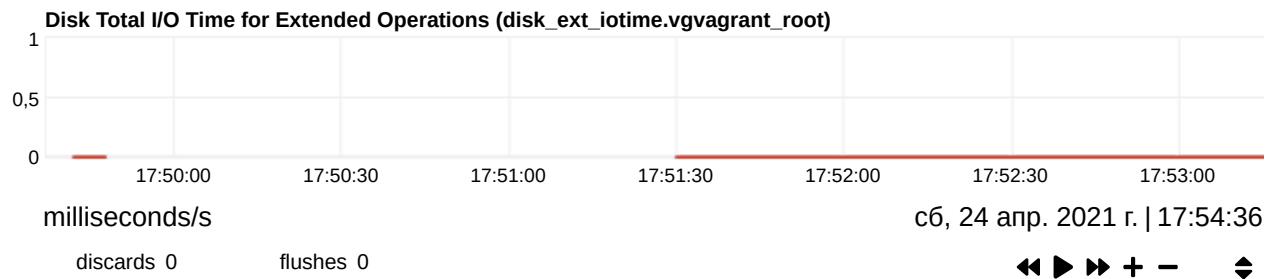
Average Service Time (disk_svctm.vgvagrant_root)



The sum of the duration of all completed I/O operations. This number can exceed the interval if the disk is able to execute I/O operations in parallel.

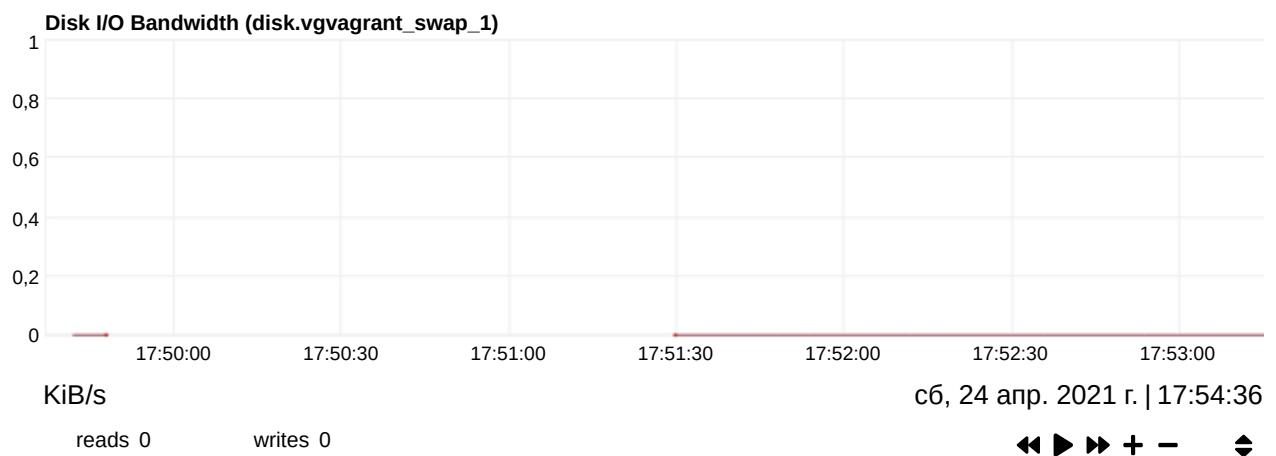
Disk Total I/O Time (disk_iotime.vgvagrant_root)



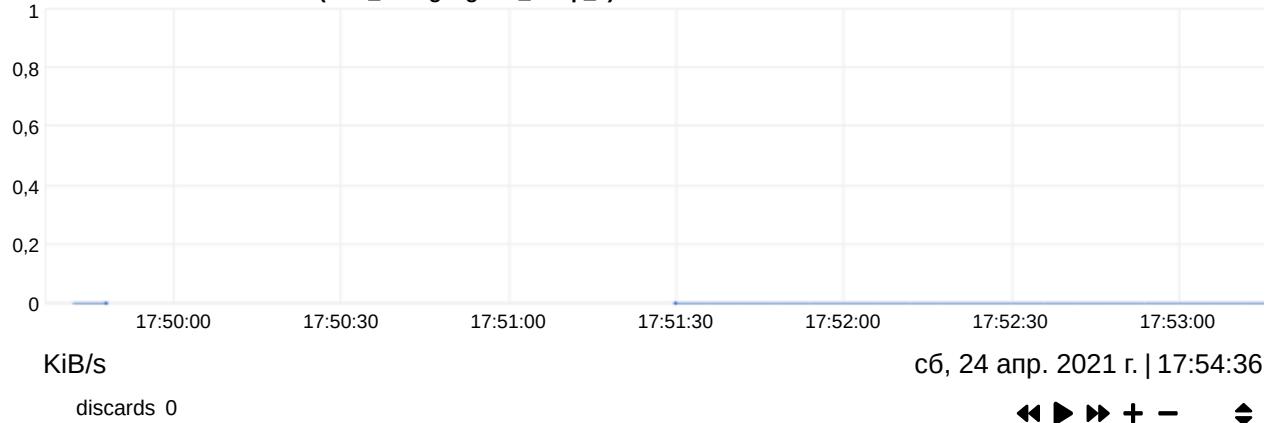


vgvagrant-swap 1

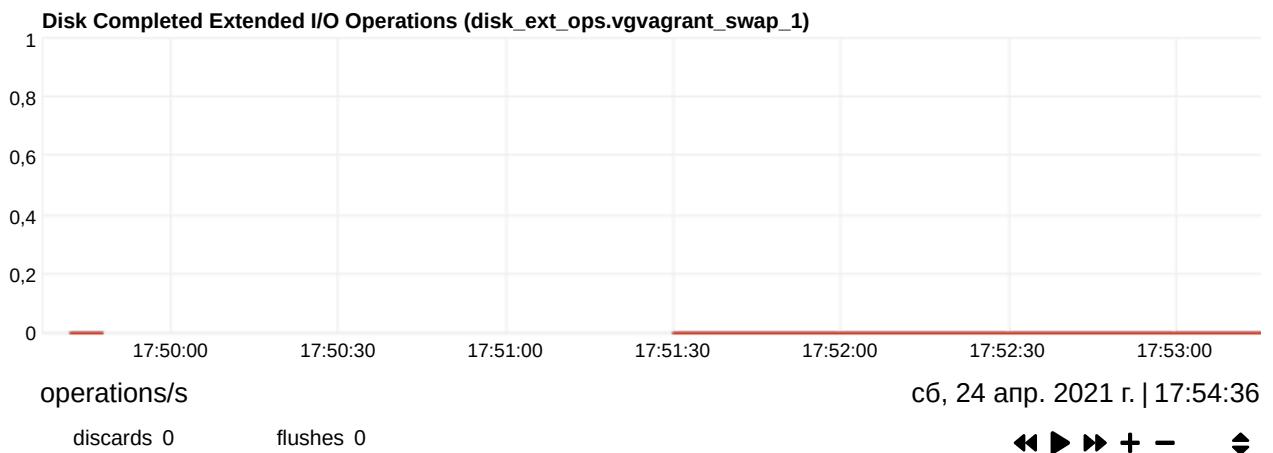
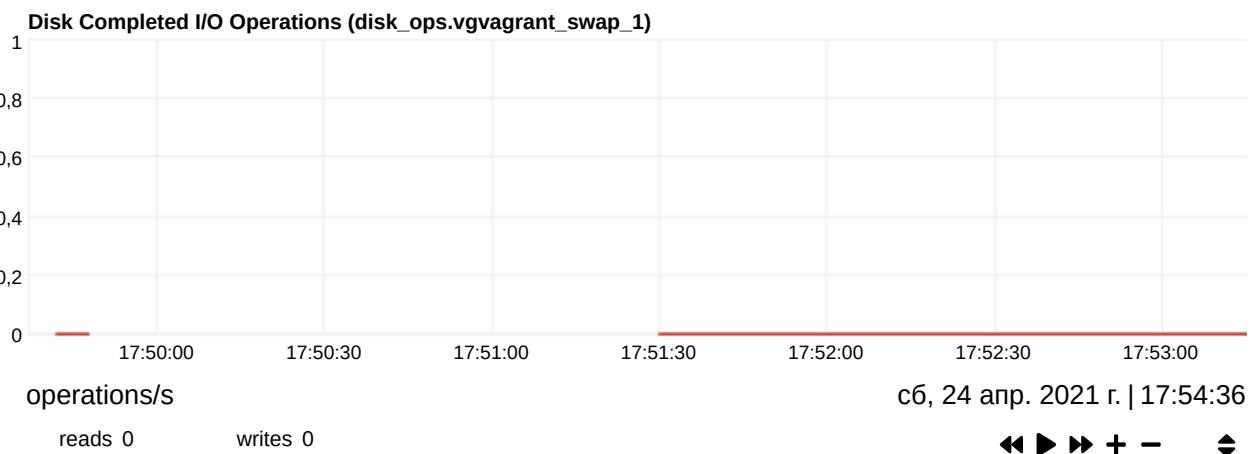
Amount of data transferred to and from disk.



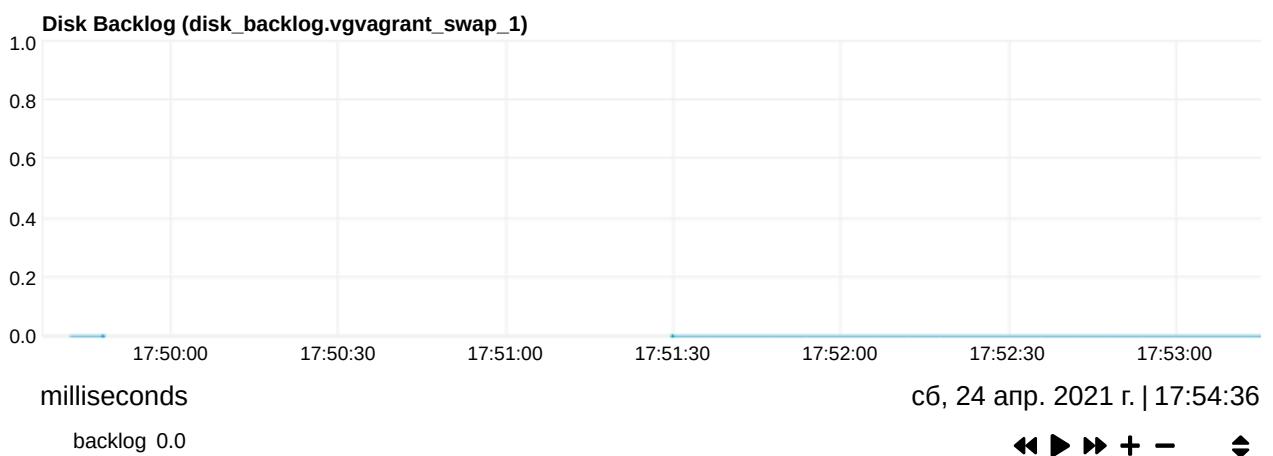
Amount of Discarded Data (disk_ext.vg vagrant_swap_1)



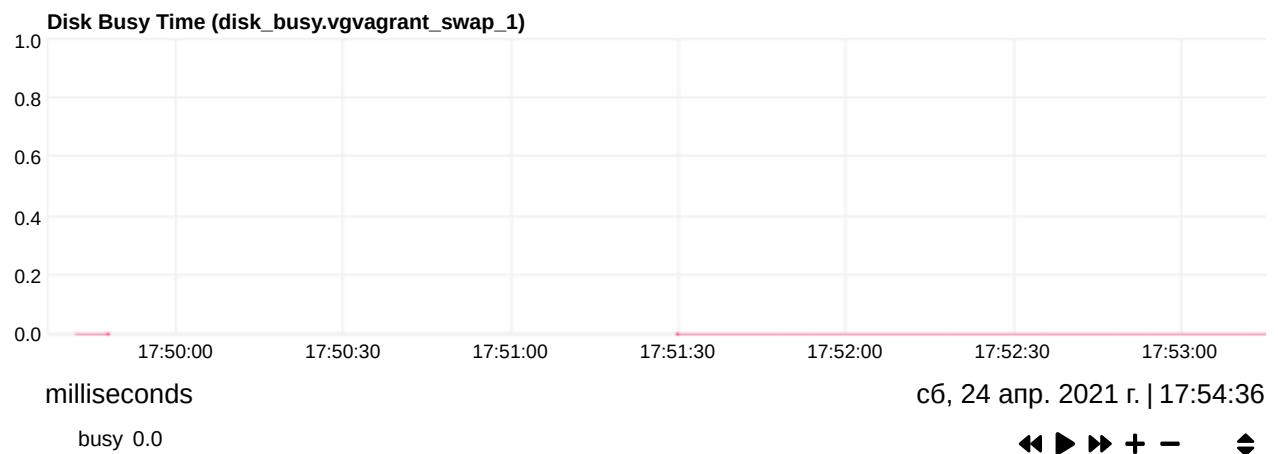
Completed disk I/O operations. Keep in mind the number of operations requested might be higher, since the system is able to merge adjacent to each other (see merged operations chart).



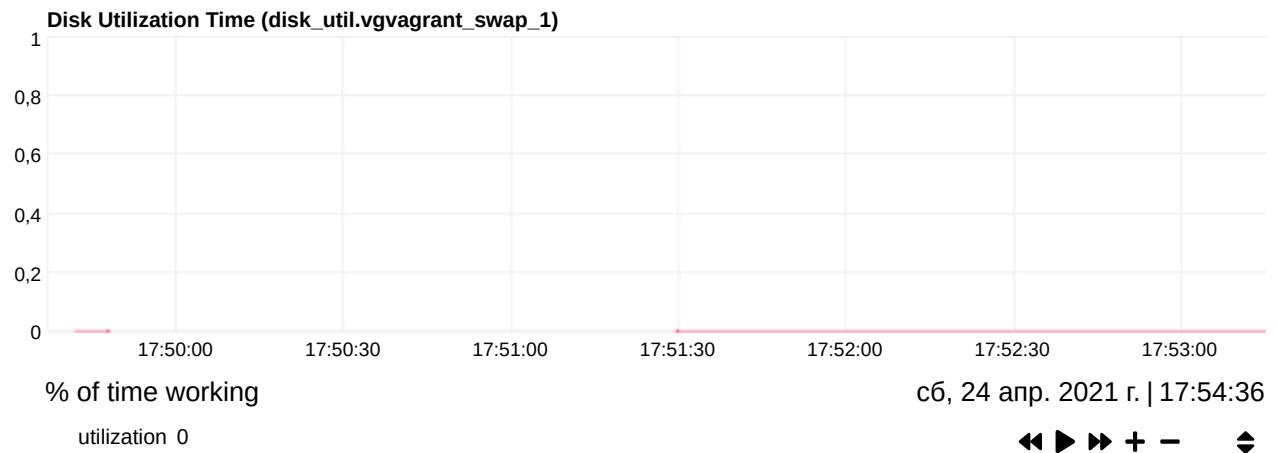
Backlog is an indication of the duration of pending disk operations. On every I/O event the system is multiplying the time spent doing I/O since the last update of this field with the number of pending operations. While not accurate, this metric can provide an indication of the expected completion time of the operations in progress.



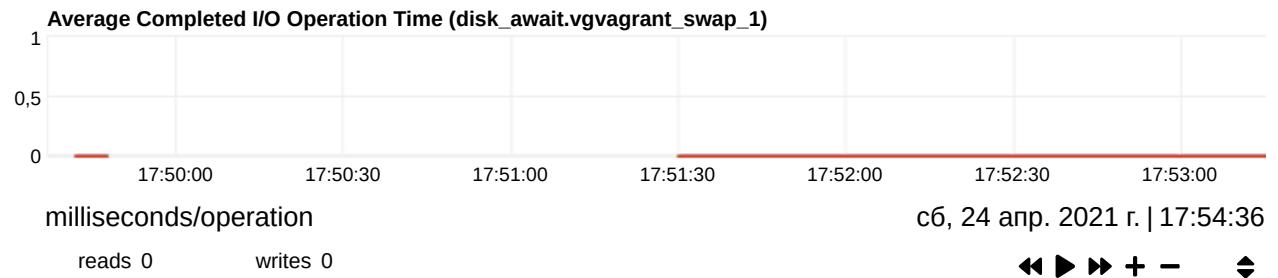
Disk Busy Time measures the amount of time the disk was busy with something.



Disk Utilization measures the amount of time the disk was busy with something. This is not related to its performance. 100% means that the system always had an outstanding operation on the disk. Keep in mind that depending on the underlying technology of the disk, 100% here may or may not be an indication of congestion.

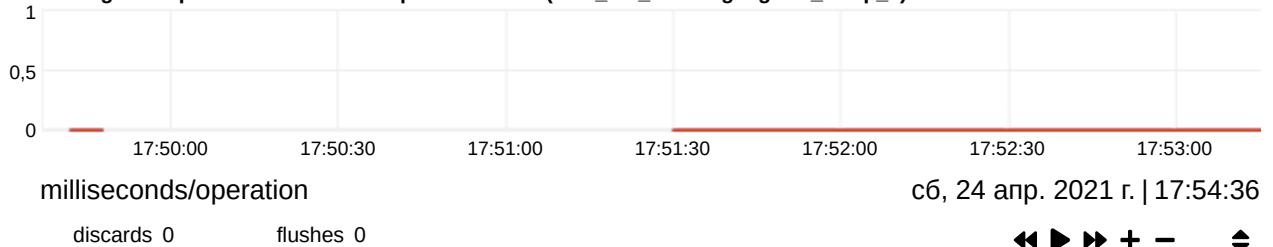


The average time for I/O requests issued to the device to be served. This includes the time spent by the requests in queue and the time spent servicing them.



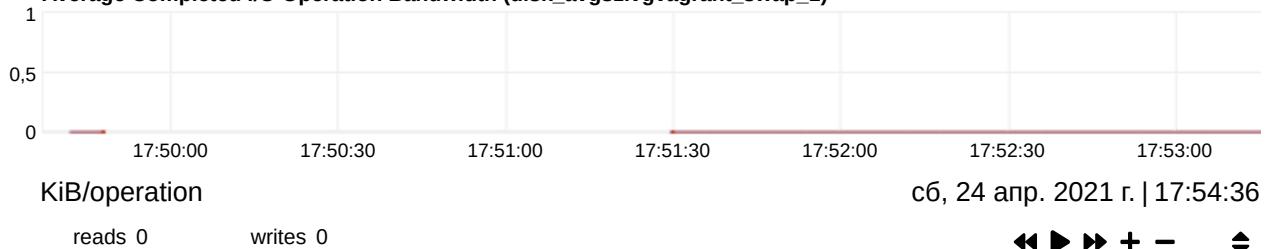
The average time for extended I/O requests issued to the device to be served. This includes the time spent by the requests in queue and the time spent servicing them.

Average Completed Extended I/O Operation Time (disk_ext_await.vgvagrant_swap_1)

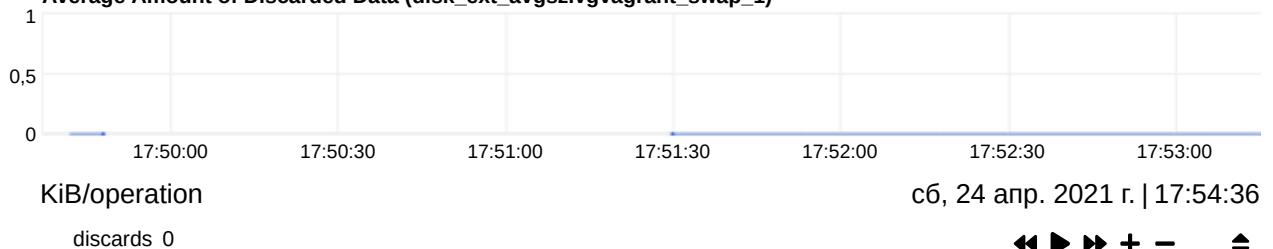


The average I/O operation size.

Average Completed I/O Operation Bandwidth (disk_avgsz.vgvagrant_swap_1)

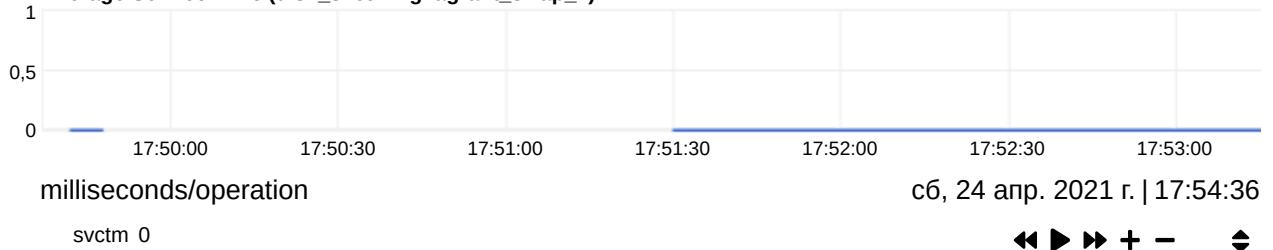


Average Amount of Discarded Data (disk_ext_avgsz.vgvagrant_swap_1)



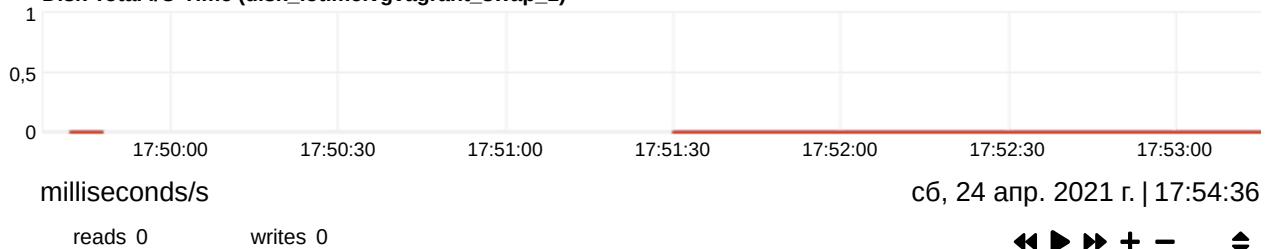
The average service time for completed I/O operations. This metric is calculated using the total busy time of the disk and the number of completed operations. If the disk is able to execute multiple parallel operations the reporting average service time will be misleading.

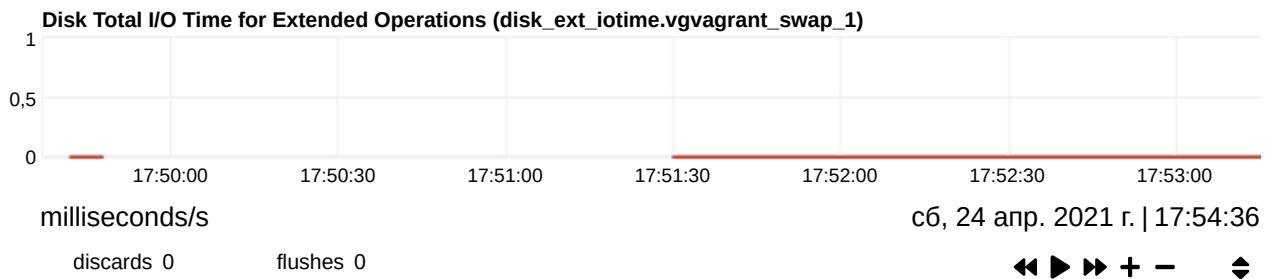
Average Service Time (disk_svctm.vgvagrant_swap_1)



The sum of the duration of all completed I/O operations. This number can exceed the interval if the disk is able to execute I/O operations in parallel.

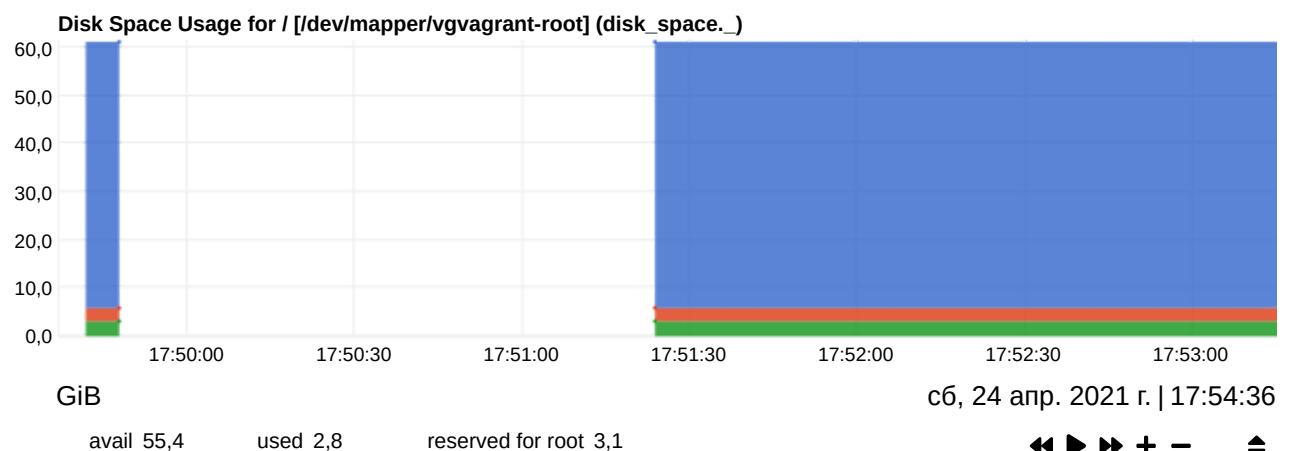
Disk Total I/O Time (disk_iotime.vgvagrant_swap_1)





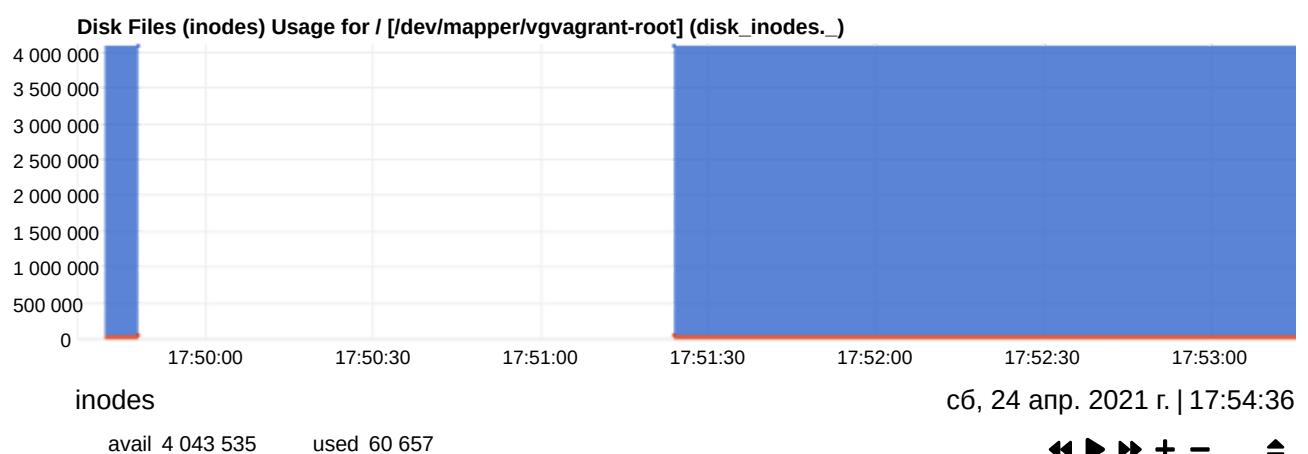
/

Disk space utilization. reserved for root is automatically reserved by the system to prevent the root user from getting out of space.



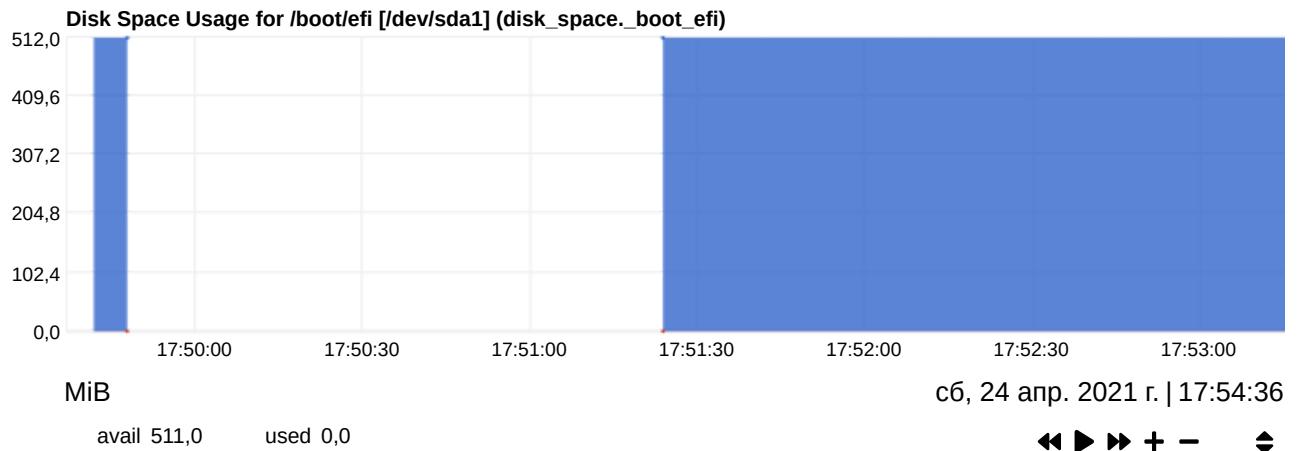
inodes (or index nodes) are filesystem objects (e.g. files and directories). On many types of file system implementations, the maximum number of inodes is fixed at filesystem creation, limiting the maximum number of files the filesystem can hold. It is possible for a device to run out of inodes.

When this happens, new files cannot be created on the device, even though there may be free space available.



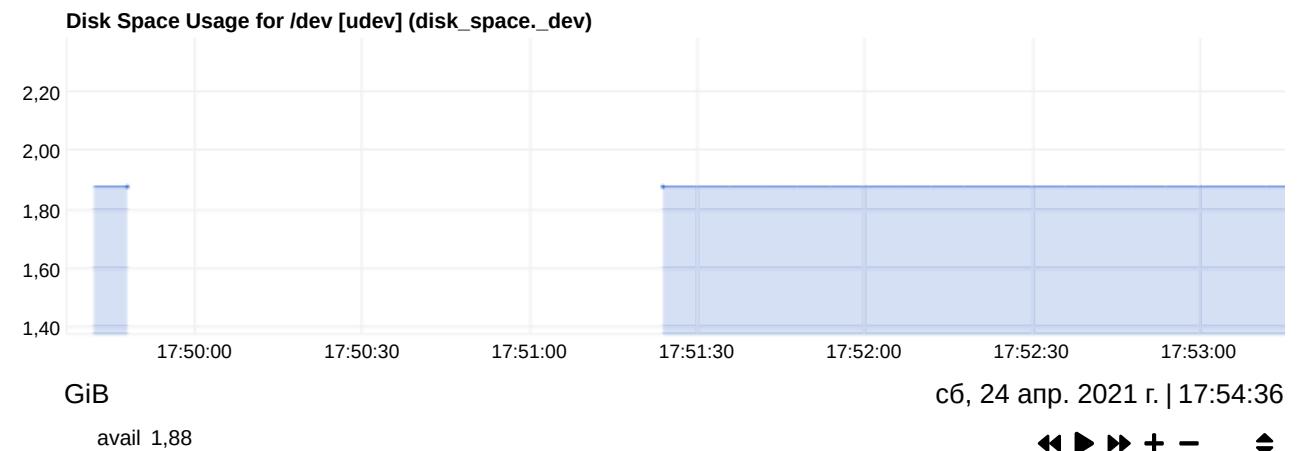
/boot/efi

Disk space utilization. reserved for root is automatically reserved by the system to prevent the root user from getting out of space.

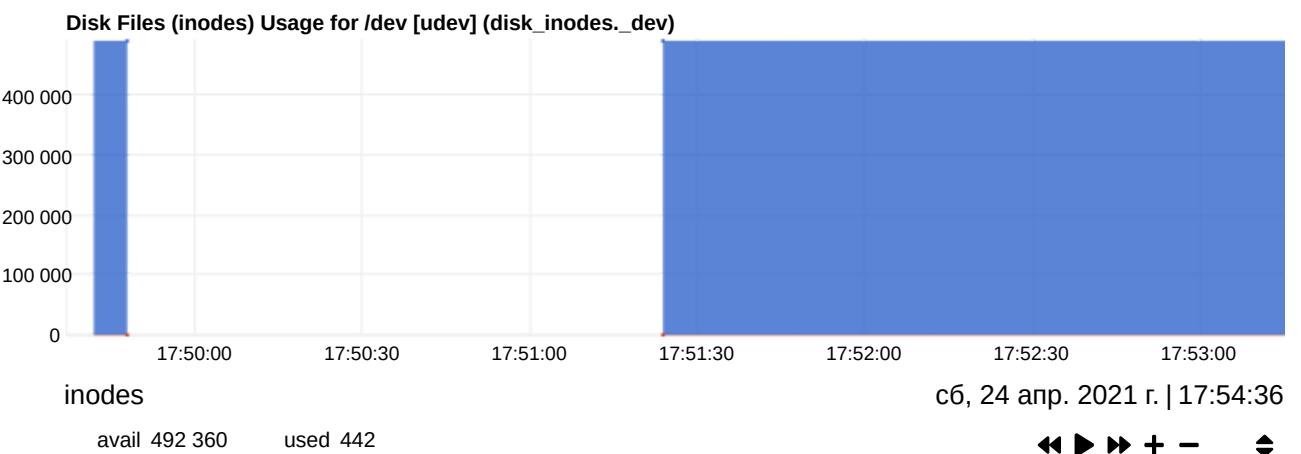


/dev

Disk space utilization. reserved for root is automatically reserved by the system to prevent the root user from getting out of space.

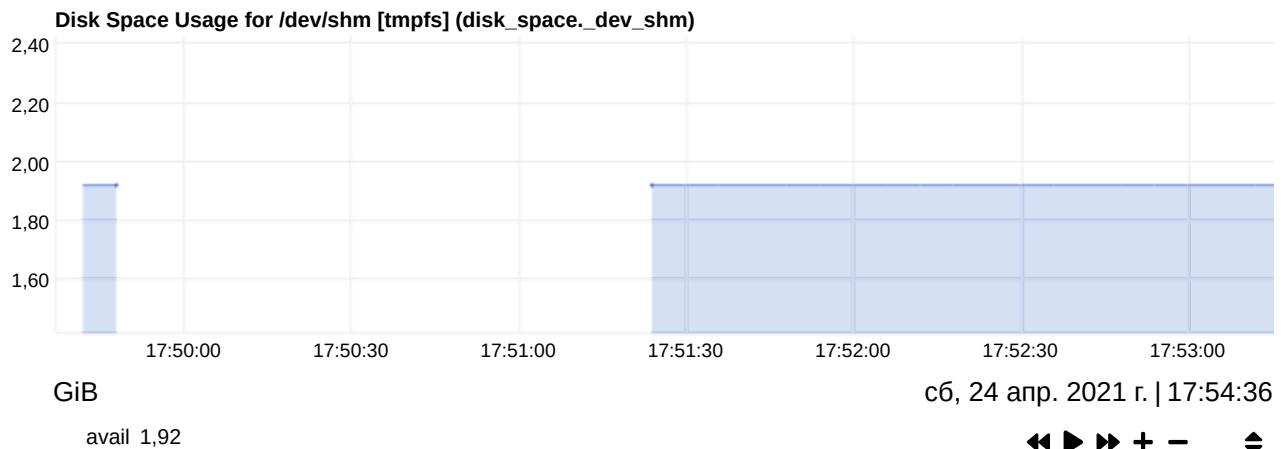


inodes (or index nodes) are filesystem objects (e.g. files and directories). On many types of file system implementations, the maximum number of inodes is fixed at filesystem creation, limiting the maximum number of files the filesystem can hold. It is possible for a device to run out of inodes. When this happens, new files cannot be created on the device, even though there may be free space available.

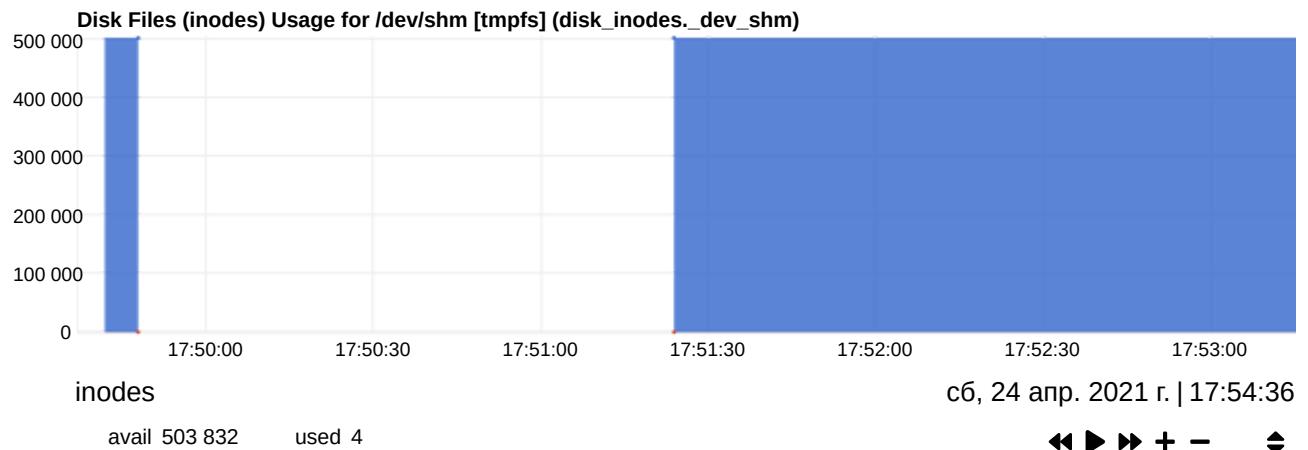


/dev/shm

Disk space utilization. reserved for root is automatically reserved by the system to prevent the root user from getting out of space.



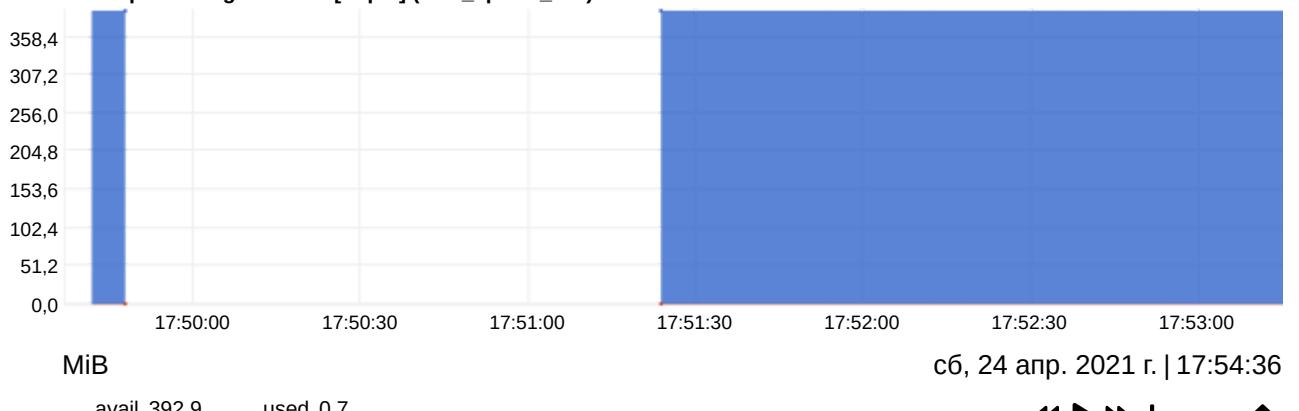
inodes (or index nodes) are filesystem objects (e.g. files and directories). On many types of file system implementations, the maximum number of inodes is fixed at filesystem creation, limiting the maximum number of files the filesystem can hold. It is possible for a device to run out of inodes. When this happens, new files cannot be created on the device, even though there may be free space available.



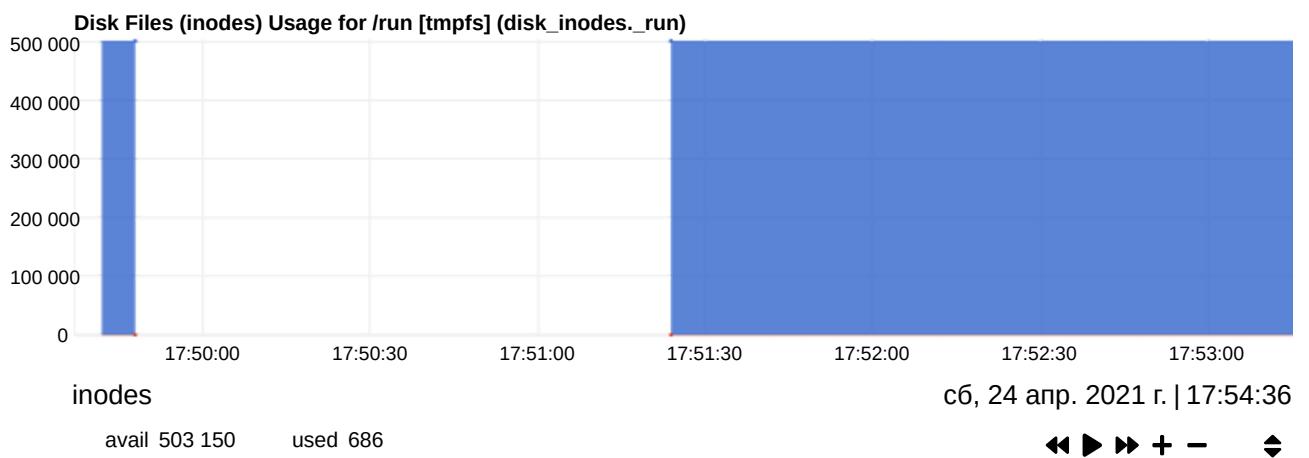
/run

Disk space utilization. reserved for root is automatically reserved by the system to prevent the root user from getting out of space.

Disk Space Usage for /run [tmpfs] (disk_space._run)



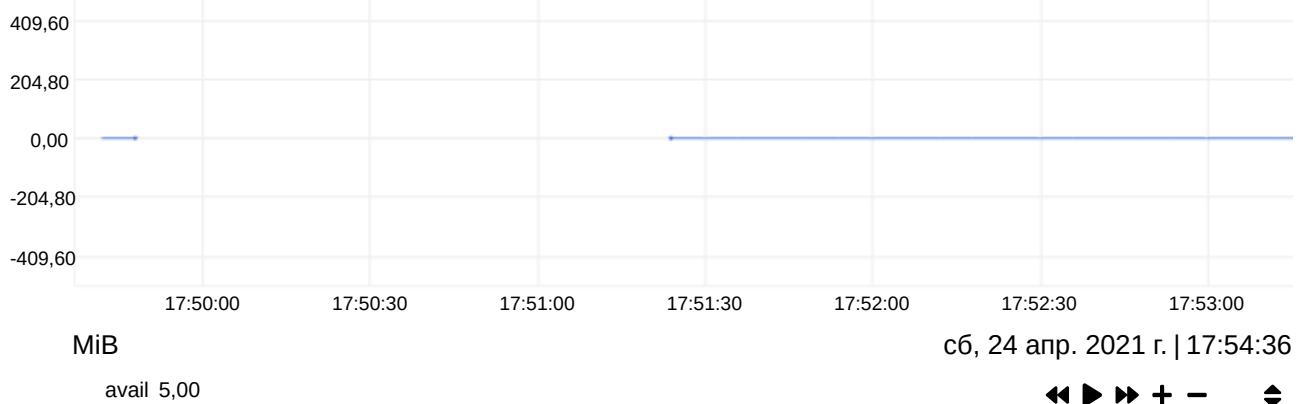
inodes (or index nodes) are filesystem objects (e.g. files and directories). On many types of file system implementations, the maximum number of inodes is fixed at filesystem creation, limiting the maximum number of files the filesystem can hold. It is possible for a device to run out of inodes. When this happens, new files cannot be created on the device, even though there may be free space available.



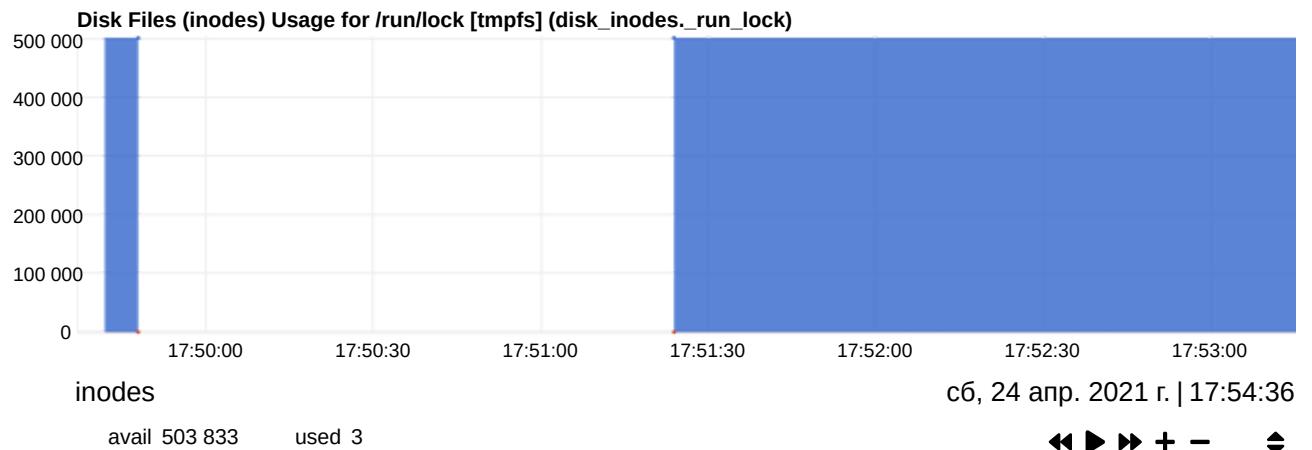
/run/lock

Disk space utilization. reserved for root is automatically reserved by the system to prevent the root user from getting out of space.

Disk Space Usage for /run/lock [tmpfs] (disk_space._run_lock)

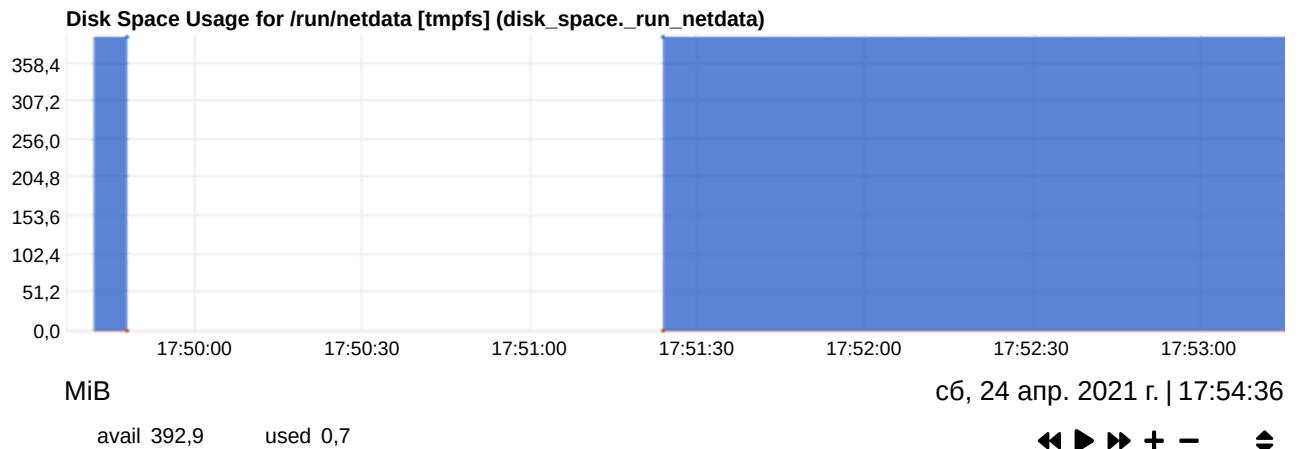


inodes (or index nodes) are filesystem objects (e.g. files and directories). On many types of file system implementations, the maximum number of inodes is fixed at filesystem creation, limiting the maximum number of files the filesystem can hold. It is possible for a device to run out of inodes. When this happens, new files cannot be created on the device, even though there may be free space available.

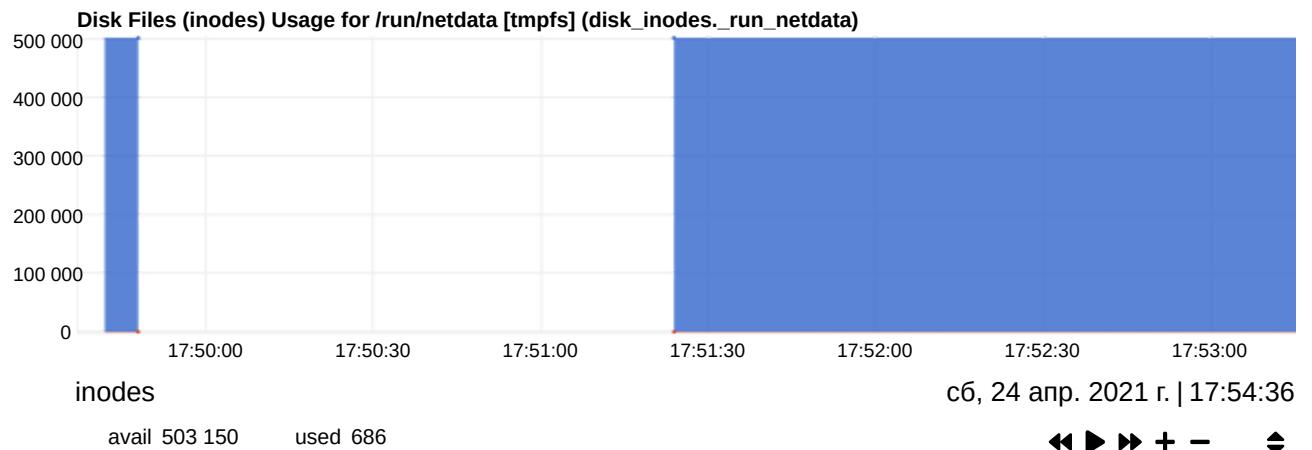


/run/netdata

Disk space utilization. reserved for root is automatically reserved by the system to prevent the root user from getting out of space.

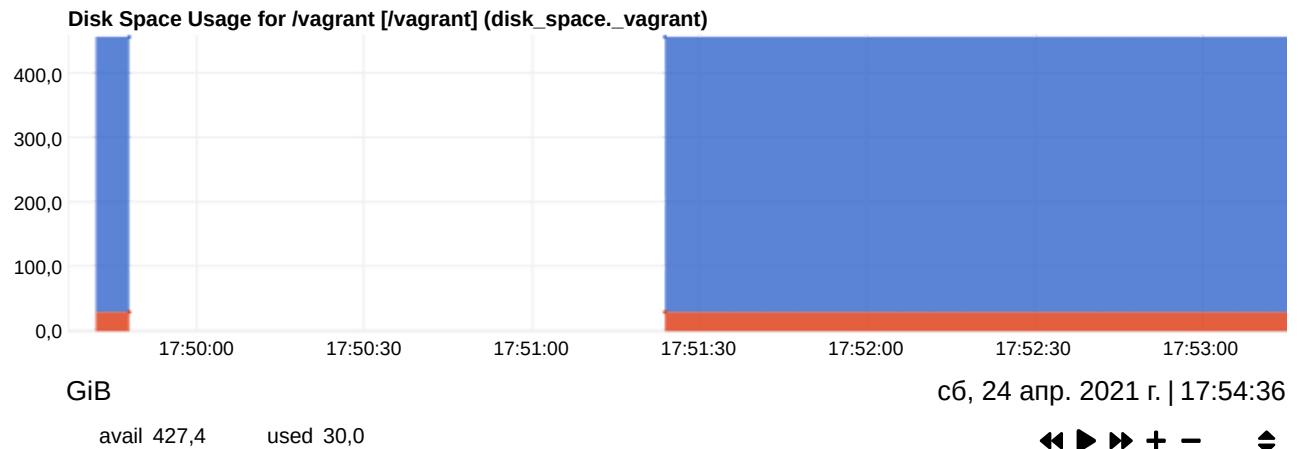


inodes (or index nodes) are filesystem objects (e.g. files and directories). On many types of file system implementations, the maximum number of inodes is fixed at filesystem creation, limiting the maximum number of files the filesystem can hold. It is possible for a device to run out of inodes. When this happens, new files cannot be created on the device, even though there may be free space available.

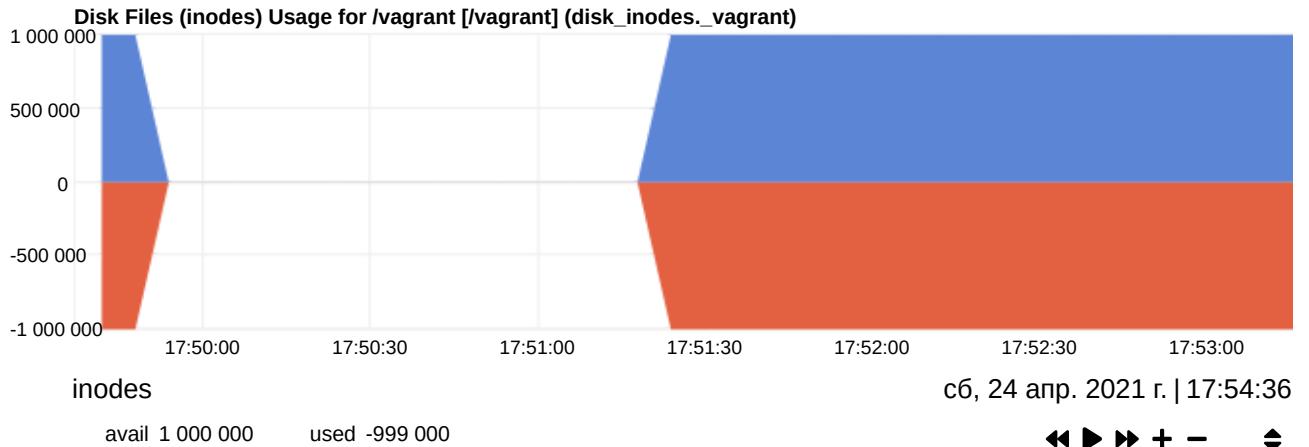


/vagrant

Disk space utilization. reserved for root is automatically reserved by the system to prevent the root user from getting out of space.



inodes (or index nodes) are filesystem objects (e.g. files and directories). On many types of file system implementations, the maximum number of inodes is fixed at filesystem creation, limiting the maximum number of files the filesystem can hold. It is possible for a device to run out of inodes. When this happens, new files cannot be created on the device, even though there may be free space available.

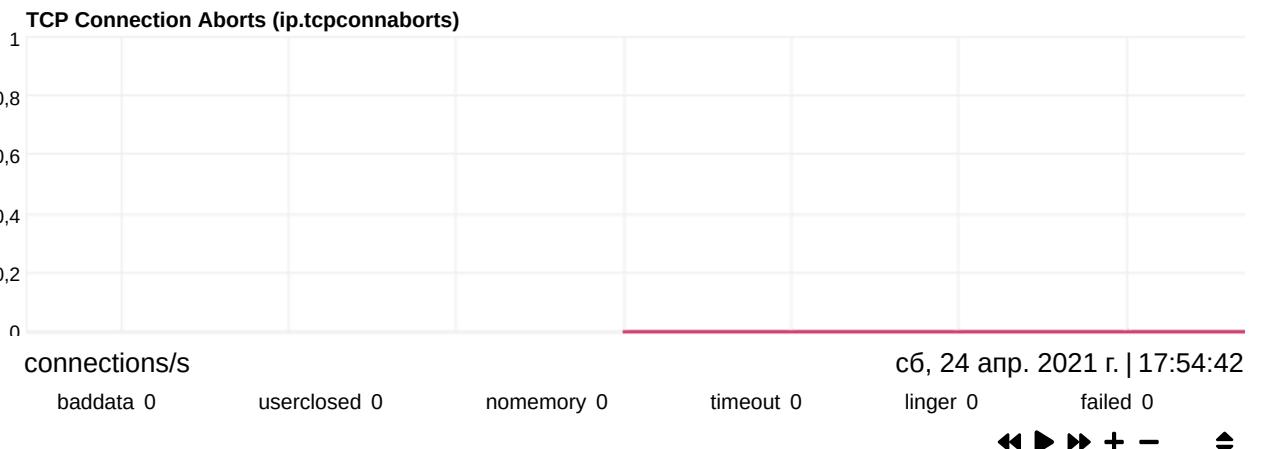


Networking Stack

Metrics for the networking stack of the system. These metrics are collected from `/proc/net/netstat`, apply to both IPv4 and IPv6 traffic and are related to operation of the kernel networking stack.

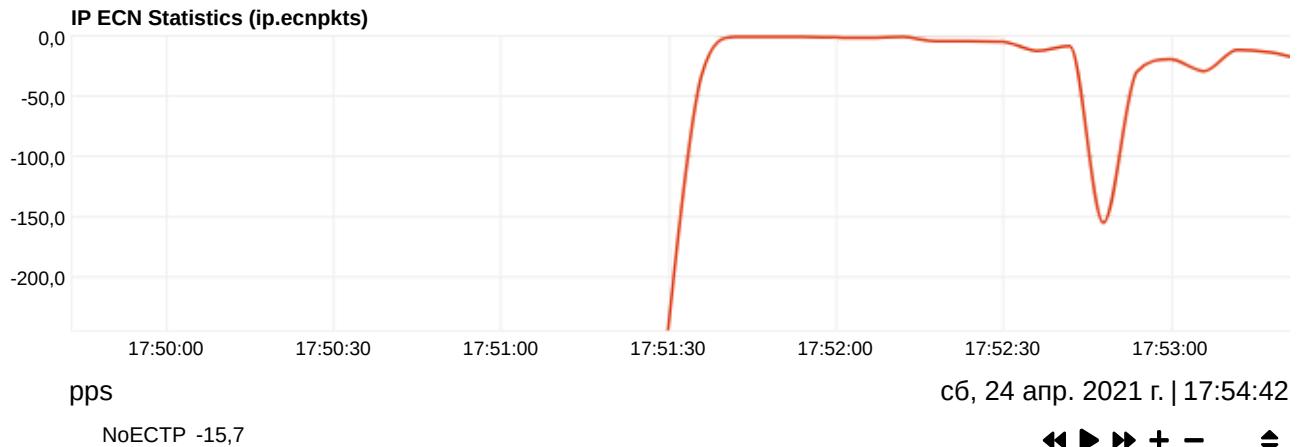
tcp

TCP connection aborts. **baddata** (`TCPAbortOnData`) happens while the connection is on `FIN_WAIT1` and the kernel receives a packet with a sequence number beyond the last one for this connection - the kernel responds with `RST` (closes the connection). **userclosed** (`TCPAbortOnClose`) happens when the kernel receives data on an already closed connection and responds with `RST`. **nomemory** (`TCPAbortOnMemory`) happens when there are too many orphaned sockets (not attached to an fd) and the kernel has to drop a connection - sometimes it will send an `RST`, sometimes it won't. **timeout** (`TCPAbortOnTimeout`) happens when a connection times out. **linger** (`TCPAbortOnLinger`) happens when the kernel killed a socket that was already closed by the application and lingered around for long enough. **failed** (`TCPAbortFailed`) happens when the kernel attempted to send an `RST` but failed because there was no memory available.



ecn

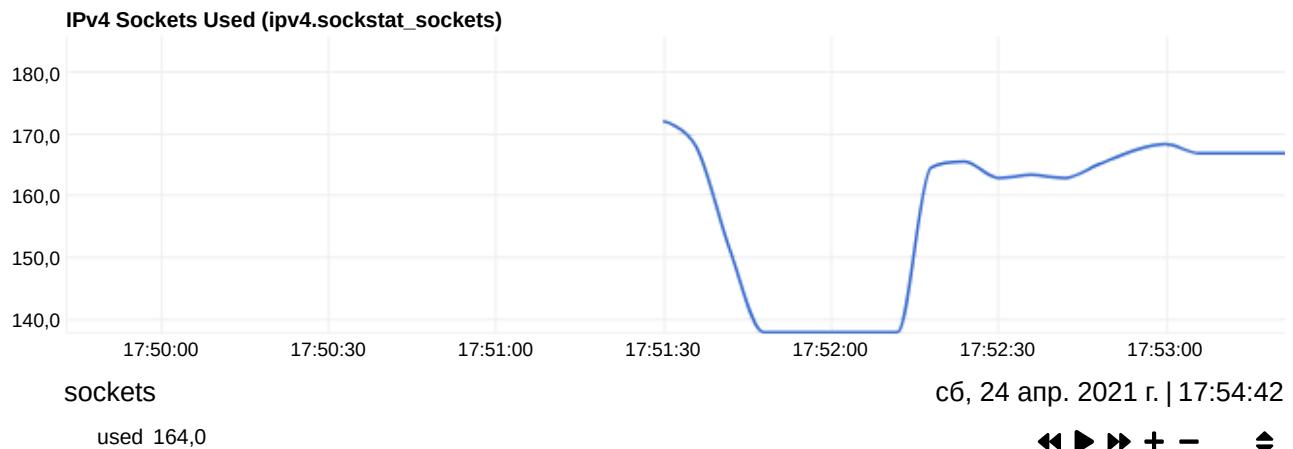
Explicit Congestion Notification (ECN) (https://en.wikipedia.org/wiki/Explicit_Congestion_Notation) is a TCP extension that allows end-to-end notification of network congestion without dropping packets. ECN is an optional feature that may be used between two ECN-enabled endpoints when the underlying network infrastructure also supports it.



IPv4 Networking

Metrics for the IPv4 stack of the system. Internet Protocol version 4 (IPv4) (<https://en.wikipedia.org/wiki/IPv4>) is the fourth version of the Internet Protocol (IP). It is one of the core protocols of standards-based internetworking methods in the Internet. IPv4 is a connectionless protocol for use on packet-switched networks. It operates on a best effort delivery model, in that it does not guarantee delivery, nor does it assure proper sequencing or avoidance of duplicate delivery. These aspects, including data integrity, are addressed by an upper layer transport protocol, such as the Transmission Control Protocol (TCP).

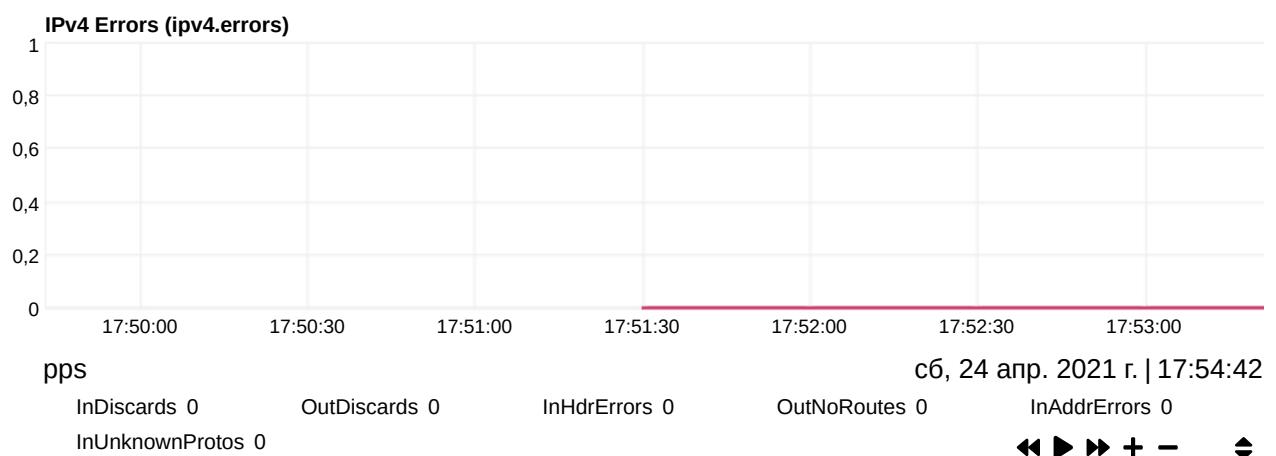
sockets



packets

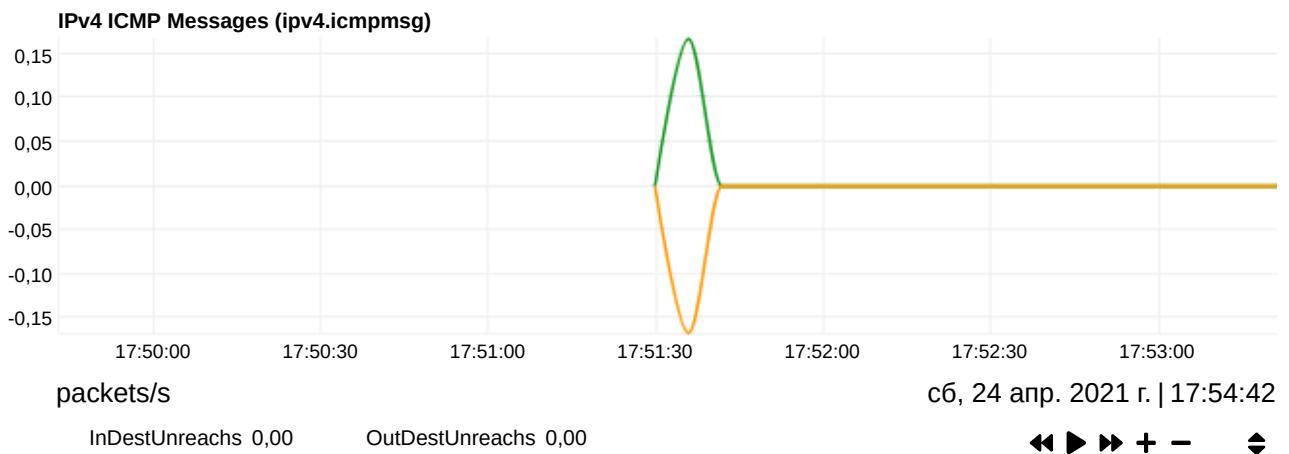
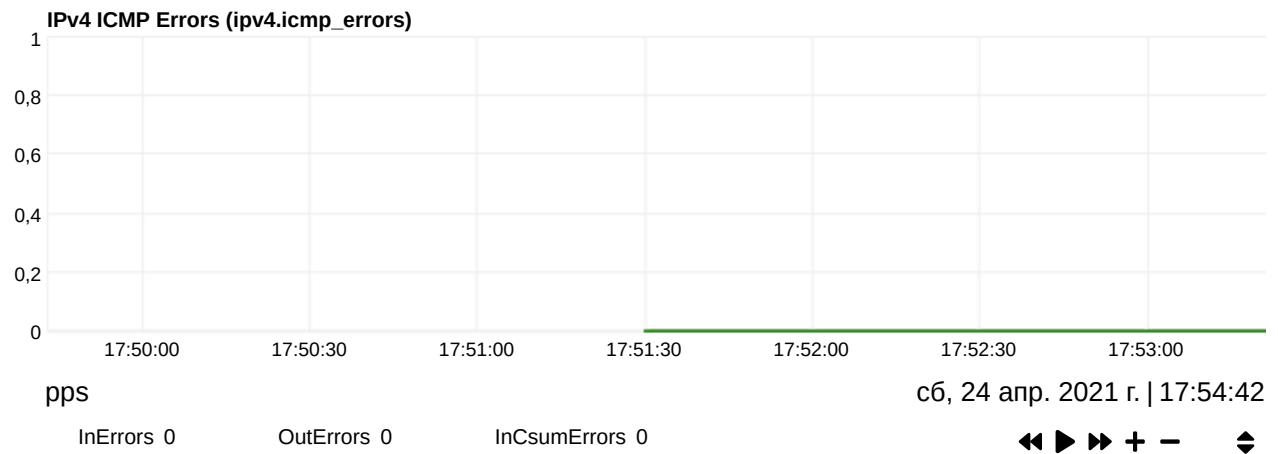


errors



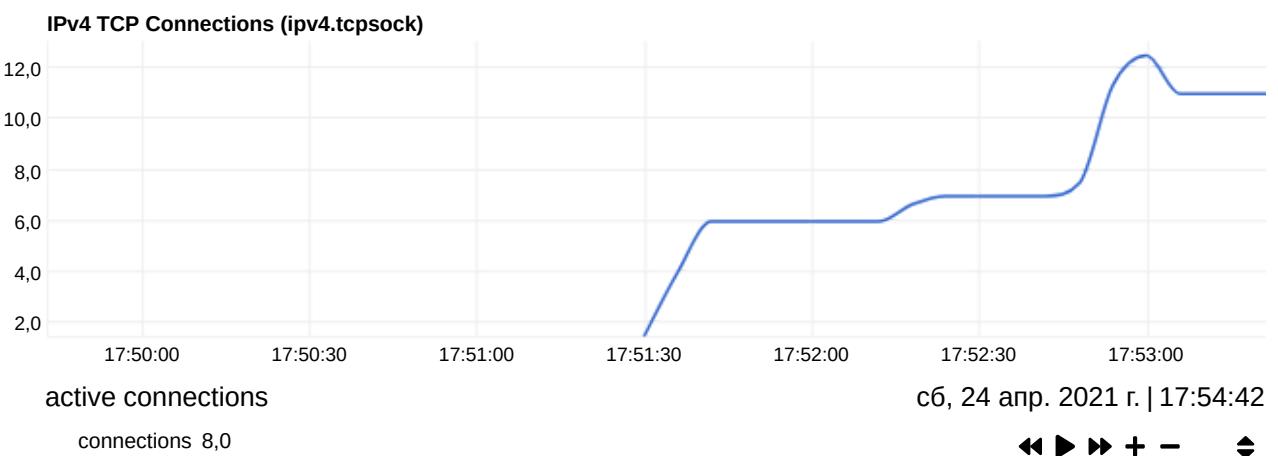
icmp



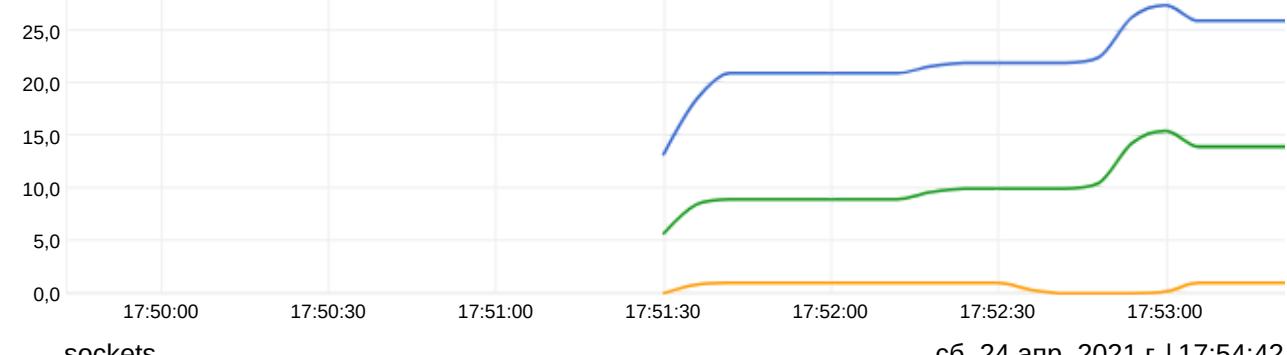


tcp

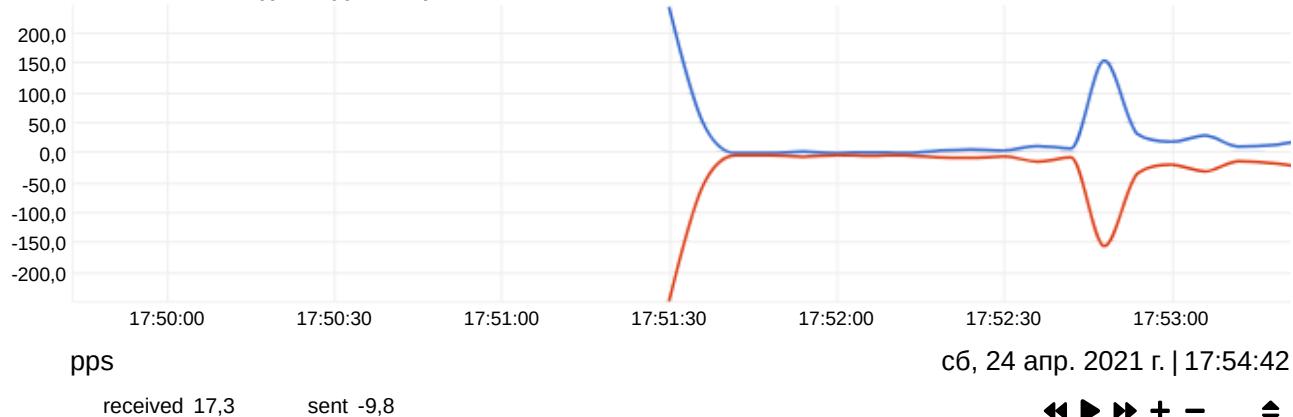
The number of established TCP connections (known as `currEstab`). This is a snapshot of the established connections at the time of measurement (i.e. a connection established and a connection disconnected within the same iteration will not affect this metric).



IPv4 TCP Sockets (ipv4.sockstat_tcp_sockets)



IPv4 TCP Packets (ipv4.tcppackets)



active or **ActiveOpens** is the number of outgoing TCP **connections attempted** by this host.

passive or **PassiveOpens** is the number of incoming TCP **connections accepted** by this host.

IPv4 TCP Opens (ipv4.tcpopens)



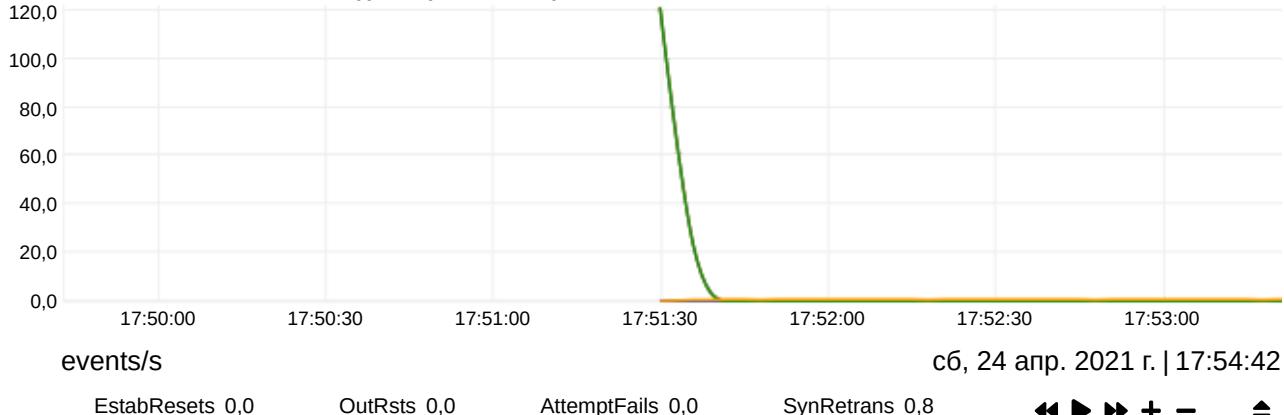
`InErrs` is the number of TCP segments received in error (including header too small, checksum errors, sequence errors, bad packets - for both IPv4 and IPv6). `InCsumErrors` is the number of TCP segments received with checksum errors (for both IPv4 and IPv6). `RetransSegs` is the number of TCP segments retransmitted.

IPv4 TCP Errors (ipv4.tcperrors)



`EstabResets` is the number of established connections resets (i.e. connections that made a direct transition from `ESTABLISHED` or `CLOSE_WAIT` to `CLOSED`). `OutRsts` is the number of TCP segments sent, with the `RST` flag set (for both IPv4 and IPv6). `AttemptFails` is the number of times TCP connections made a direct transition from either `SYN_SENT` or `SYN_RECV` to `CLOSED`, plus the number of times TCP connections made a direct transition from the `SYN_RECV` to `LISTEN`. `TCPSynRetrans` shows retries for new outbound TCP connections, which can indicate general connectivity issues or backlog on the remote host.

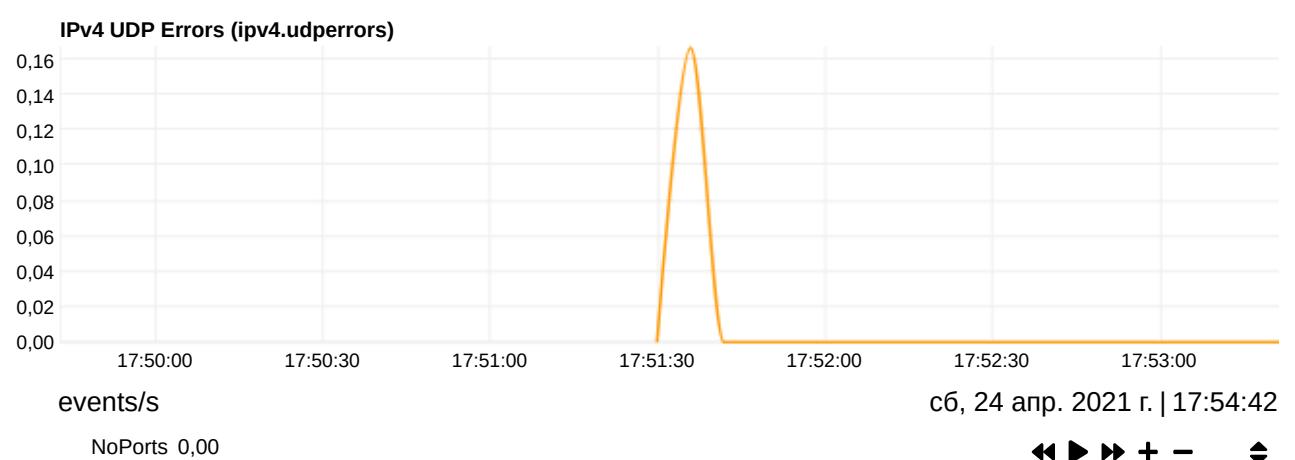
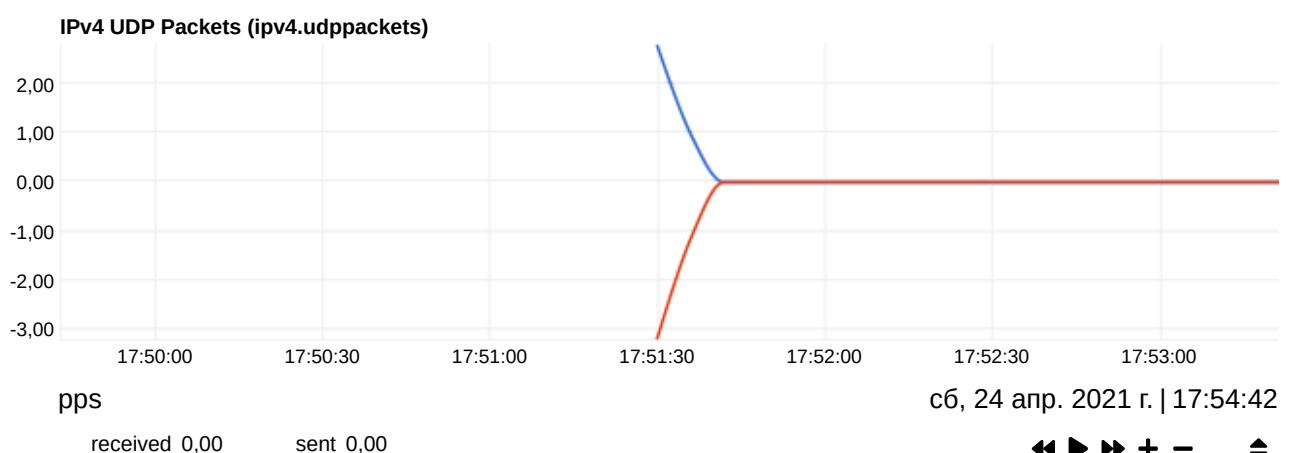
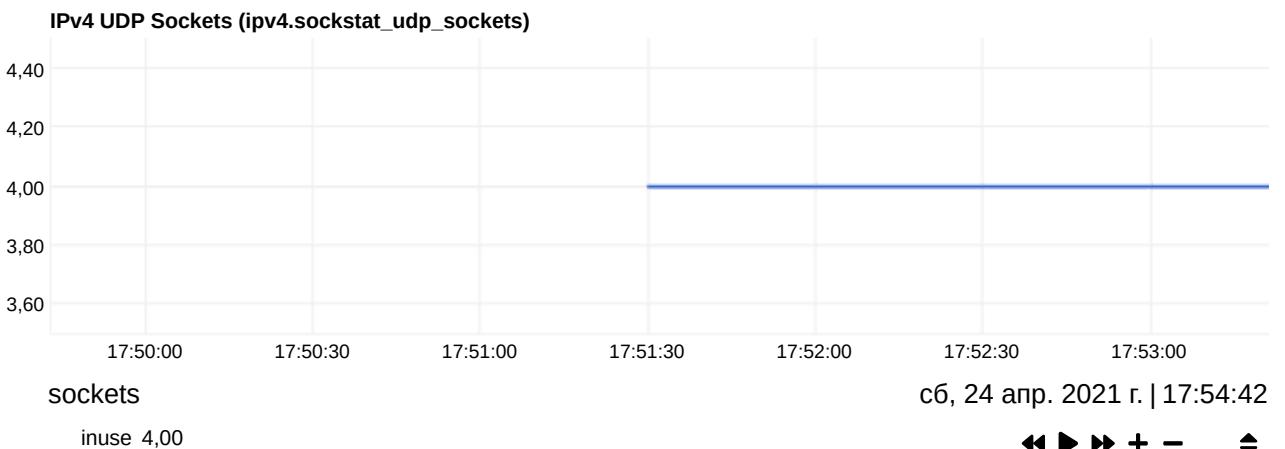
IPv4 TCP Handshake Issues (ipv4.tcphandshake)

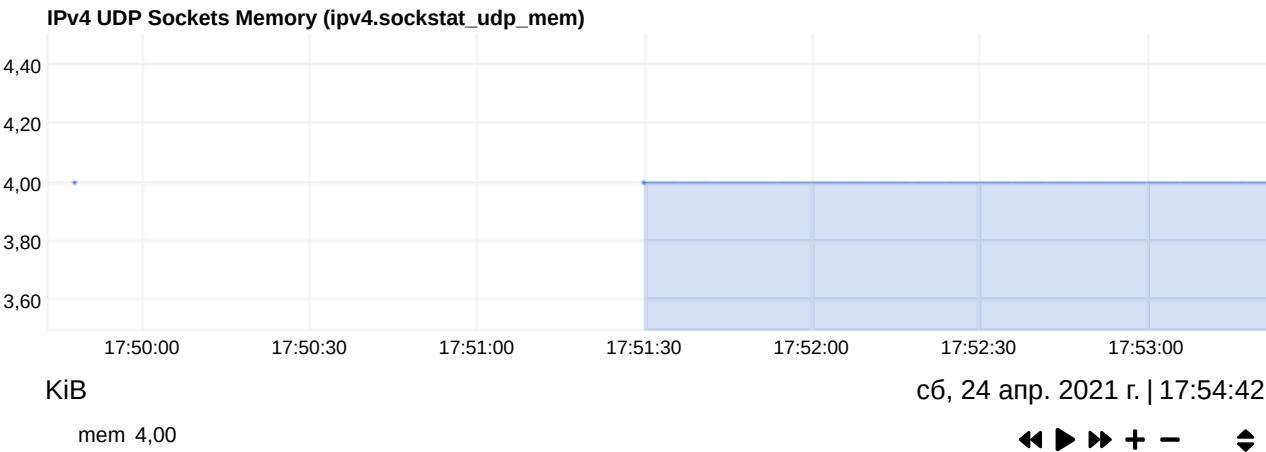


IPv4 TCP Sockets Memory (ipv4.sockstat_tcp_mem)



udp





IPv6 Networking

Metrics for the IPv6 stack of the system. Internet Protocol version 6 (IPv6) (<https://en.wikipedia.org/wiki/IPv6>) is the most recent version of the Internet Protocol (IP), the communications protocol that provides an identification and location system for computers on networks and routes traffic across the Internet. IPv6 was developed by the Internet Engineering Task Force (IETF) to deal with the long-anticipated problem of IPv4 address exhaustion. IPv6 is intended to replace IPv4.

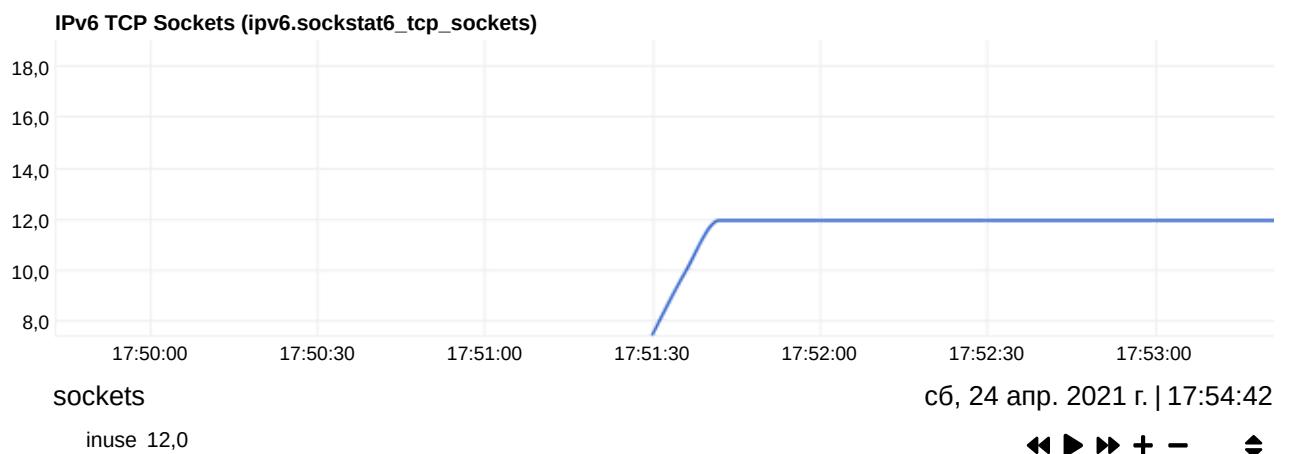
packets



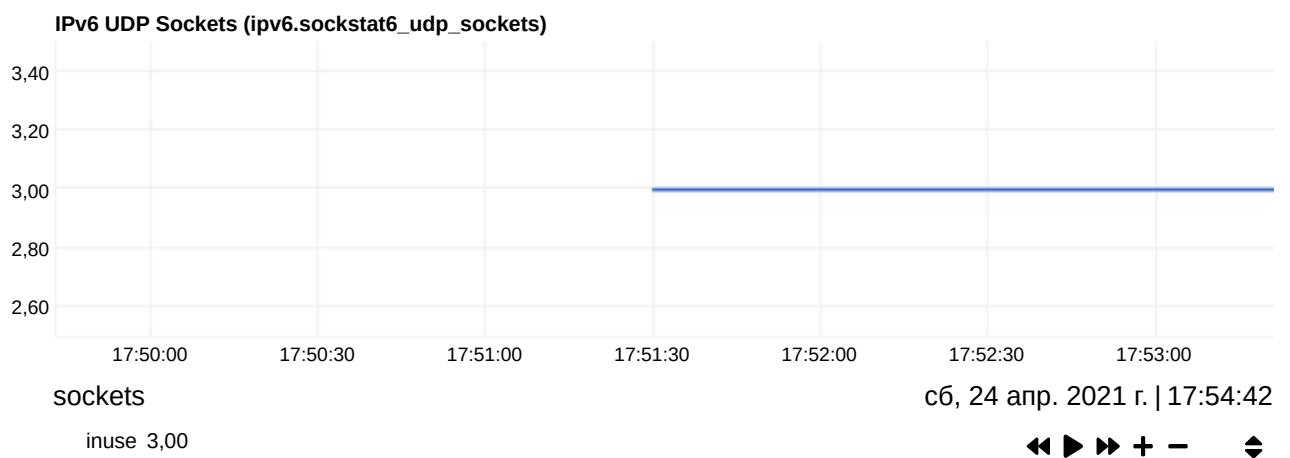
IPv6 ECT Packets (ipv6.ect)

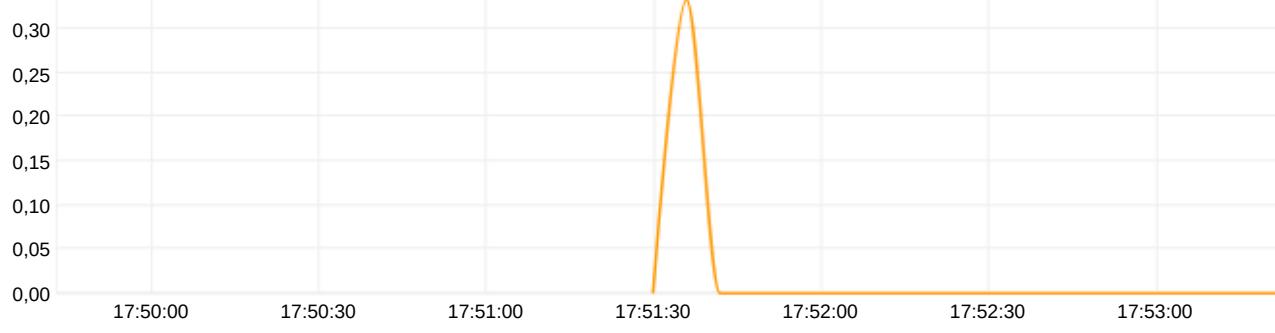
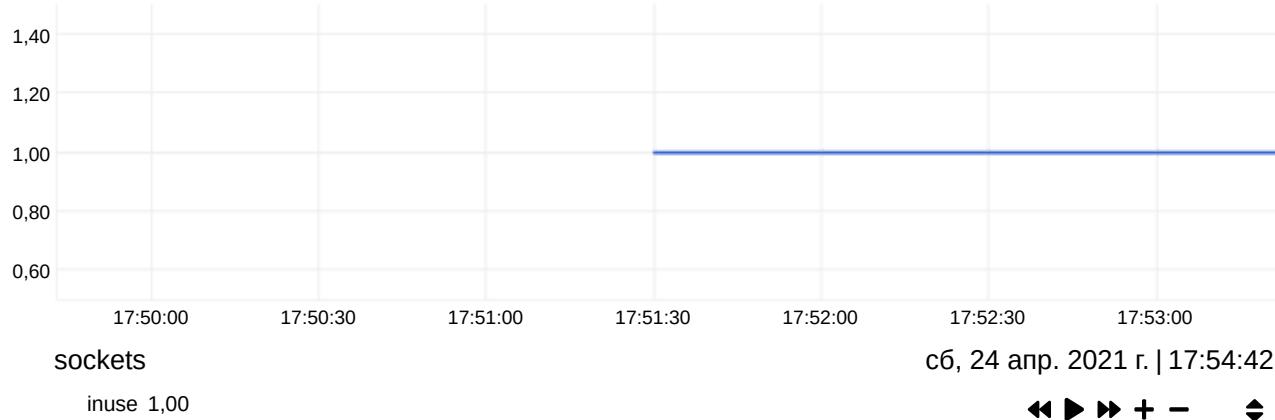


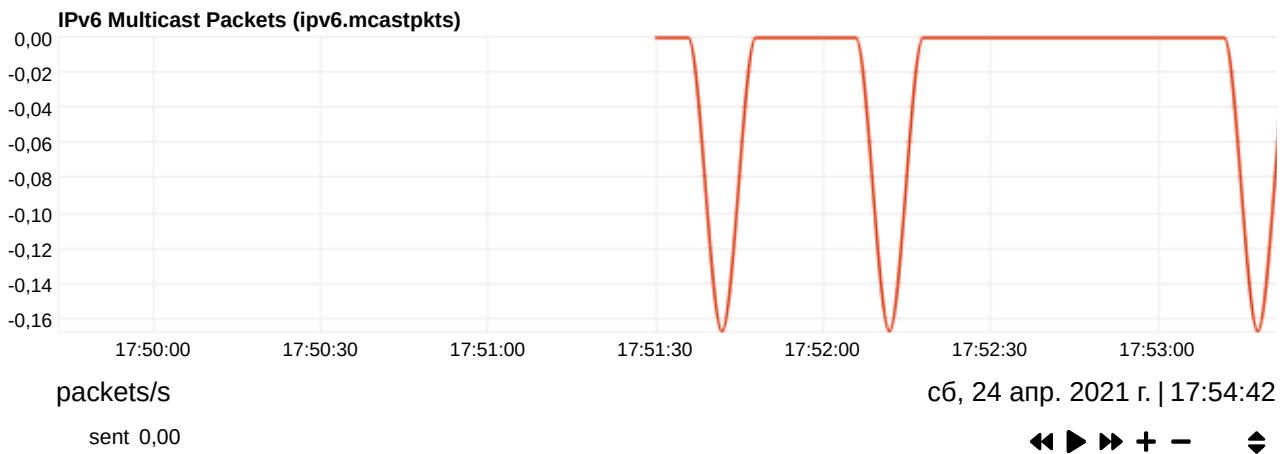
tcp6



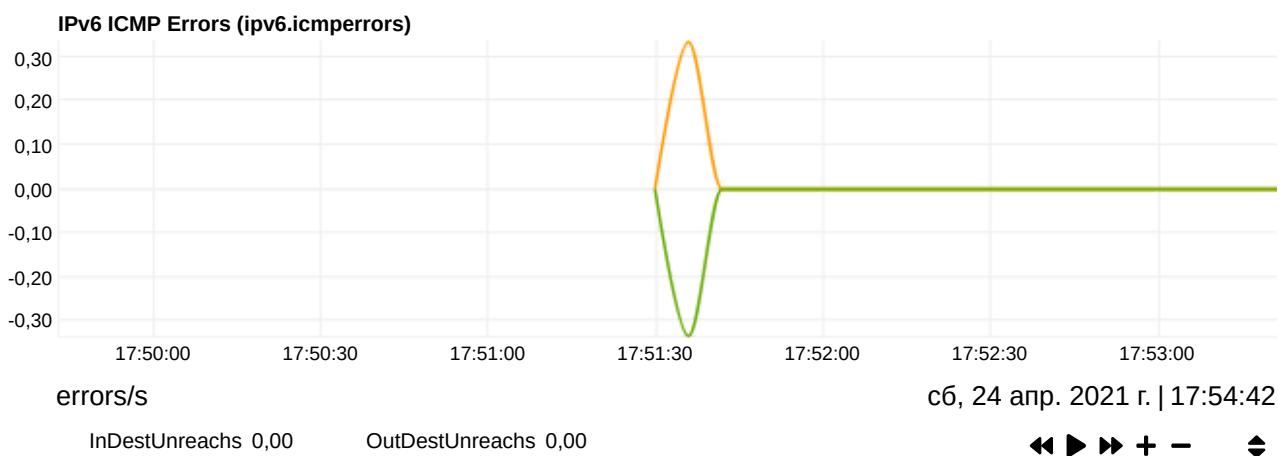
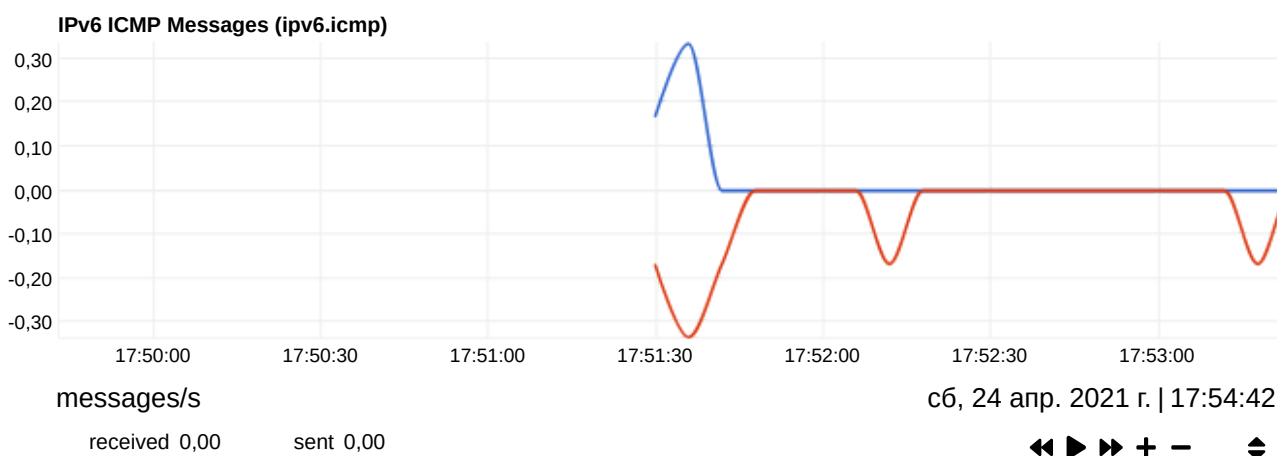
udp6

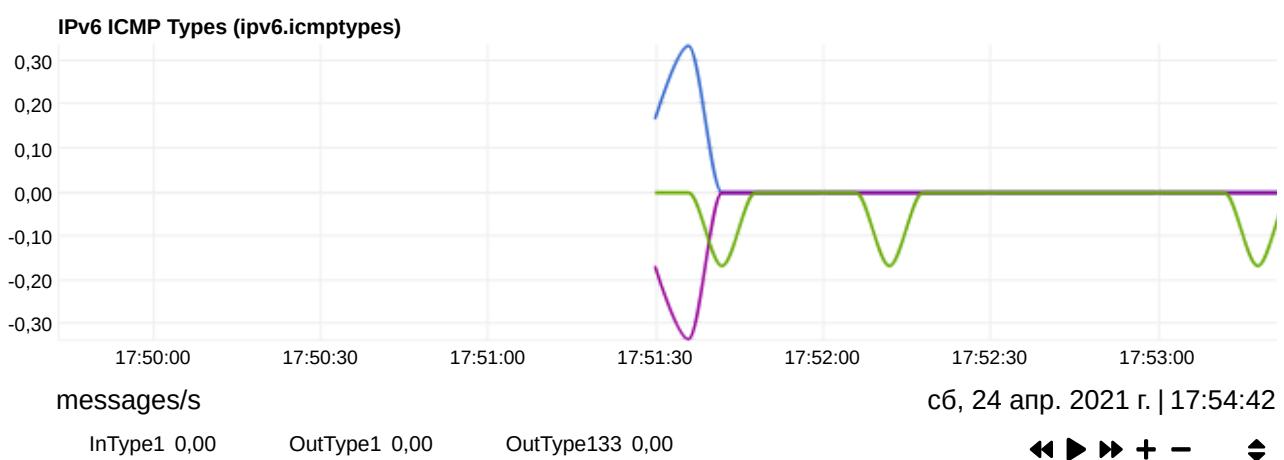
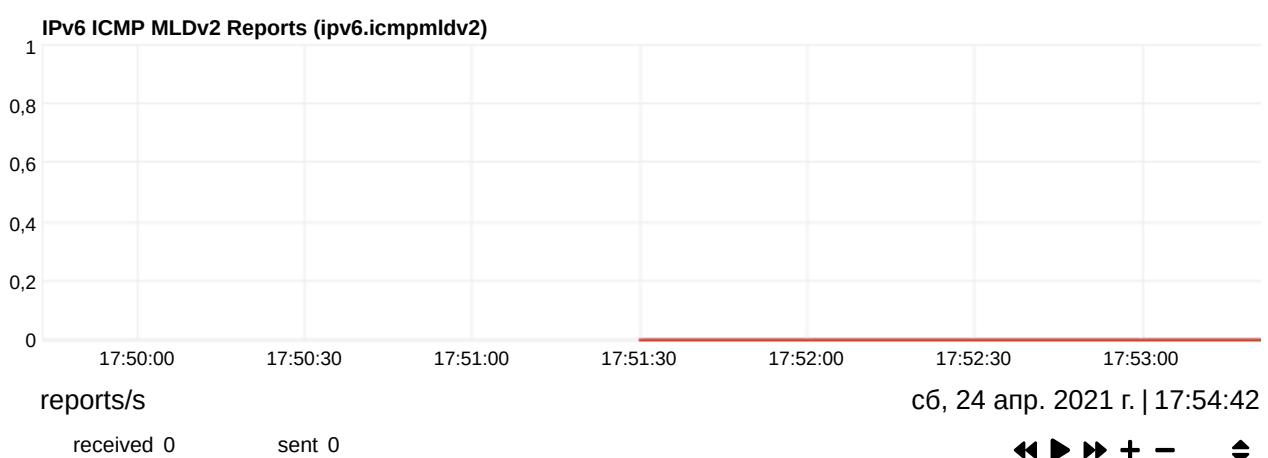
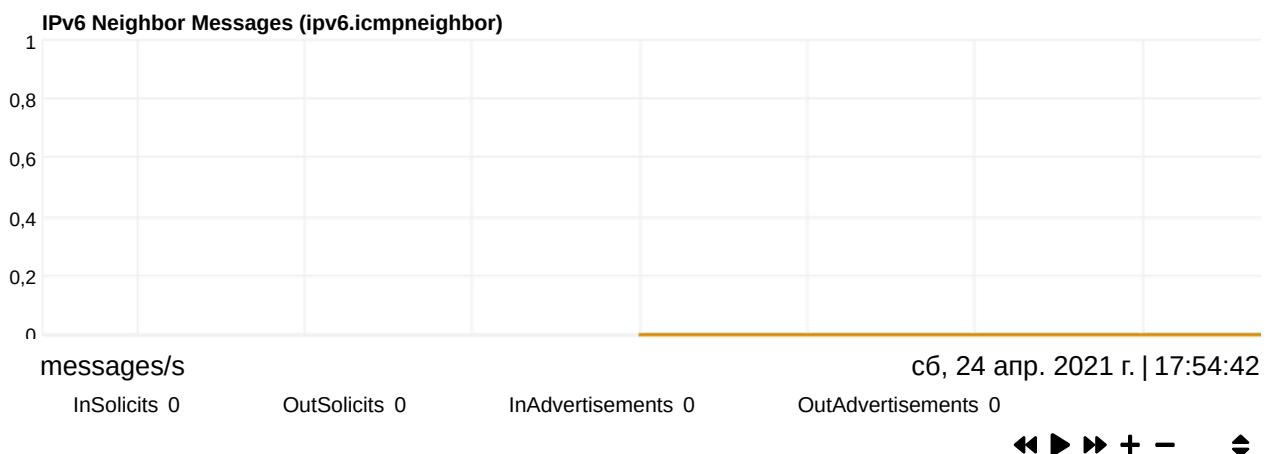
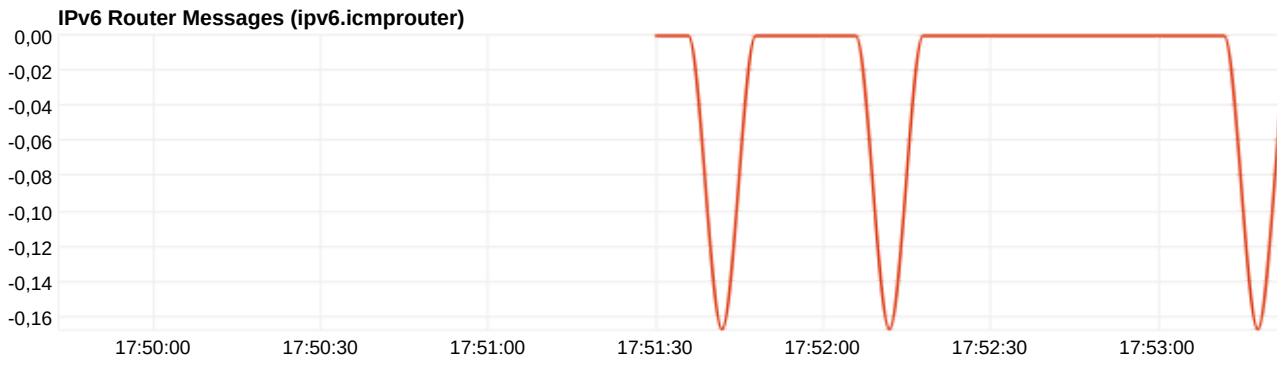


IPv6 UDP Errors (ipv6.udperrors)**raw6****IPv6 RAW Sockets (ipv6.sockstat6_raw_sockets)****multicast6****IPv6 Multicast Bandwidth (ipv6.mcast)**



icmp6





Network Interfaces

Performance metrics for network interfaces.

eth0



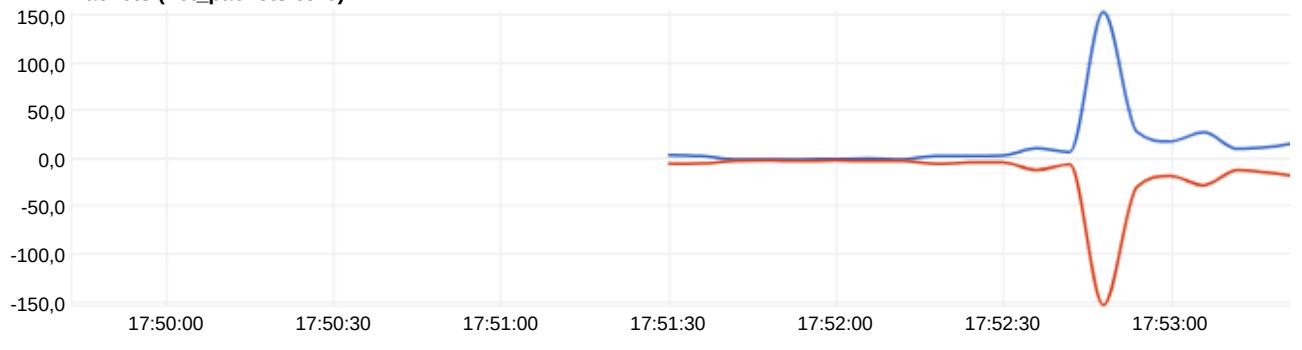
megabits/s

received 0,07 sent -0,06

сб, 24 апр. 2021 г. | 17:54:42

◀▶▶+−◆

Packets (net_packets.eth0)



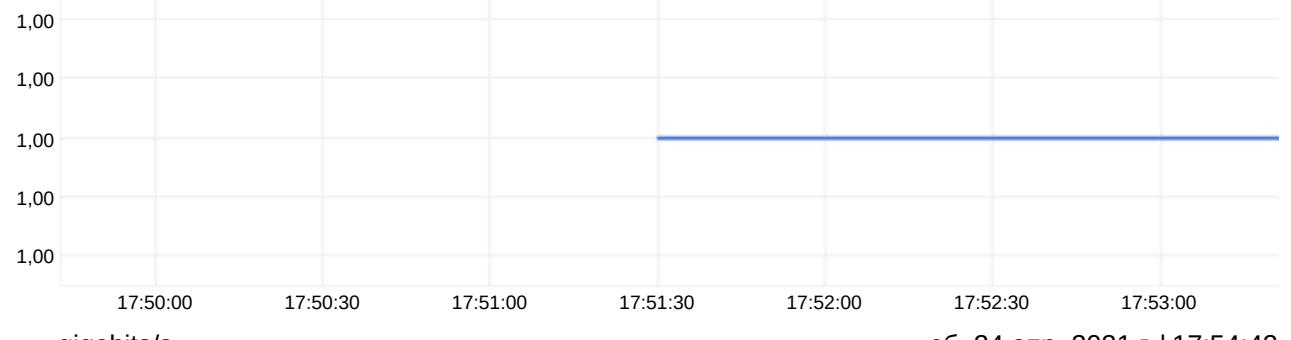
pps

received 15,7 sent -9,0

сб, 24 апр. 2021 г. | 17:54:42

◀▶▶+−◆

Interface Speed (net_speed.eth0)



gigabits/s

speed 1,00

сб, 24 апр. 2021 г. | 17:54:42

◀▶▶+−◆

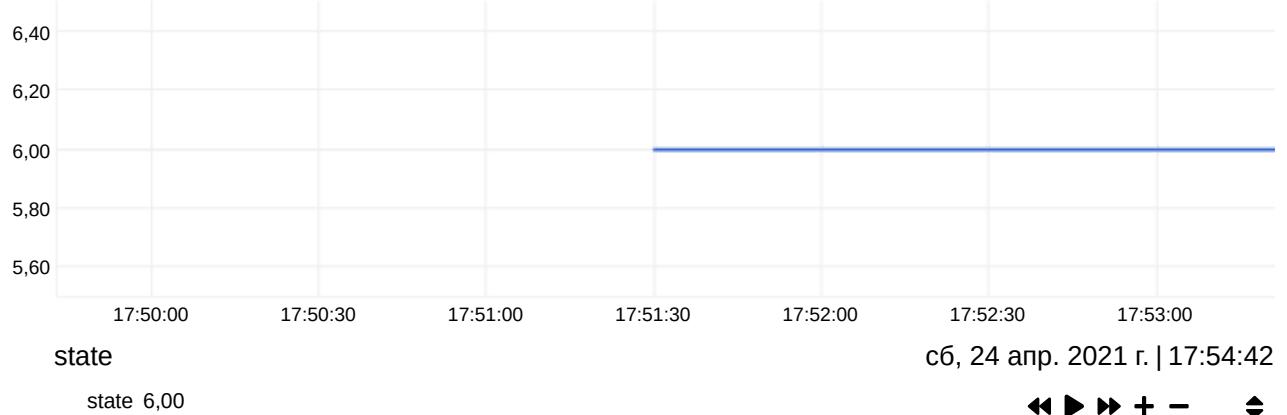
State map: 0 - unknown, 1 - half duplex, 2 - full duplex

Interface Duplex State (net_duplex.eth0)



State map: 0 - unknown, 1 - notpresent, 2 - down, 3 - lowerlayerdown, 4 - testing, 5 - dormant, 6 - up

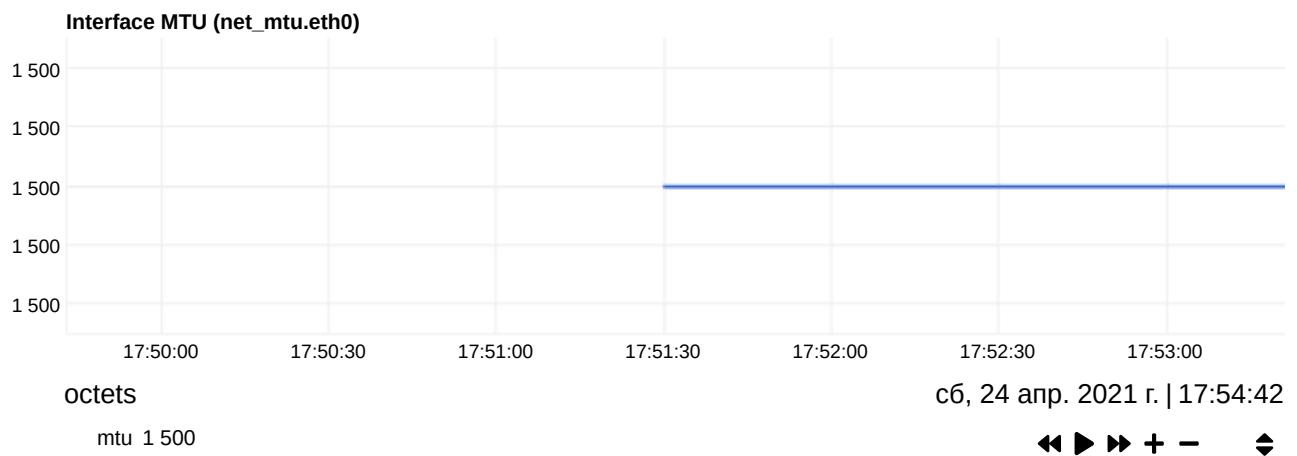
Interface Operational State (net_operstate.eth0)



State map: 0 - down, 1 - up

Interface Physical Link State (net_carrier.eth0)



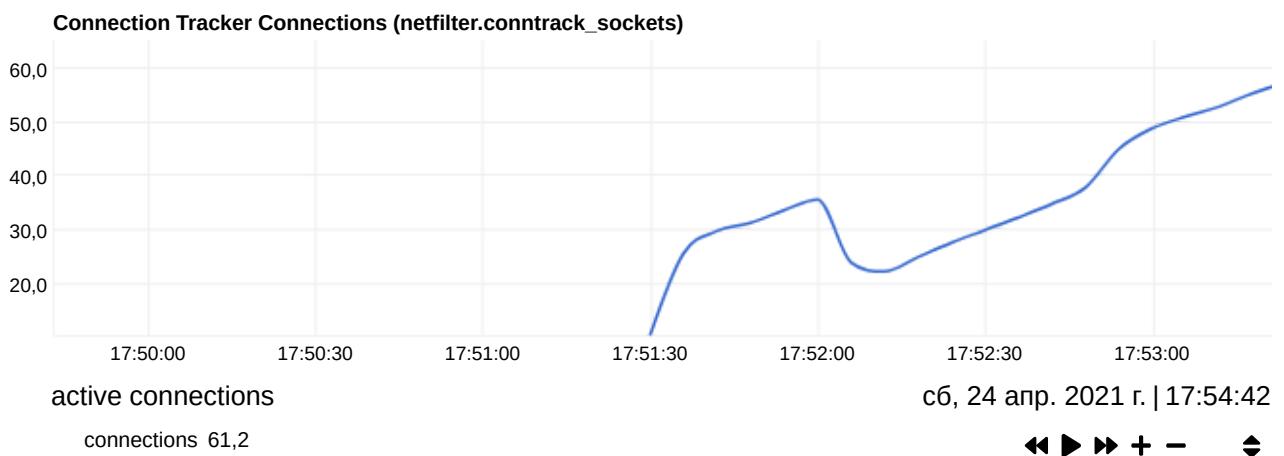


Firewall (netfilter)

Performance metrics of the netfilter components.

connection tracker

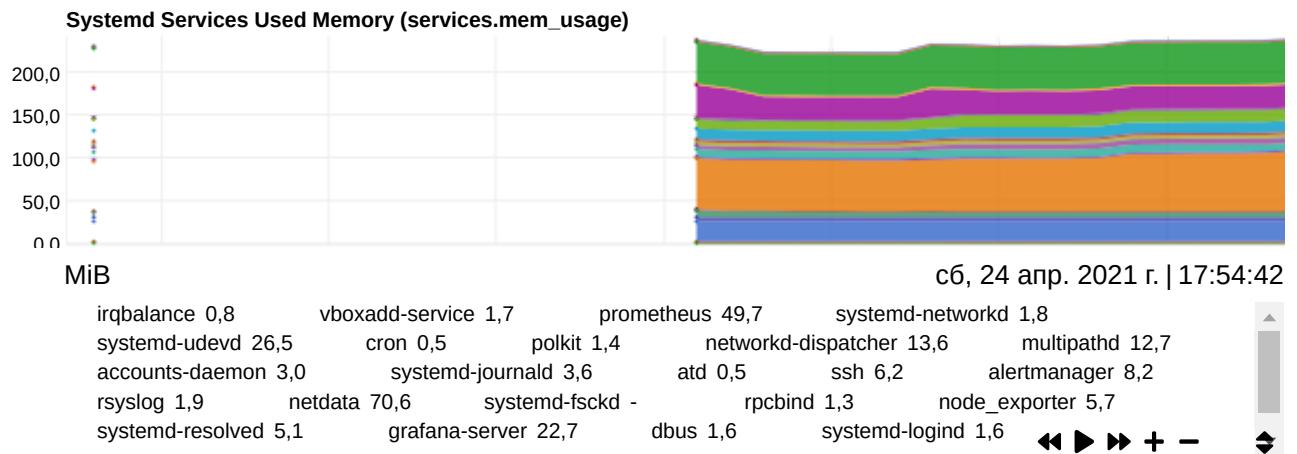
Netfilter Connection Tracker performance metrics. The connection tracker keeps track of all connections of the machine, inbound and outbound. It works by keeping a database with all open connections, tracking network and address translation and connection expectations.



systemd Services

Resources utilization of systemd services. netdata monitors all systemd services via CGROUPS (the resources accounting used by containers).

mem

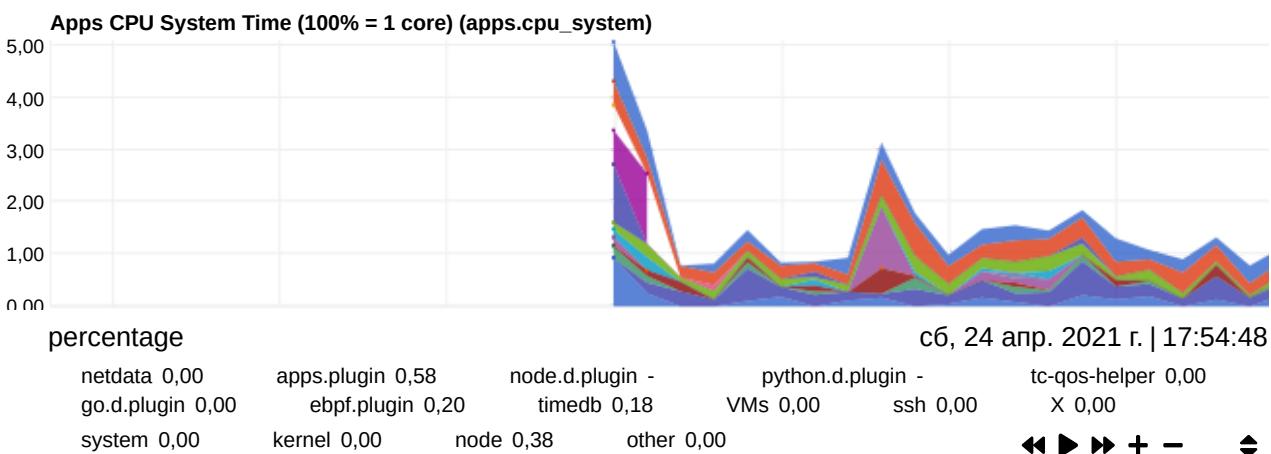
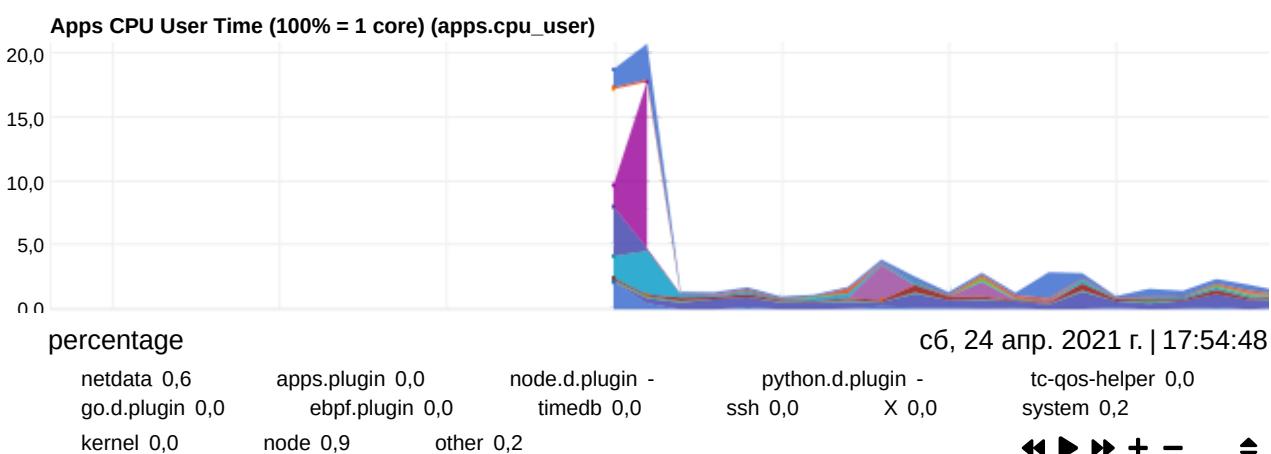
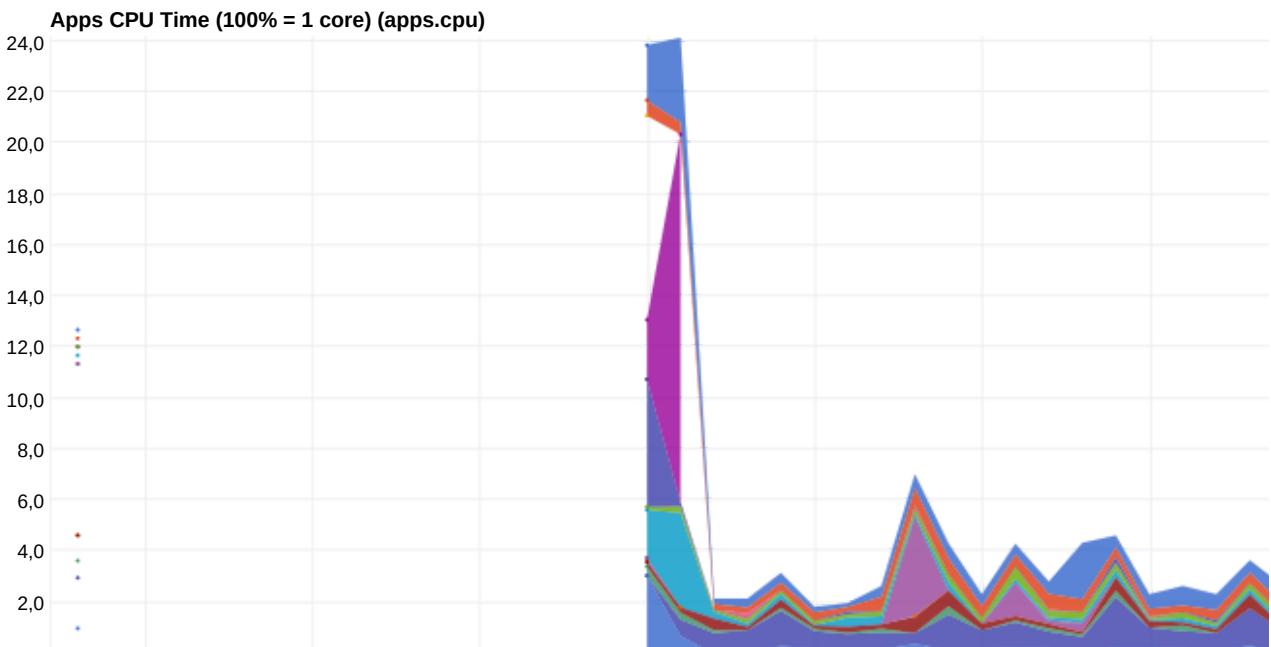


❤️ Applications

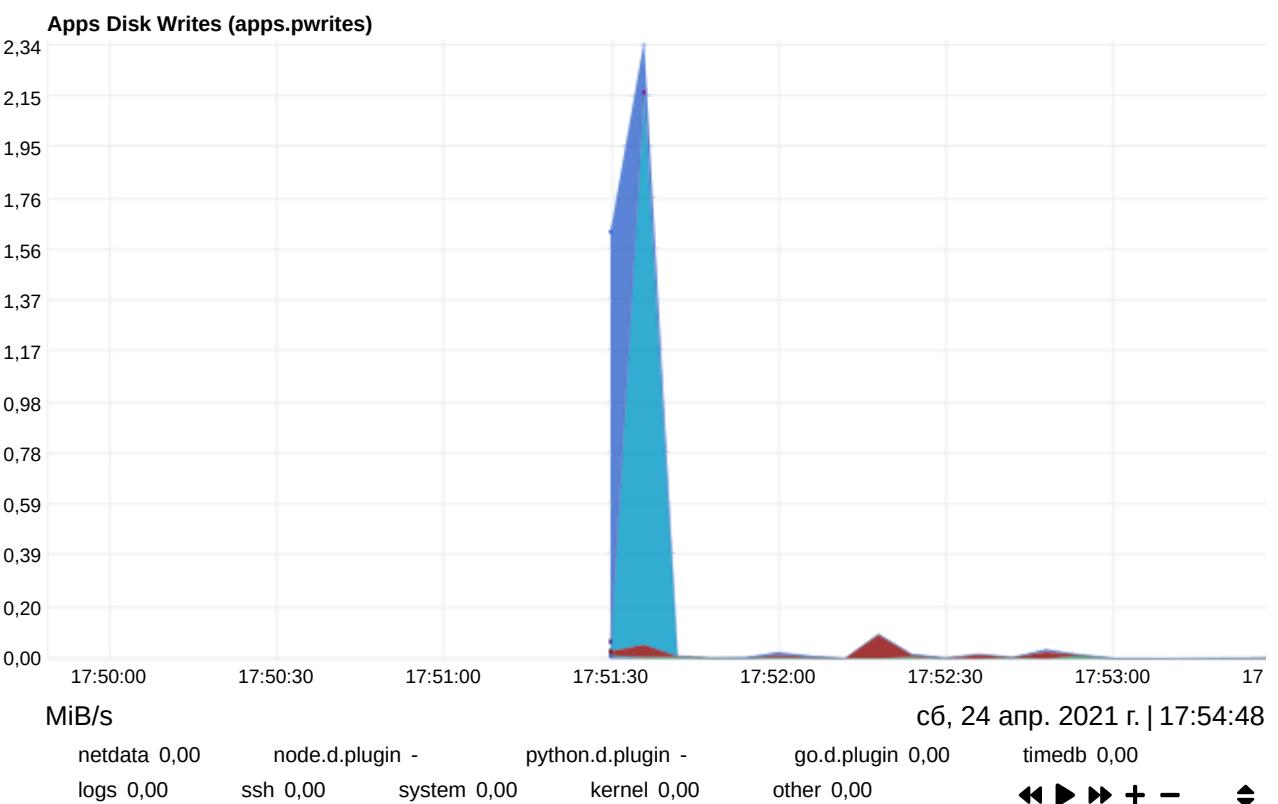
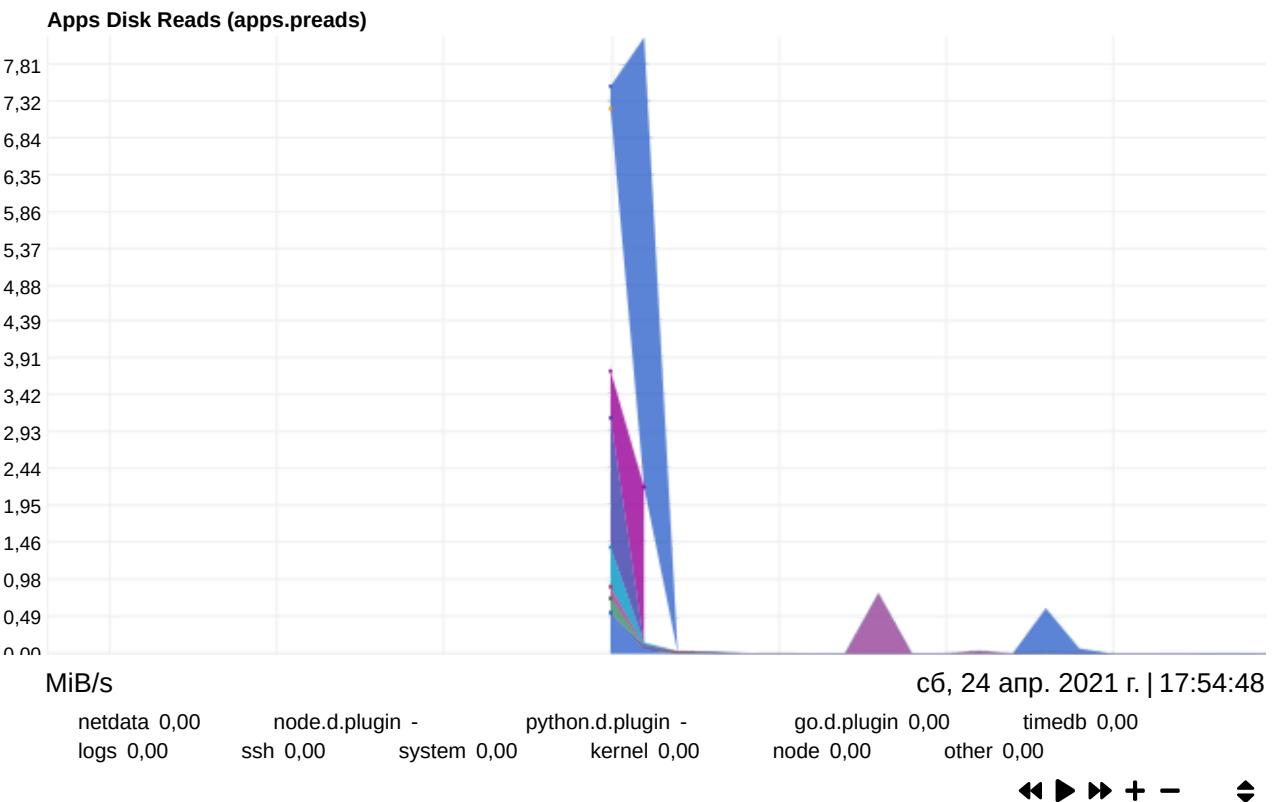
Per application statistics are collected using netdata's `apps.plugin`. This plugin walks through all processes and aggregates statistics for applications of interest, defined in

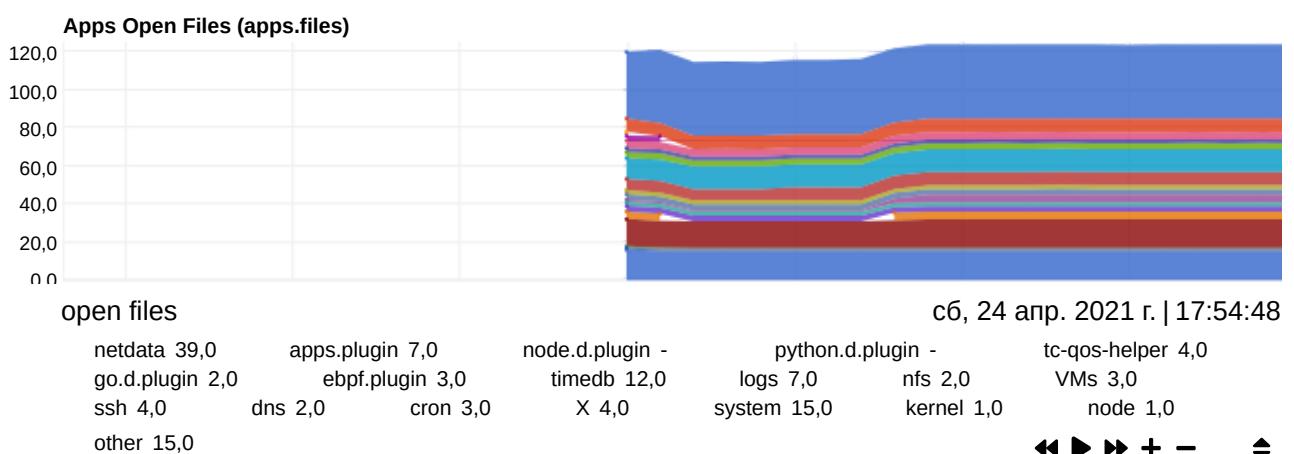
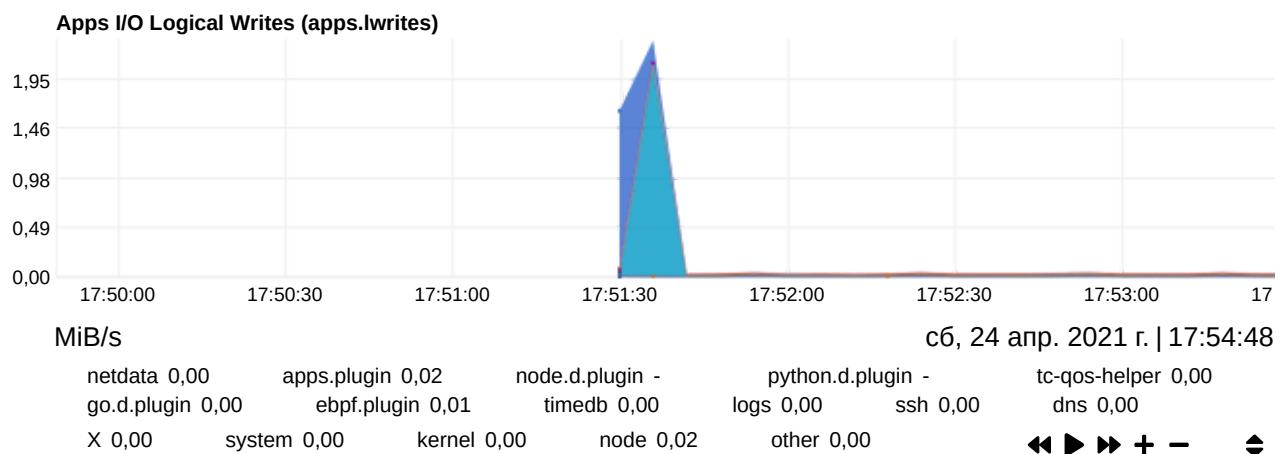
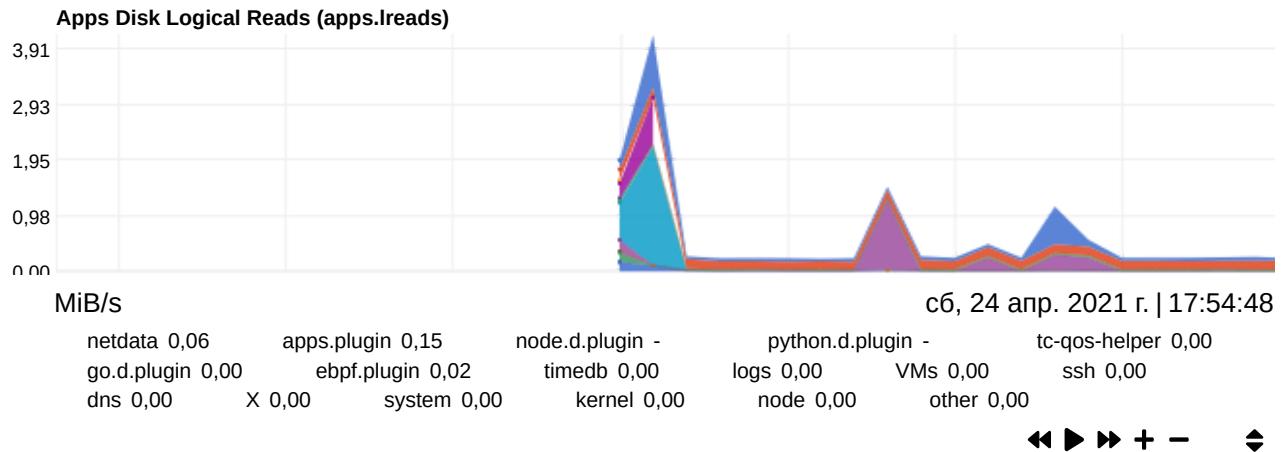
`/etc/netdata/apps_groups.conf`, which can be edited by running `$ /etc/netdata/edit-config apps_groups.conf` (the default is here (https://github.com/netdata/netdata/blob/master/collectors/apps.plugin/apps_groups.conf)). The plugin internally builds a process tree (much like `ps fax` does), and groups processes together (evaluating both child and parent processes) so that the result is always a chart with a predefined set of dimensions (of course, only application groups found running are reported). The reported values are compatible with `top`, although the netdata plugin counts also the resources of exited children (unlike `top` which shows only the resources of the currently running processes). So for processes like shell scripts, the reported values include the resources used by the commands these scripts run within each timeframe.

cpu



disk

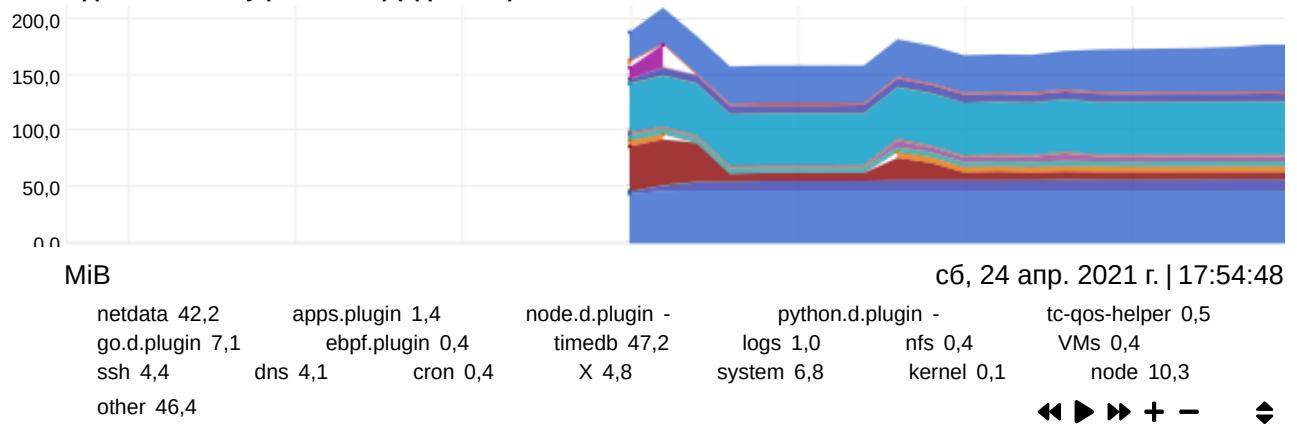




mem

Real memory (RAM) used by applications. This does not include shared memory.

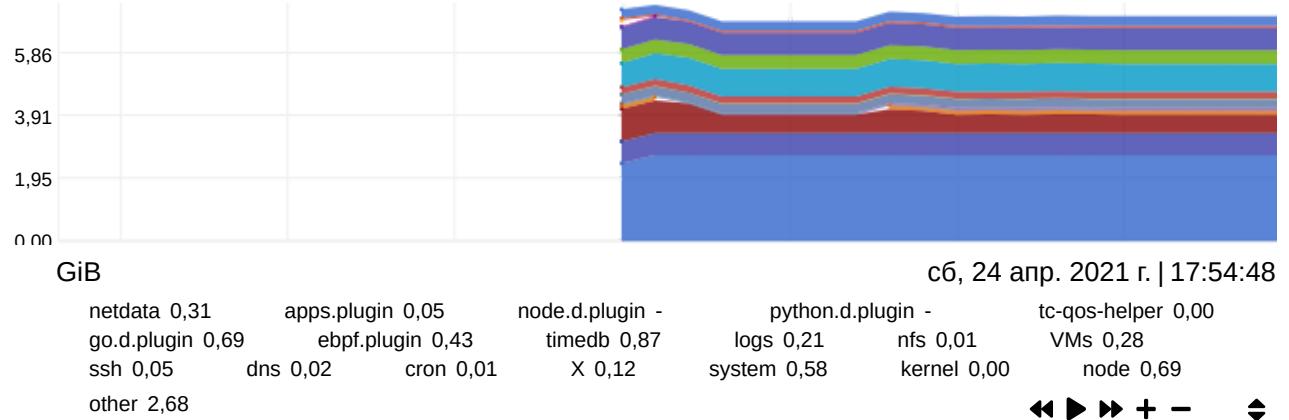
Apps Real Memory (w/o shared) (apps.mem)



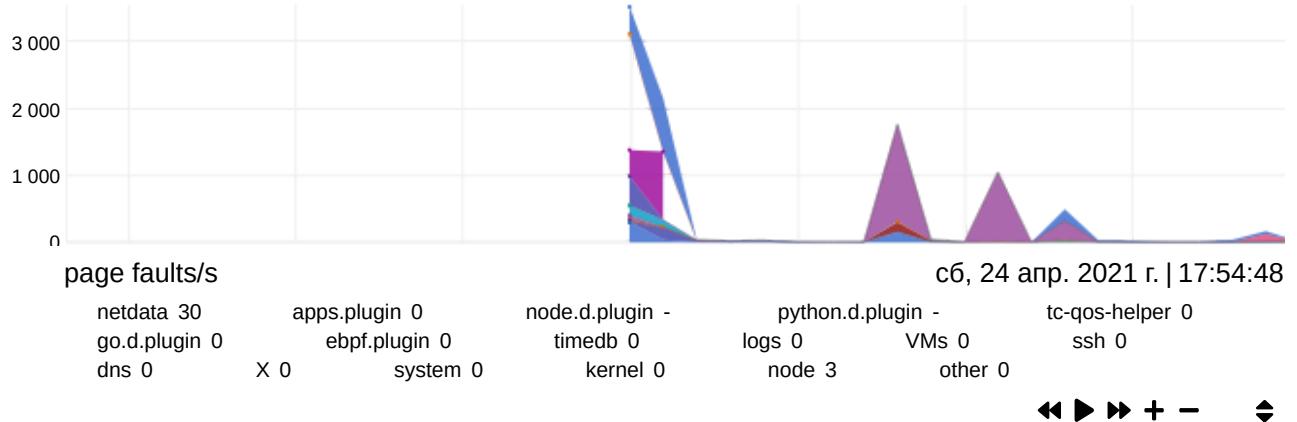
Virtual memory allocated by applications. Please check this article

(<https://github.com/netdata/netdata/tree/master/daemon#virtual-memory>) for more information.

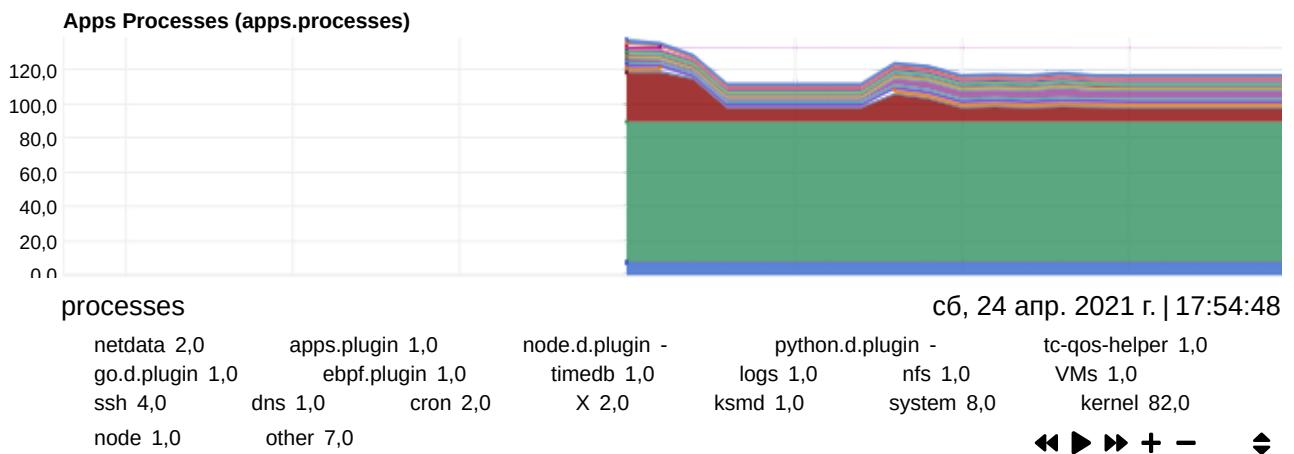
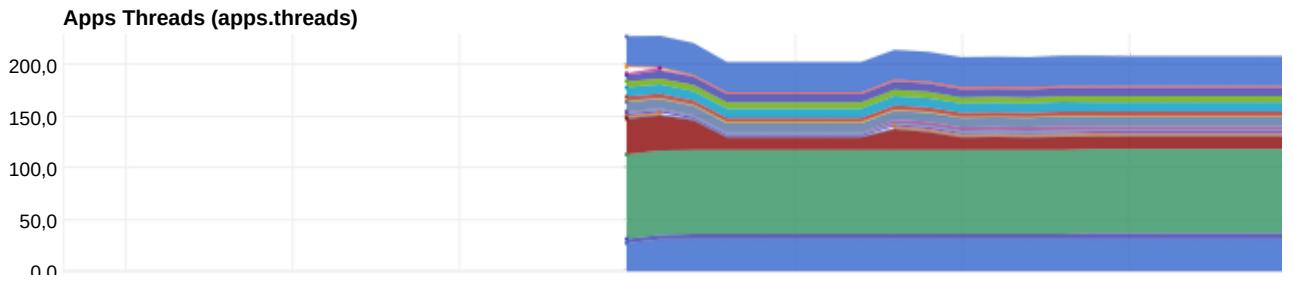
Apps Virtual Memory Size (apps.vmem)



Apps Minor Page Faults (apps.minor_faults)

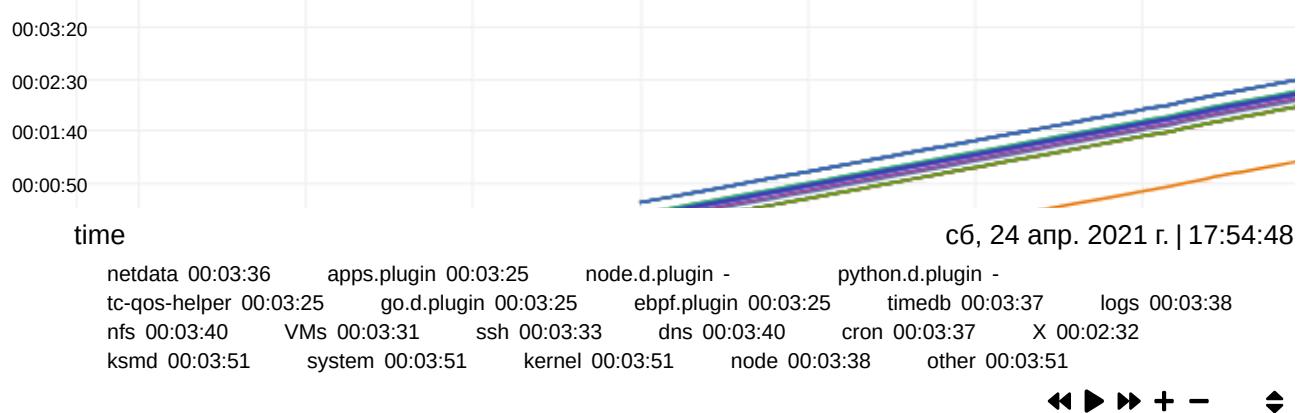


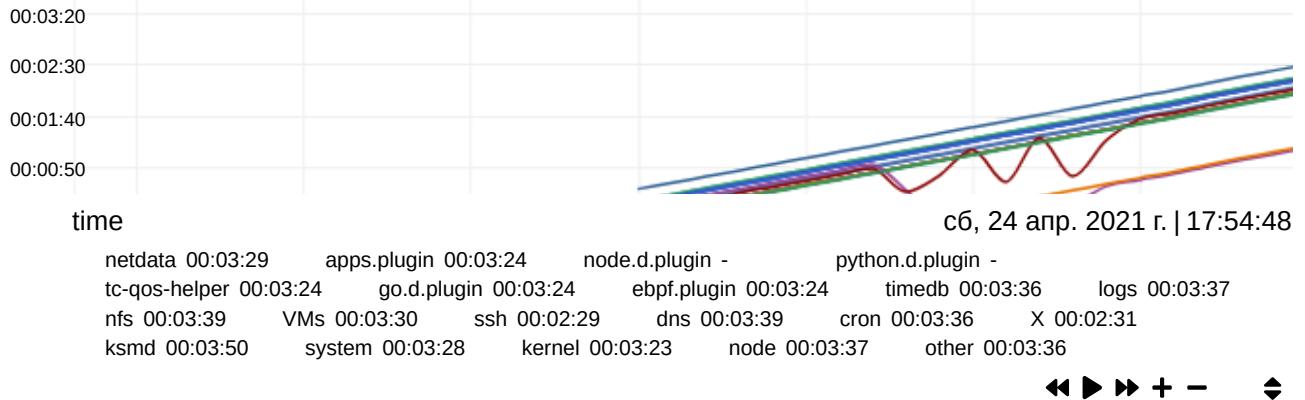
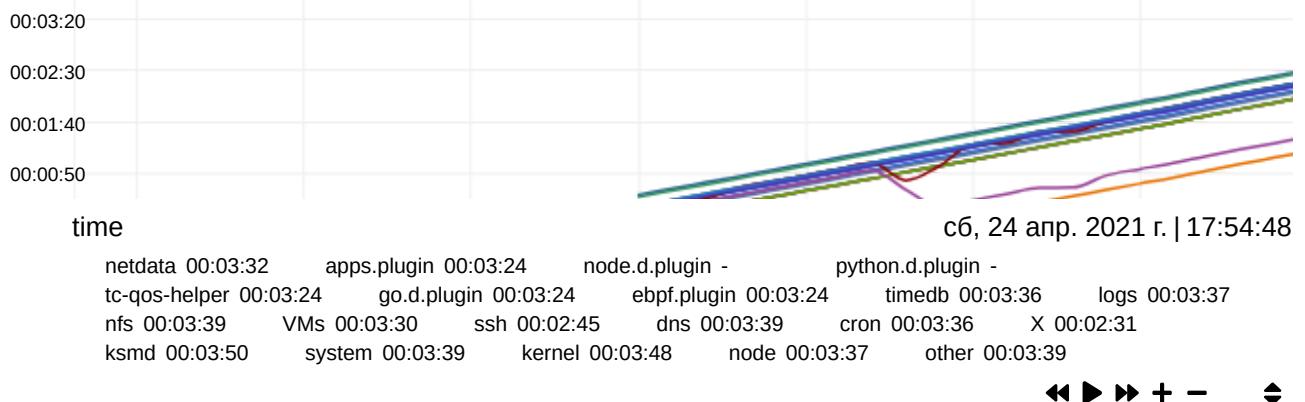
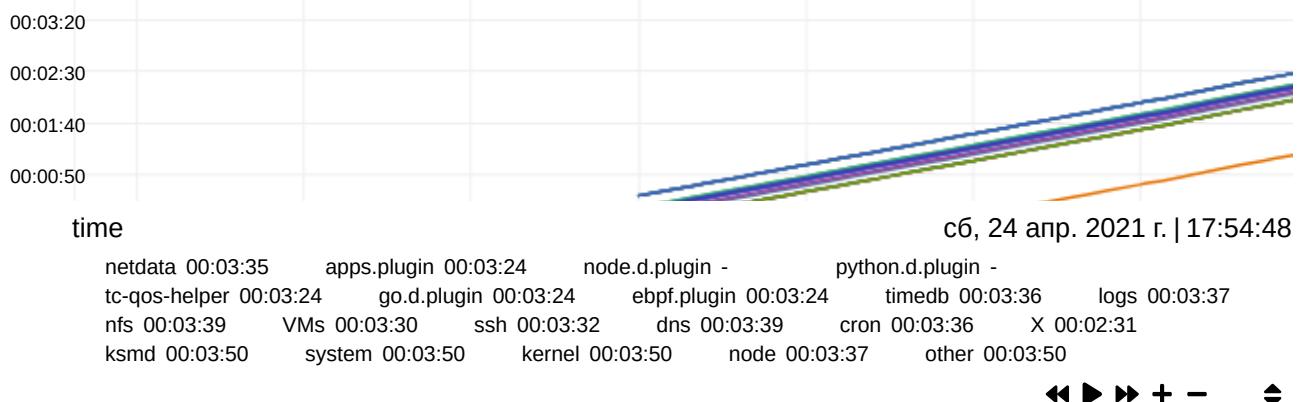
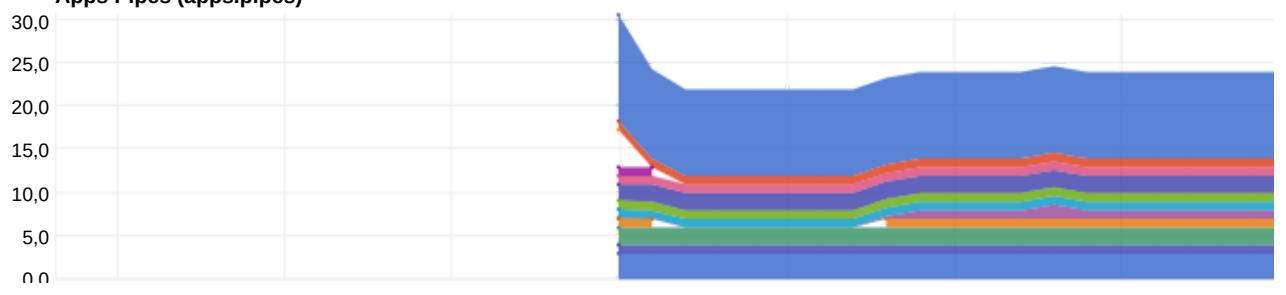
processes



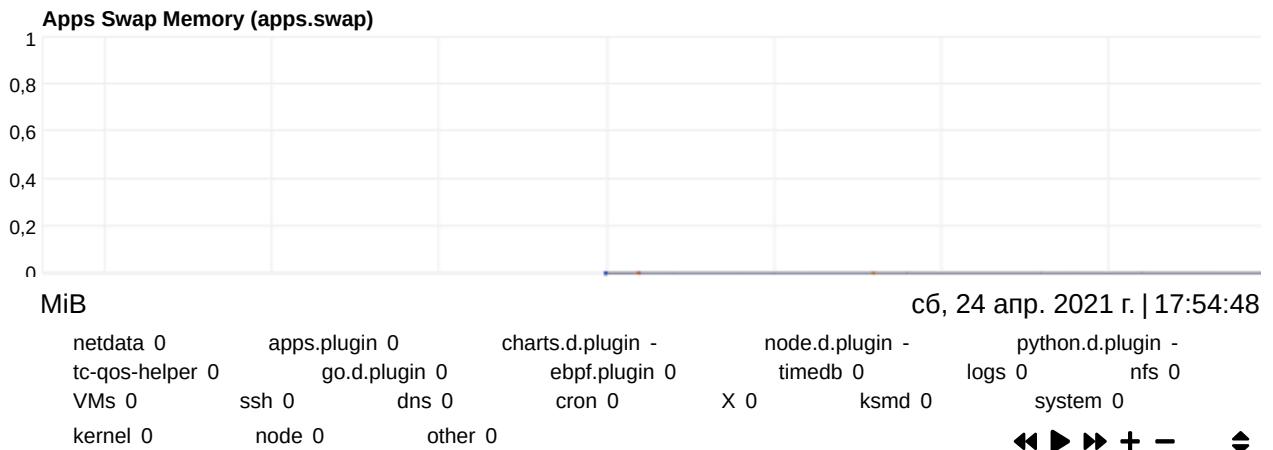
Carried over process group uptime since the Netdata restart. The period of time within which at least one process in the group was running.

Apps Carried Over Uptime (apps.uptime)

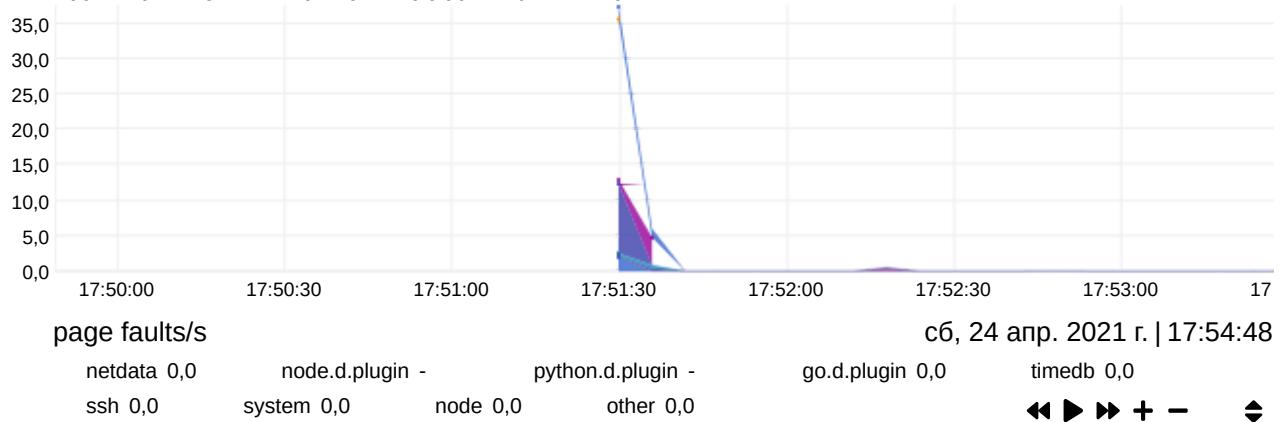


Apps Minimum Uptime (apps.uptime_min)**Apps Average Uptime (apps.uptime_avg)****Apps Maximum Uptime (apps.uptime_max)****Apps Pipes (apps.pipes)**

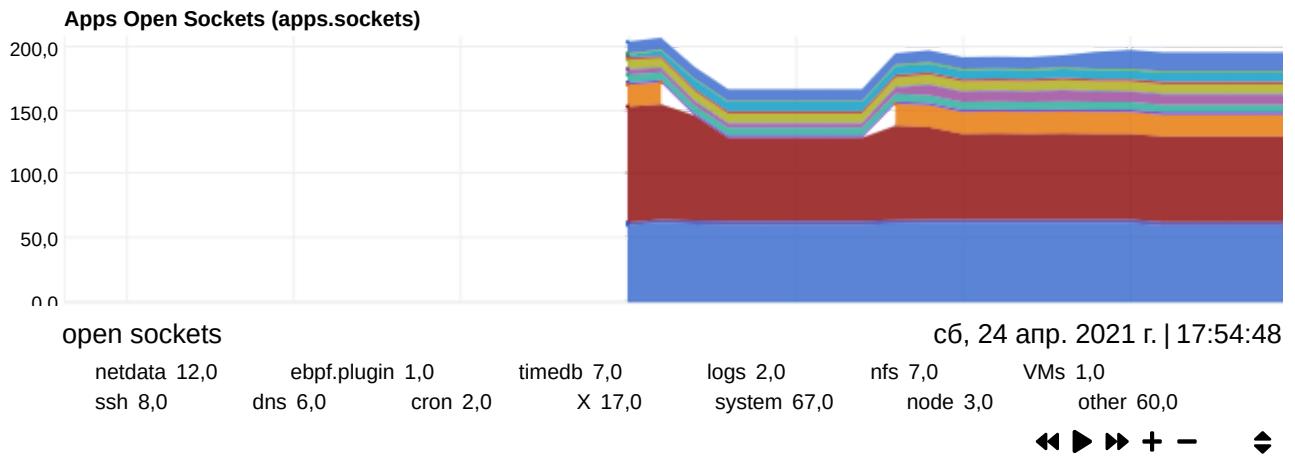
Swap



Apps Major Page Faults (swap read) (apps.major_faults)

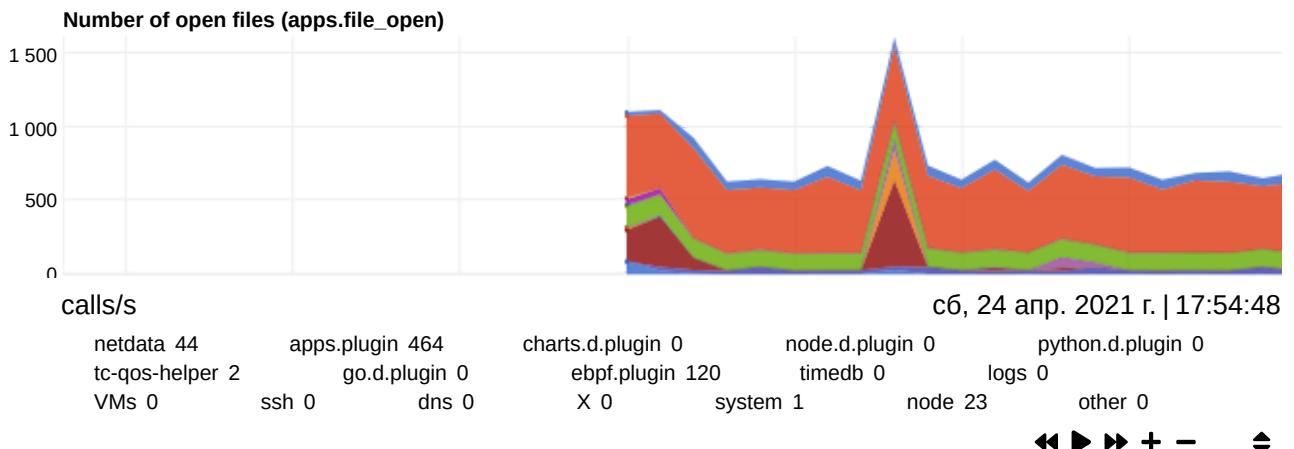


net

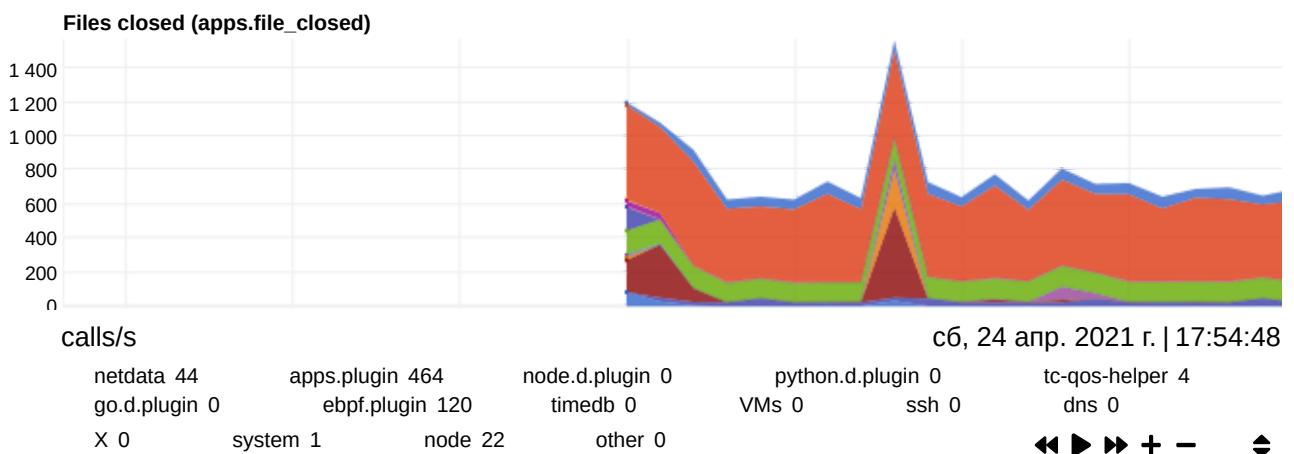


file (eBPF)

Calls to the internal function `do_sys_open` (For kernels newer than 5.5.19 we add a kprobe to `do_sys_openat2`), which is the common function called from `open(2)` (<https://www.man7.org/linux/man-pages/man2/open.2.html>) and `openat(2)` (<https://www.man7.org/linux/man-pages/man2/openat.2.html>).



Calls to the internal function `__close_fd` (<https://elixir.bootlin.com/linux/v5.10/source/fs/file.c#L665>) or `close_fd` (<https://elixir.bootlin.com/linux/v5.11/source/fs/file.c#L617>) according to your kernel version, which is called from `close(2)` (<https://www.man7.org/linux/man-pages/man2/close.2.html>).

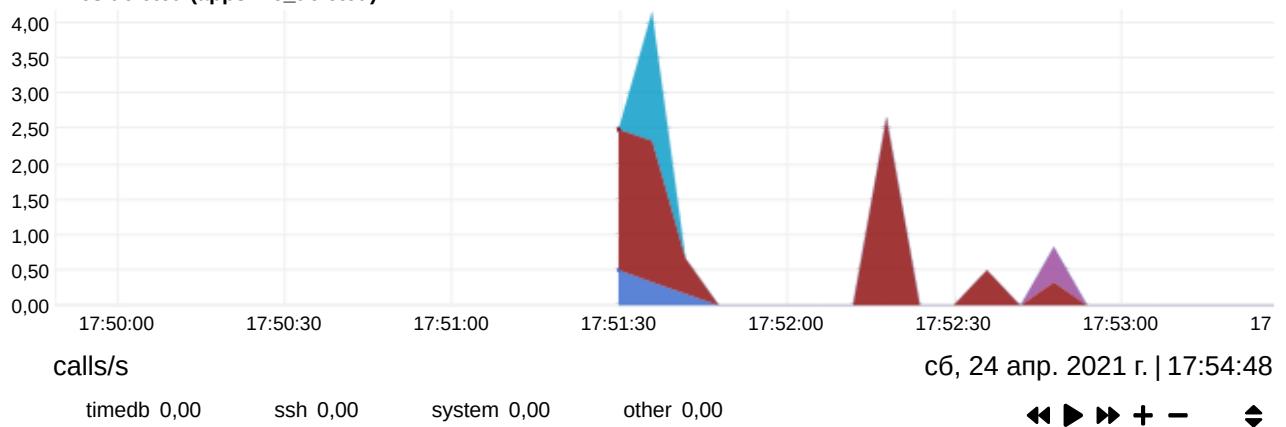


vfs (eBPF)

Calls to the function `vfs_unlink` (<https://www.kernel.org/doc/html/docs/filesystems/API-vfs-unlink.html>).

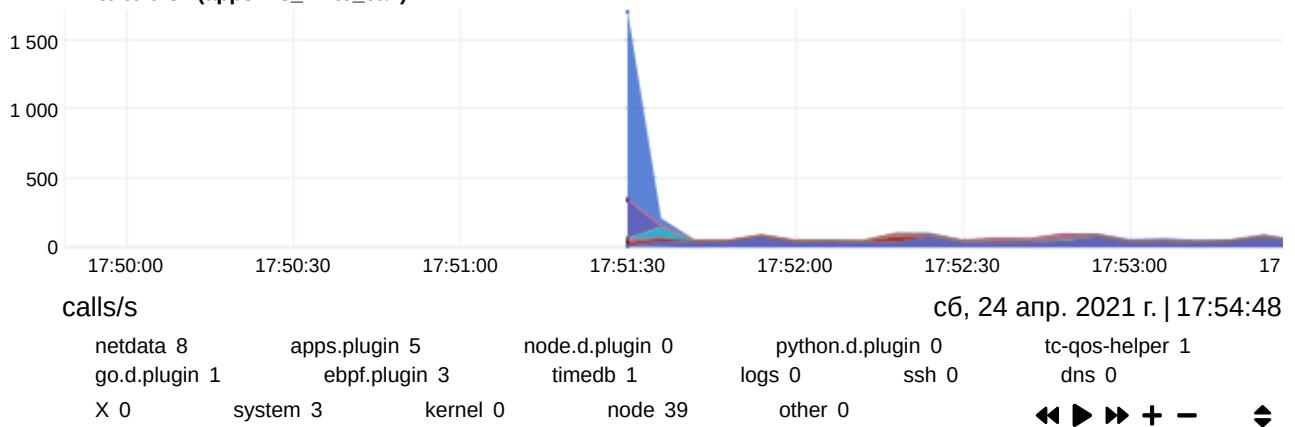
This chart does not show all events that remove files from the filesystem, because filesystems can create their own functions to remove files.

Files deleted (apps.file_deleted)



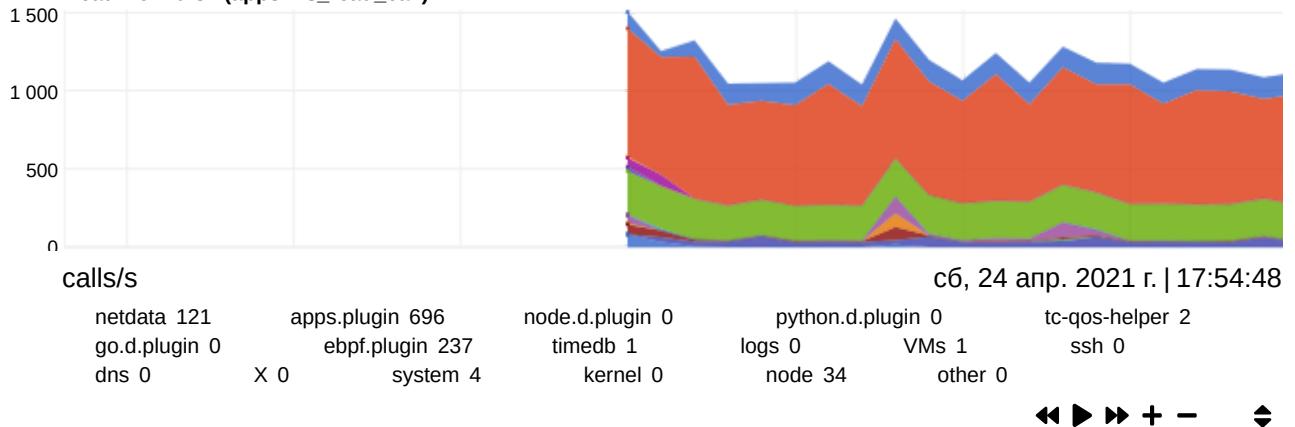
Successful calls to the function `vfs_write` (https://topic.alibabacloud.com/a/kernel-state-file-operation-work-information-kernel_8_8_20287135.html). This chart may not show all filesystem events if it uses other functions to store data on disk.

Write to disk (apps.vfs_write_call)

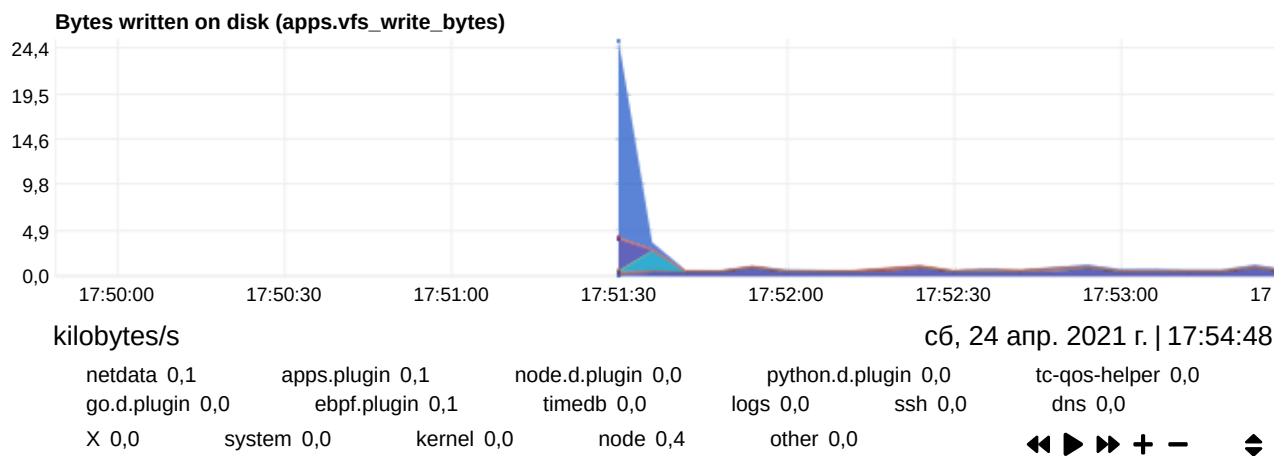


Successful calls to the function `vfs_read` (https://topic.alibabacloud.com/a/kernel-state-file-operation-work-information-kernel_8_8_20287135.html). This chart may not show all filesystem events if it uses other functions to store data on disk.

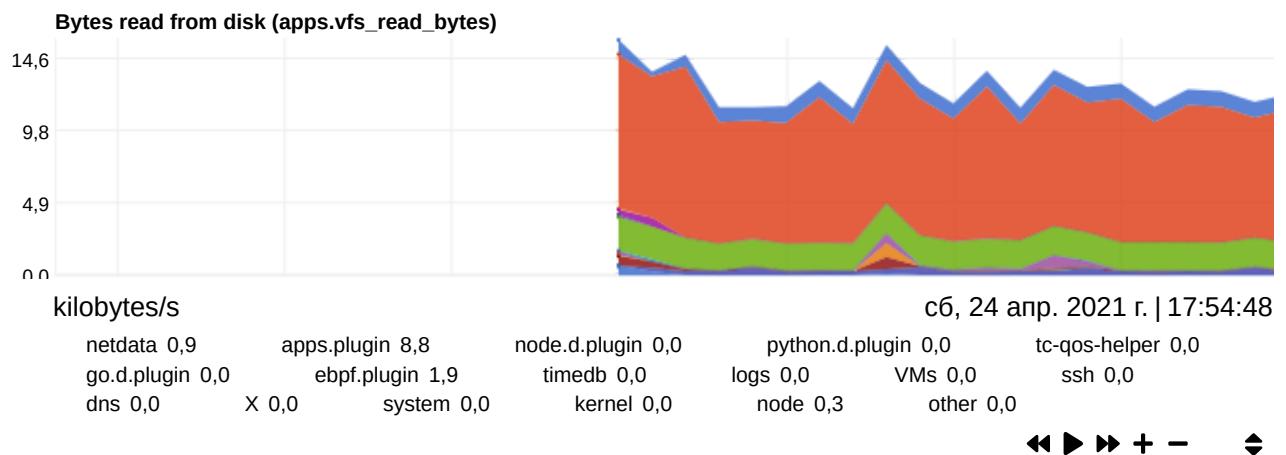
Read from disk (apps.vfs_read_call)



Total of bytes successfully written using the function vfs_write
https://topic.alibabacloud.com/a/kernel-state-file-operation-__-work-information-kernel_8_8_20287135.html.



Total of bytes successfully read using the function vfs_read (https://topic.alibabacloud.com/a/kernel-state-file-operation-__-work-information-kernel_8_8_20287135.html).

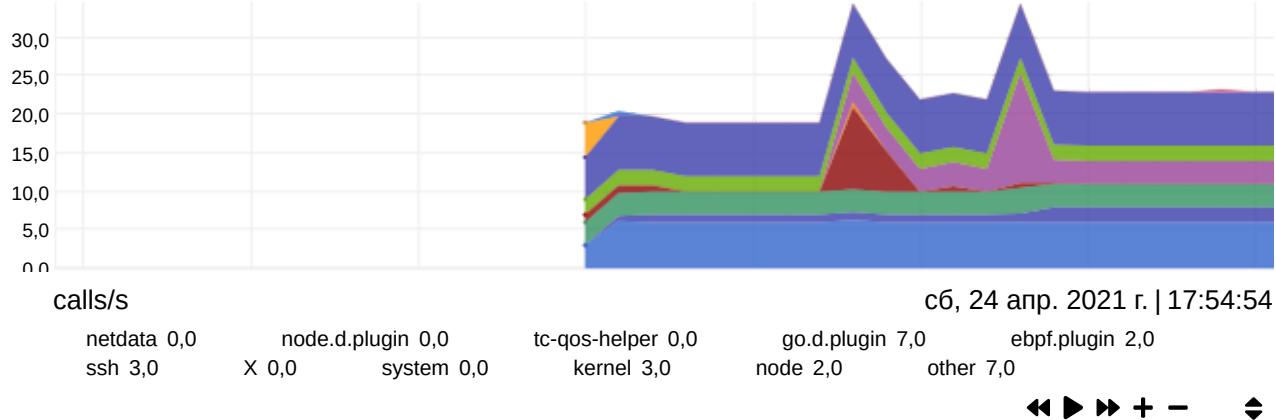


process (eBPF)

Calls to either do_fork

(<https://www.ece.uic.edu/~yshi1/linux/lkse/node4.html#SECTION00421000000000000000>), or `kernel_clone` if you are running kernel newer than 5.9.16, to create a new task, which is the common name used to define process and tasks inside the kernel. Netdata identifies the process by counting the number of calls to `sys_clone` (<https://linux.die.net/man/2/clone>) that do not have the flag `CLONE_THREAD` set.

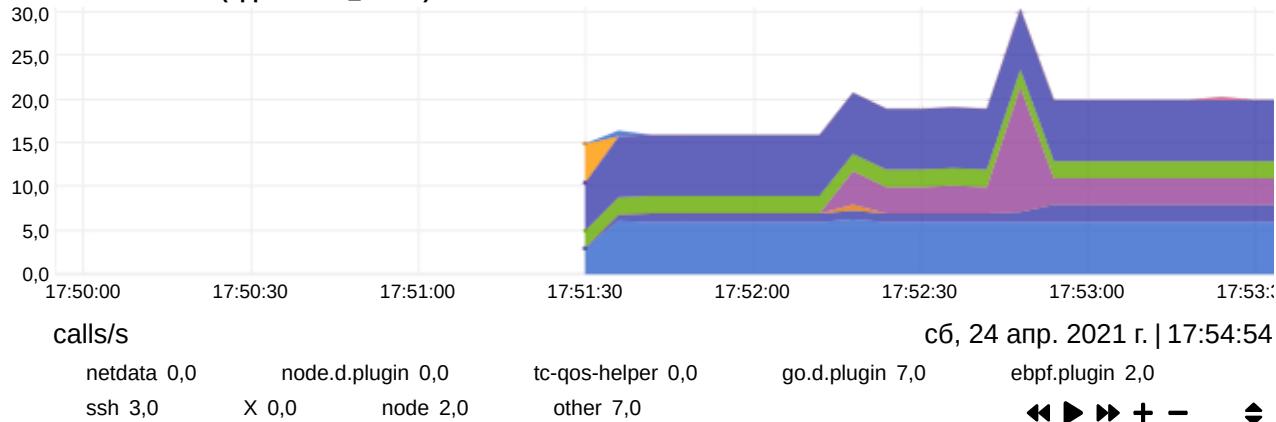
Process started (apps.process_create)



Calls to either do_fork

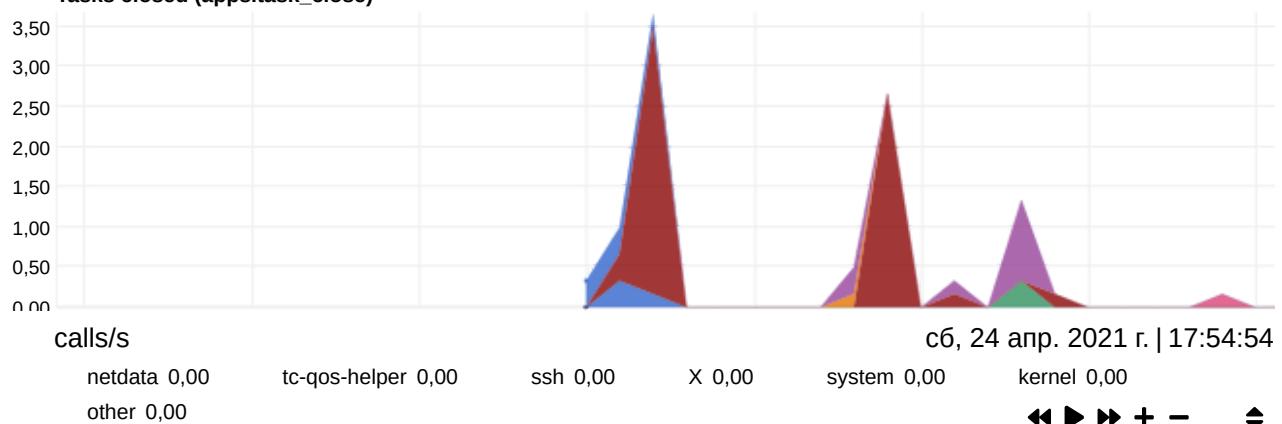
(<https://www.ece.uic.edu/~yshi1/linux/lkse/node4.html#SECTION00421000000000000000>), or `kernel_clone` if you are running kernel newer than 5.9.16, to create a new task, which is the common name used to define process and tasks inside the kernel. Netdata identifies the threads by counting the number of calls to `sys_clone` (<https://linux.die.net/man/2/clone>) that have the flag `CLONE_THREAD` set.

Threads started (apps.thread_create)



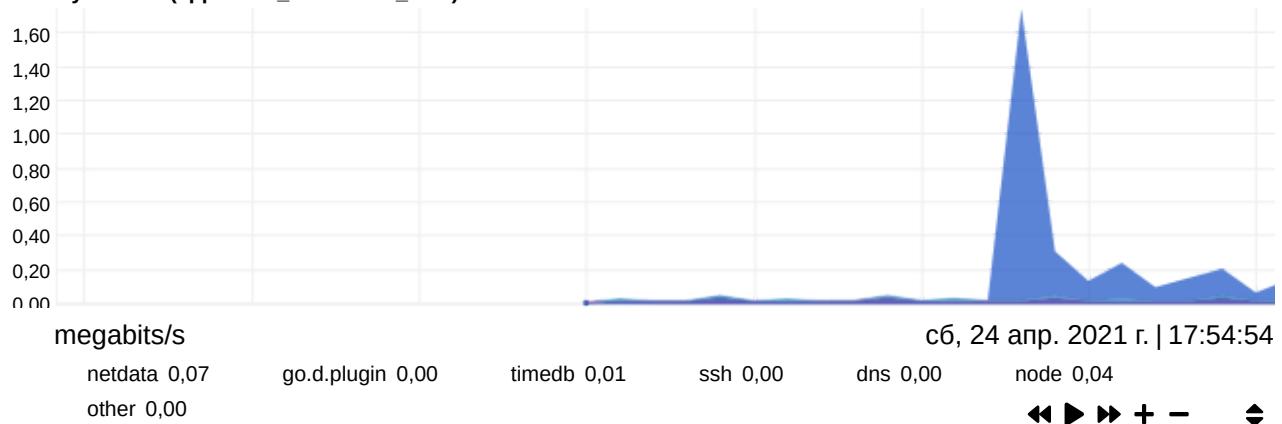
Calls to the functions responsible for closing (do_exit (<https://www.informit.com/articles/article.aspx?p=370047&seqNum=4>)) and releasing (release_task (<https://www.informit.com/articles/article.aspx?p=370047&seqNum=4>)) tasks.

Tasks closed (apps.task_close)

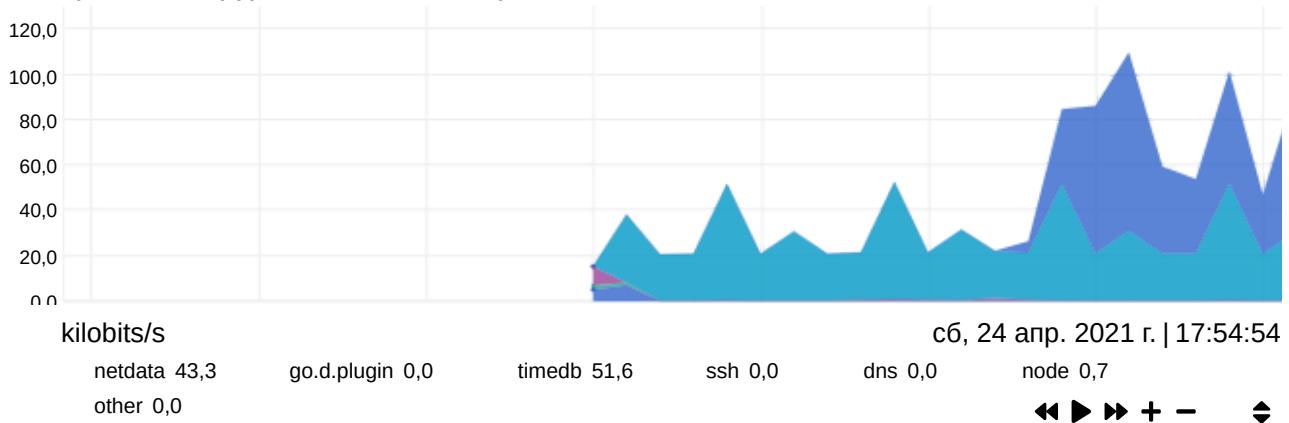


net (eBPF)

Bytes sent (apps.total_bandwidth_sent)

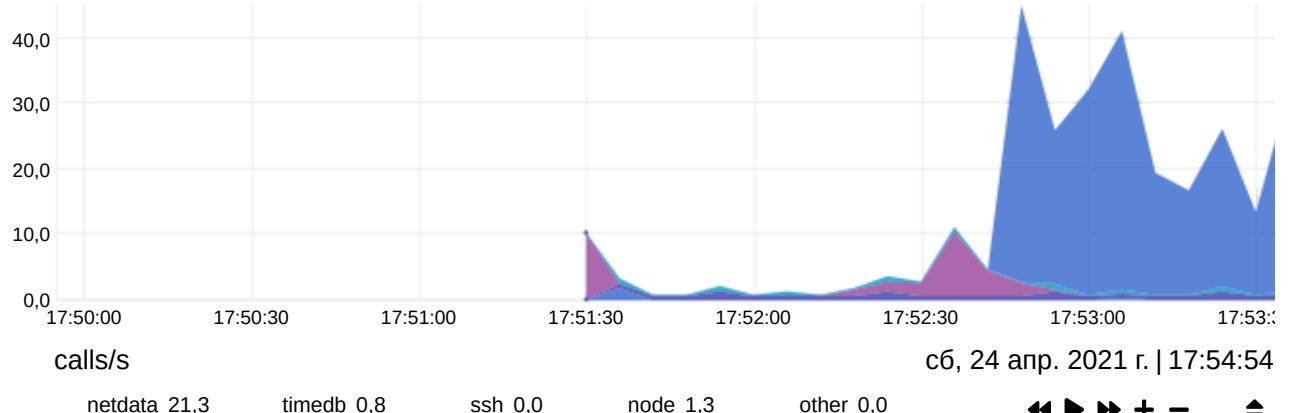


bytes received (apps.total_bandwidth_recv)



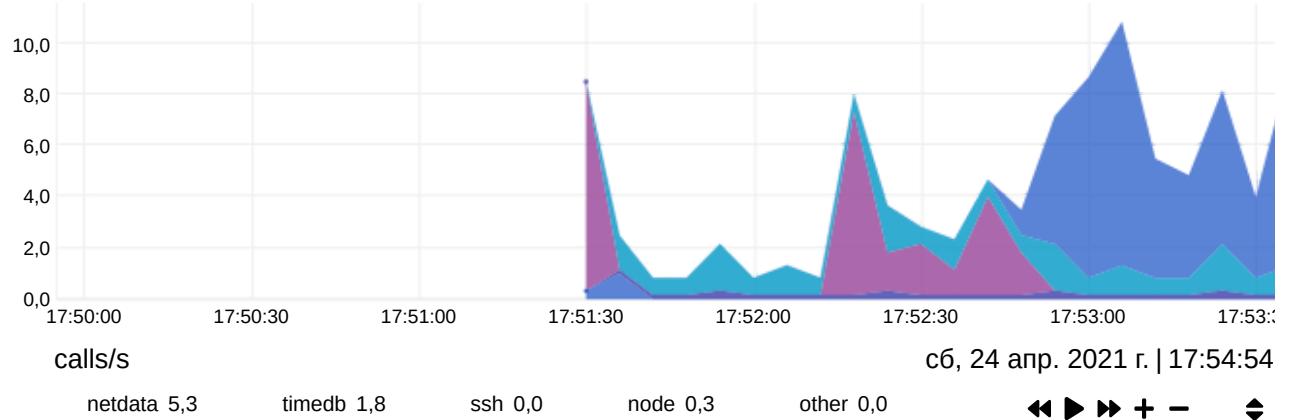
Calls for function `tcp_sendmsg`.

Calls for `tcp_sendmsg` (apps.bandwidth_tcp_send)



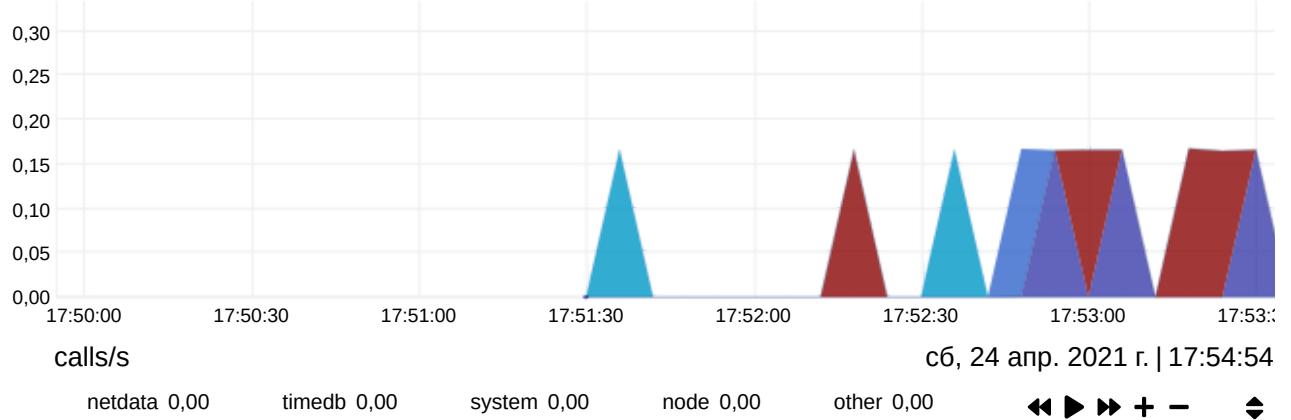
Calls for functions `tcp_cleanup_rbuf`. We use `tcp_cleanup_rbuf` instead `tcp_recvmsg`, because this last misses `tcp_read_sock()` traffic and we would also need to have more probes to get the socket and package size.

Calls for `tcp_cleanup_rbuf` (apps.bandwidth_tcp_recv)



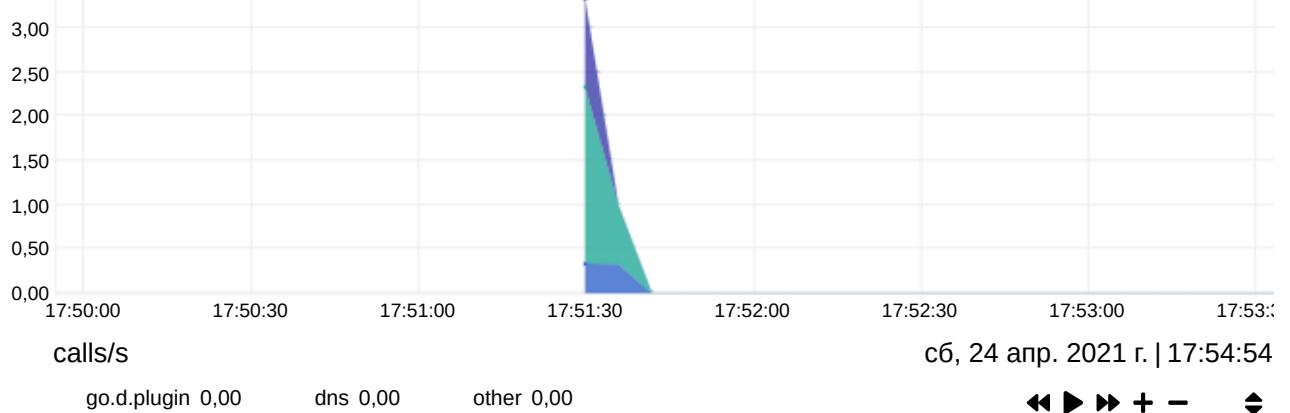
Calls for functions `tcp_retransmit_skb`.

Calls for `tcp_retransmit` (apps.bandwidth_tcp_retransmit)



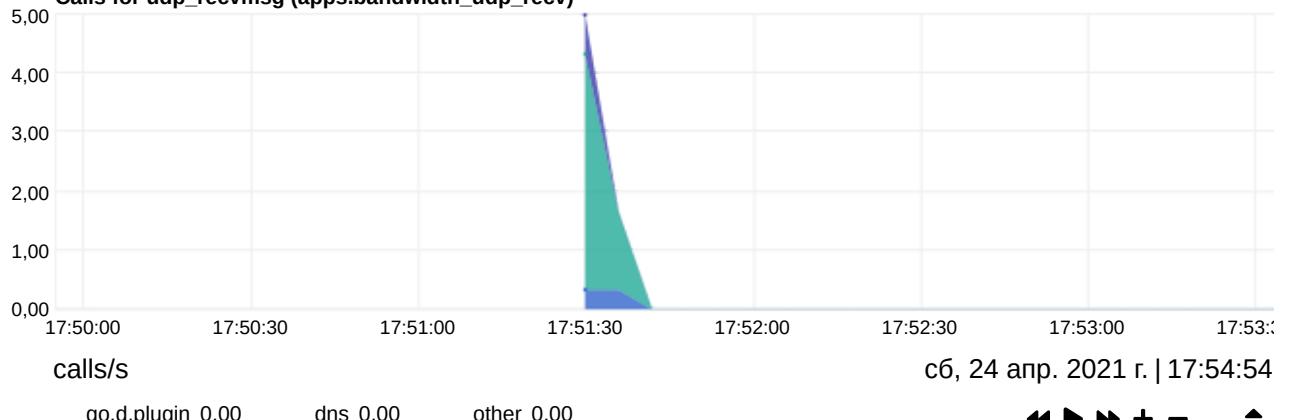
Calls for function `udp_sendmsg`.

Calls for `udp_sendmsg (apps.bandwidth_udp_send)`



Calls for function `udp_recvmsg`.

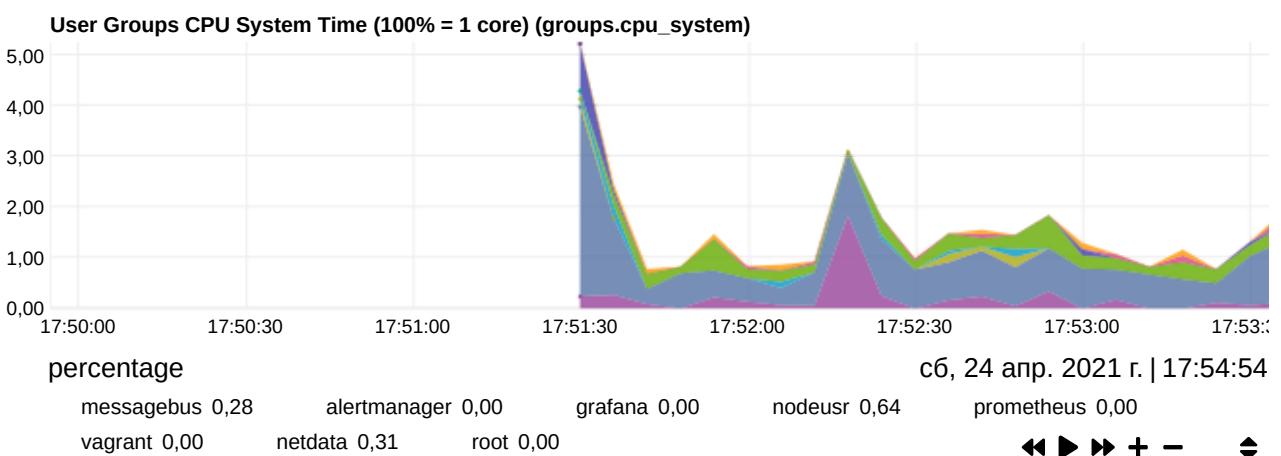
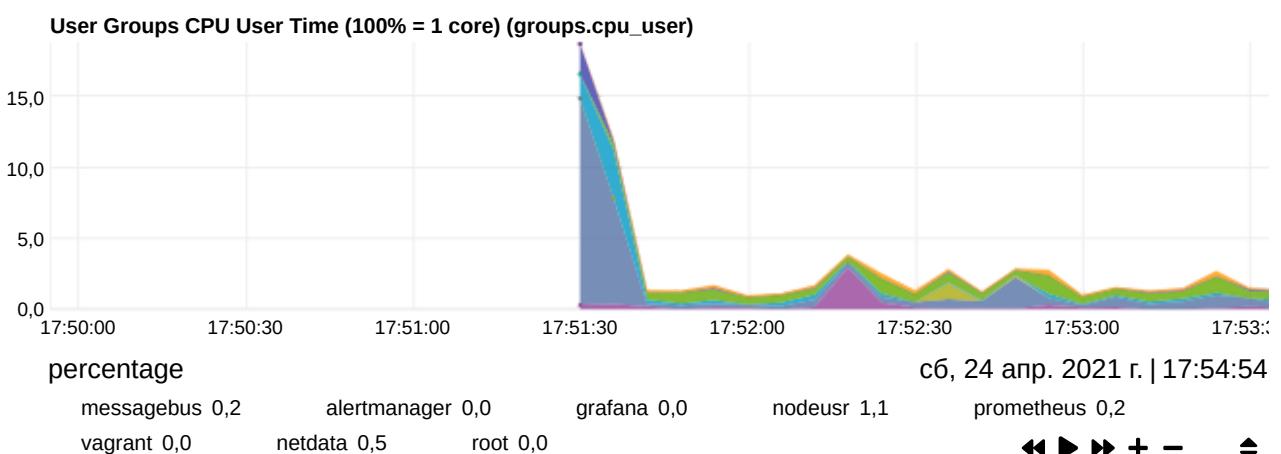
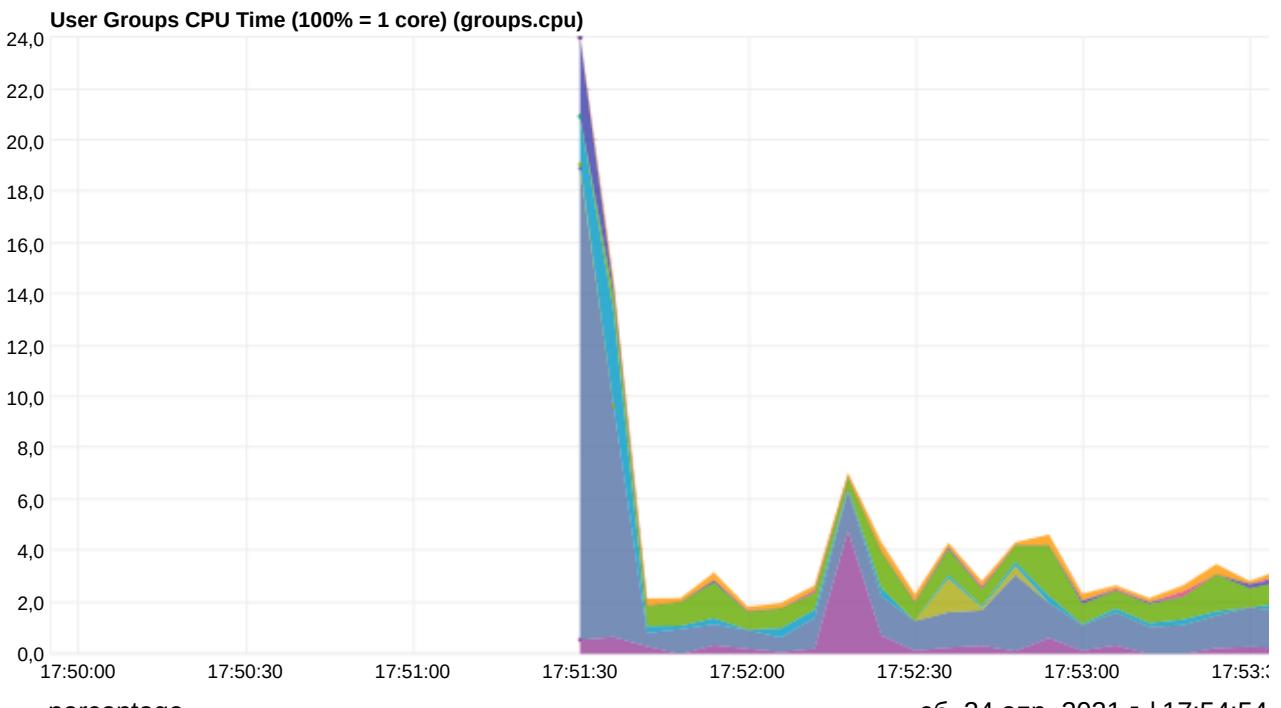
Calls for `udp_recvmsg (apps.bandwidth_udp_recv)`



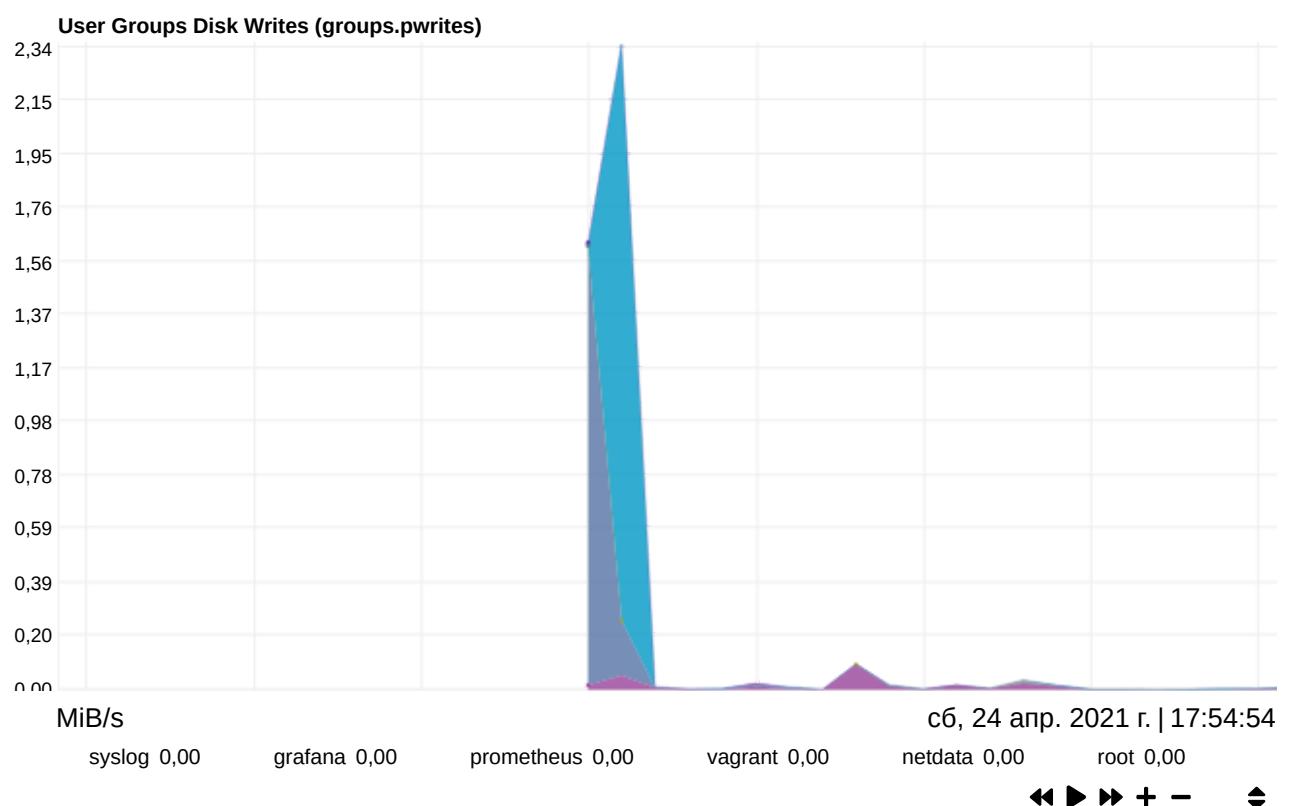
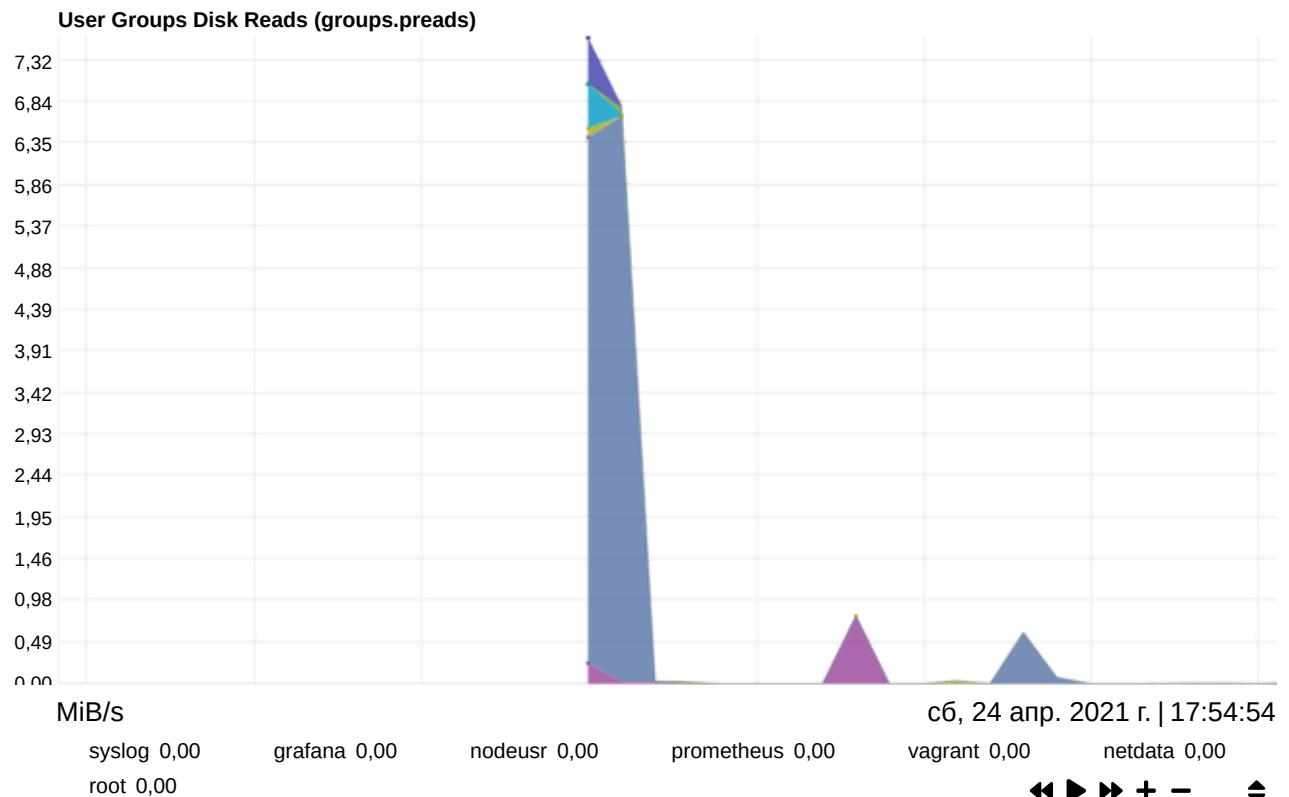
User Groups

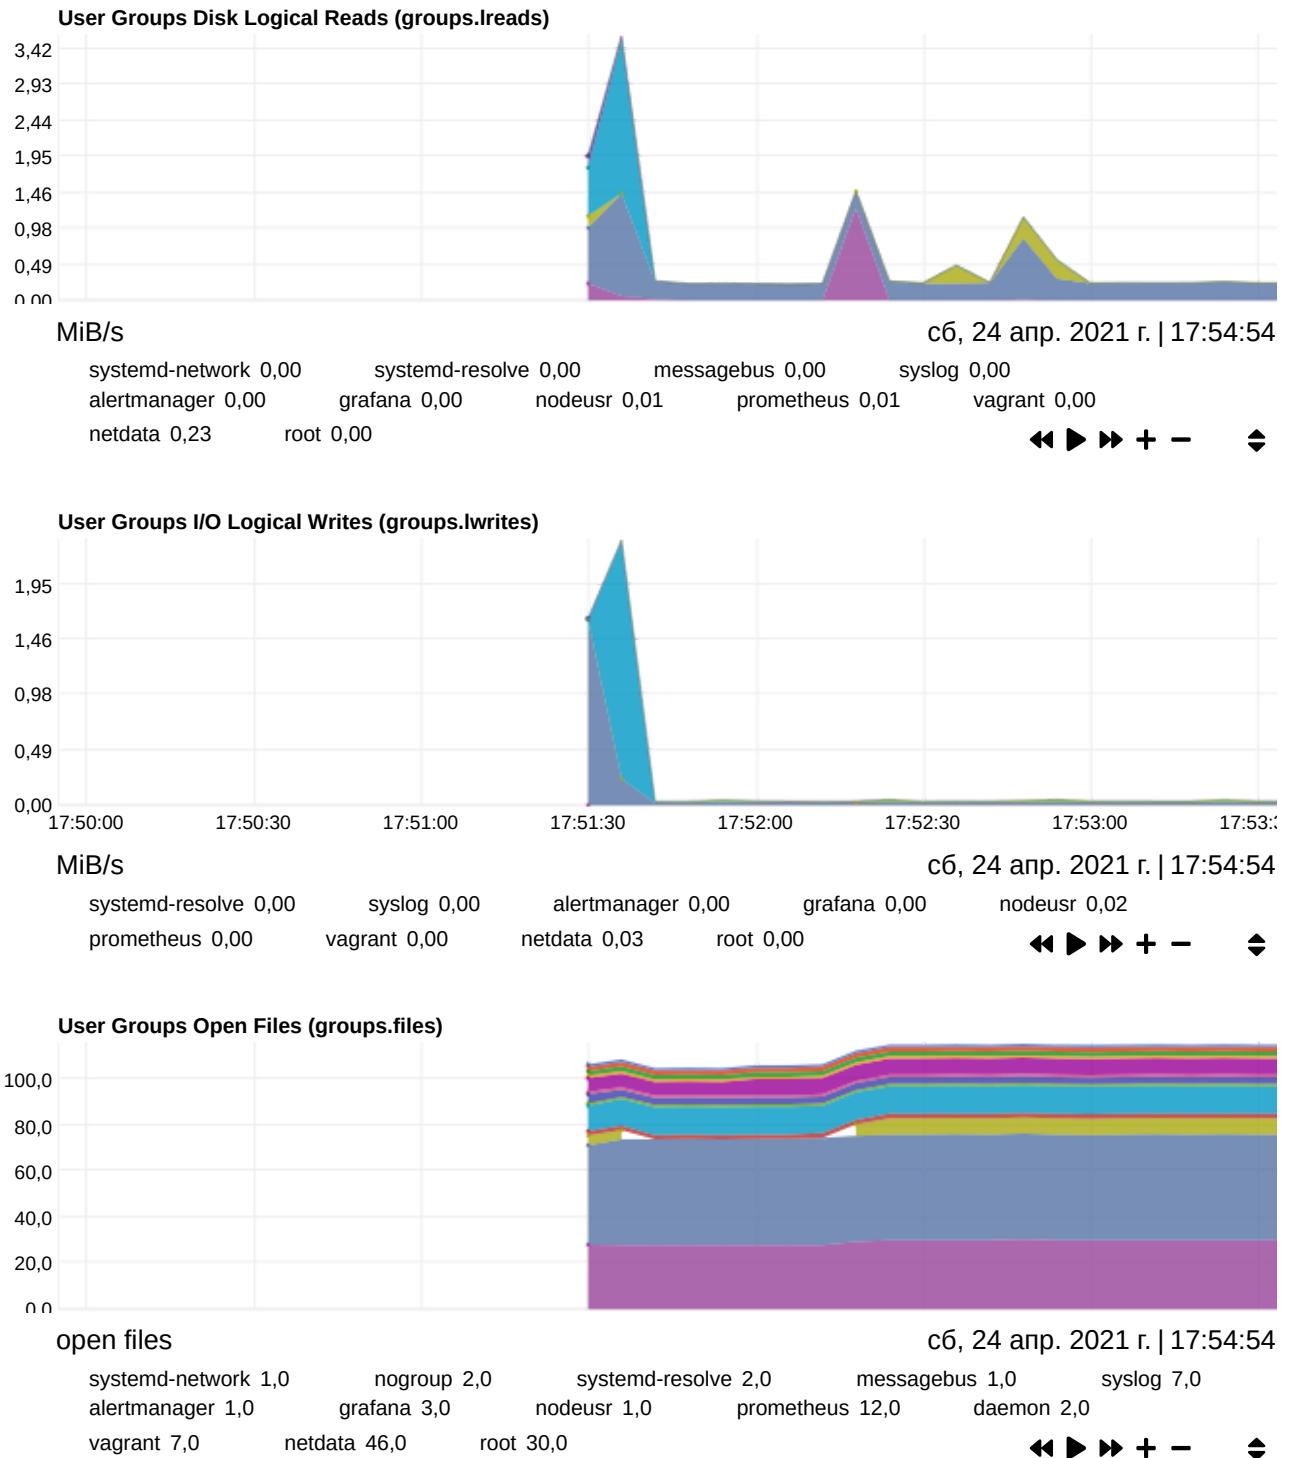
Per user group statistics are collected using netdata's `apps.plugin`. This plugin walks through all processes and aggregates statistics per user group. The reported values are compatible with `top`, although the netdata plugin counts also the resources of exited children (unlike `top` which shows only the resources of the currently running processes). So for processes like shell scripts, the reported values include the resources used by the commands these scripts run within each timeframe.

cpu



disk

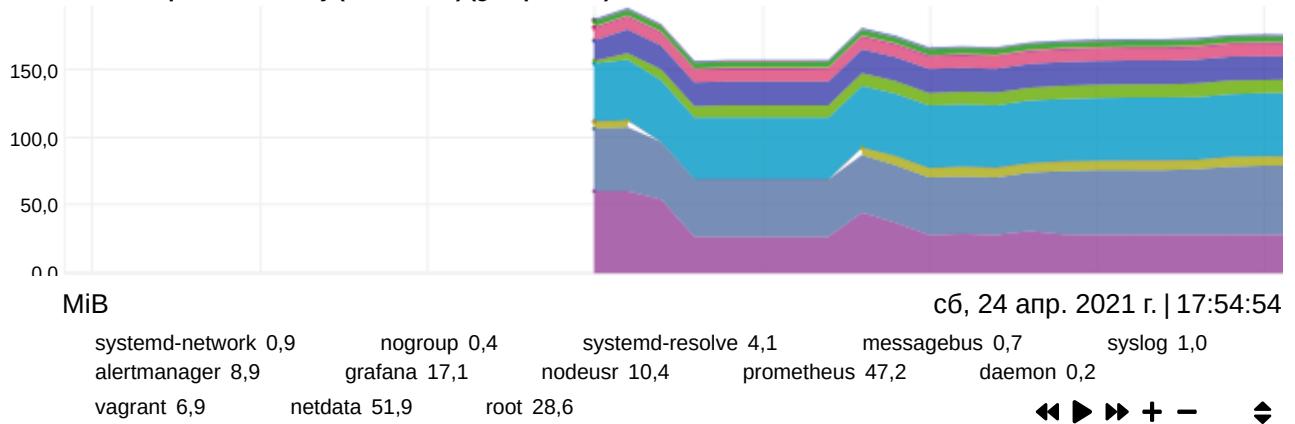




mem

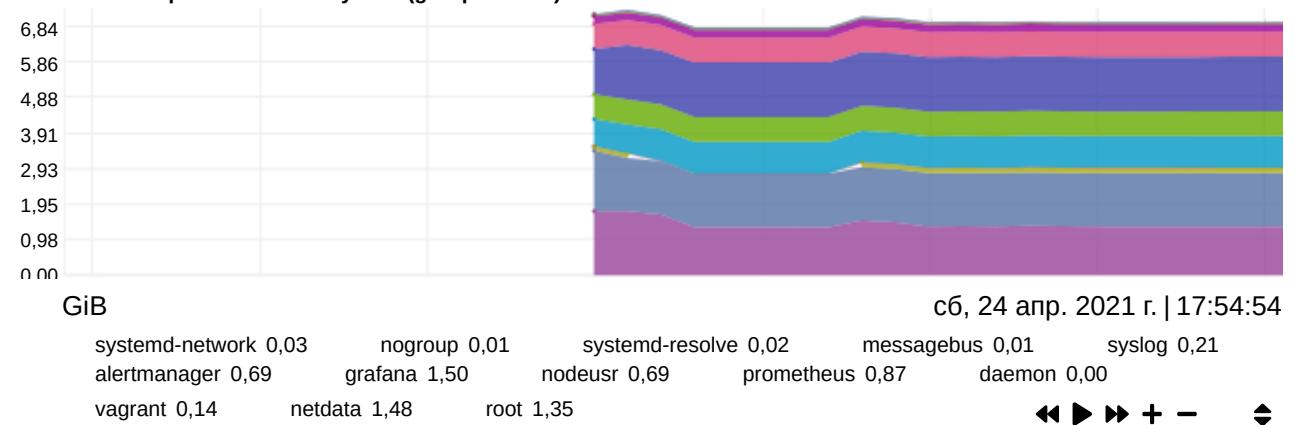
Real memory (RAM) used per user group. This does not include shared memory.

User Groups Real Memory (w/o shared) (groups.mem)

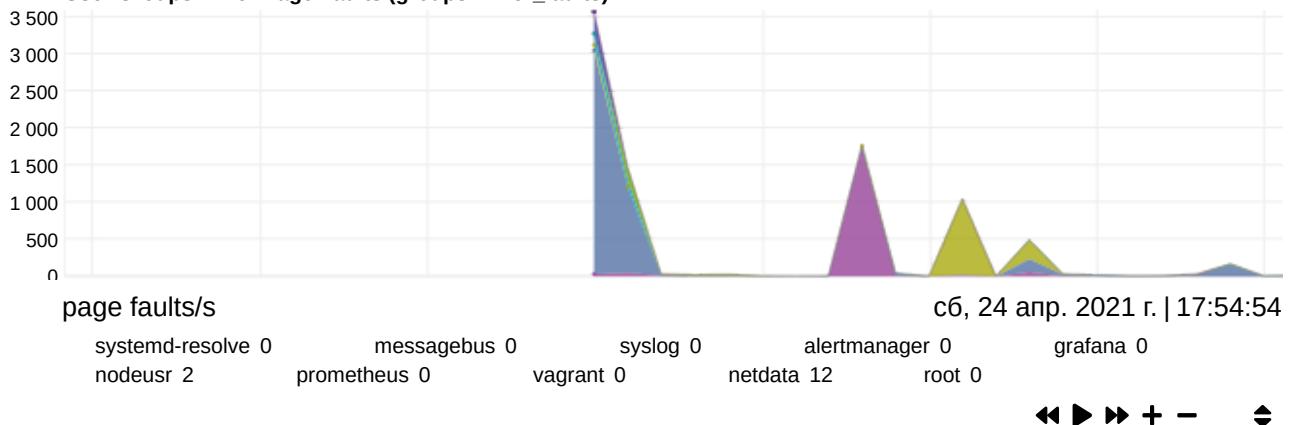


Virtual memory allocated per user group since the Netdata restart. Please check this article (<https://github.com/netdata/netdata/tree/master/daemon#virtual-memory>) for more information.

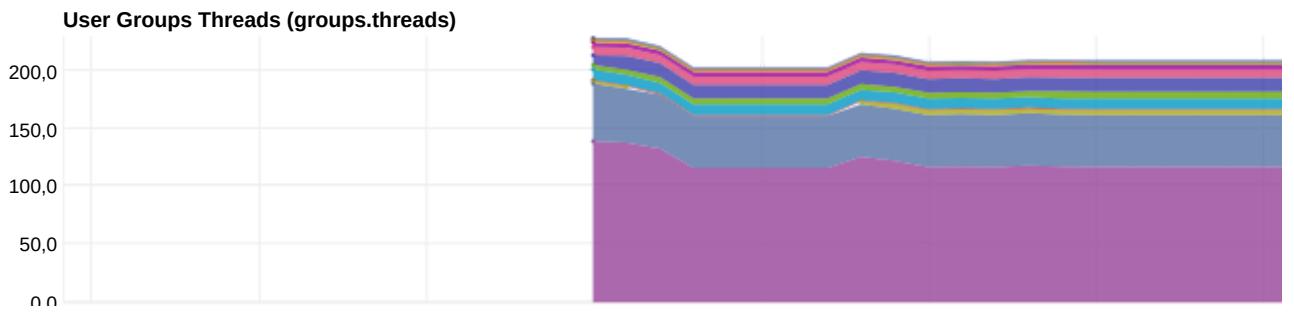
User Groups Virtual Memory Size (groups.vmem)



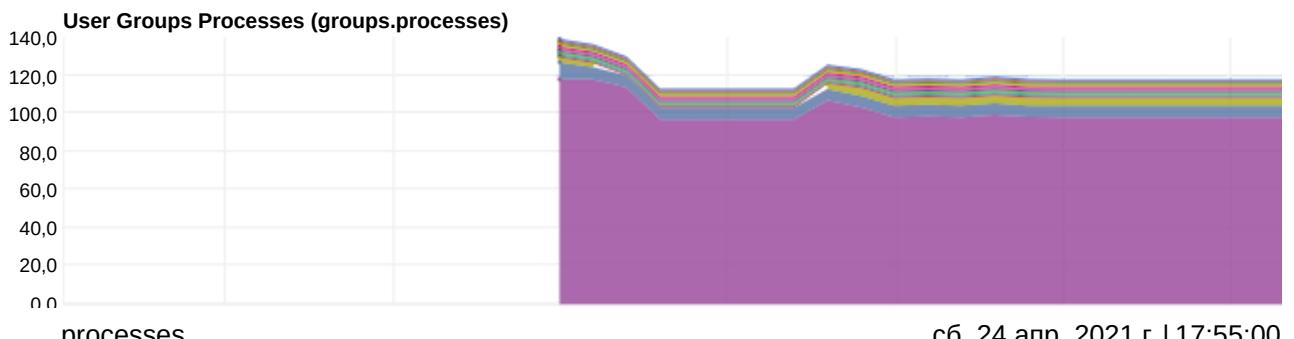
User Groups Minor Page Faults (groups.minor_faults)



processes



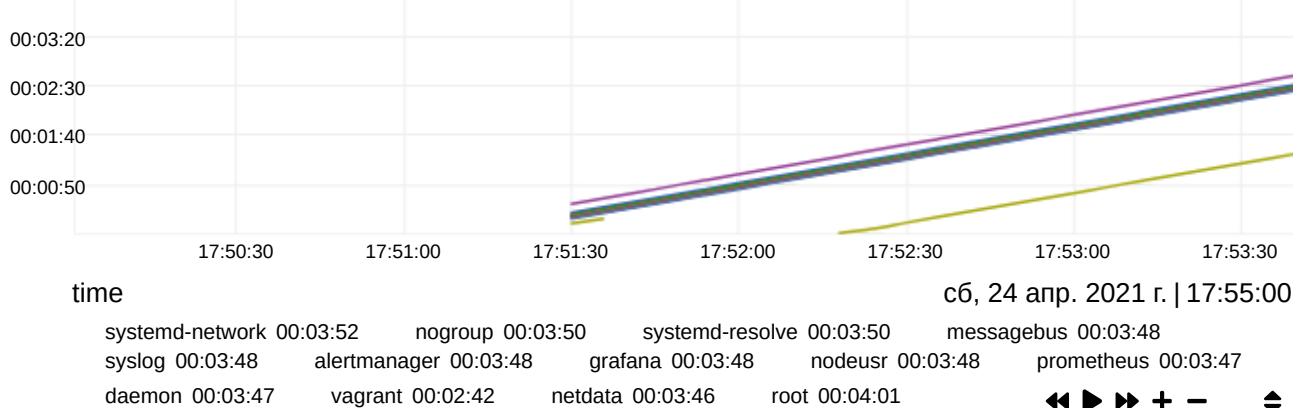
◀ ▶ ▷ + - ⌂



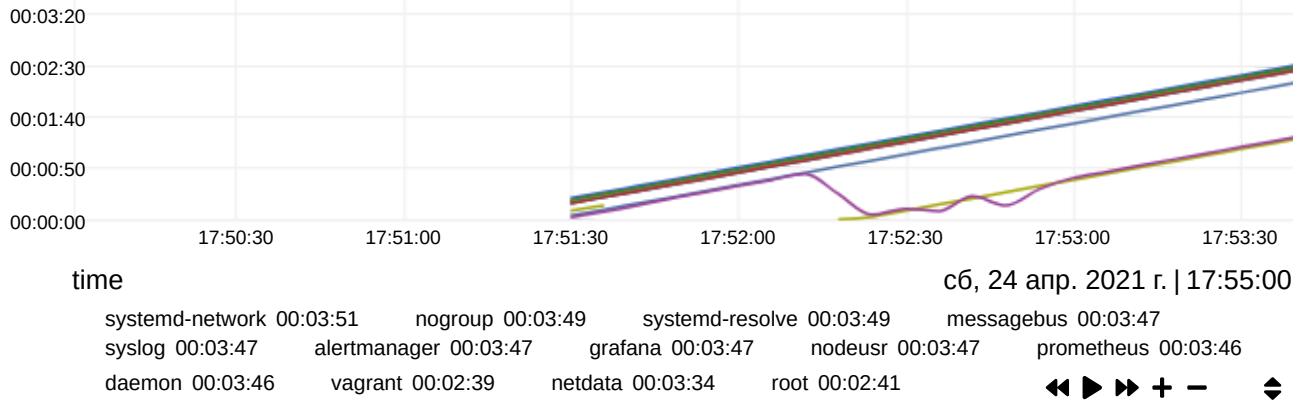
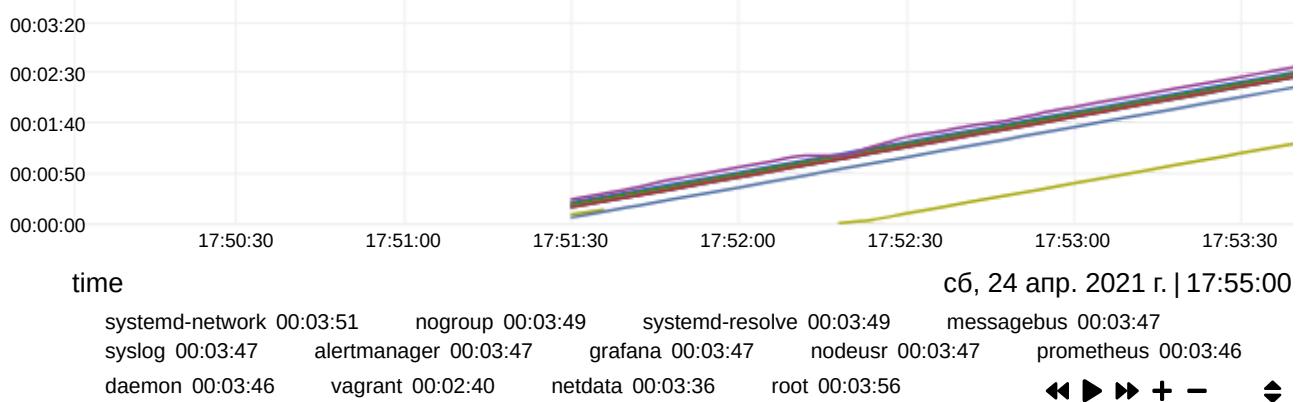
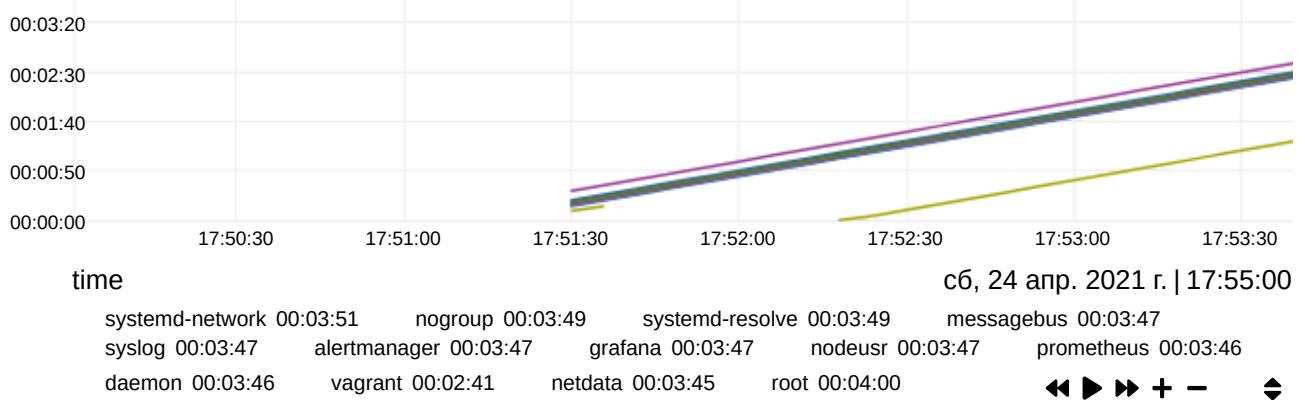
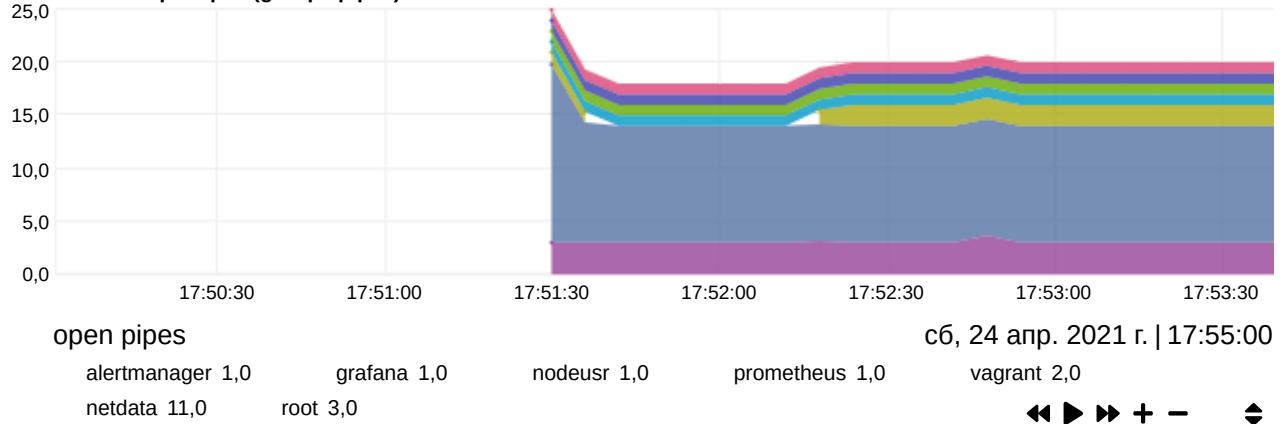
◀ ▶ ▷ + - ⌂

Carried over process group uptime. The period of time within which at least one process in the group was running.

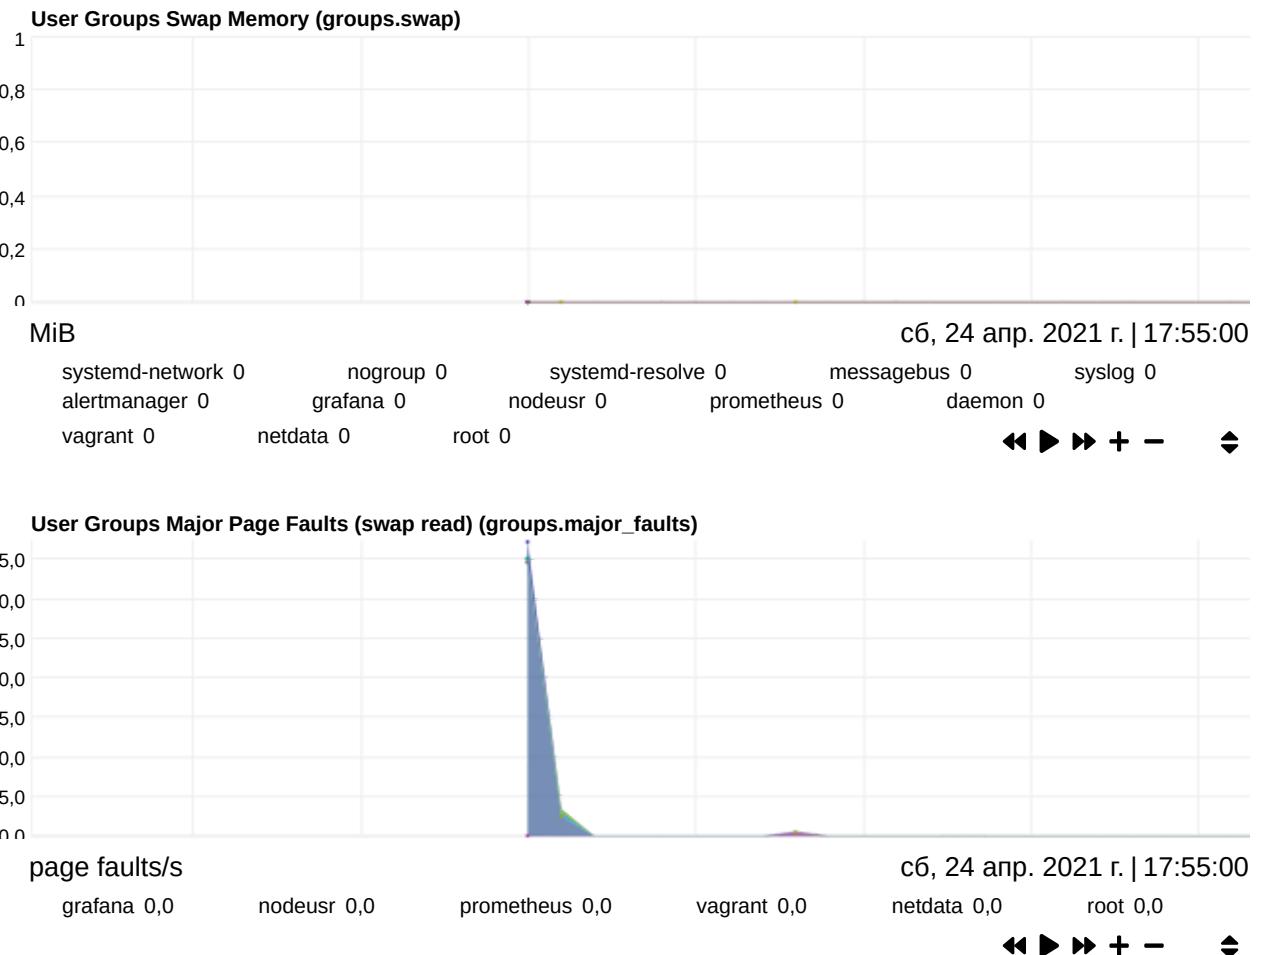
User Groups Carried Over Uptime (groups.uptime)



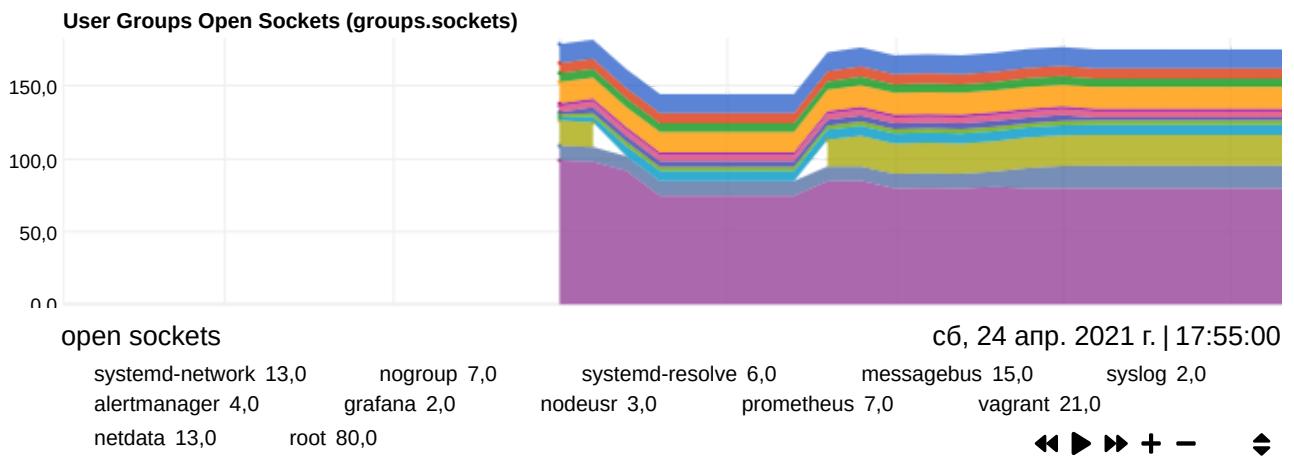
◀ ▶ ▷ + - ⌂

User Groups Minimum Uptime (groups.uptime_min)**User Groups Average Uptime (groups.uptime_avg)****User Groups Maximum Uptime (groups.uptime_max)****User Groups Pipes (groups.pipes)**

Swap



net

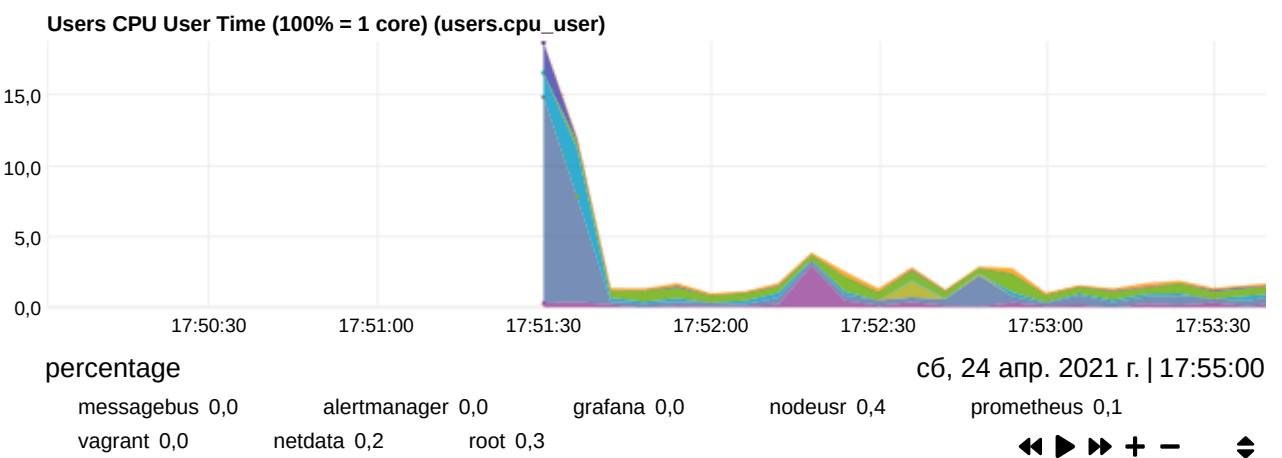
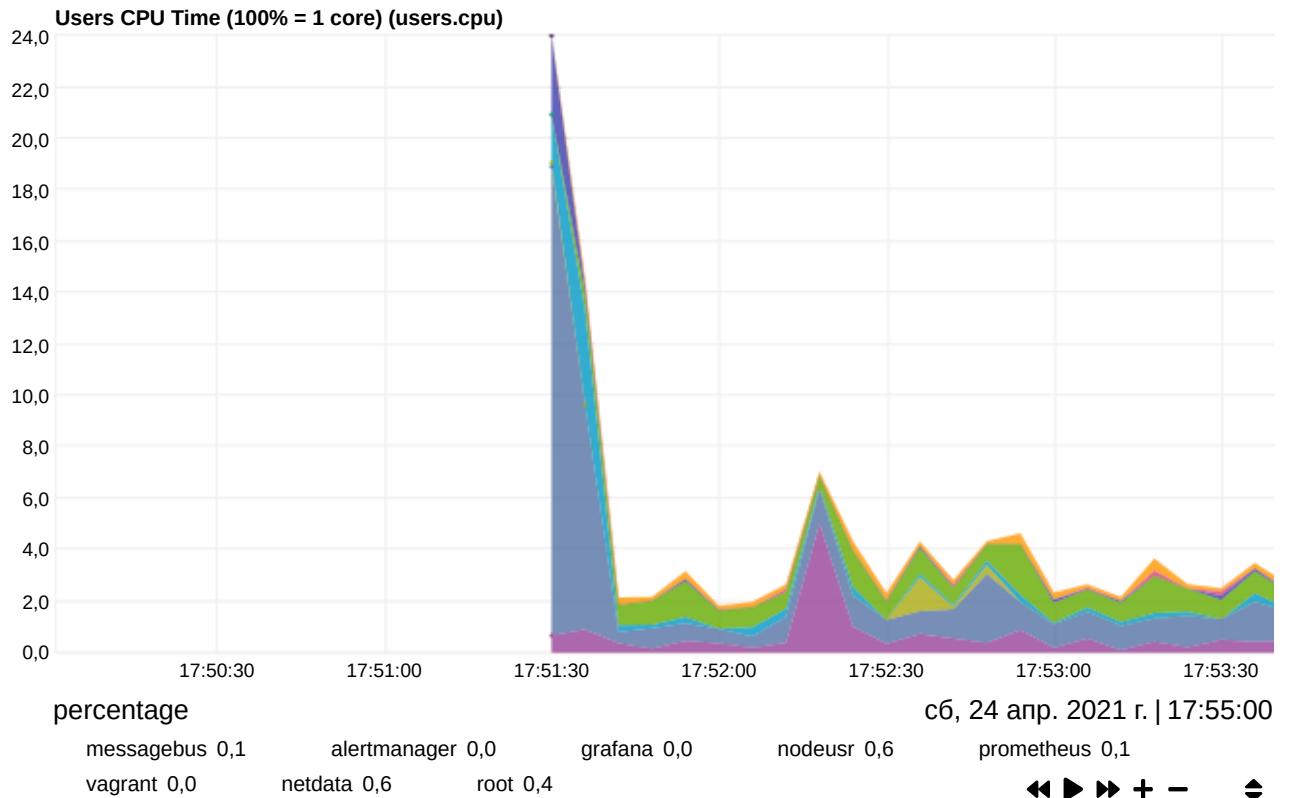


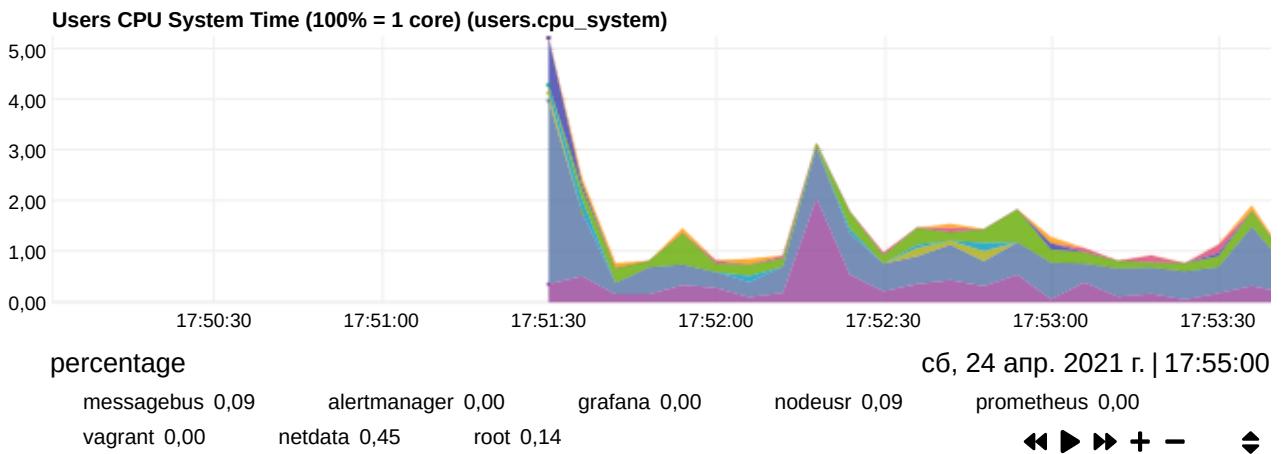
Users

Per user statistics are collected using netdata's `apps.plugin`. This plugin walks through all processes and aggregates statistics per user. The reported values are compatible with `top`, although the netdata plugin counts also the resources of exited children (unlike `top` which shows only the resources of the

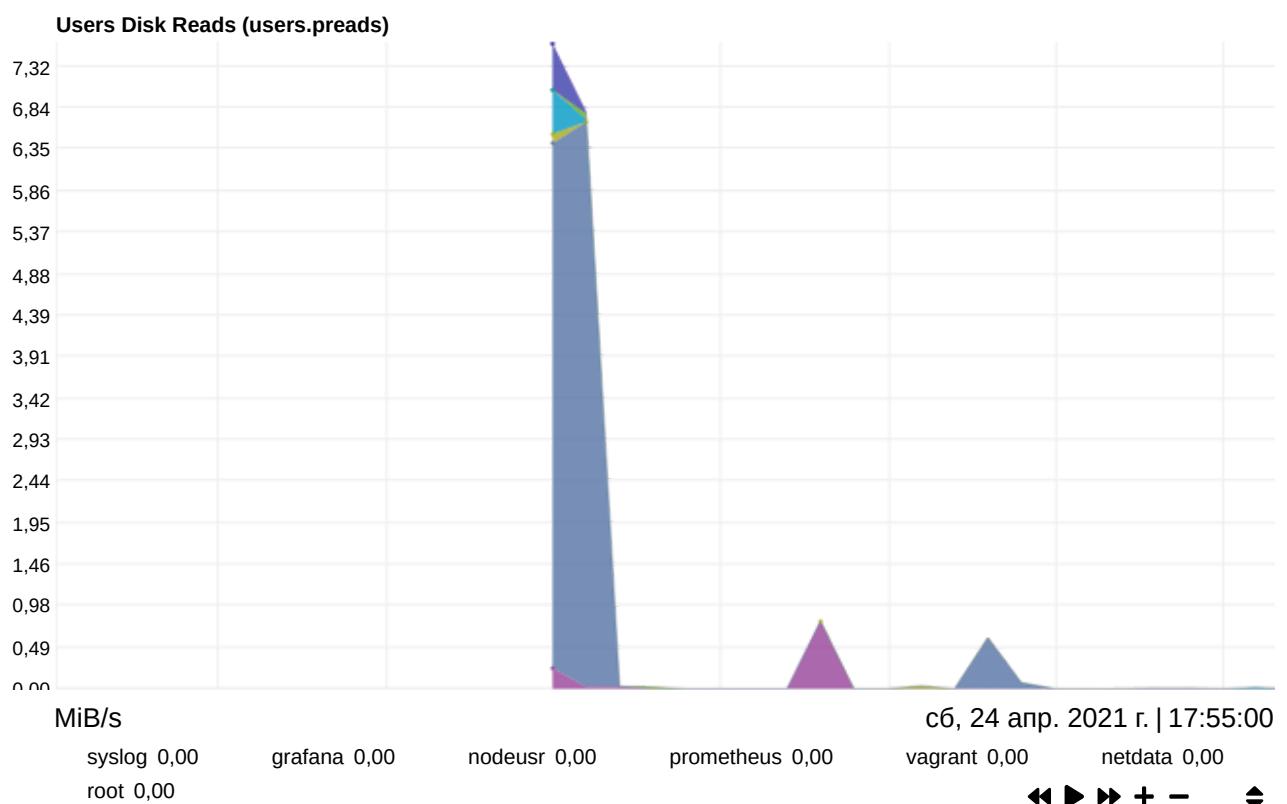
currently running processes). So for processes like shell scripts, the reported values include the resources used by the commands these scripts run within each timeframe.

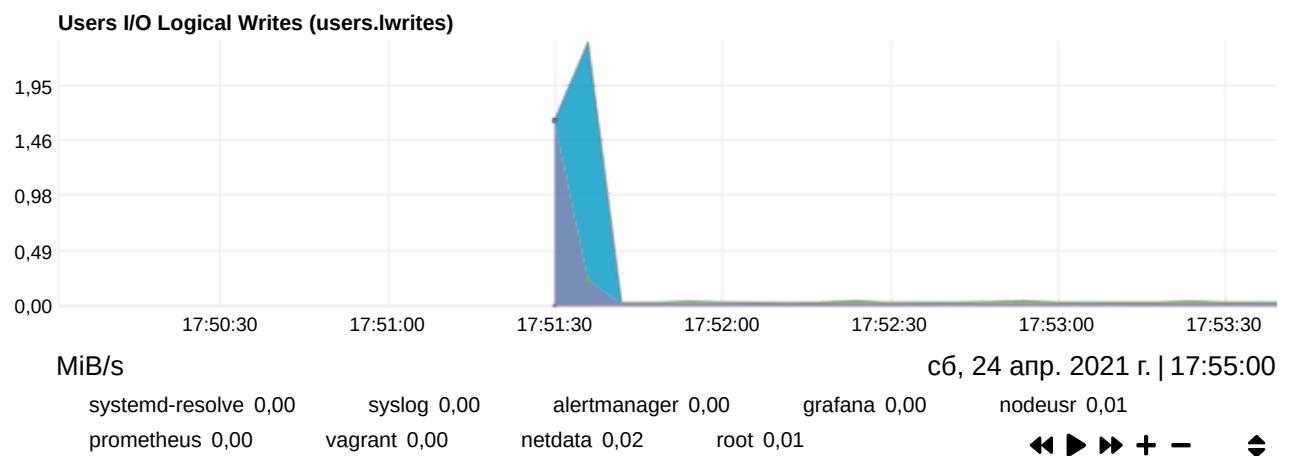
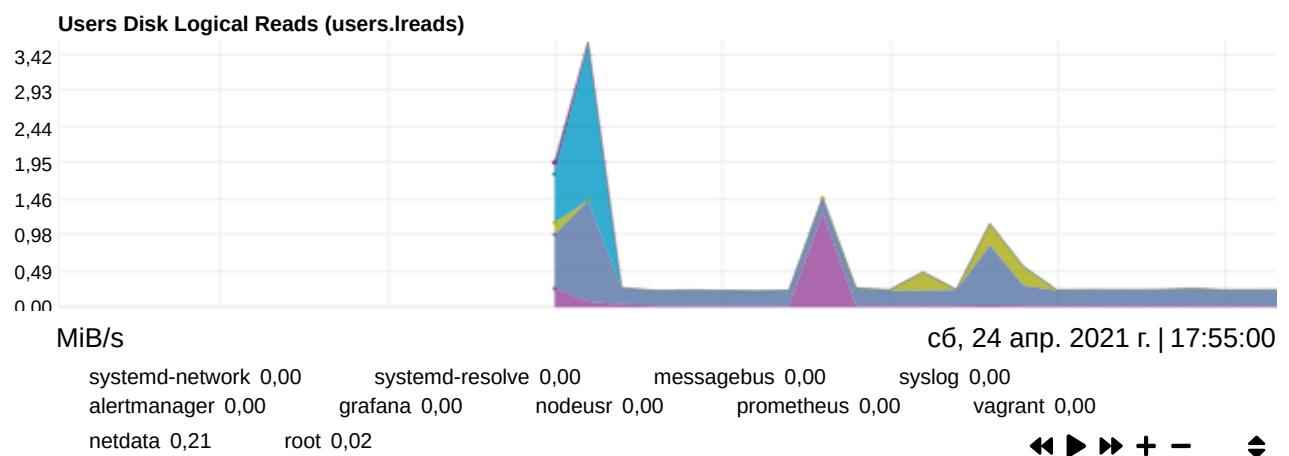
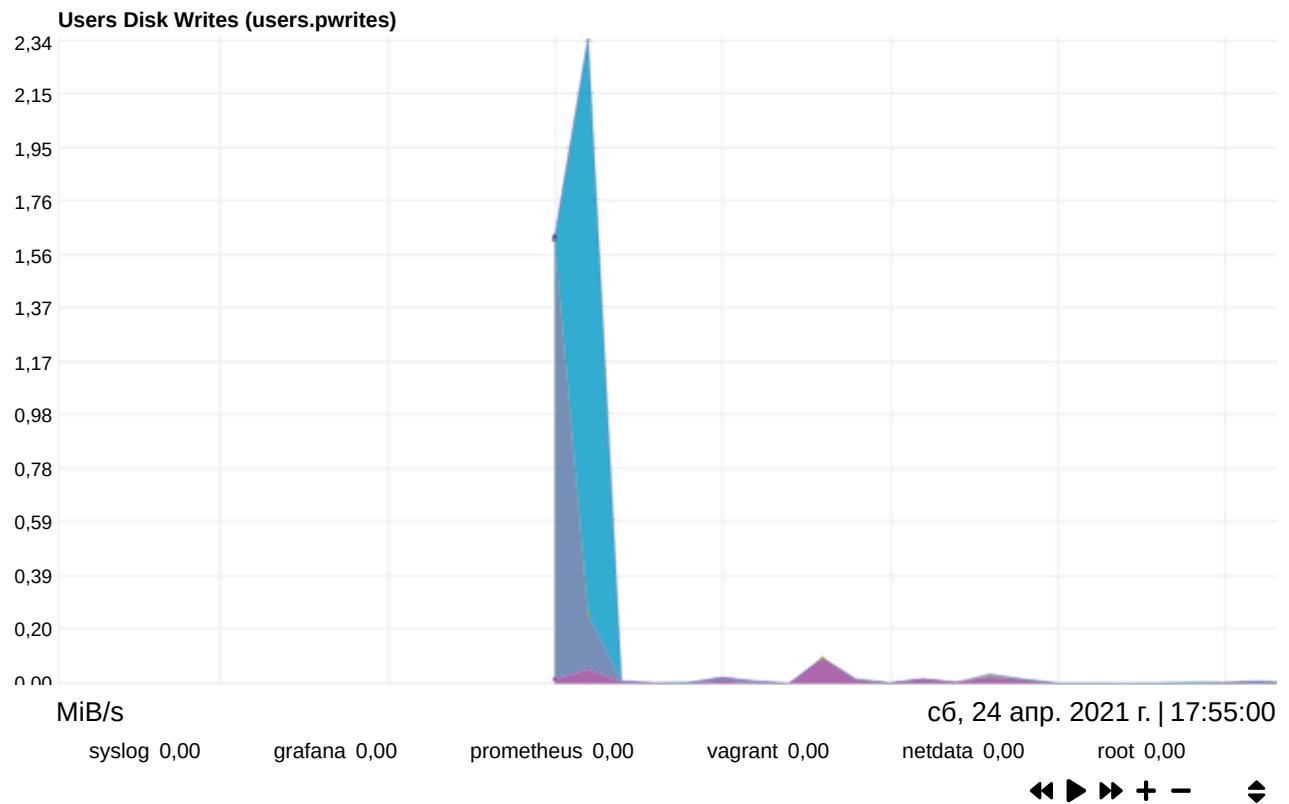
cpu

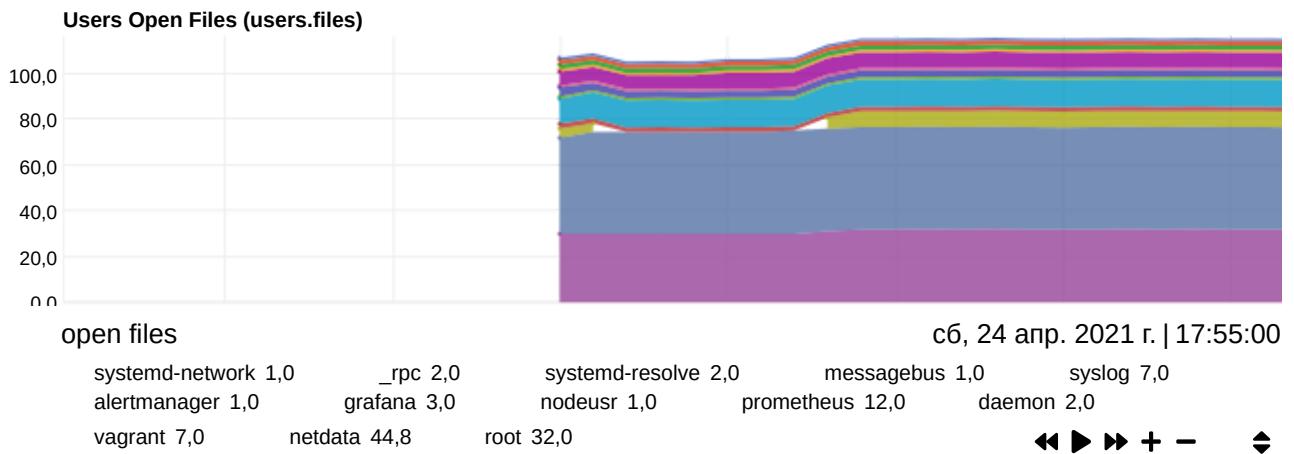




disk

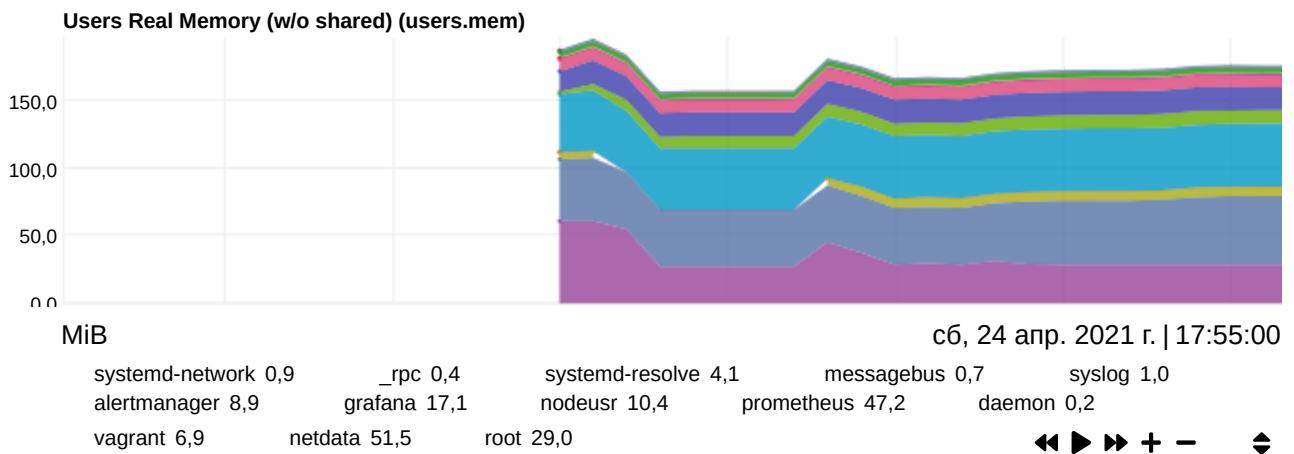






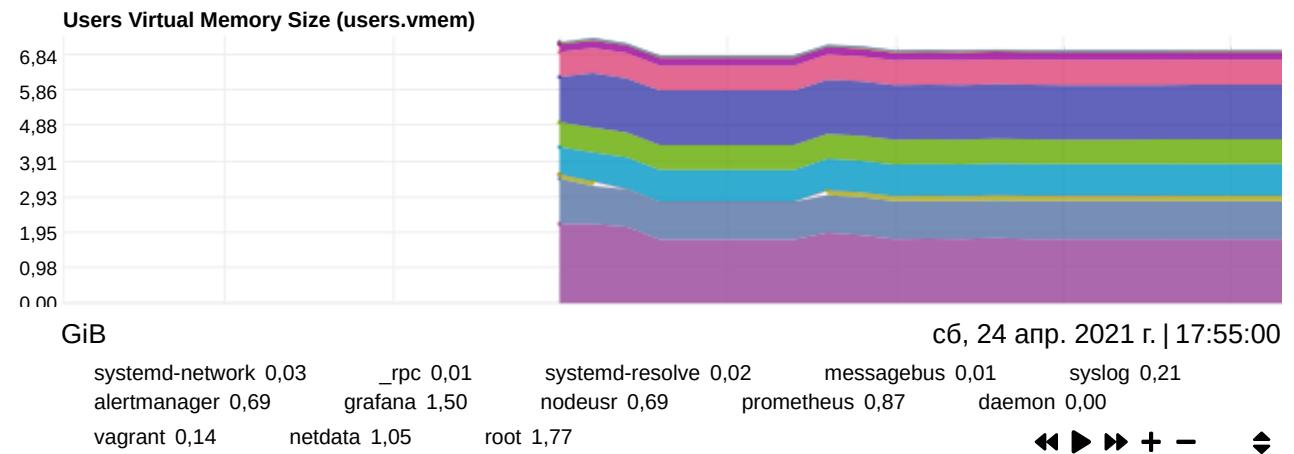
mem

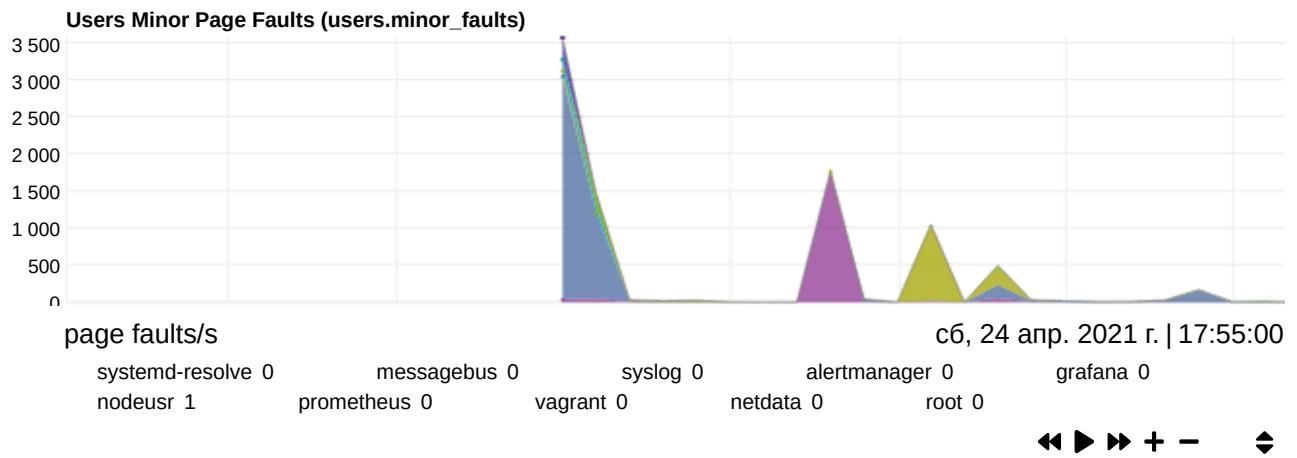
Real memory (RAM) used per user. This does not include shared memory.



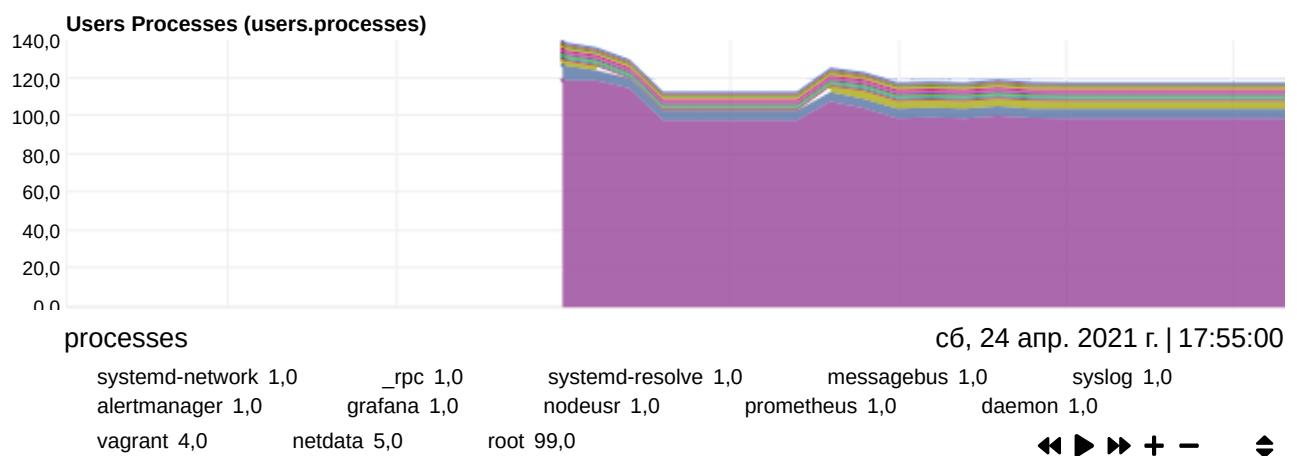
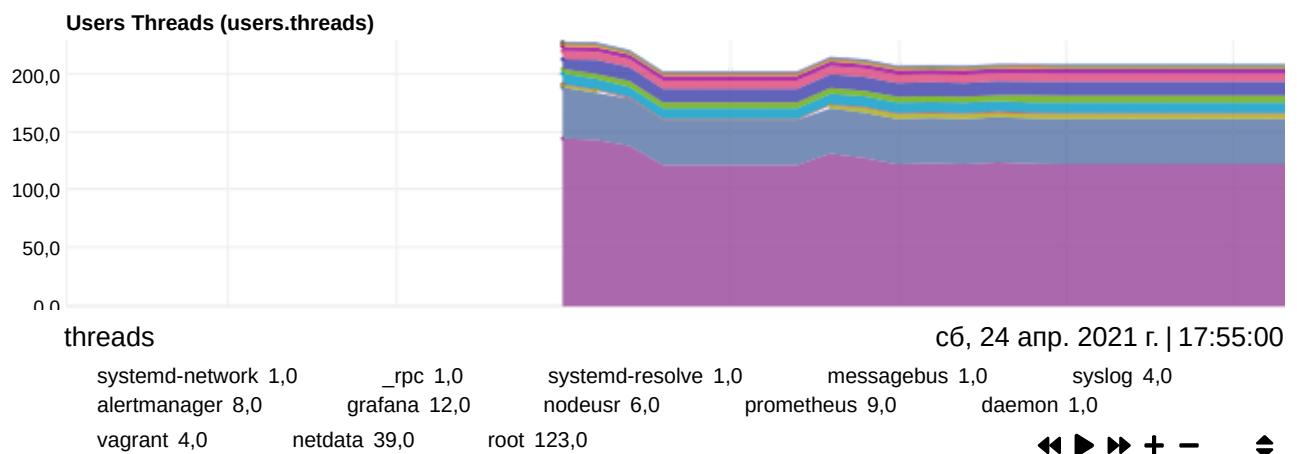
Virtual memory allocated per user. Please check this article

(<https://github.com/netdata/netdata/tree/master/daemon#virtual-memory>) for more information.



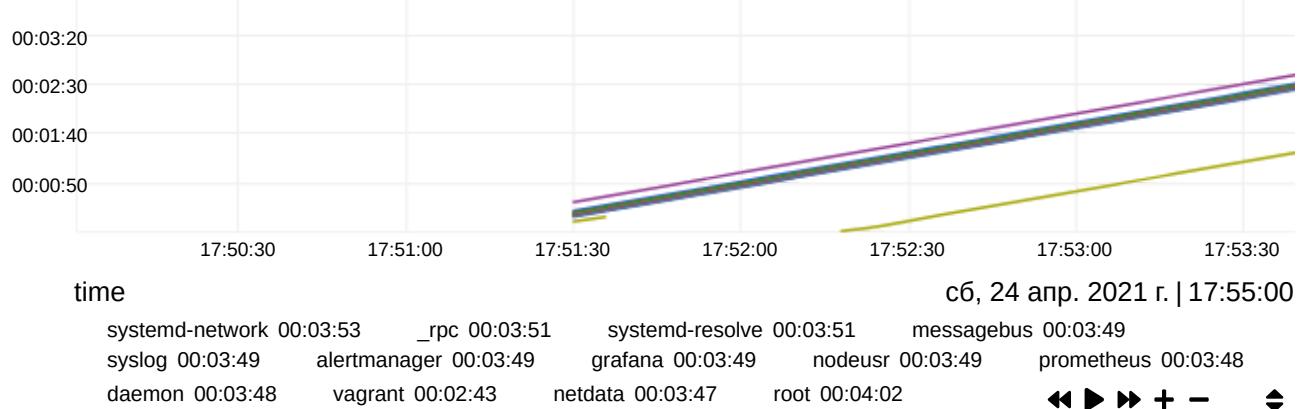


processes

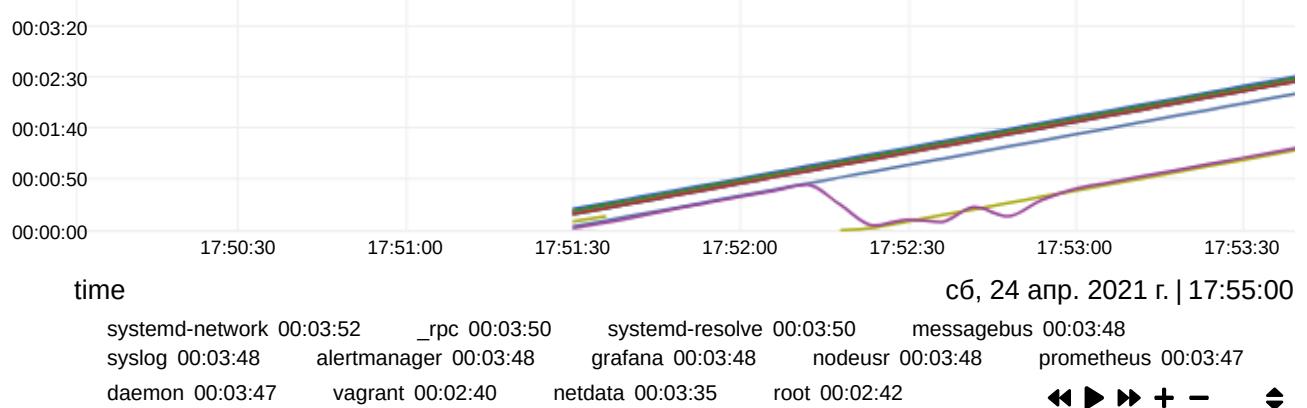


Carried over process group uptime since the Netdata restart. The period of time within which at least one process in the group was running.

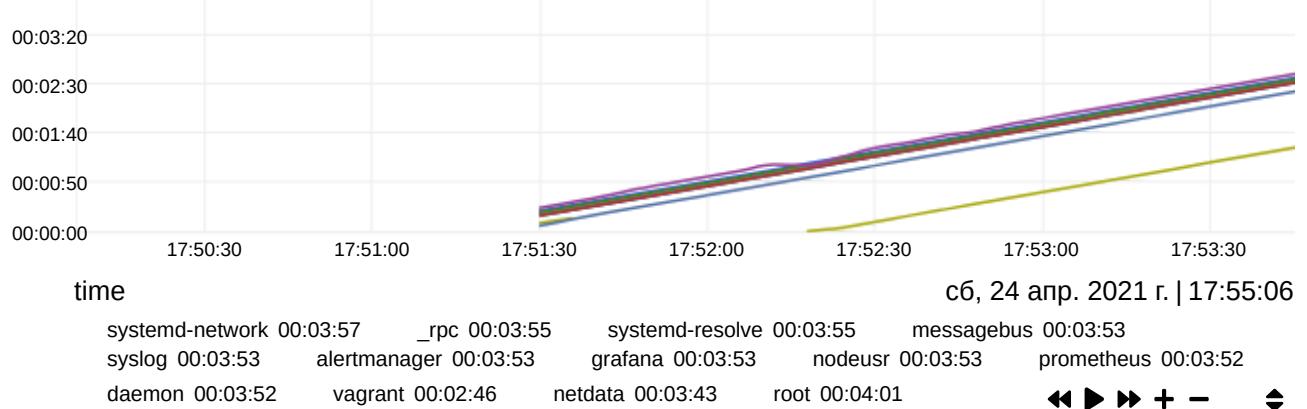
Users Carried Over Uptime (users.uptime)

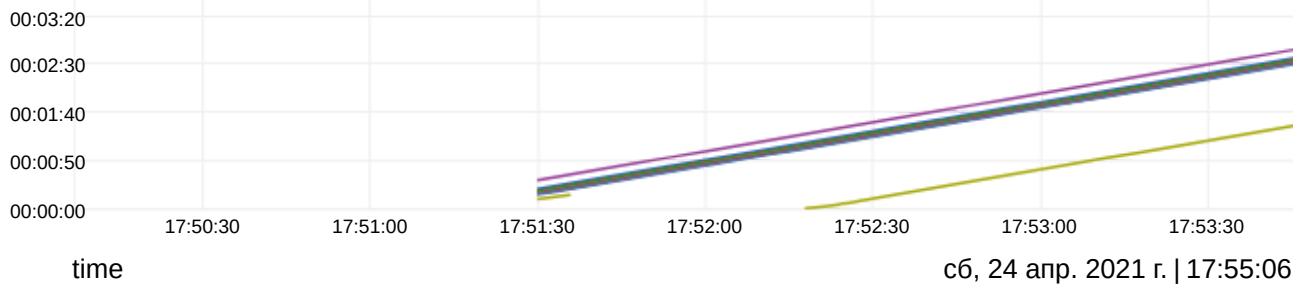


Users Minimum Uptime (users.uptime_min)

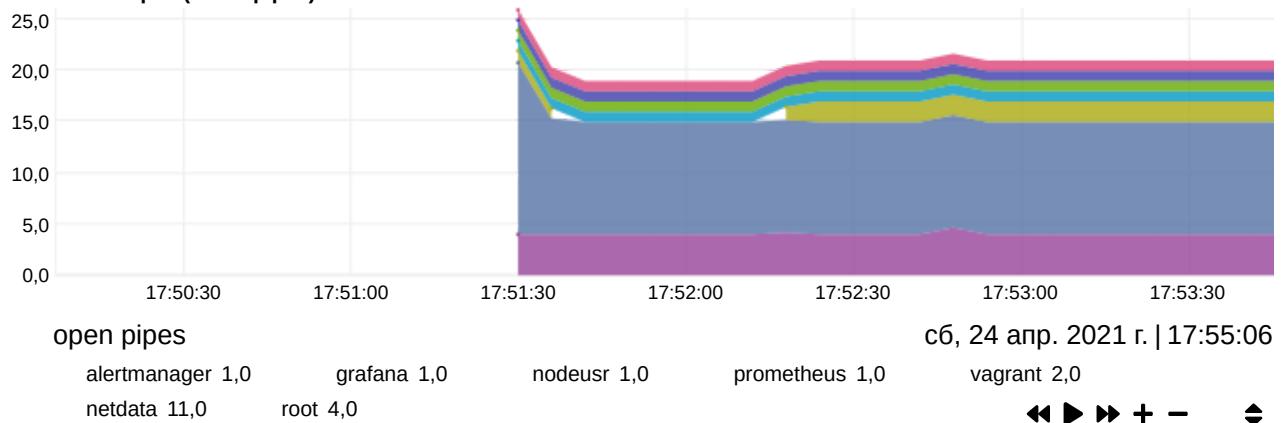


Users Average Uptime (users.uptime_avg)



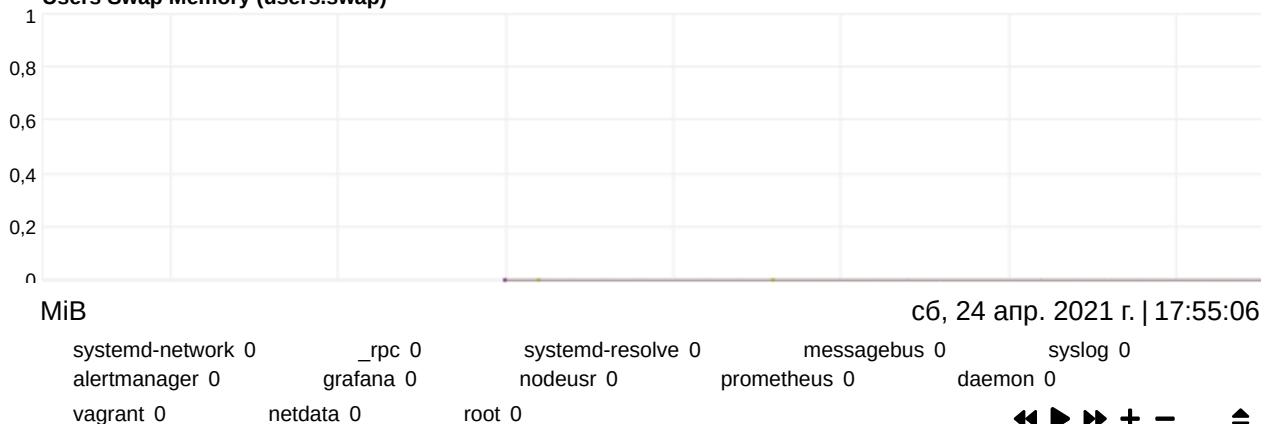
Users Maximum Uptime (users.uptime_max)

сб, 24 апр. 2021 г. | 17:55:06

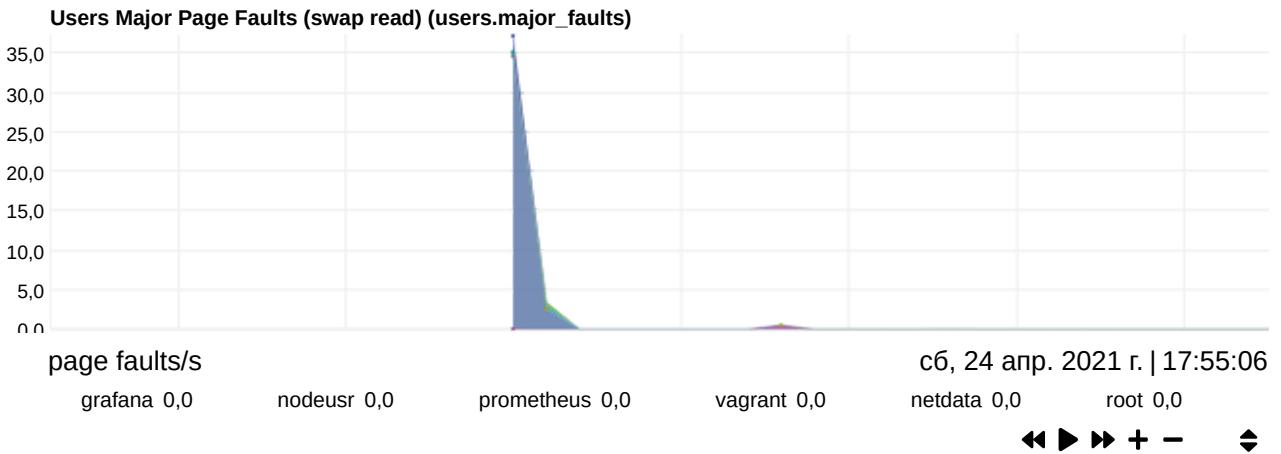
Users Pipes (users.pipes)

сб, 24 апр. 2021 г. | 17:55:06

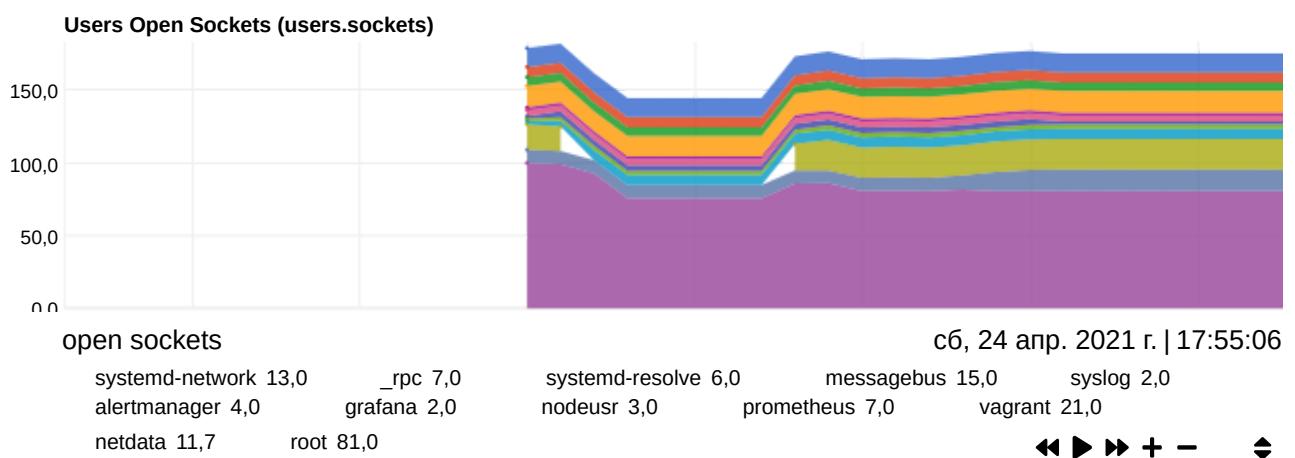
Swap

Users Swap Memory (users.swap)

сб, 24 апр. 2021 г. | 17:55:06



net

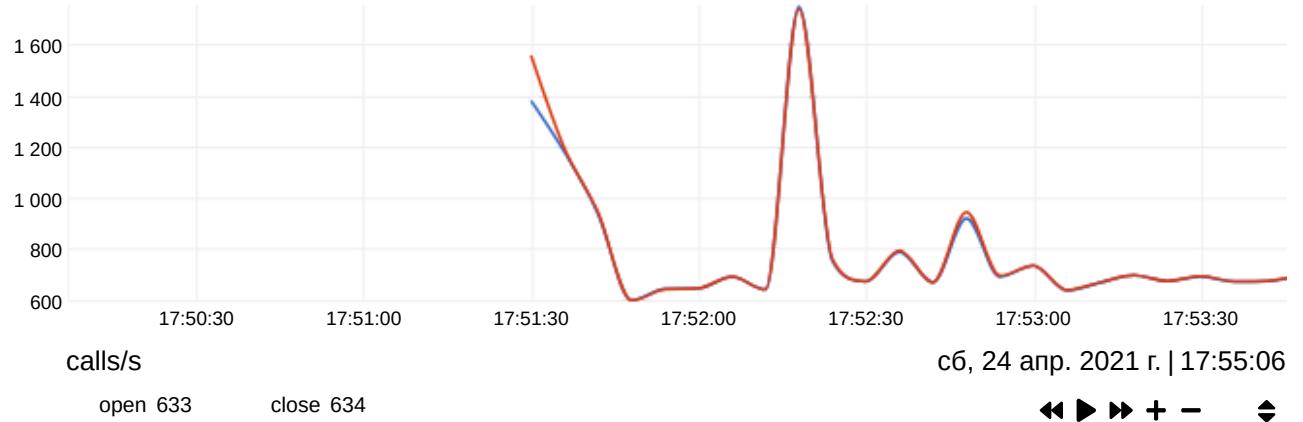


Monitor system calls, internal functions, bytes read, bytes written and errors using eBPF .

File

Calls for internal functions on Linux kernel. The open dimension is attached to the kernel internal function `do_sys_open` (For kernels newer than 5.5.19 we add a kprobe to `do_sys_openat2`.), which is the common function called from `open(2)` (<https://www.man7.org/linux/man-pages/man2/open.2.html>) and `openat(2)` (<https://www.man7.org/linux/man-pages/man2/openat.2.html>). The close dimension is attached to the function `__close_fd` or `close_fd` according to your kernel version, which is called from system call `close(2)` (<https://www.man7.org/linux/man-pages/man2/close.2.html>).

Open and close calls (ebpf.file_descriptor)



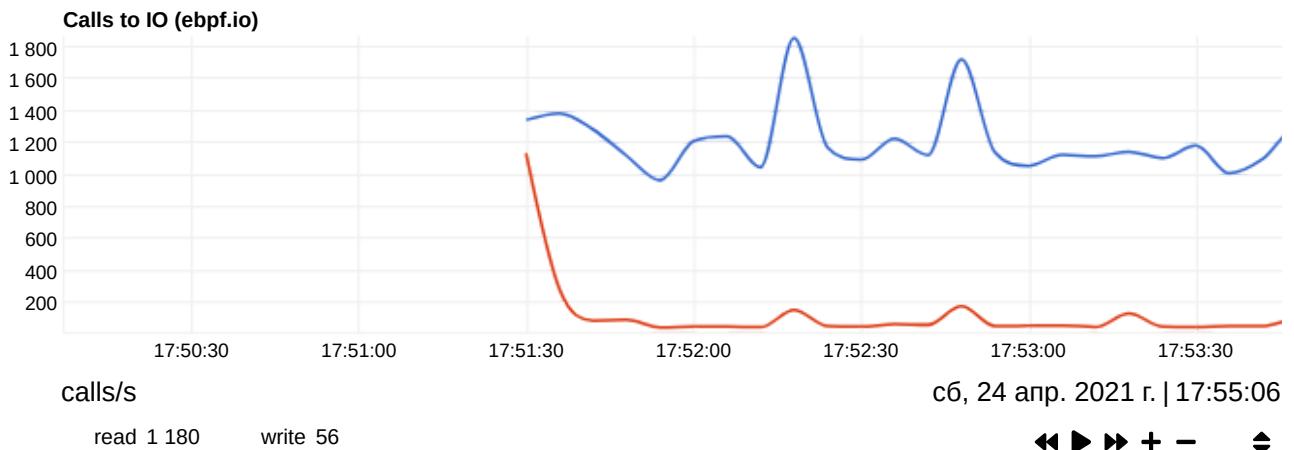
VFS

This chart does not show all events that remove files from the file system, because file systems can create their own functions to remove files, it shows calls for the function `vfs_unlink` (<https://www.kernel.org/doc/htmldocs/filesystems/API-vfs-unlink.html>).

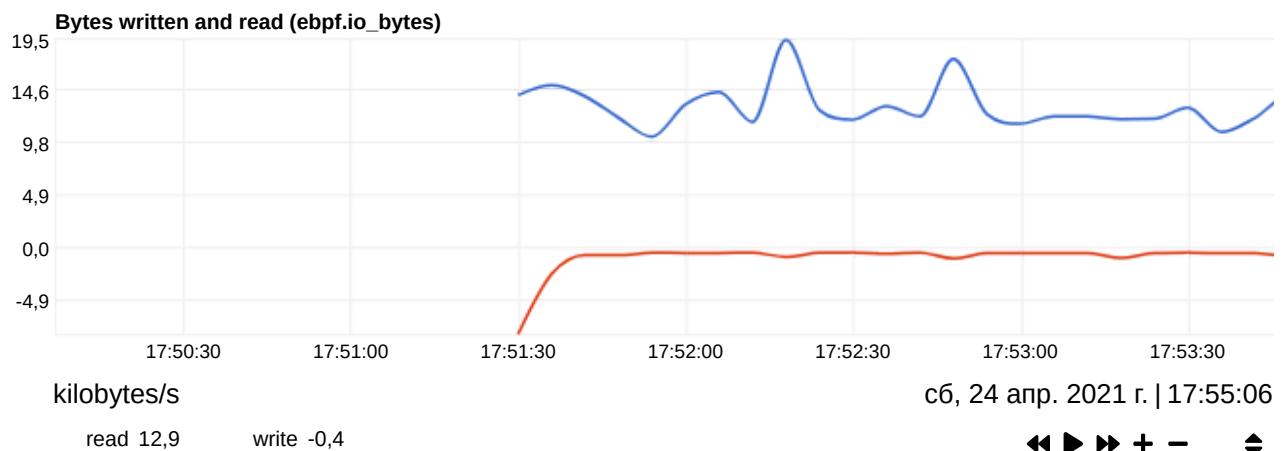
Remove files (ebpf.deleted_objects)



Successful or failed calls to functions vfs_read (https://topic.alibabacloud.com/a/kernel-state-file-operation-__-work-information-kernel_8_8_20287135.html) and vfs_write (https://topic.alibabacloud.com/a/kernel-state-file-operation-__-work-information-kernel_8_8_20287135.html). This chart may not show all file system events if it uses other functions to store data on disk.



Total of bytes read or written with success using the functions vfs_read (https://topic.alibabacloud.com/a/kernel-state-file-operation-__-work-information-kernel_8_8_20287135.html) and vfs_write (https://topic.alibabacloud.com/a/kernel-state-file-operation-__-work-information-kernel_8_8_20287135.html).



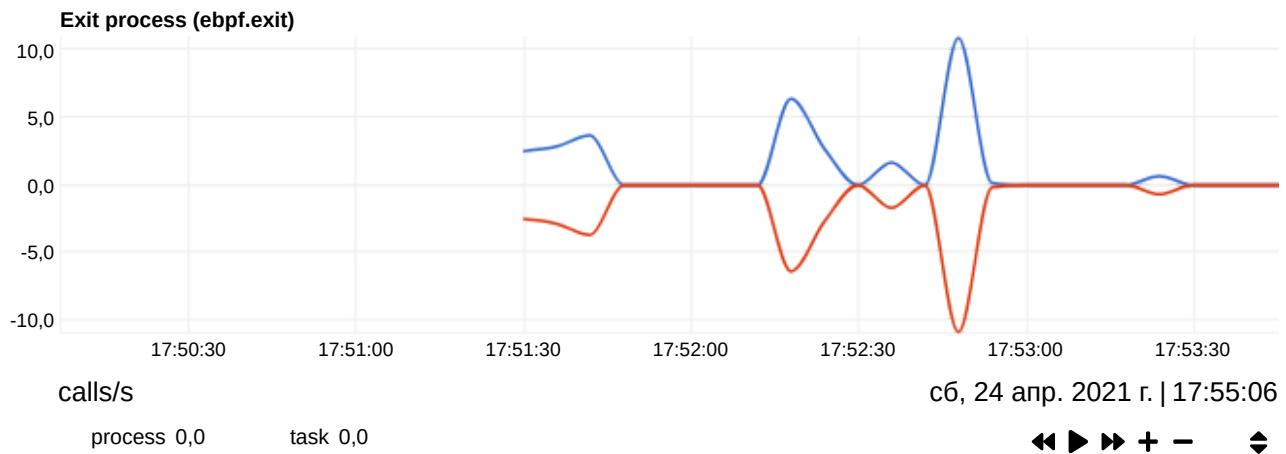
Process

Number of times that either do fork

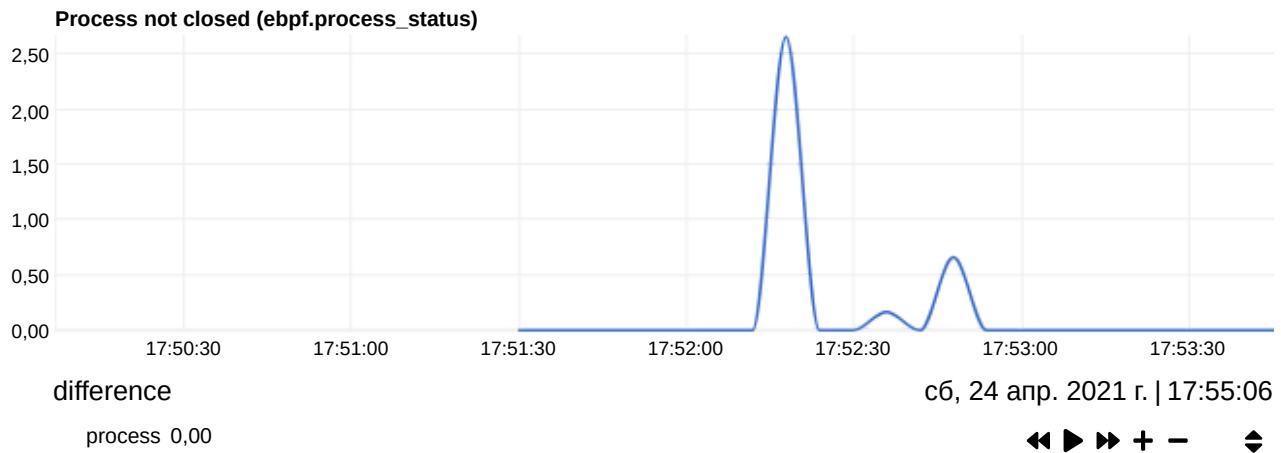
(<https://www.ece.uic.edu/~yshi1/linux/lkse/node4.html#SECTION00421000000000000000>), or `kernel_clone` if you are running kernel newer than 5.9.16, is called to create a new task, which is the common name used to define process and tasks inside the kernel. Netdata identifies the threads by counting the number of calls for `sys_clone` (<https://linux.die.net/man/2/clone>) that has the flag `CLONE_THREAD` set.



Calls for the functions responsible for closing (`do_exit` (<https://www.informit.com/articles/article.aspx?p=370047&seqNum=4>)) and releasing (`release_task` (<https://www.informit.com/articles/article.aspx?p=370047&seqNum=4>)) tasks.



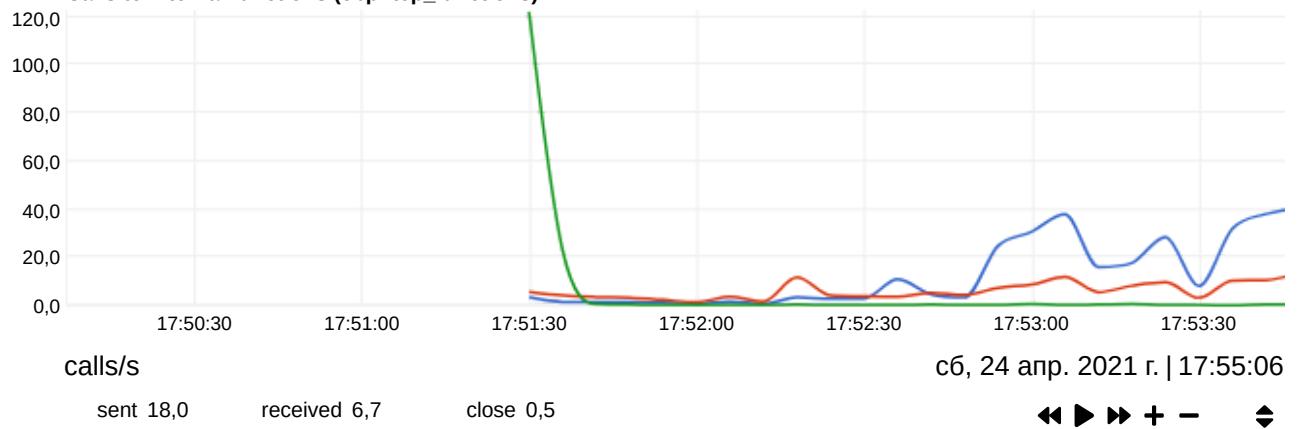
Difference between the number of process created and the number of threads created per period(process dimension), it also shows the number of possible zombie process running on system.



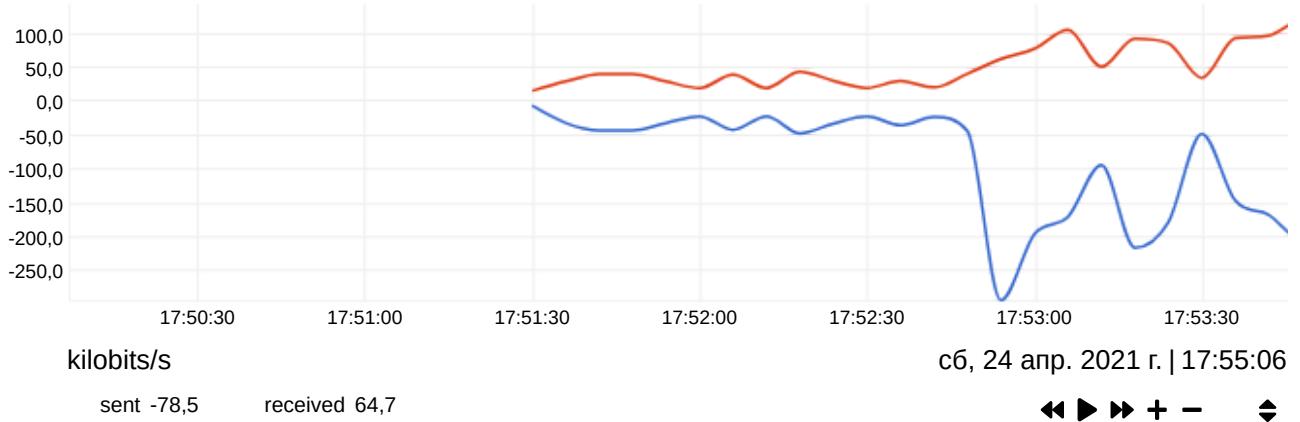
Socket

Successful or failed calls to functions `tcp_sendmsg`, `tcp_cleanup_rbuf` and `tcp_close`.

Calls to internal functions (`ebpf.tcp_functions`)



TCP bandwidth (`ebpf.total_tcp_bandwidth`)



Number of packets retransmitted for function `tcp_retransmit_skb`.

Packages retransmitted (`ebpf.tcp_retransmit`)



Successful or failed calls to functions `udp_sendmsg` and `udp_recvmsg`.

UDP calls (ebpf.udp_functions)



UDP bandwidth (ebpf.total_udp_bandwidth)

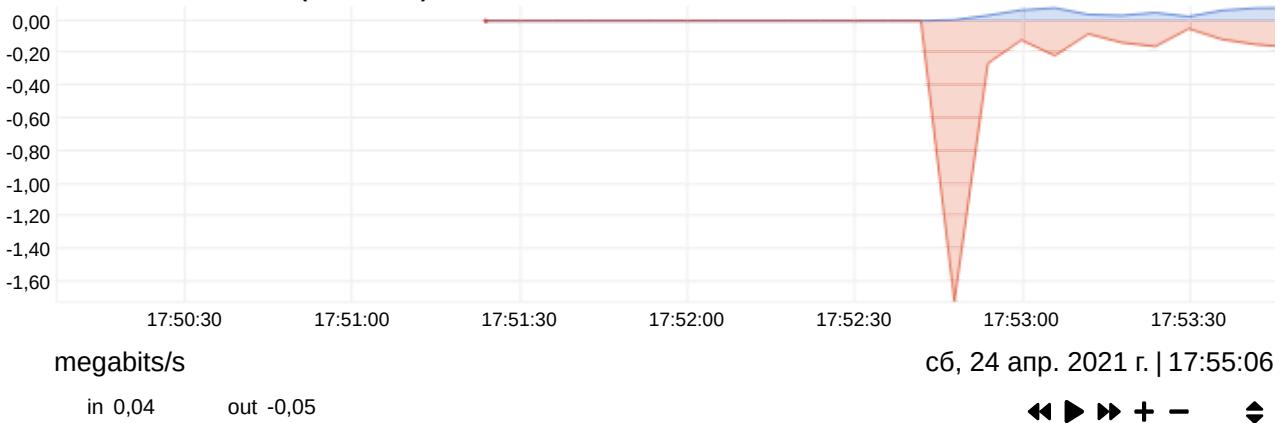


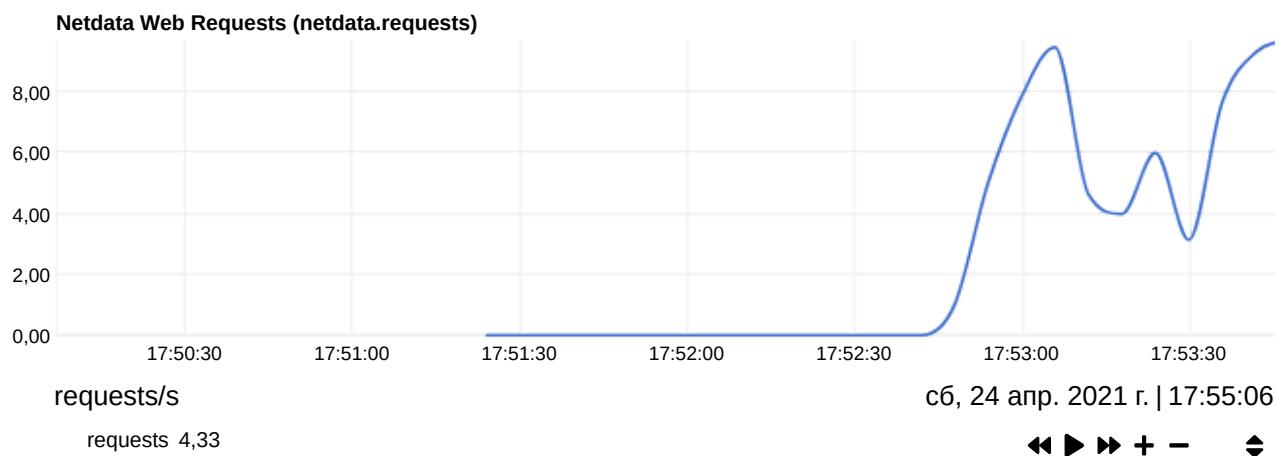
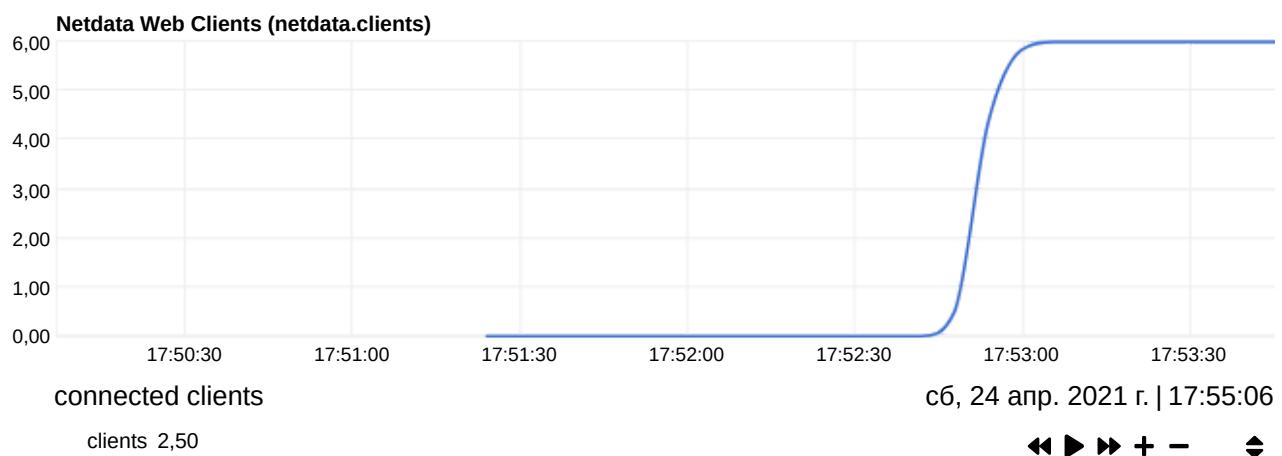
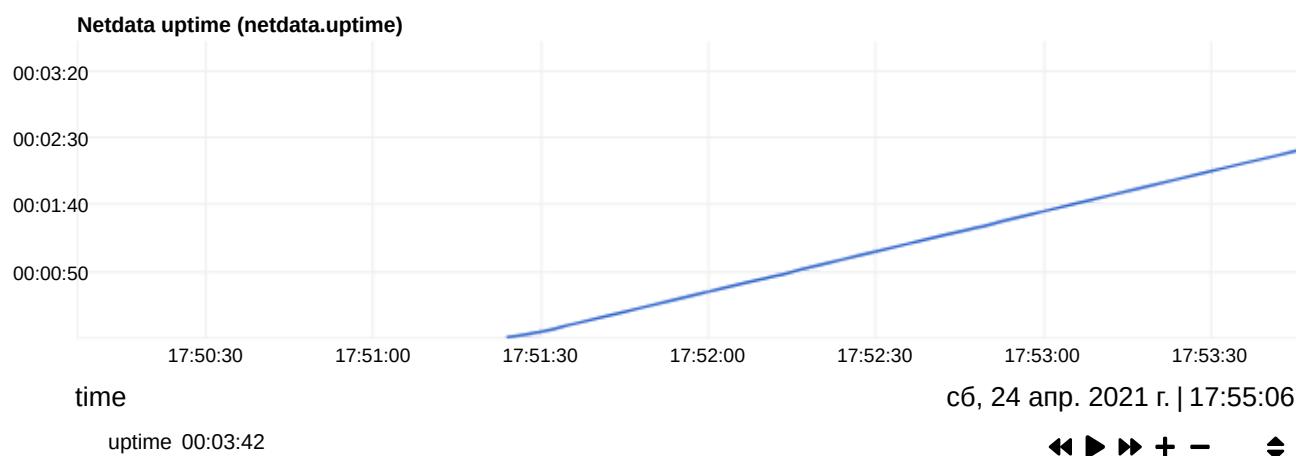
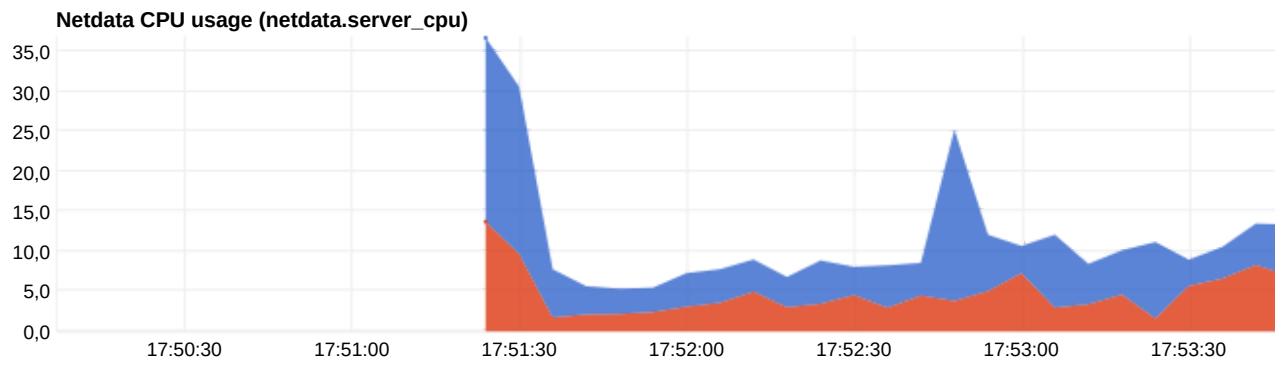
Netdata Monitoring

Performance metrics for the operation of netdata itself and its plugins.

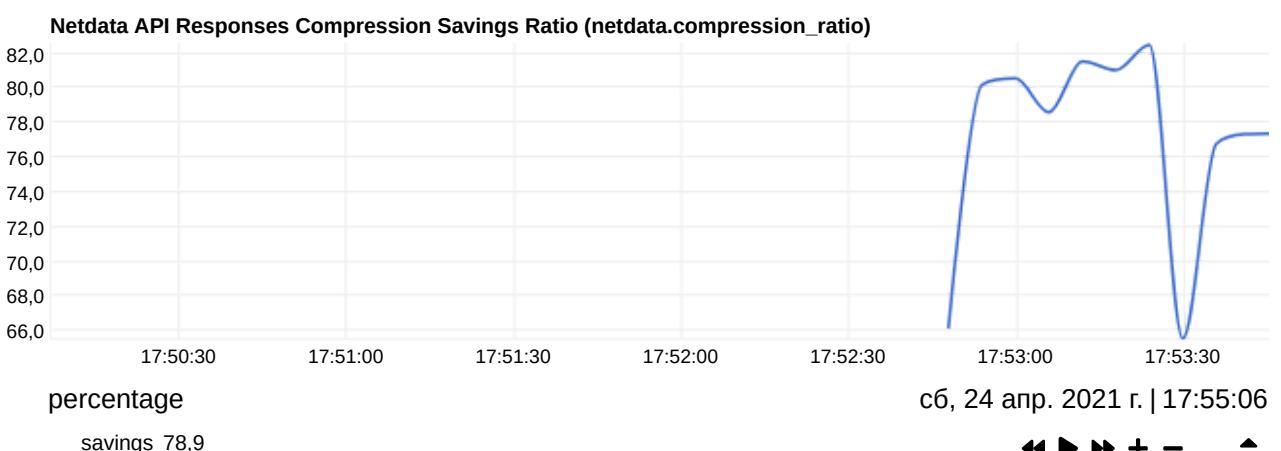
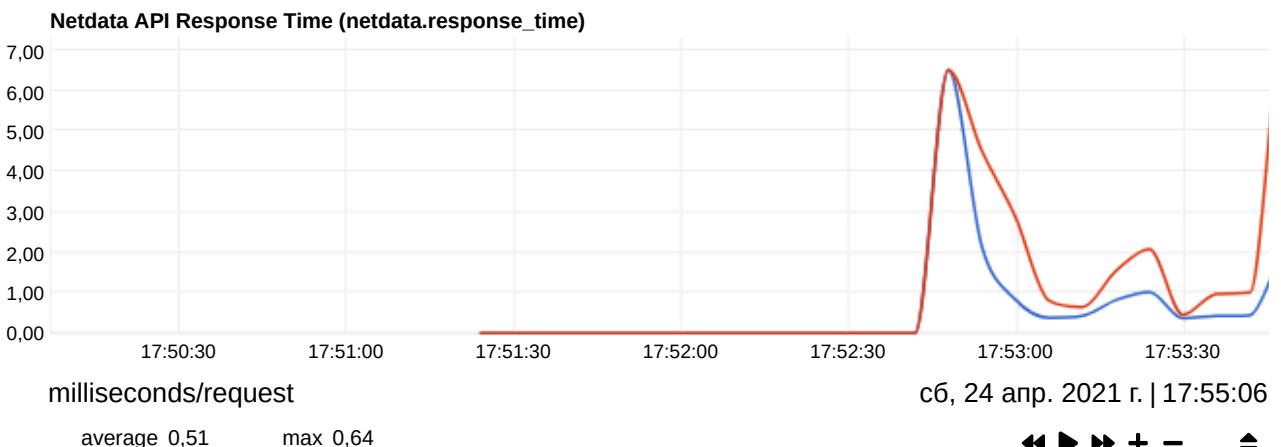
netdata

Netdata Network Traffic (netdata.net)

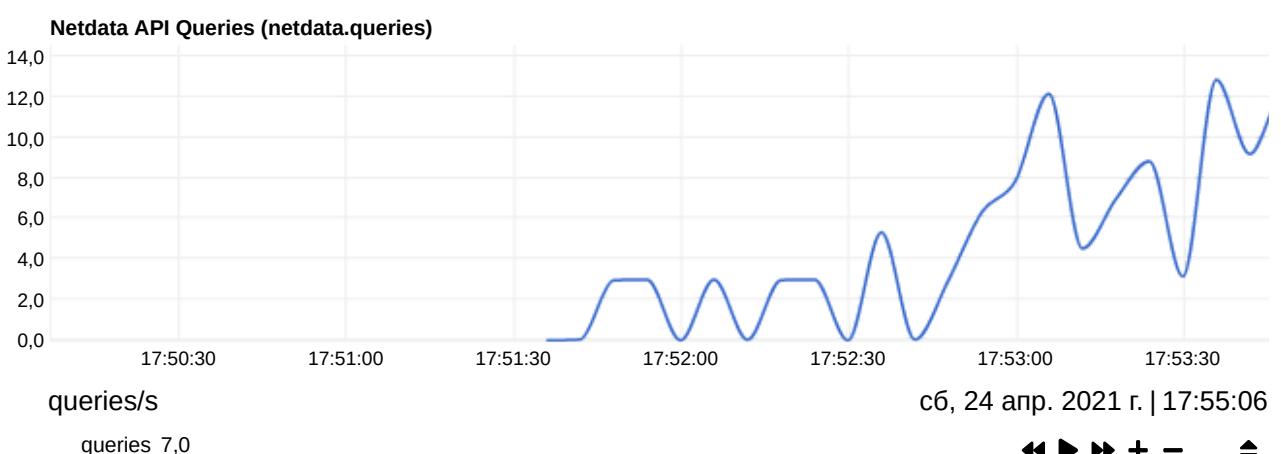


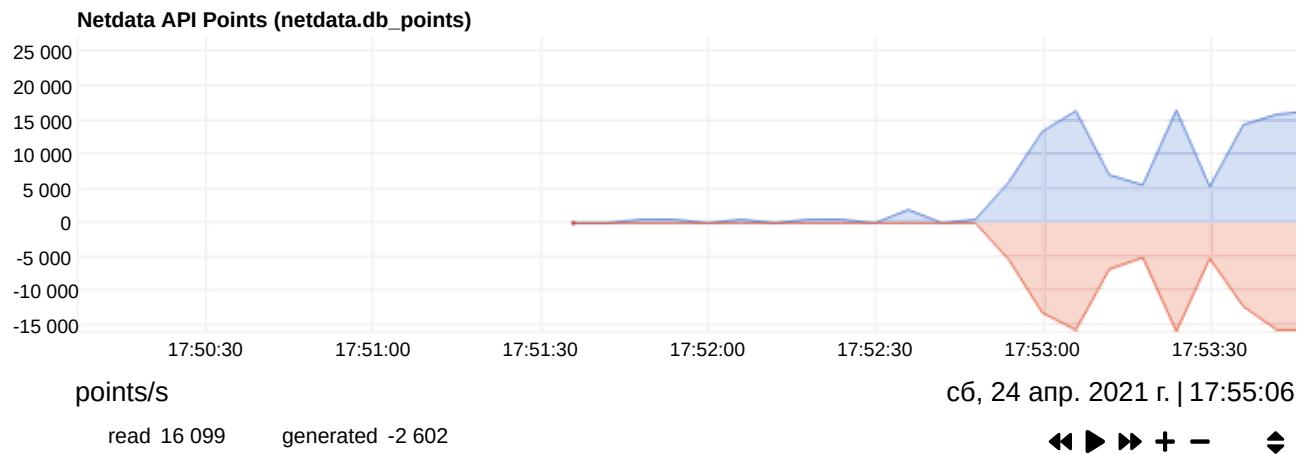


The netdata API response time measures the time netdata needed to serve requests. This time includes everything, from the reception of the first byte of a request, to the dispatch of the last byte of its reply, therefore it includes all network latencies involved (i.e. a client over a slow network will influence these metrics).

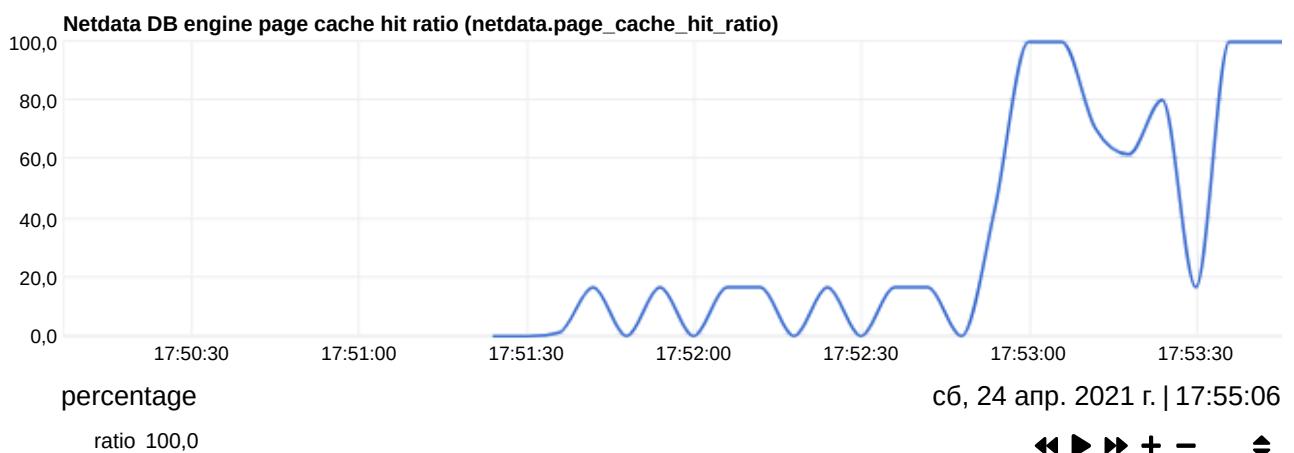
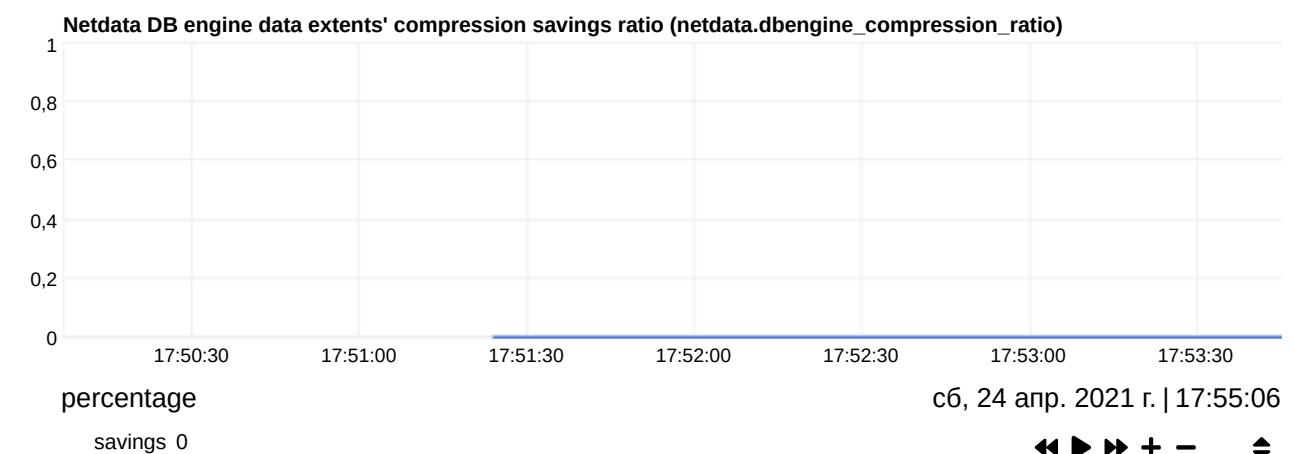


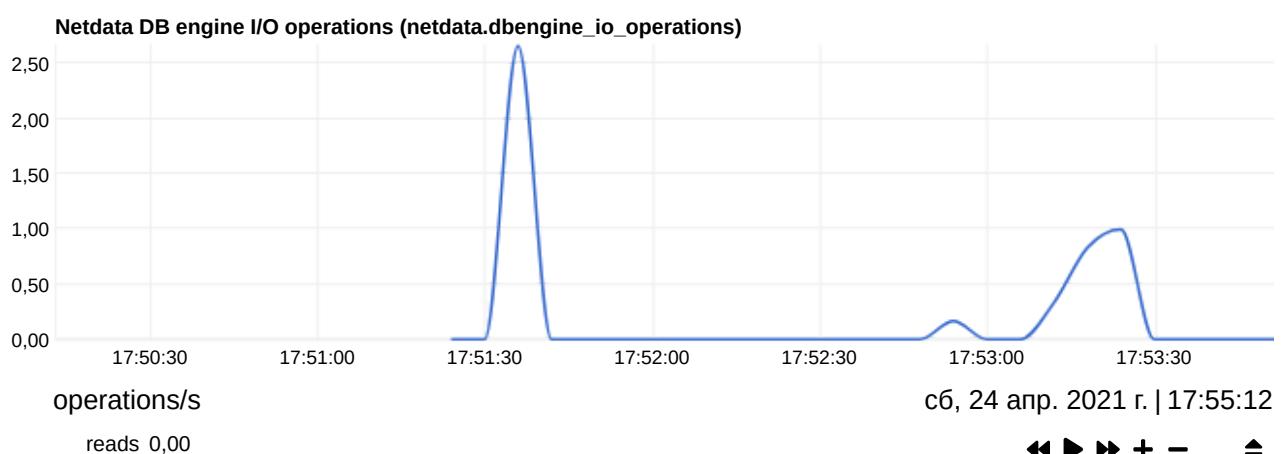
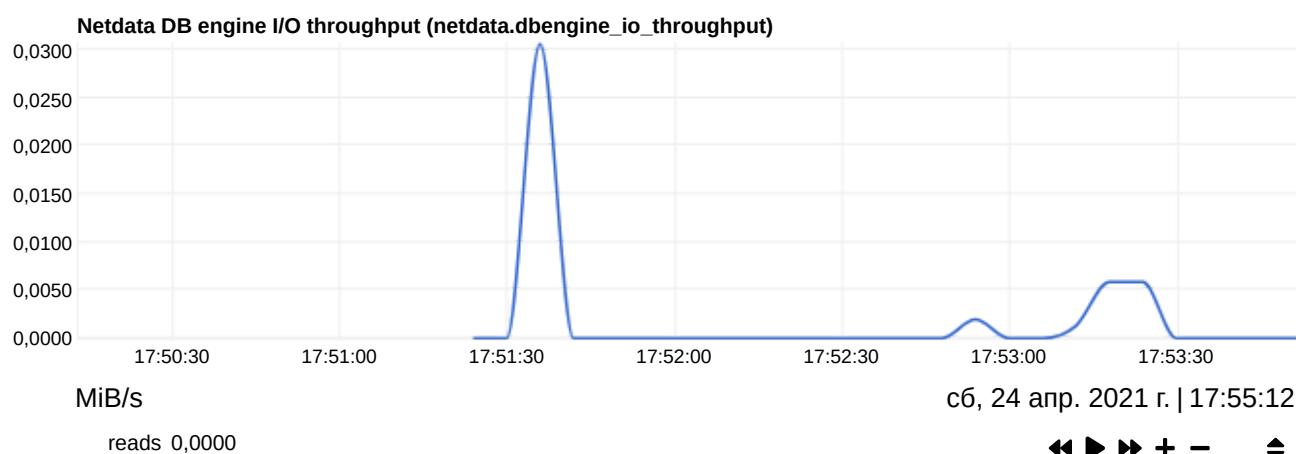
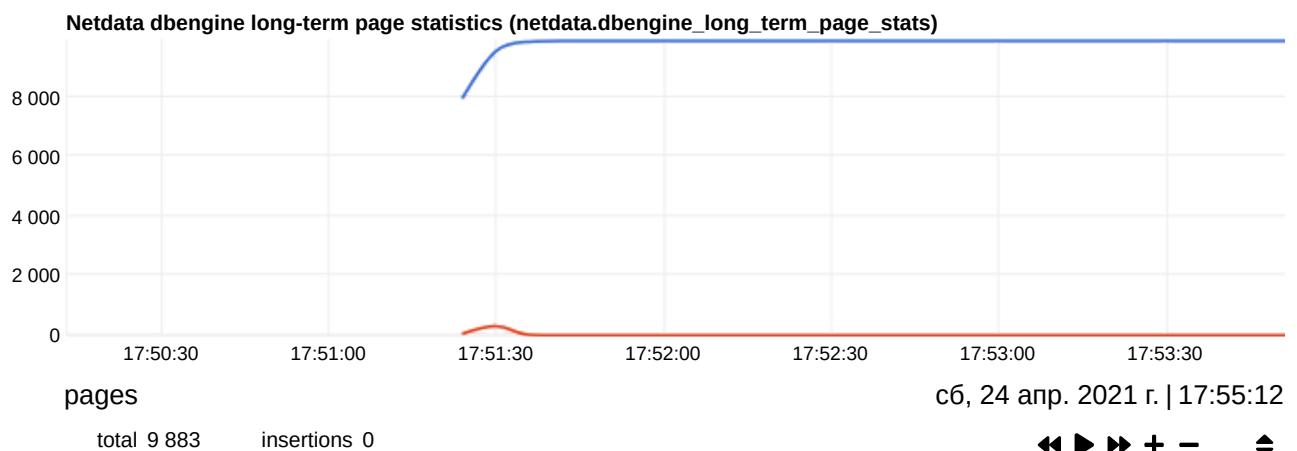
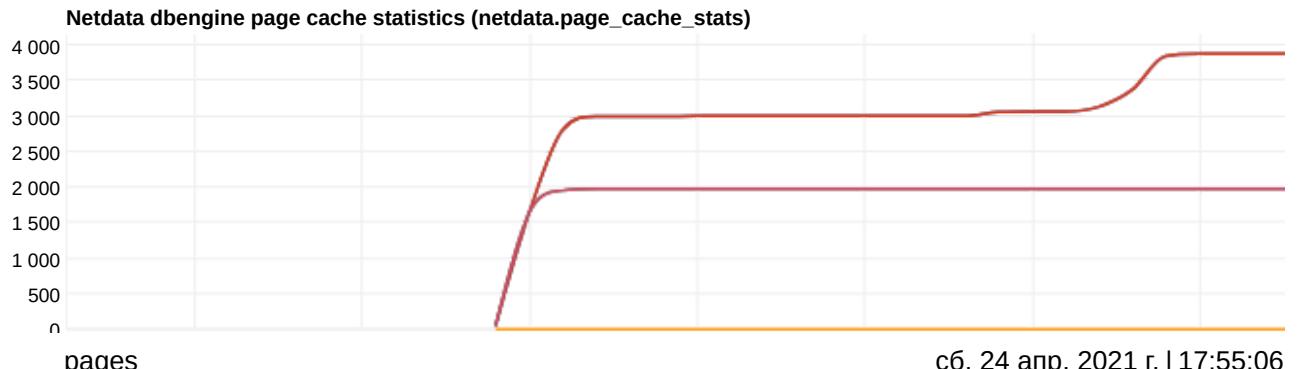
queries

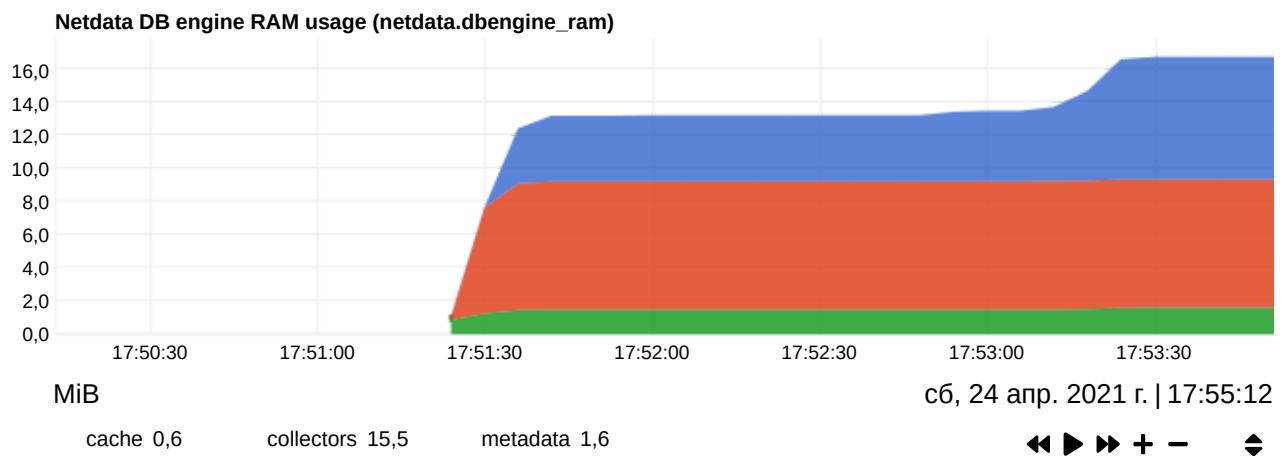
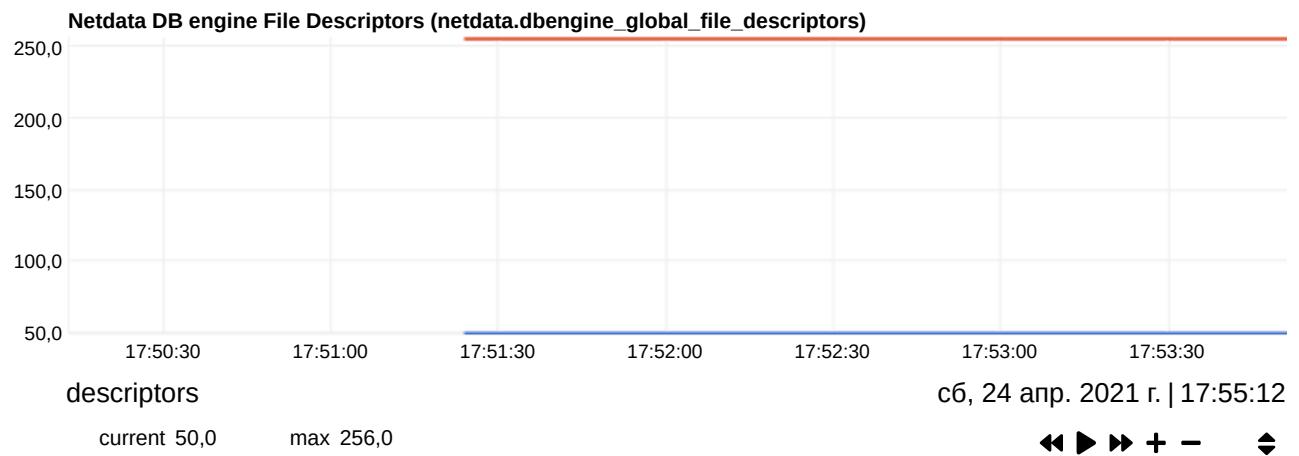
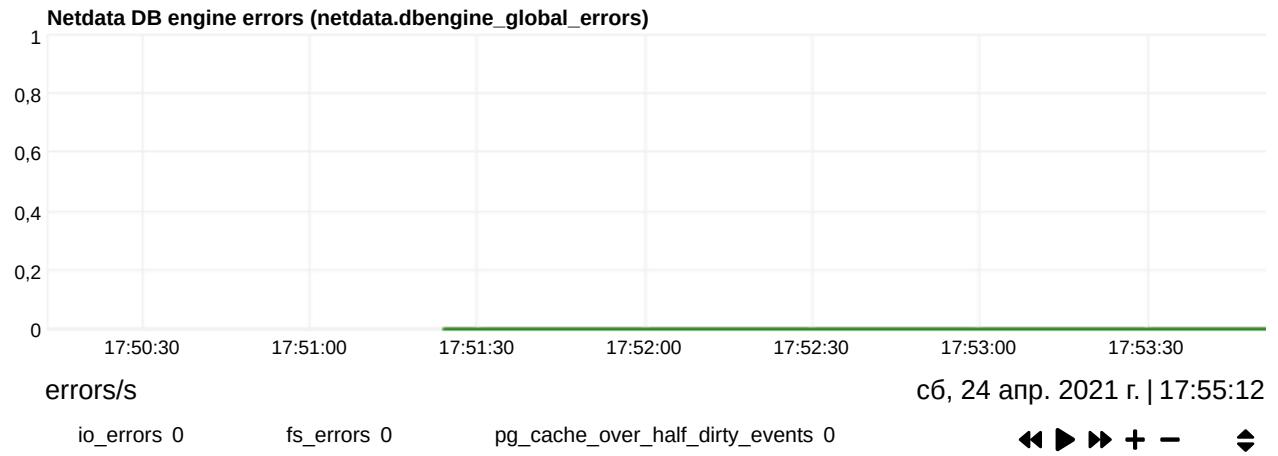




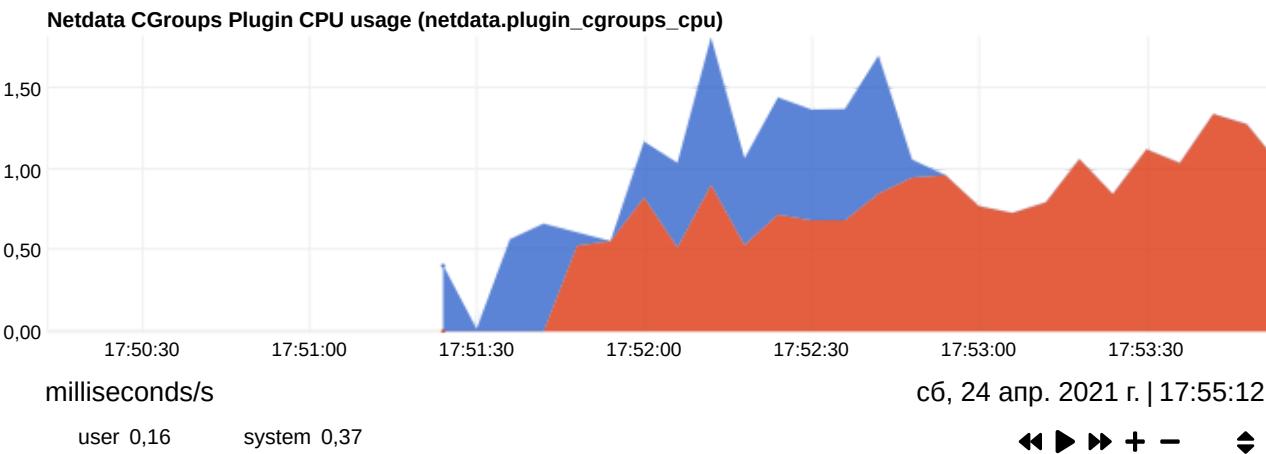
dbengine



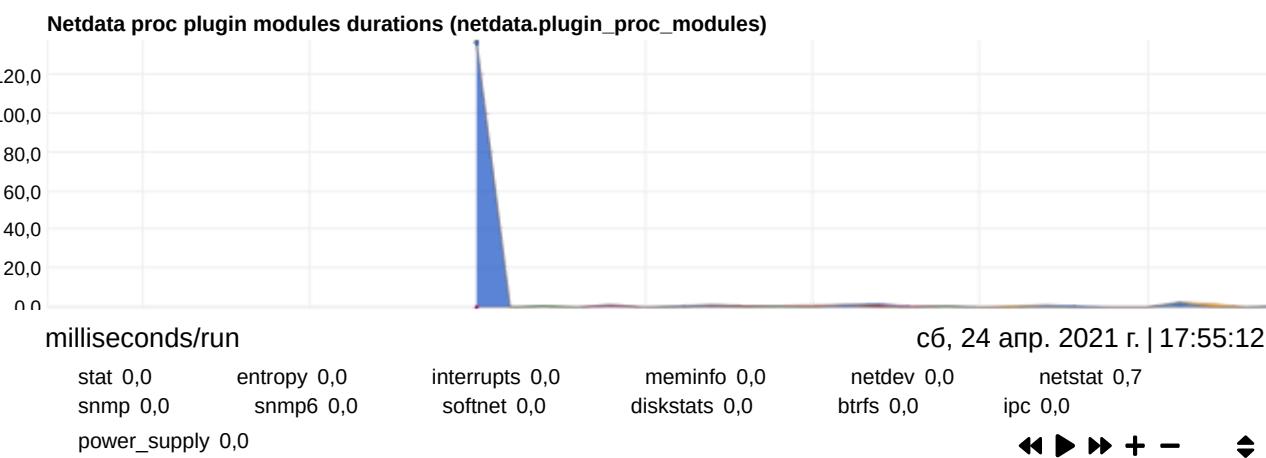
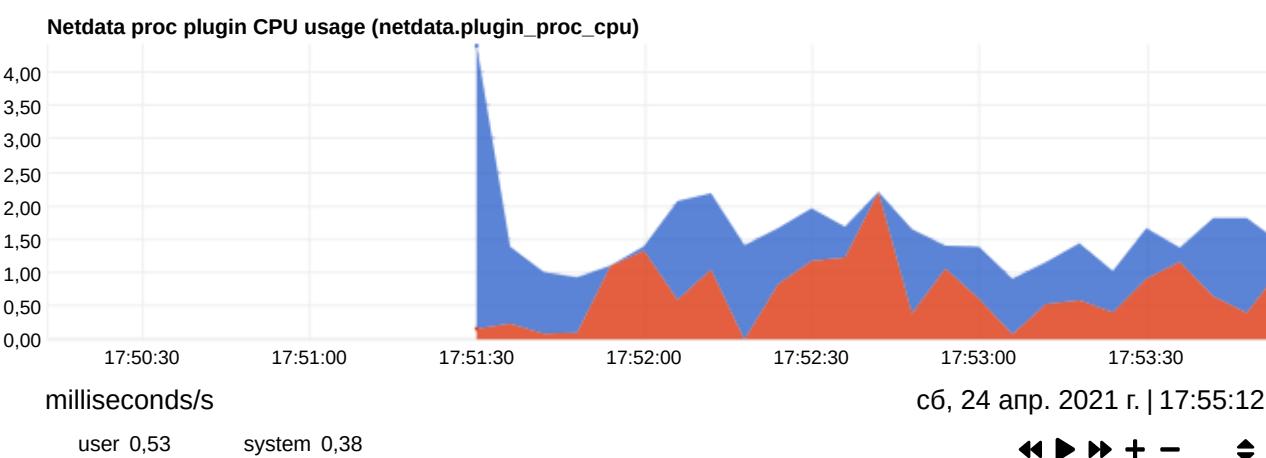




cgroups



proc

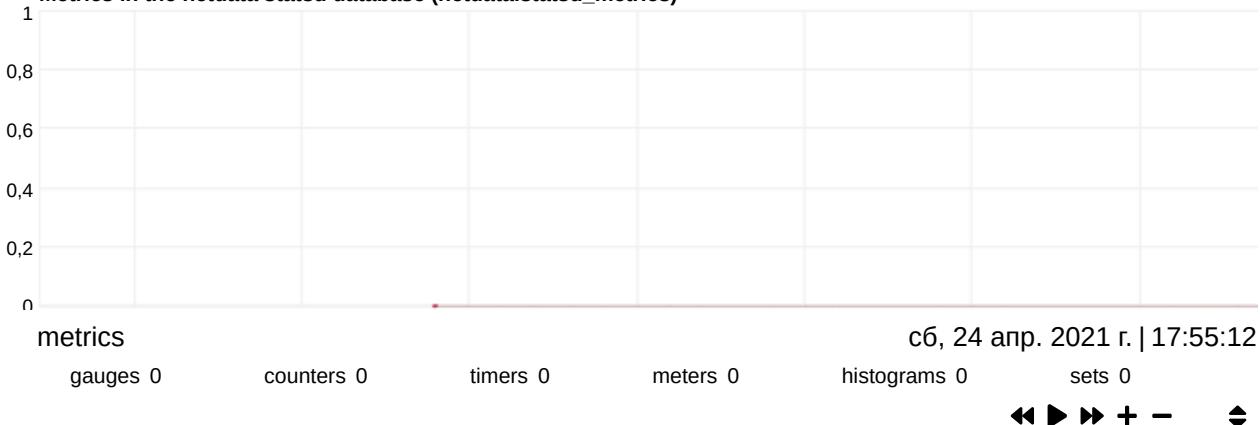
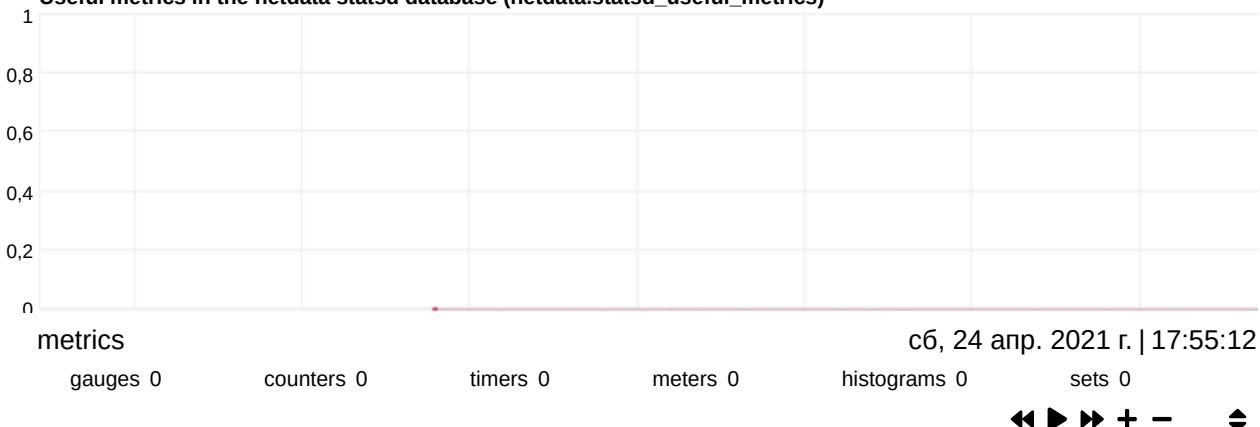
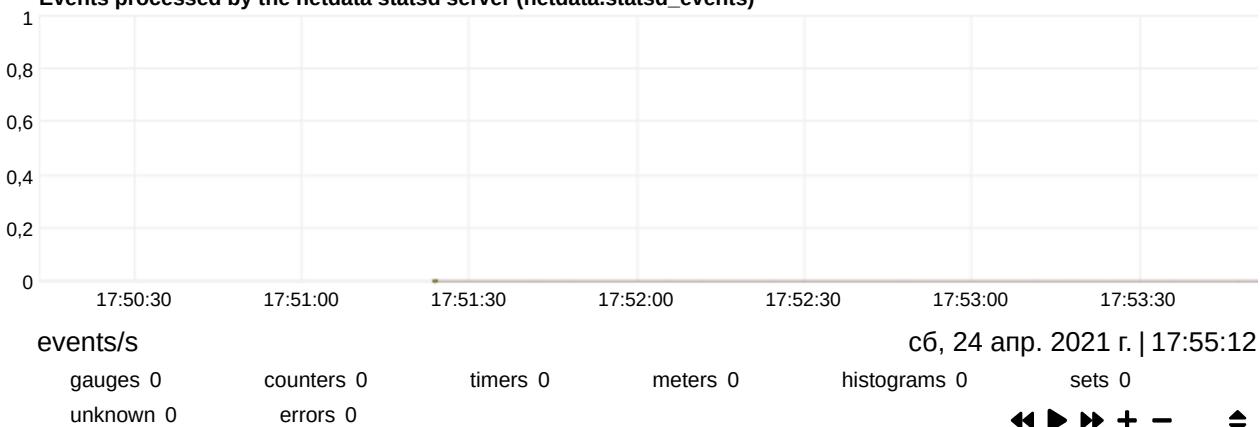


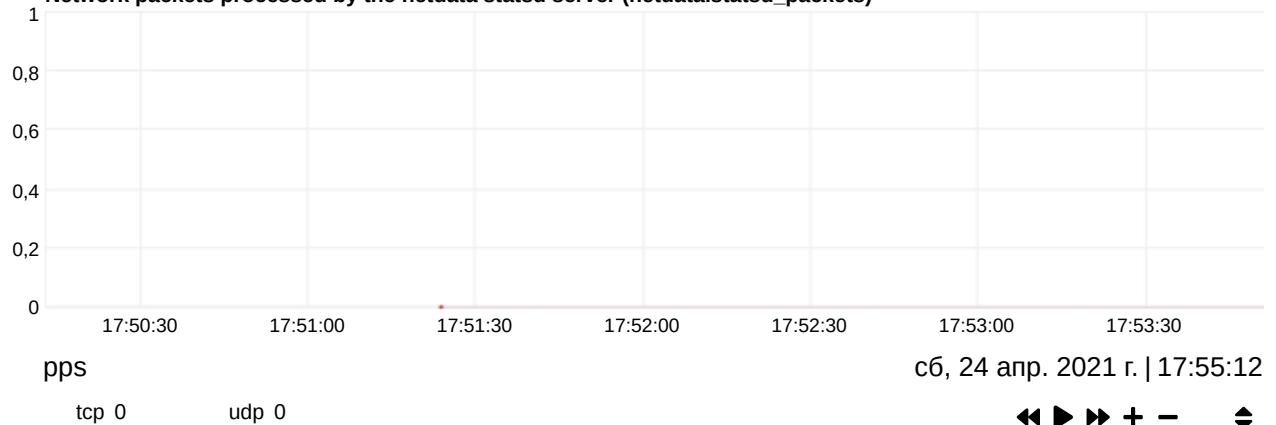
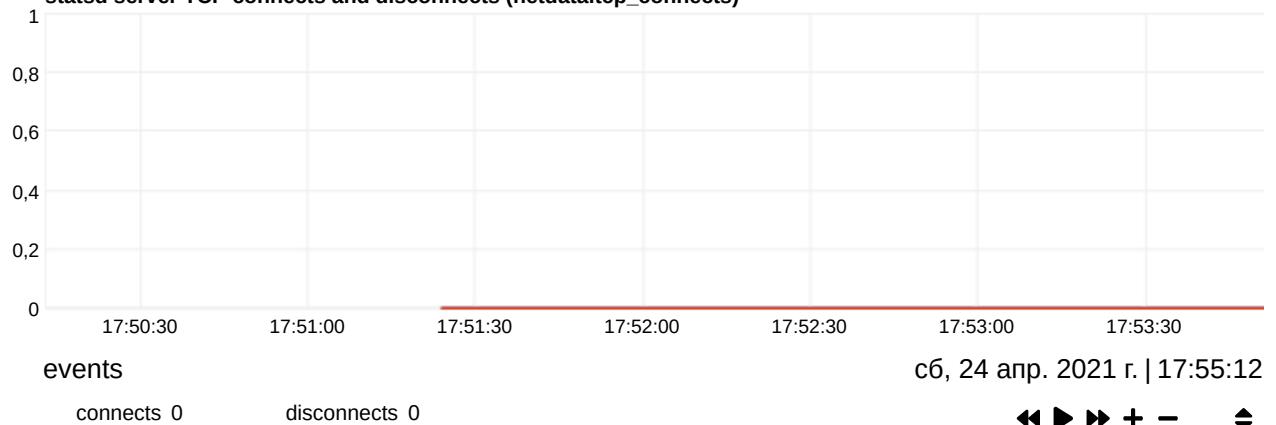
web

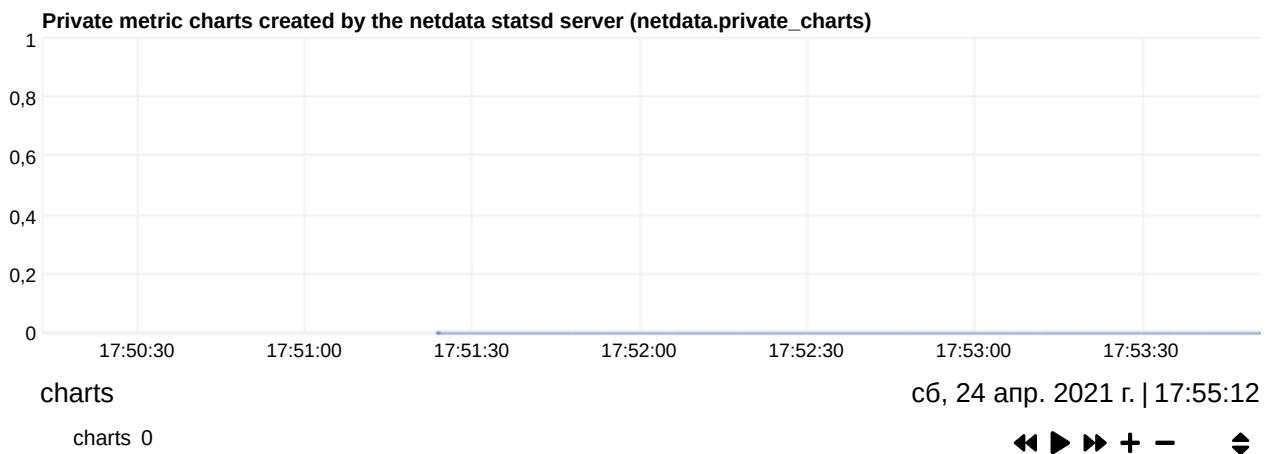
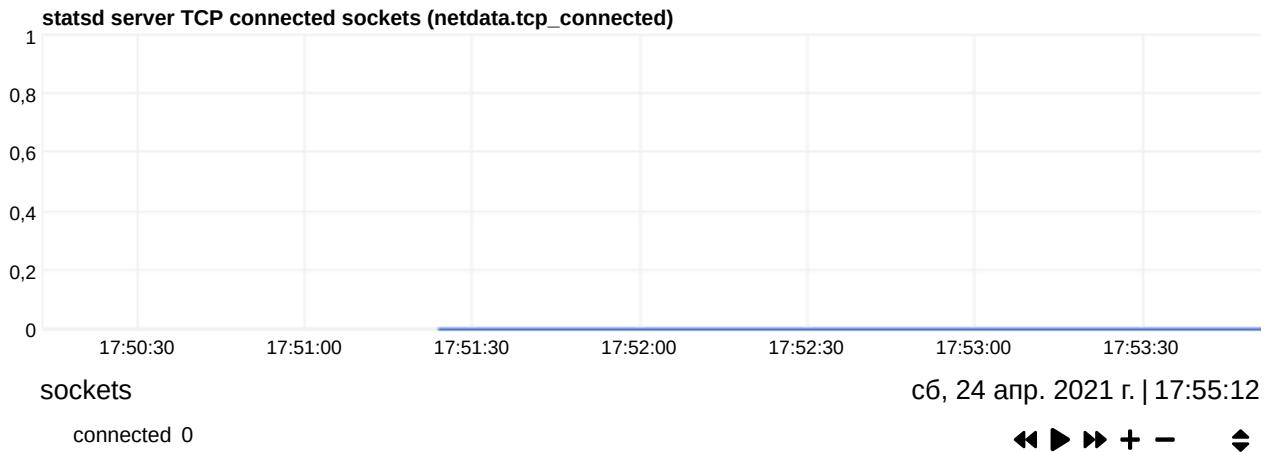
Netdata web server thread No 1 CPU usage (netdata.web_thread1_cpu)**Netdata web server thread No 2 CPU usage (netdata.web_thread2_cpu)**

statsd

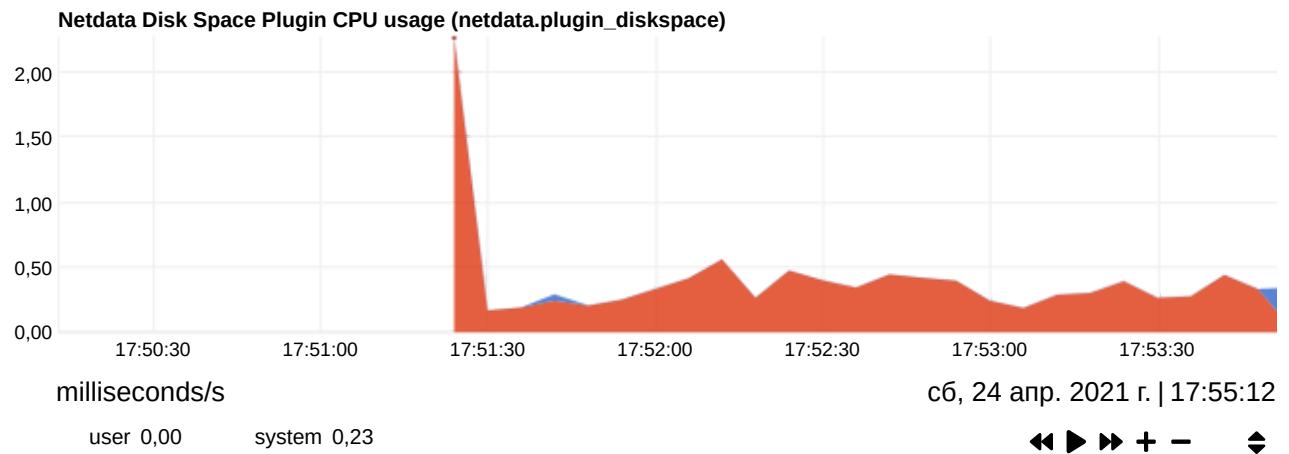
Netdata statsd charting thread CPU usage (netdata.plugin_statsd_charting_cpu)

NetData statsd collector thread No 1 CPU usage (netdata.plugin_statsd_collector1_cpu)**Metrics in the netdata statsd database (netdata.statsd_metrics)****Useful metrics in the netdata statsd database (netdata.statsd_useful_metrics)****Events processed by the netdata statsd server (netdata.statsd_events)**

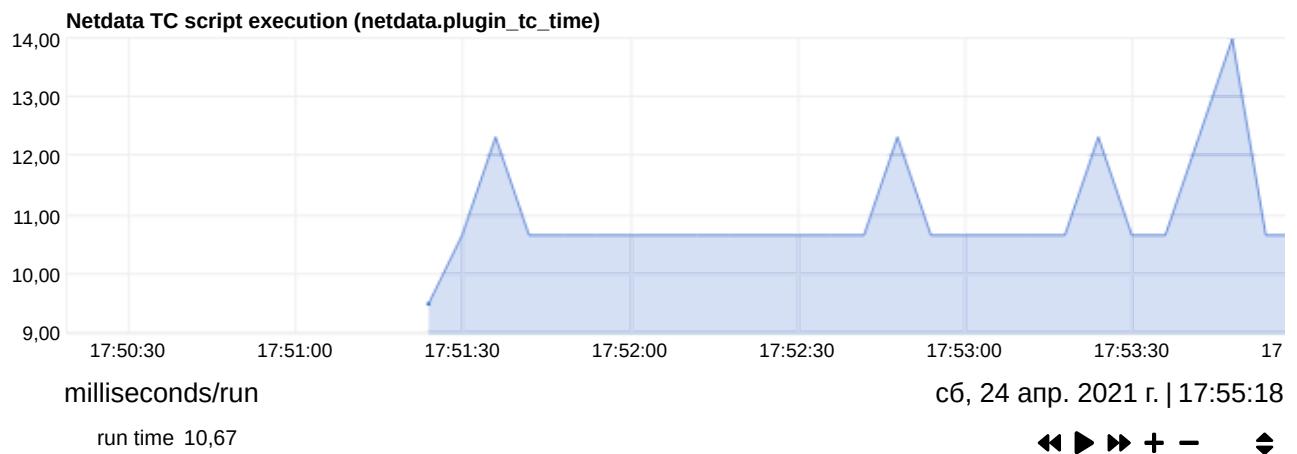
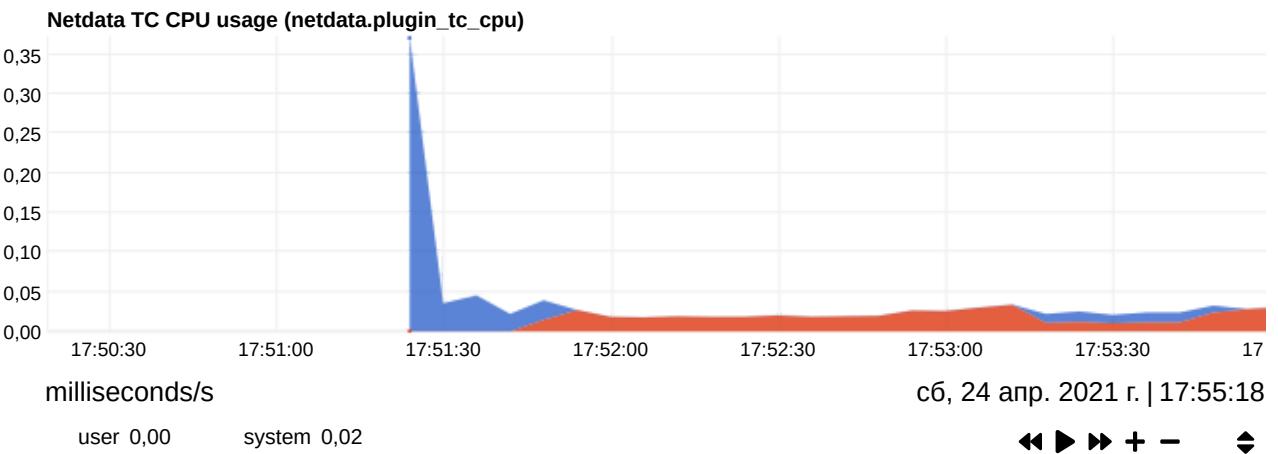
Read operations made by the netdata statsd server (netdata.statsd_reads)**Bytes read by the netdata statsd server (netdata.statsd_bytes)****Network packets processed by the netdata statsd server (netdata.statsd_packets)****statsd server TCP connects and disconnects (netdata.tcp_connects)**



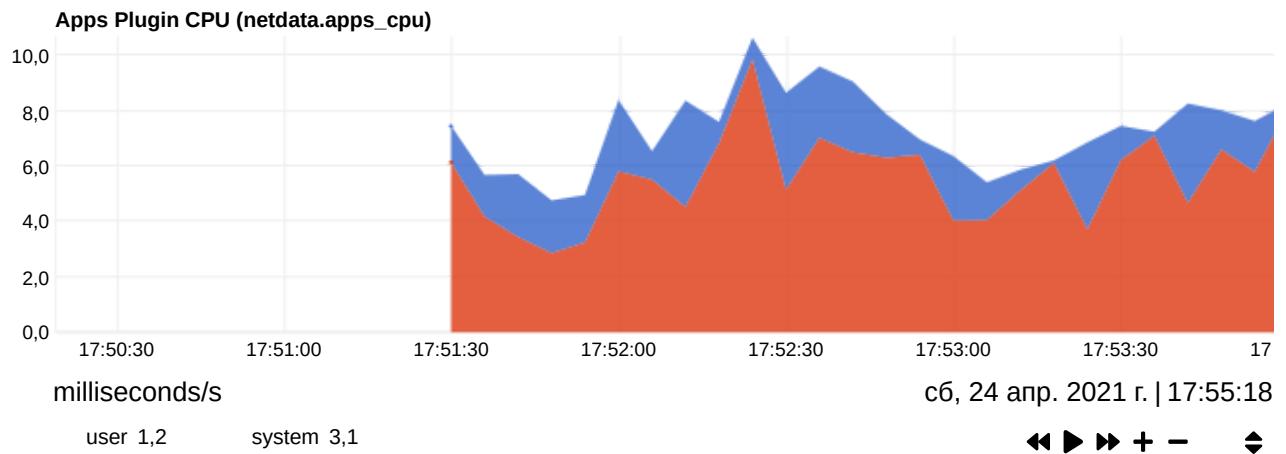
diskspace

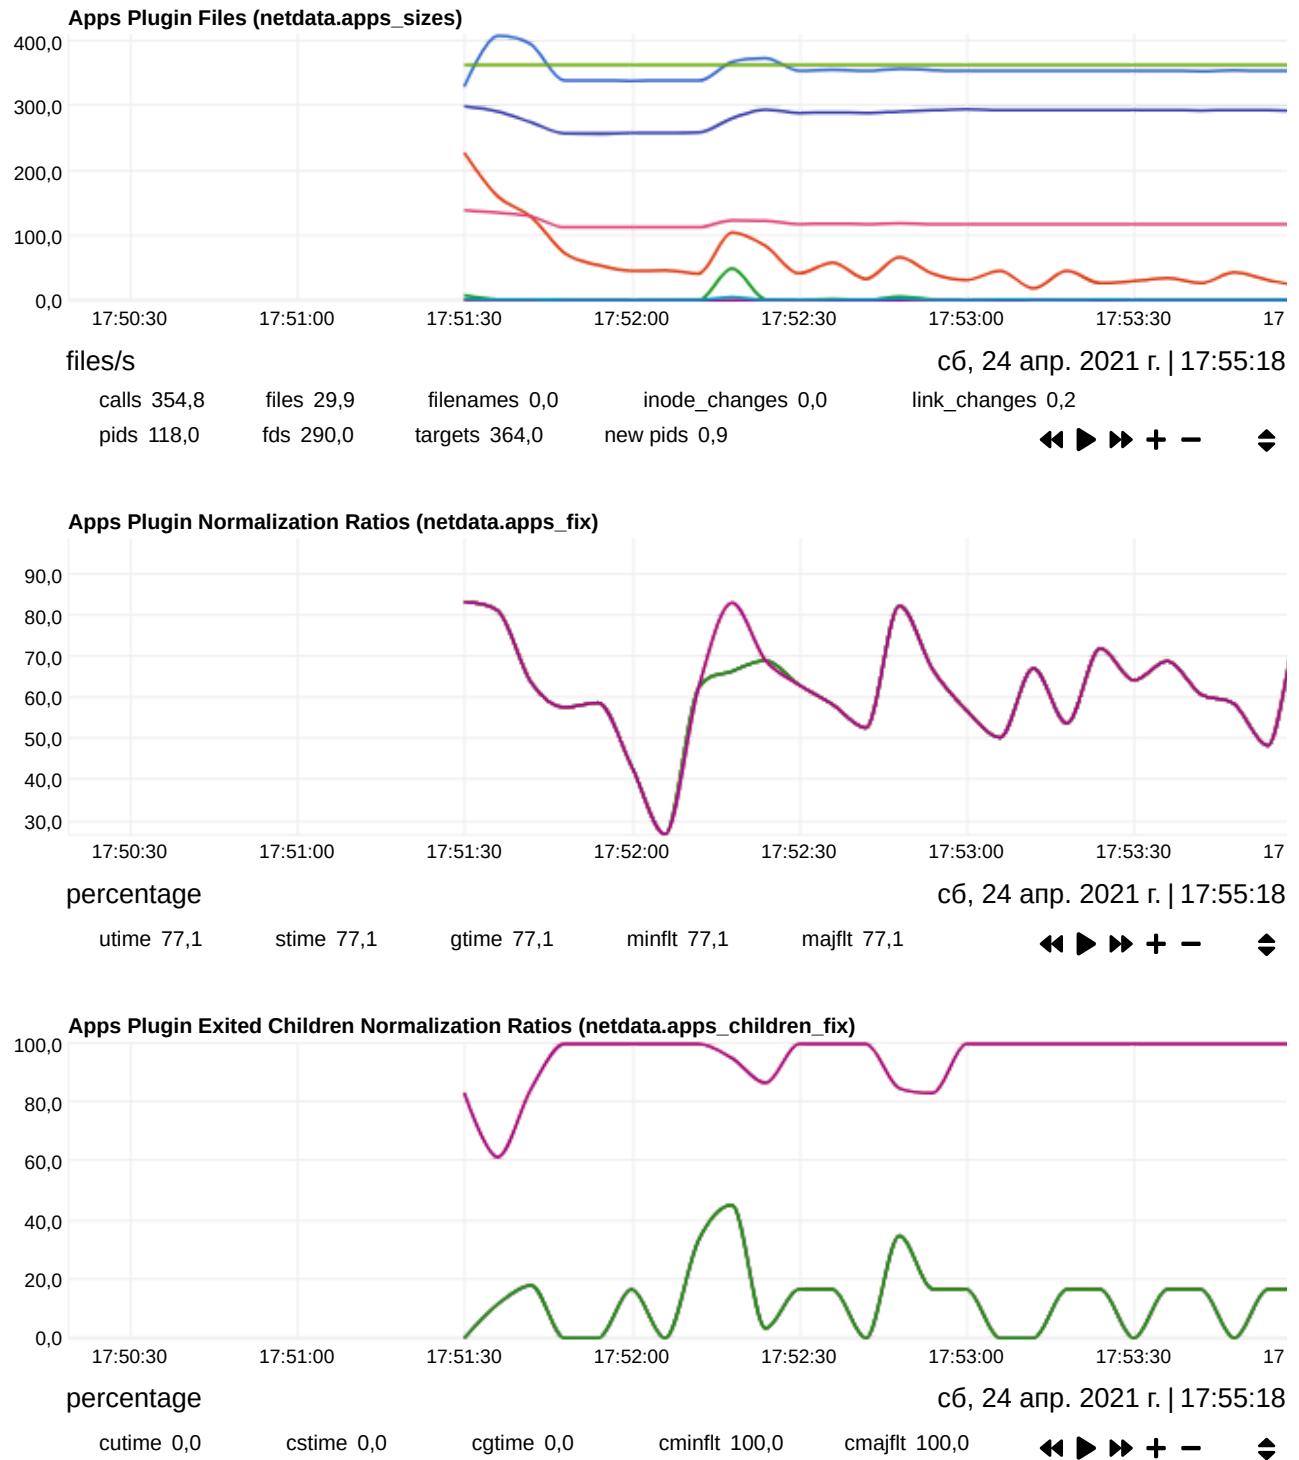


Netdata Disk Space Plugin Duration (netdata.plugin_diskspace_dt)**timex****Netdata Timex Plugin CPU usage (netdata.plugin_timex)****Netdata Timex Plugin Duration (netdata.plugin_timex_dt)****tc.helper**



apps.plugin





aclk

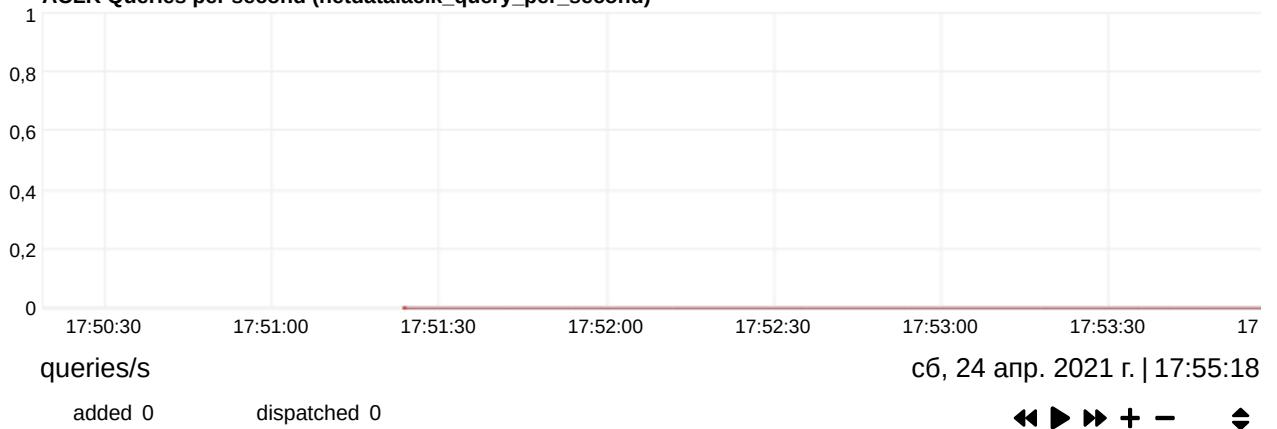
This chart shows if ACLK was online during entirety of the sample duration.

ACLK/Cloud connection status (netdata.aclk_status)

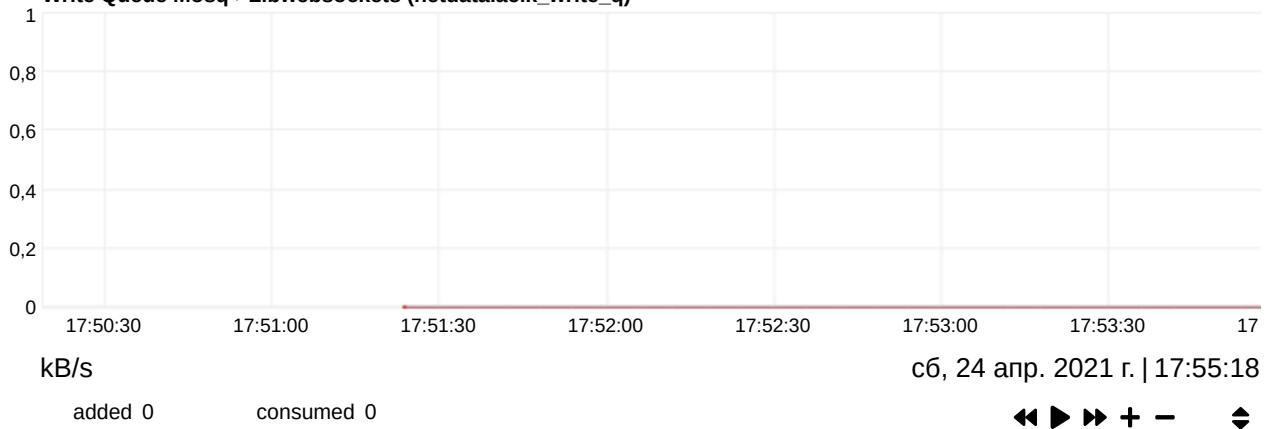


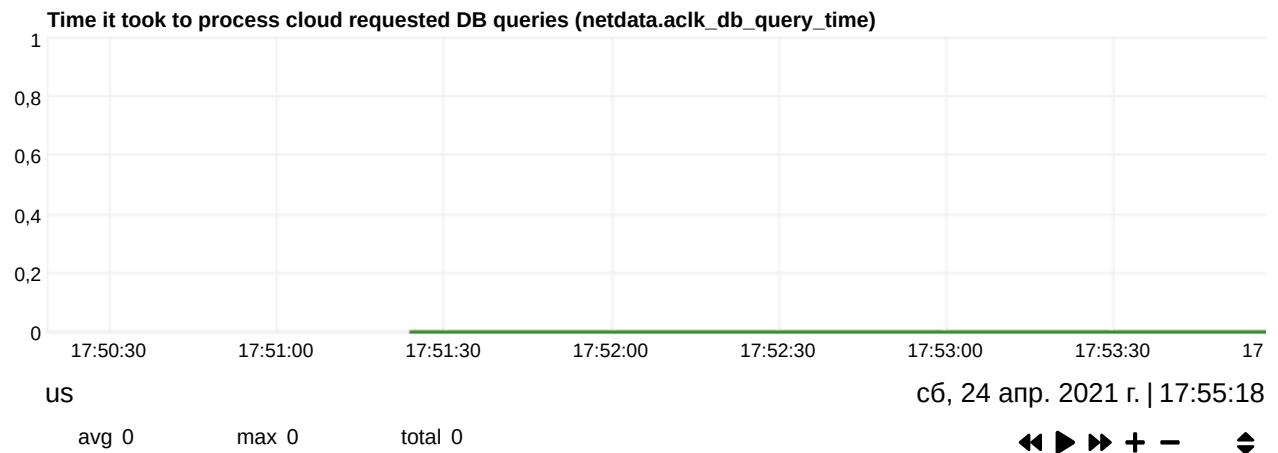
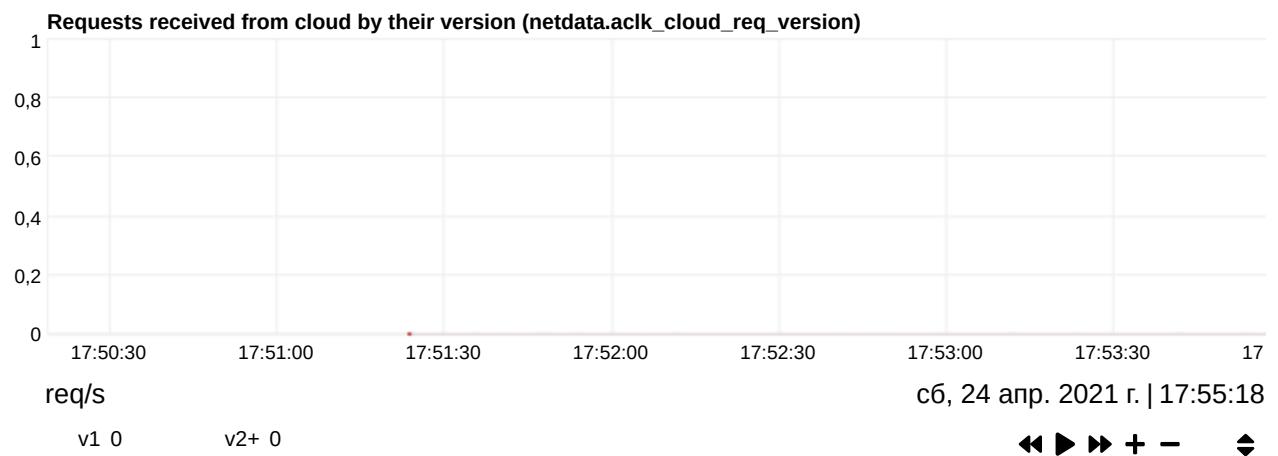
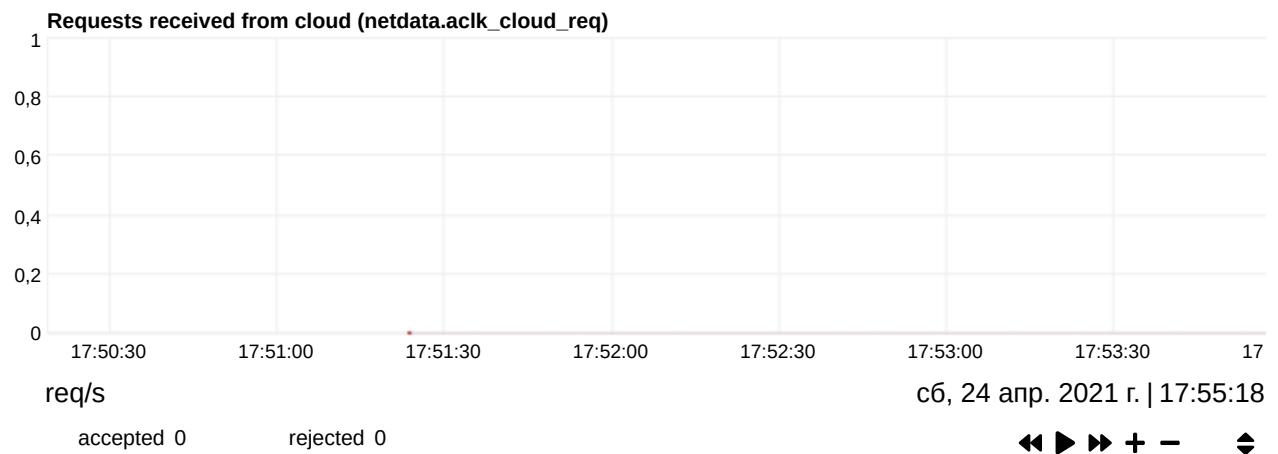
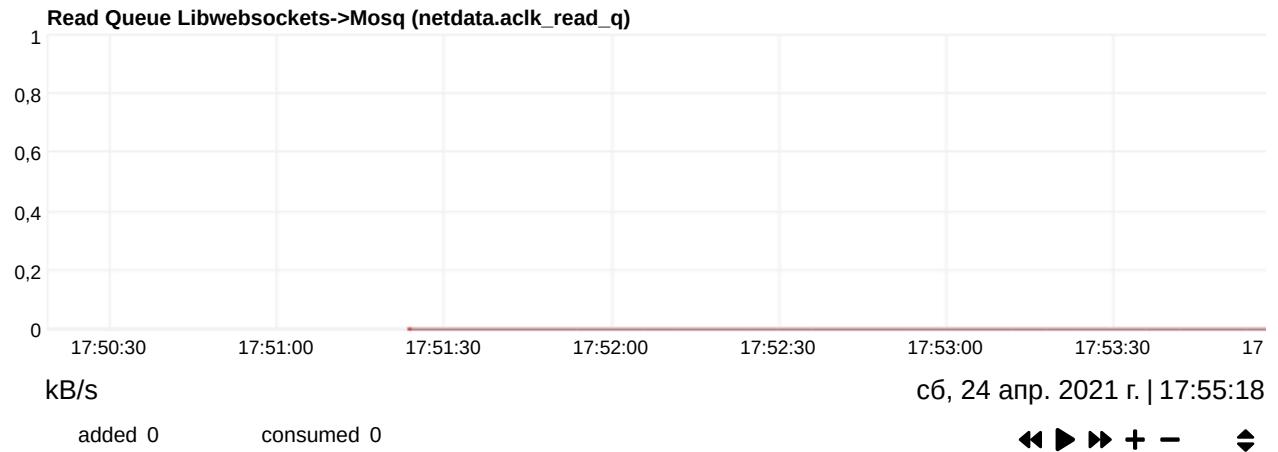
This chart shows how many queries were added for ACLK_query thread to process and how many it was actually able to process.

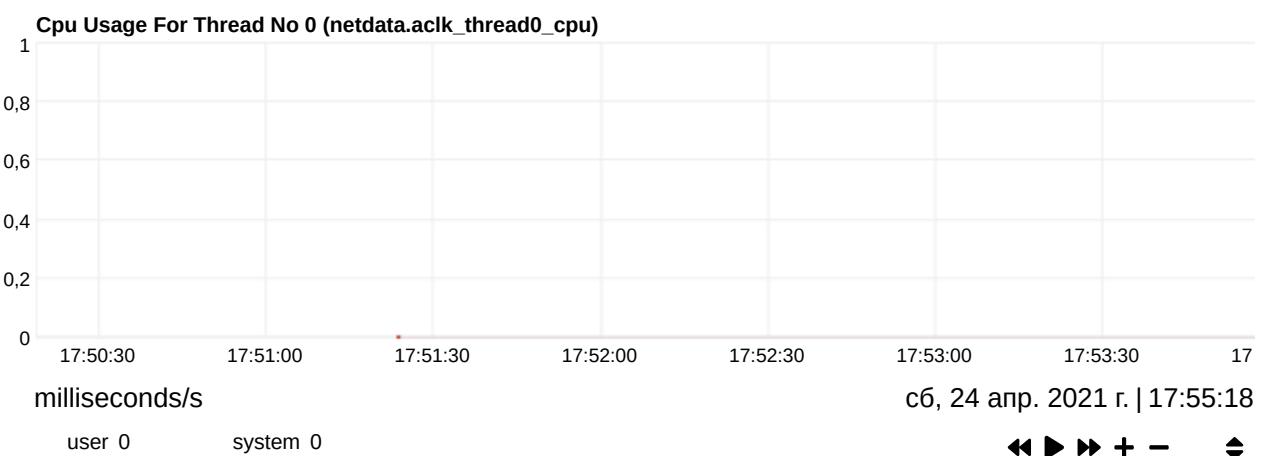
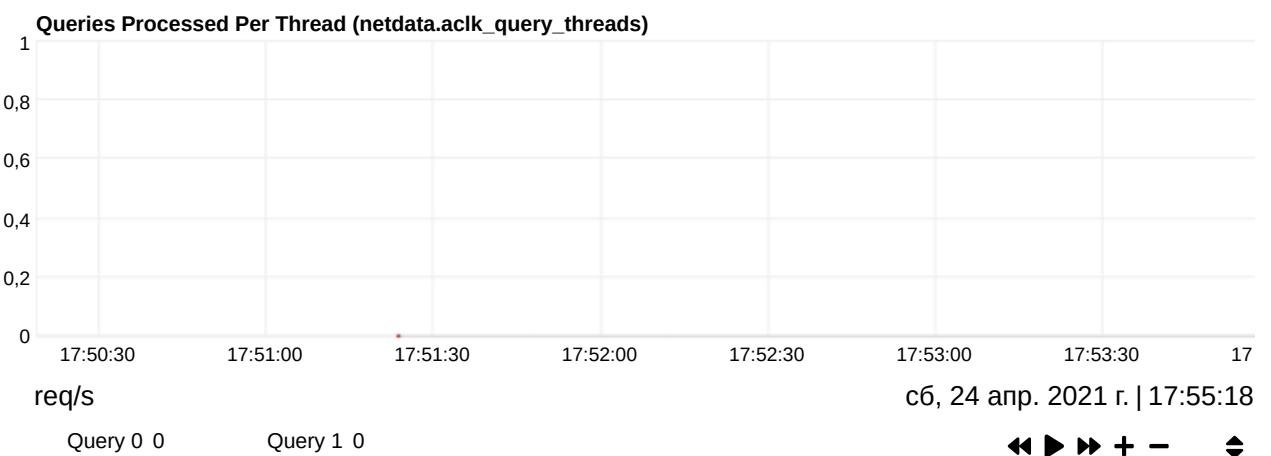
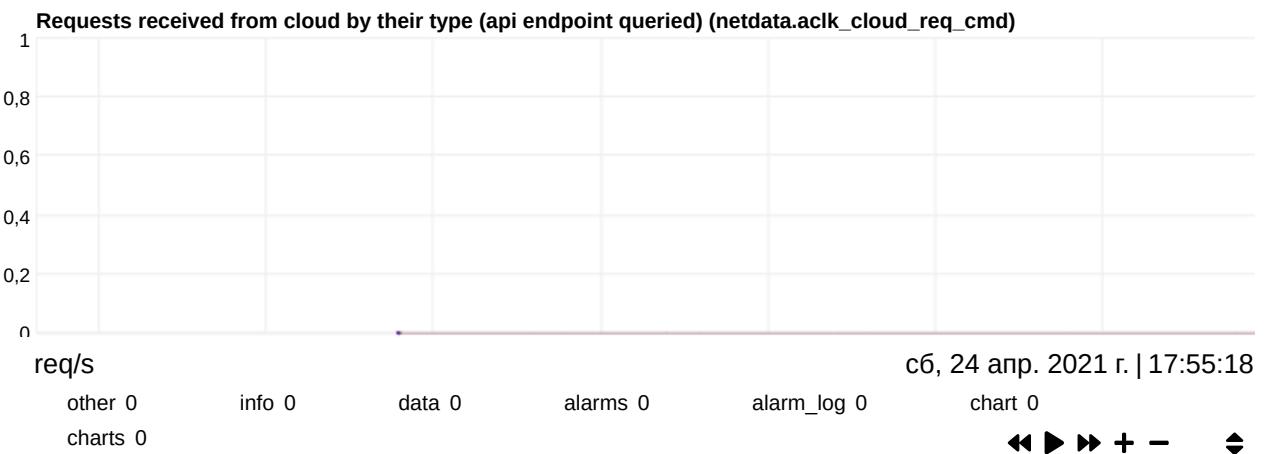
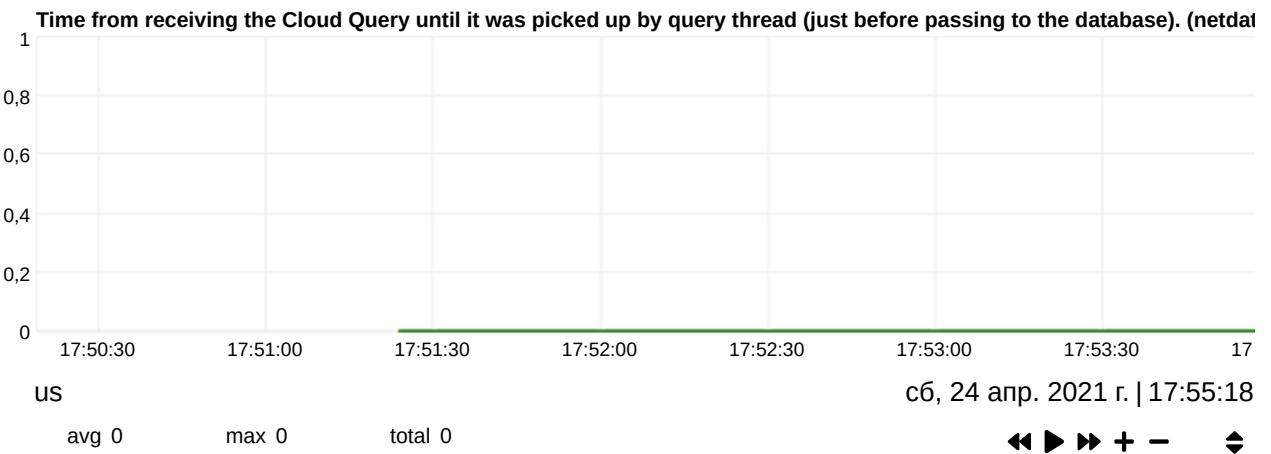
ACLK Queries per second (netdata.aclk_query_per_second)

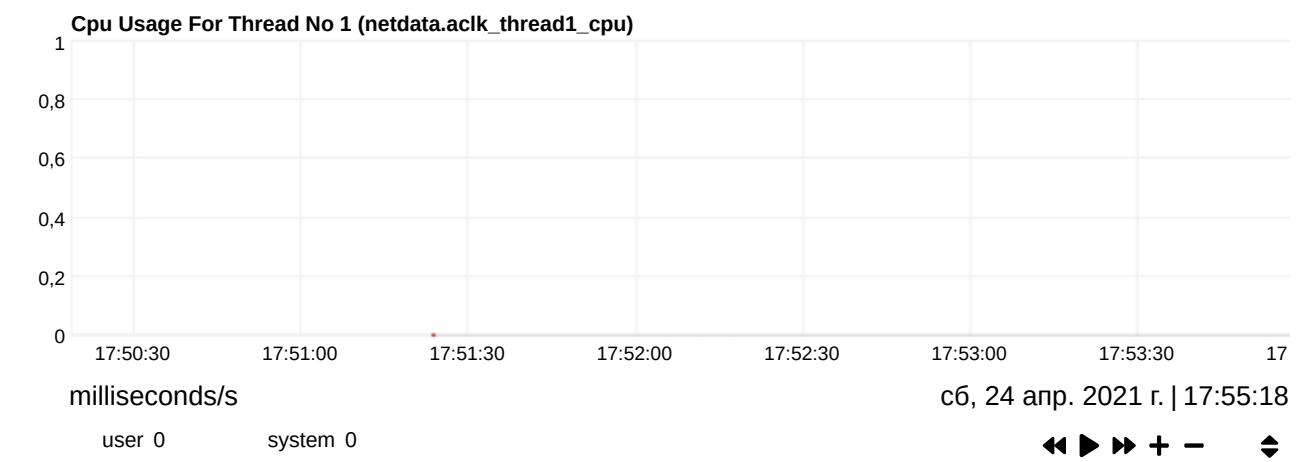


Write Queue Mosq->Libwebsockets (netdata.aclk_write_q)









Netdata (<https://github.com/netdata/netdata/wiki>)

© Copyright 2020, Netdata, Inc (<mailto:info@netdata.cloud>).

[Terms and conditions](https://www.netdata.cloud/terms/) (<https://www.netdata.cloud/terms/>).

[Privacy Policy](https://www.netdata.cloud/privacy/) (<https://www.netdata.cloud/privacy/>).

Released under GPL v3 or later (<http://www.gnu.org/licenses/gpl-3.0.en.html>). Netdata uses third party tools (<https://github.com/netdata/netdata/blob/master/REDISTRIBUTED.md>).